



US011860760B1

(12) **United States Patent**
Agarwal et al.(10) **Patent No.:** US 11,860,760 B1
(45) **Date of Patent:** Jan. 2, 2024(54) **AGGREGATING METRICS FOR WORKFLOWS ASSOCIATED WITH A REAL USER SESSION**(71) Applicant: **Splunk Inc.**, San Francisco, CA (US)(72) Inventors: **Mayank Agarwal**, Mountain View, CA (US); **Shashwat Sehgal**, Belmont, CA (US); **Justin Smith**, San Francisco, CA (US); **David Tagatac**, San Diego, CA (US); **Rashmi Kalyani Vasudevan**, San Francisco, CA (US); **John Bennett Wundes**, Santa Rosa, CA (US)(73) Assignee: **SPLUNK Inc.**, San Francisco, CA (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 322 days.

(21) Appl. No.: **17/390,614**(22) Filed: **Jul. 30, 2021**(51) **Int. Cl.**

G06F 16/00 (2019.01)
G06F 11/34 (2006.01)
G06F 11/32 (2006.01)
G06F 16/23 (2019.01)

(52) **U.S. Cl.**

CPC **G06F 11/3495** (2013.01); **G06F 11/328** (2013.01); **G06F 11/3438** (2013.01); **G06F 16/2365** (2019.01)

(58) **Field of Classification Search**

CPC G06F 16/2365
See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

7,937,344 B2 5/2011 Baum et al.
8,112,425 B2 2/2012 Baum et al.
8,751,529 B2 6/2014 Zhang et al.

8,788,525 B2 7/2014 Neels et al.
9,215,240 B2 12/2015 Merza et al.
9,286,413 B1 3/2016 Coates et al.
10,127,258 B2 11/2018 Lamas et al.
2019/0098106 A1 3/2019 Mungel et al.
2020/0334416 A1* 10/2020 Vianu G06N 3/044

OTHER PUBLICATIONS

Splunk Enterprise 8.0.0 Overview, available online, retrieved May 20, 2020 from docs.splunk.com.

Splunk Cloud 8.0.2004 User Manual, available online, retrieved May 20, 2020 from docs.splunk.com.

Splunk Quick Reference Guide, updated 2019, available online at <https://www.splunk.com/pdfs/solution-guides/splunk-quick-reference-guide.pdf>, retrieved May 20, 2020.

Carraso, David, "Exploring Splunk," published by CITO Research, New York, NY, Apr. 2012.

Bitincka, Ledion et al., "Optimizing Data Analysis with a Semi-structured Time Series Database," self-published, first presented at "Workshop on Managing Systems via Log Analysis and Machine Learning Techniques (SLAML)", Vancouver, British Columbia, Oct. 3, 2010.

* cited by examiner

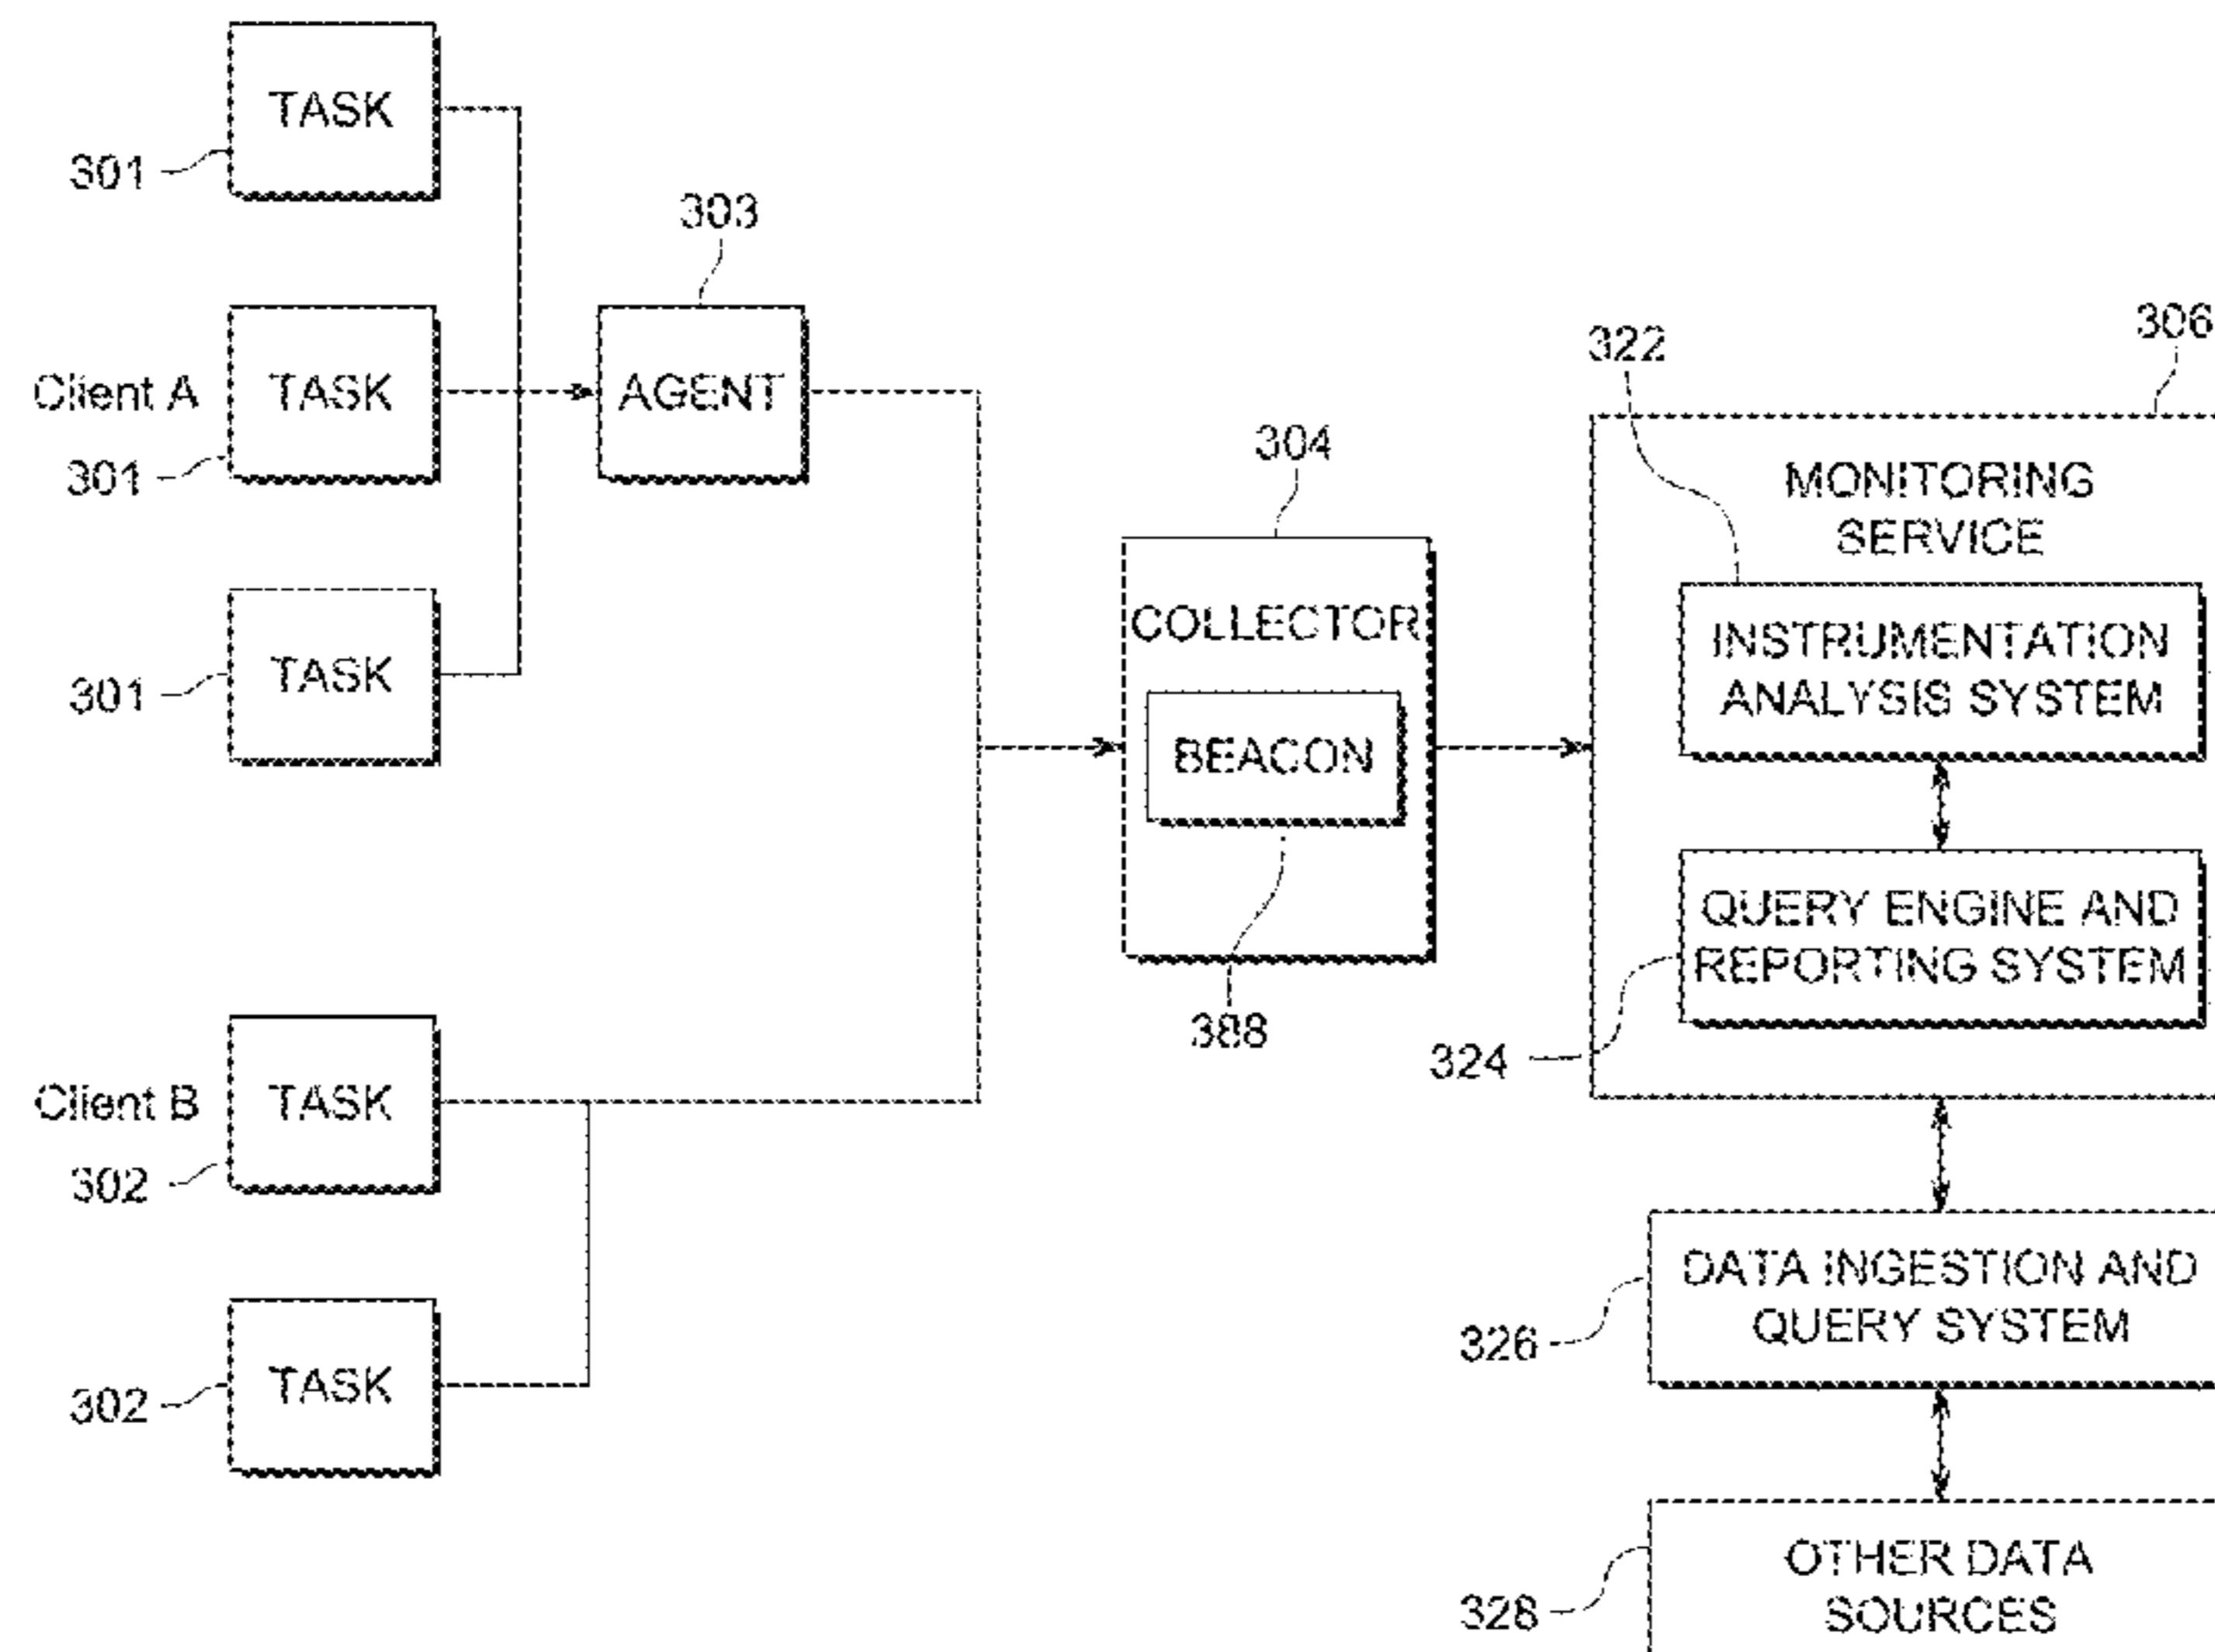
Primary Examiner — Ajith Jacob

(74) *Attorney, Agent, or Firm* — Kilpatrick Townsend & Stockton LLP

(57) **ABSTRACT**

A method of aggregating metrics associated with a user interaction during a real user session comprises identifying a span comprising a tag associated with a workflow from ingested spans associated with the real user session, where the workflow comprises spans generated in response to the user interaction. The method also comprises identifying other spans associated with the workflow from the ingested spans. The method further comprises grouping the other spans associated with the workflow with the tagged span and aggregating metrics for the workflow over a duration of time.

20 Claims, 35 Drawing Sheets



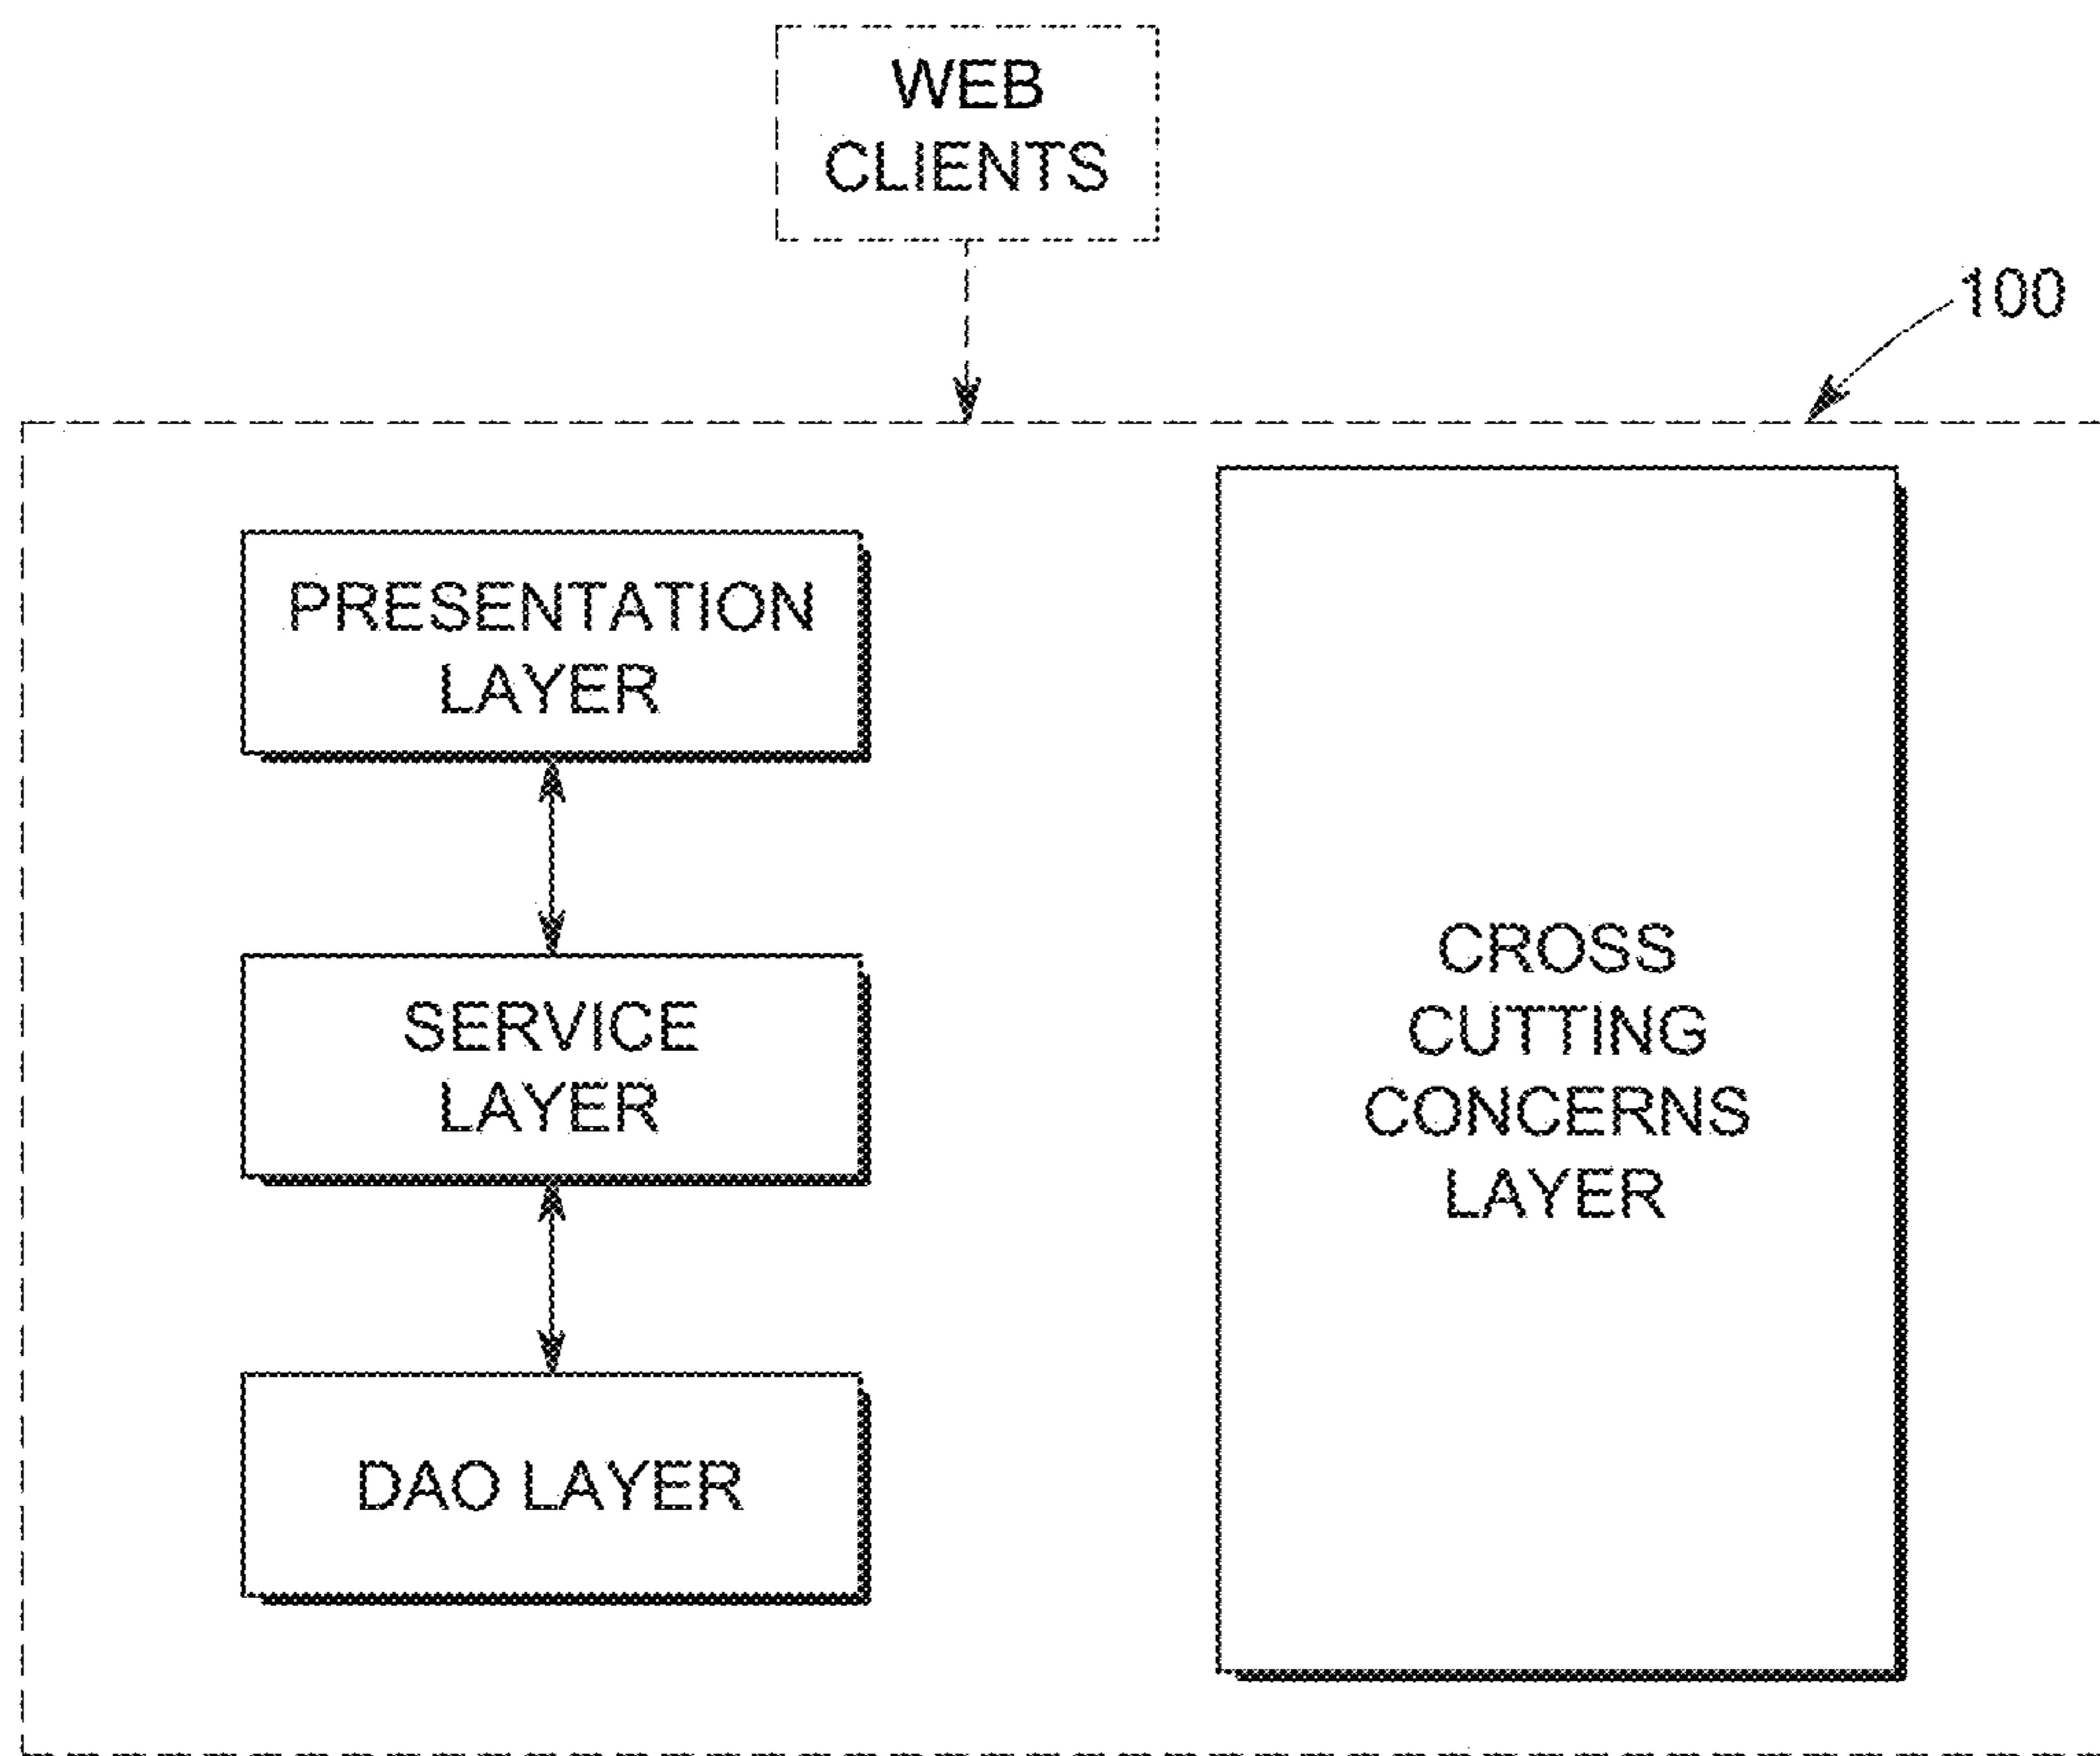


FIG. 1A

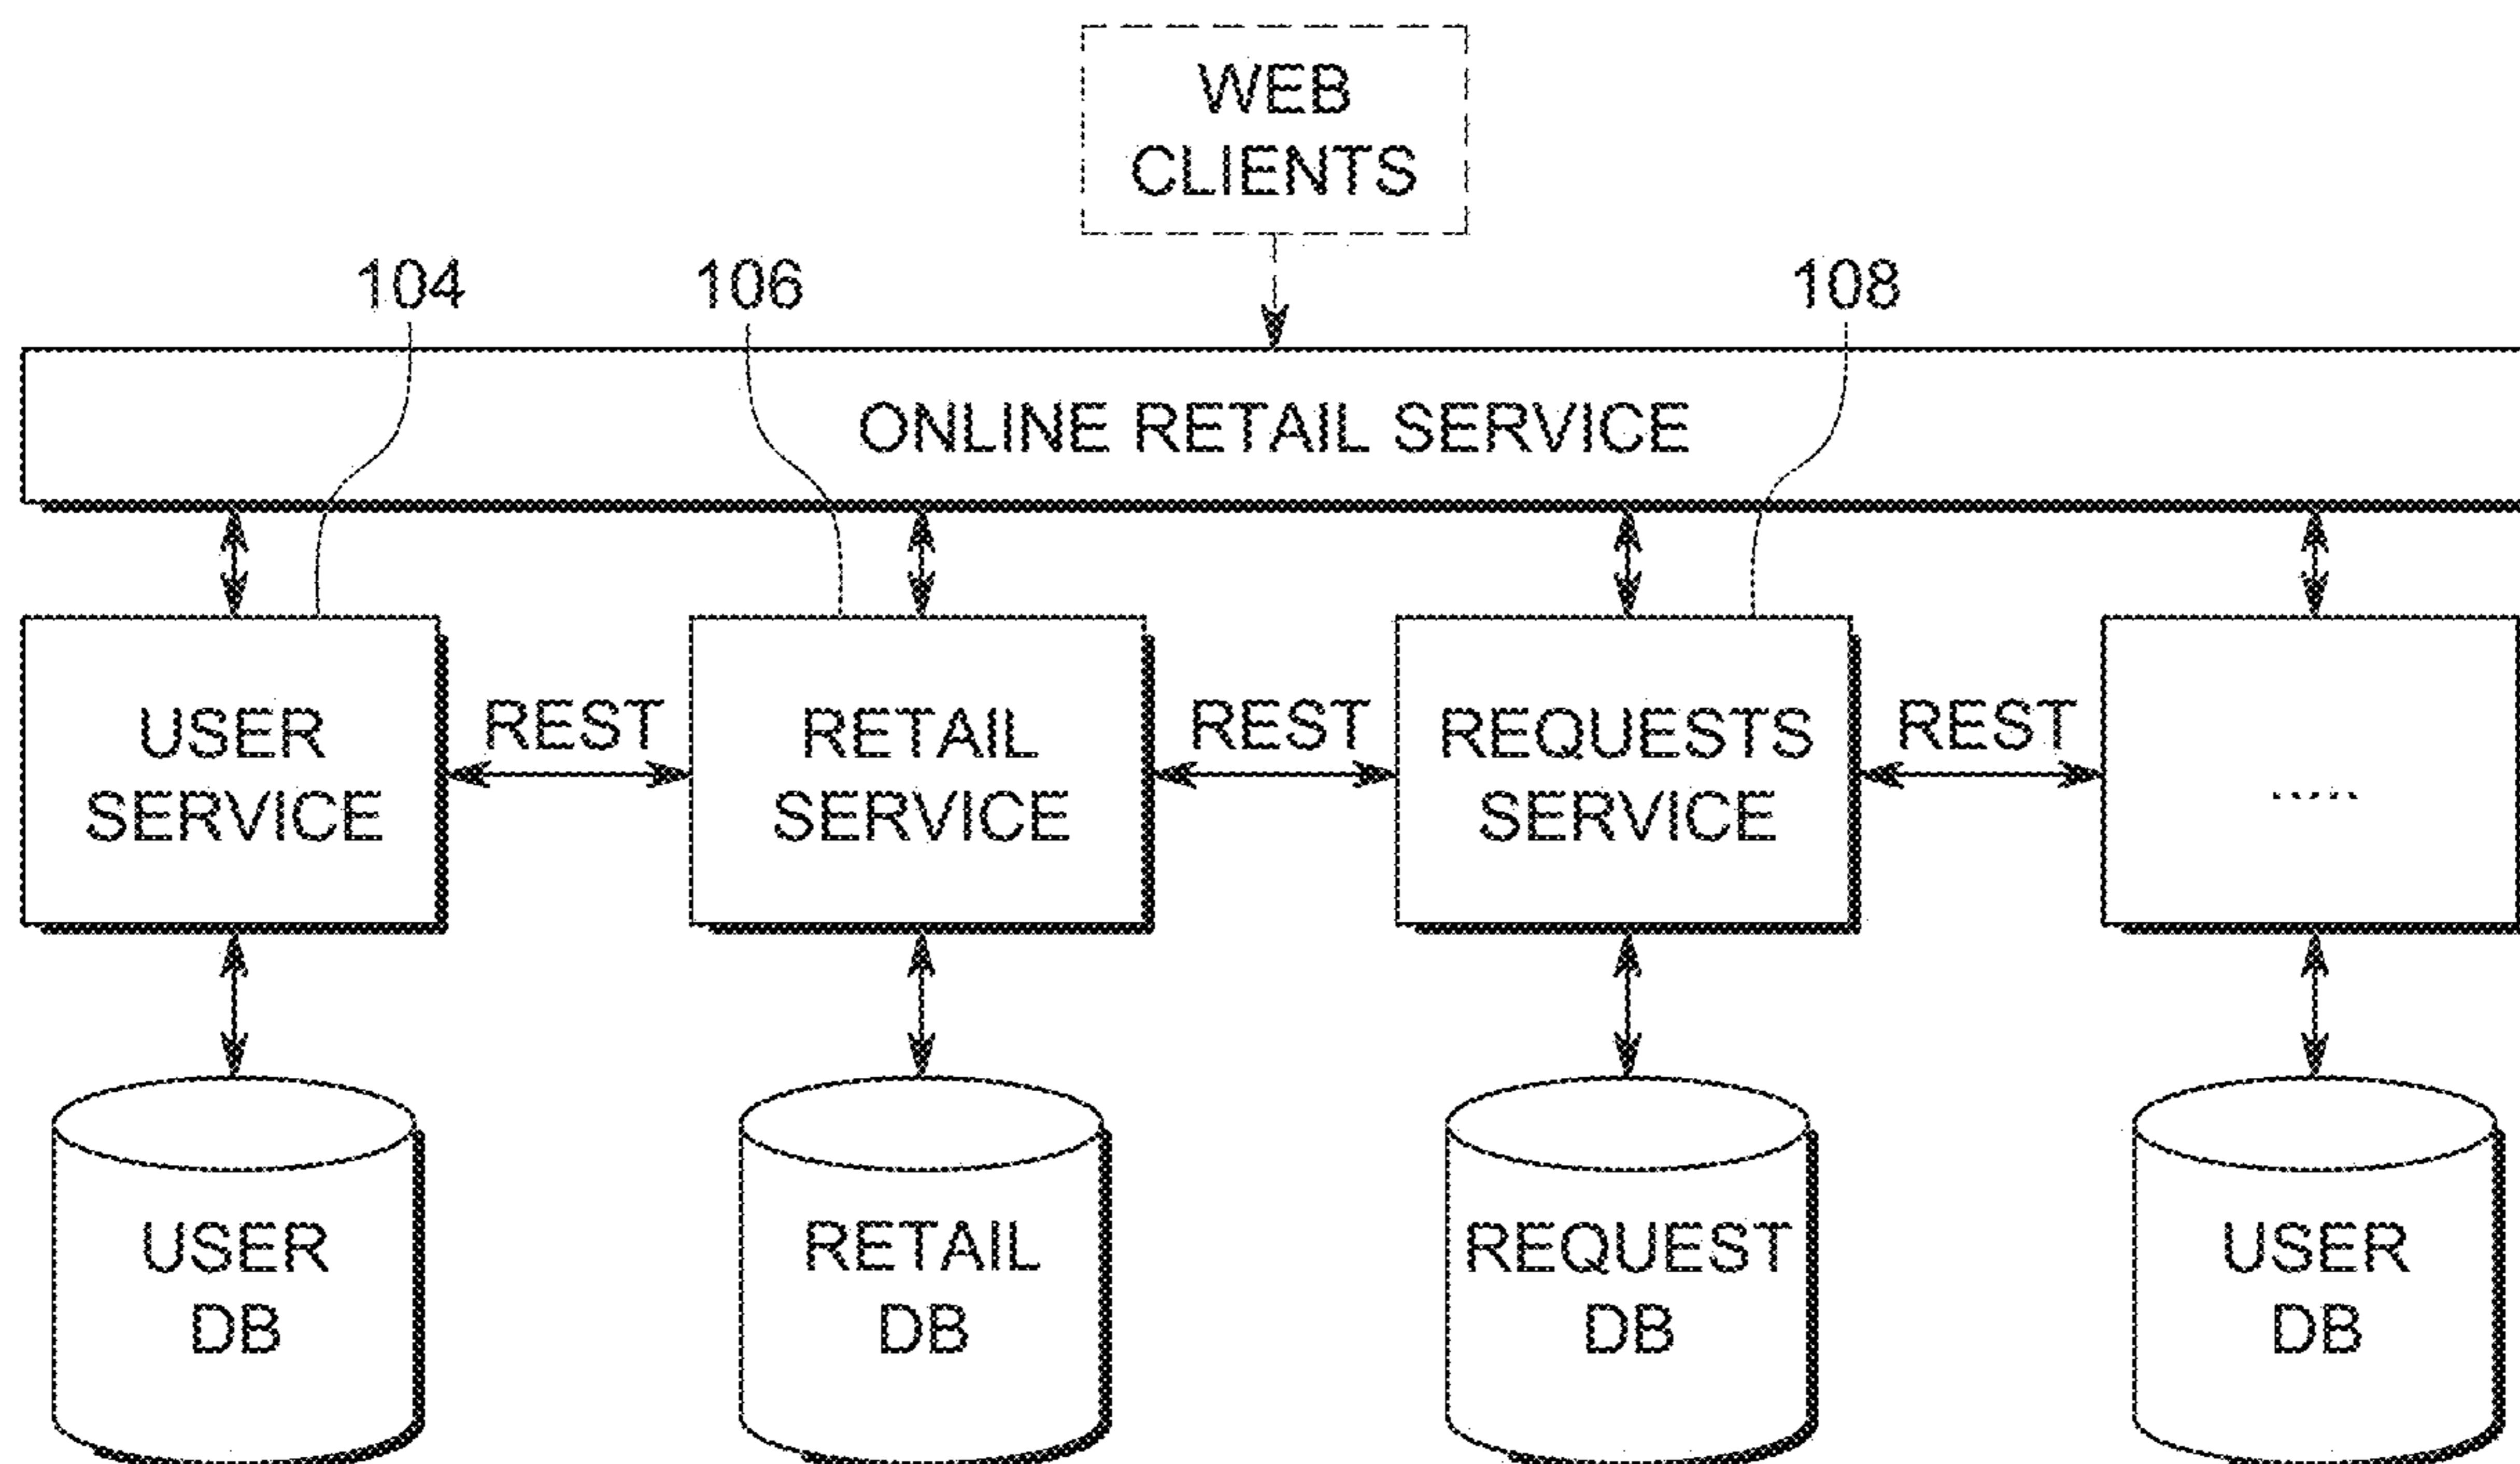


FIG. 1B

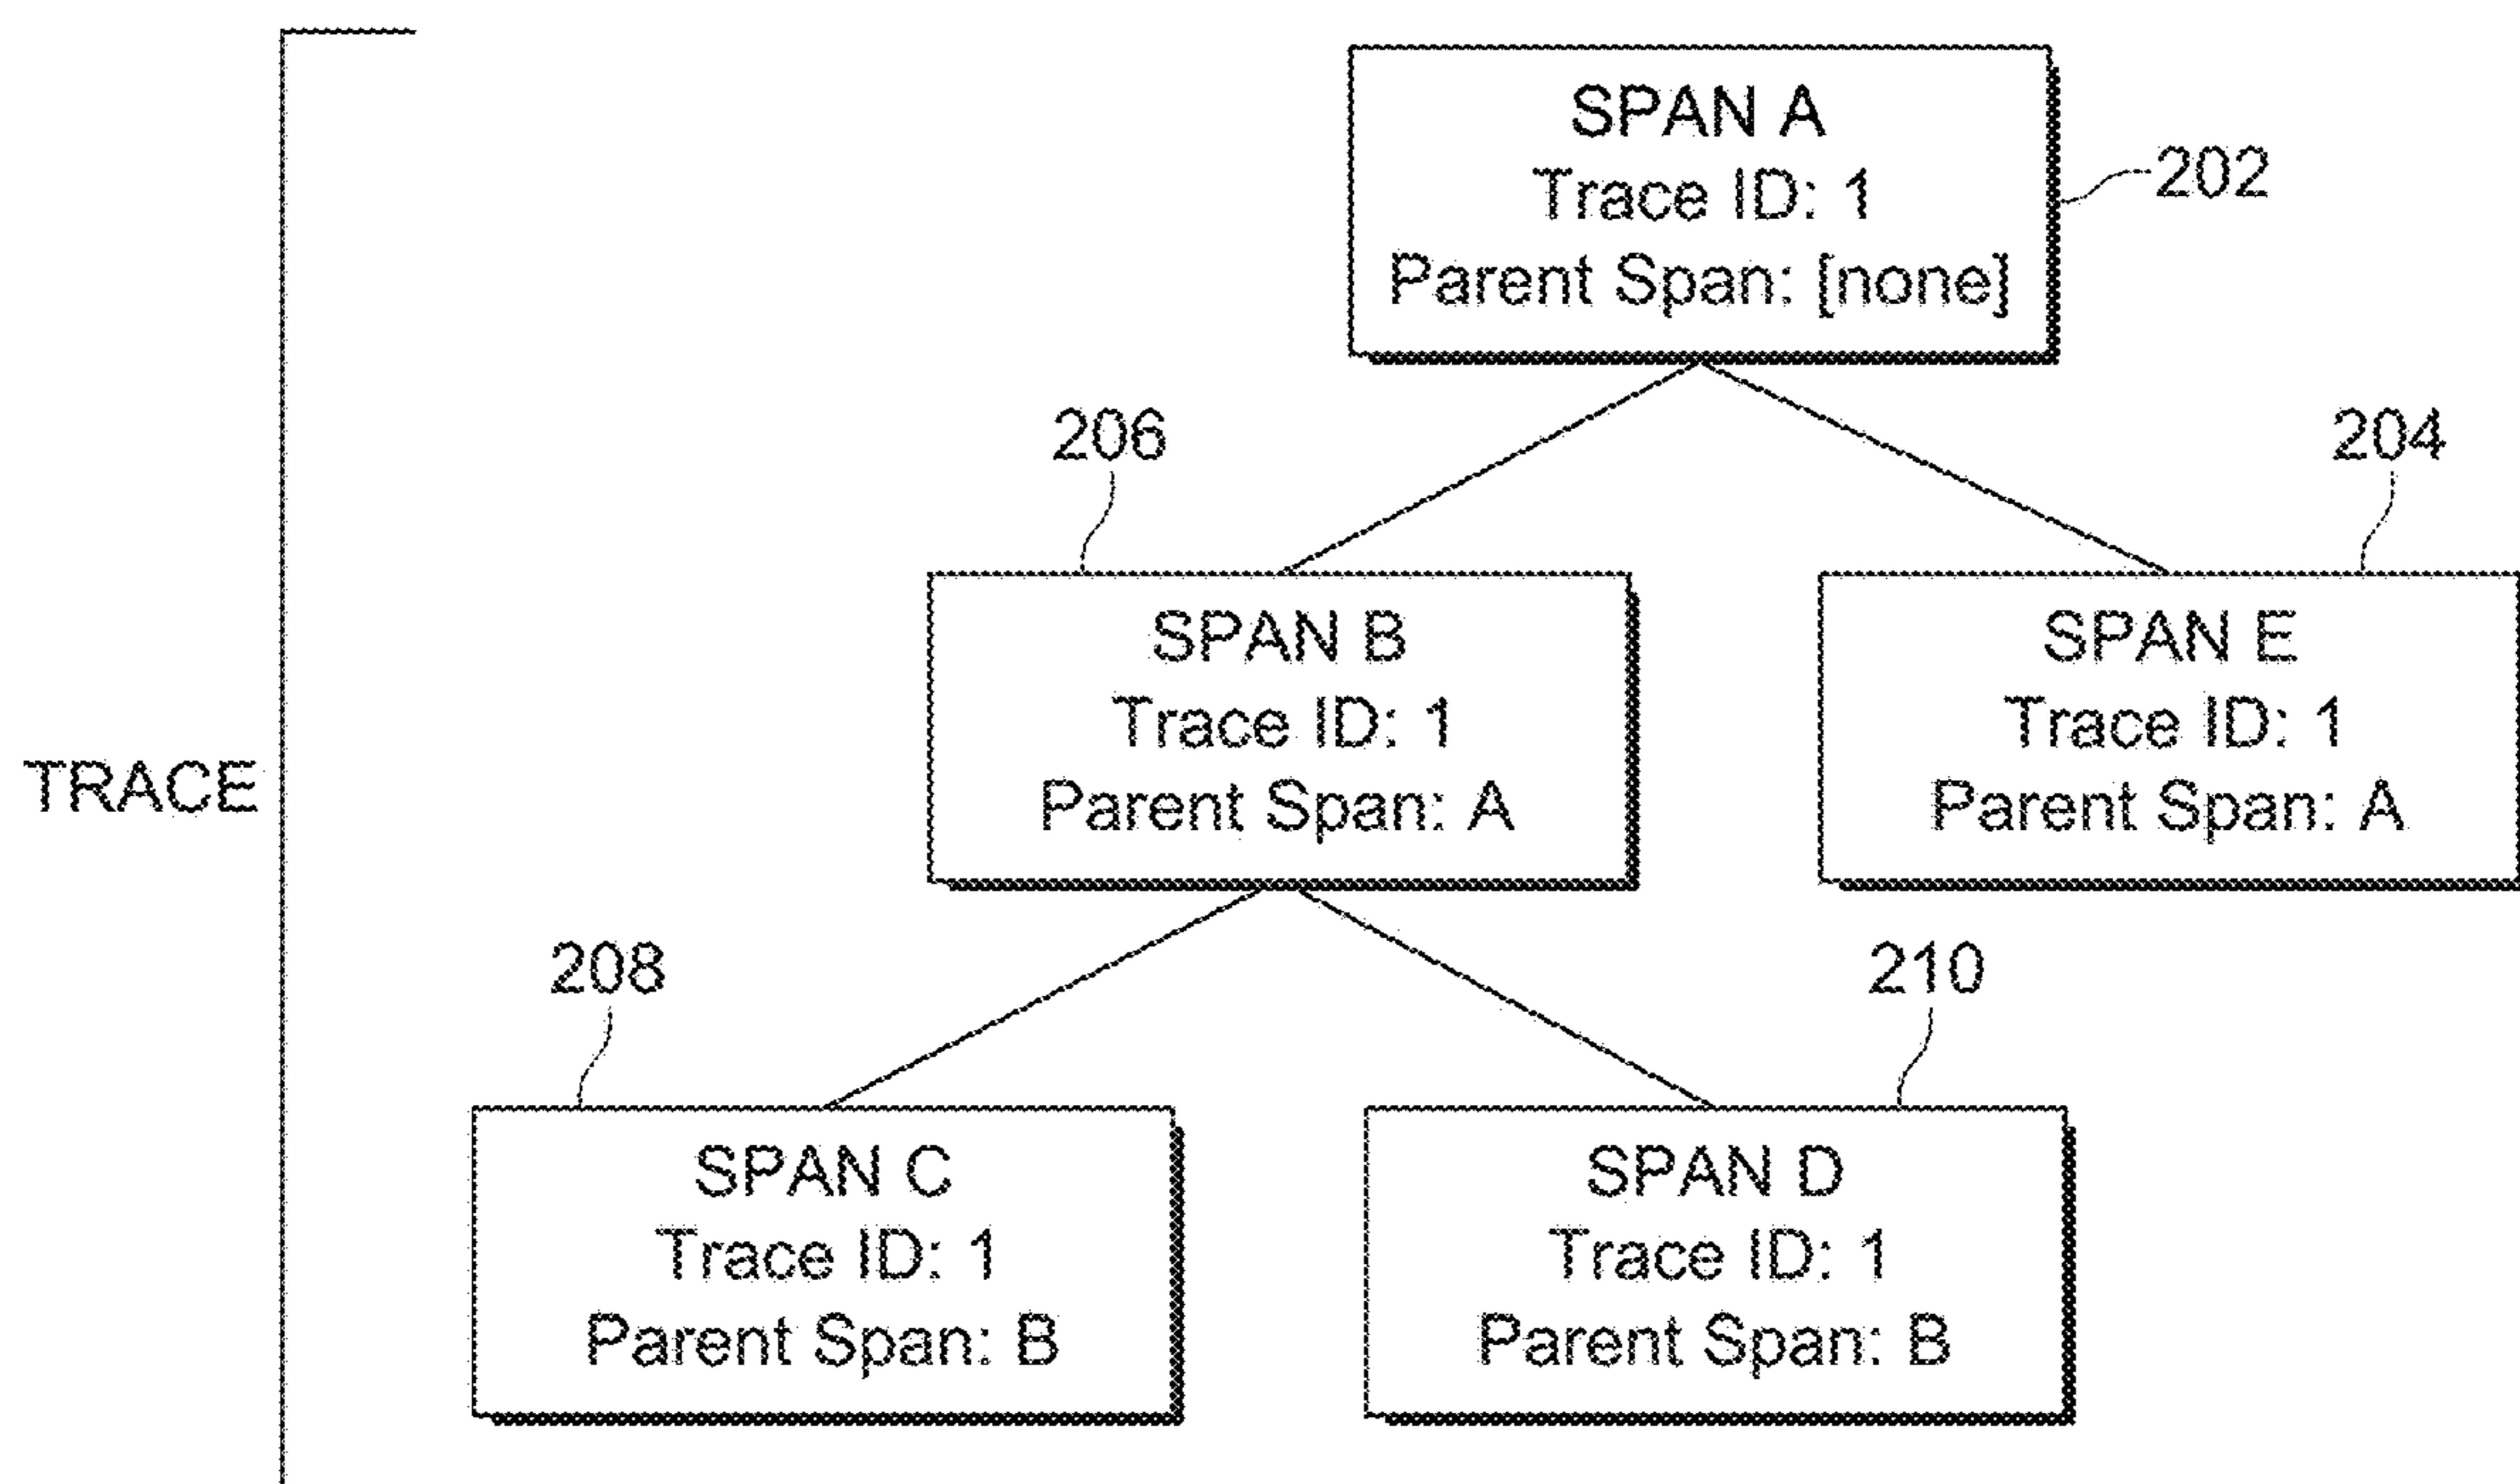


FIG. 2A

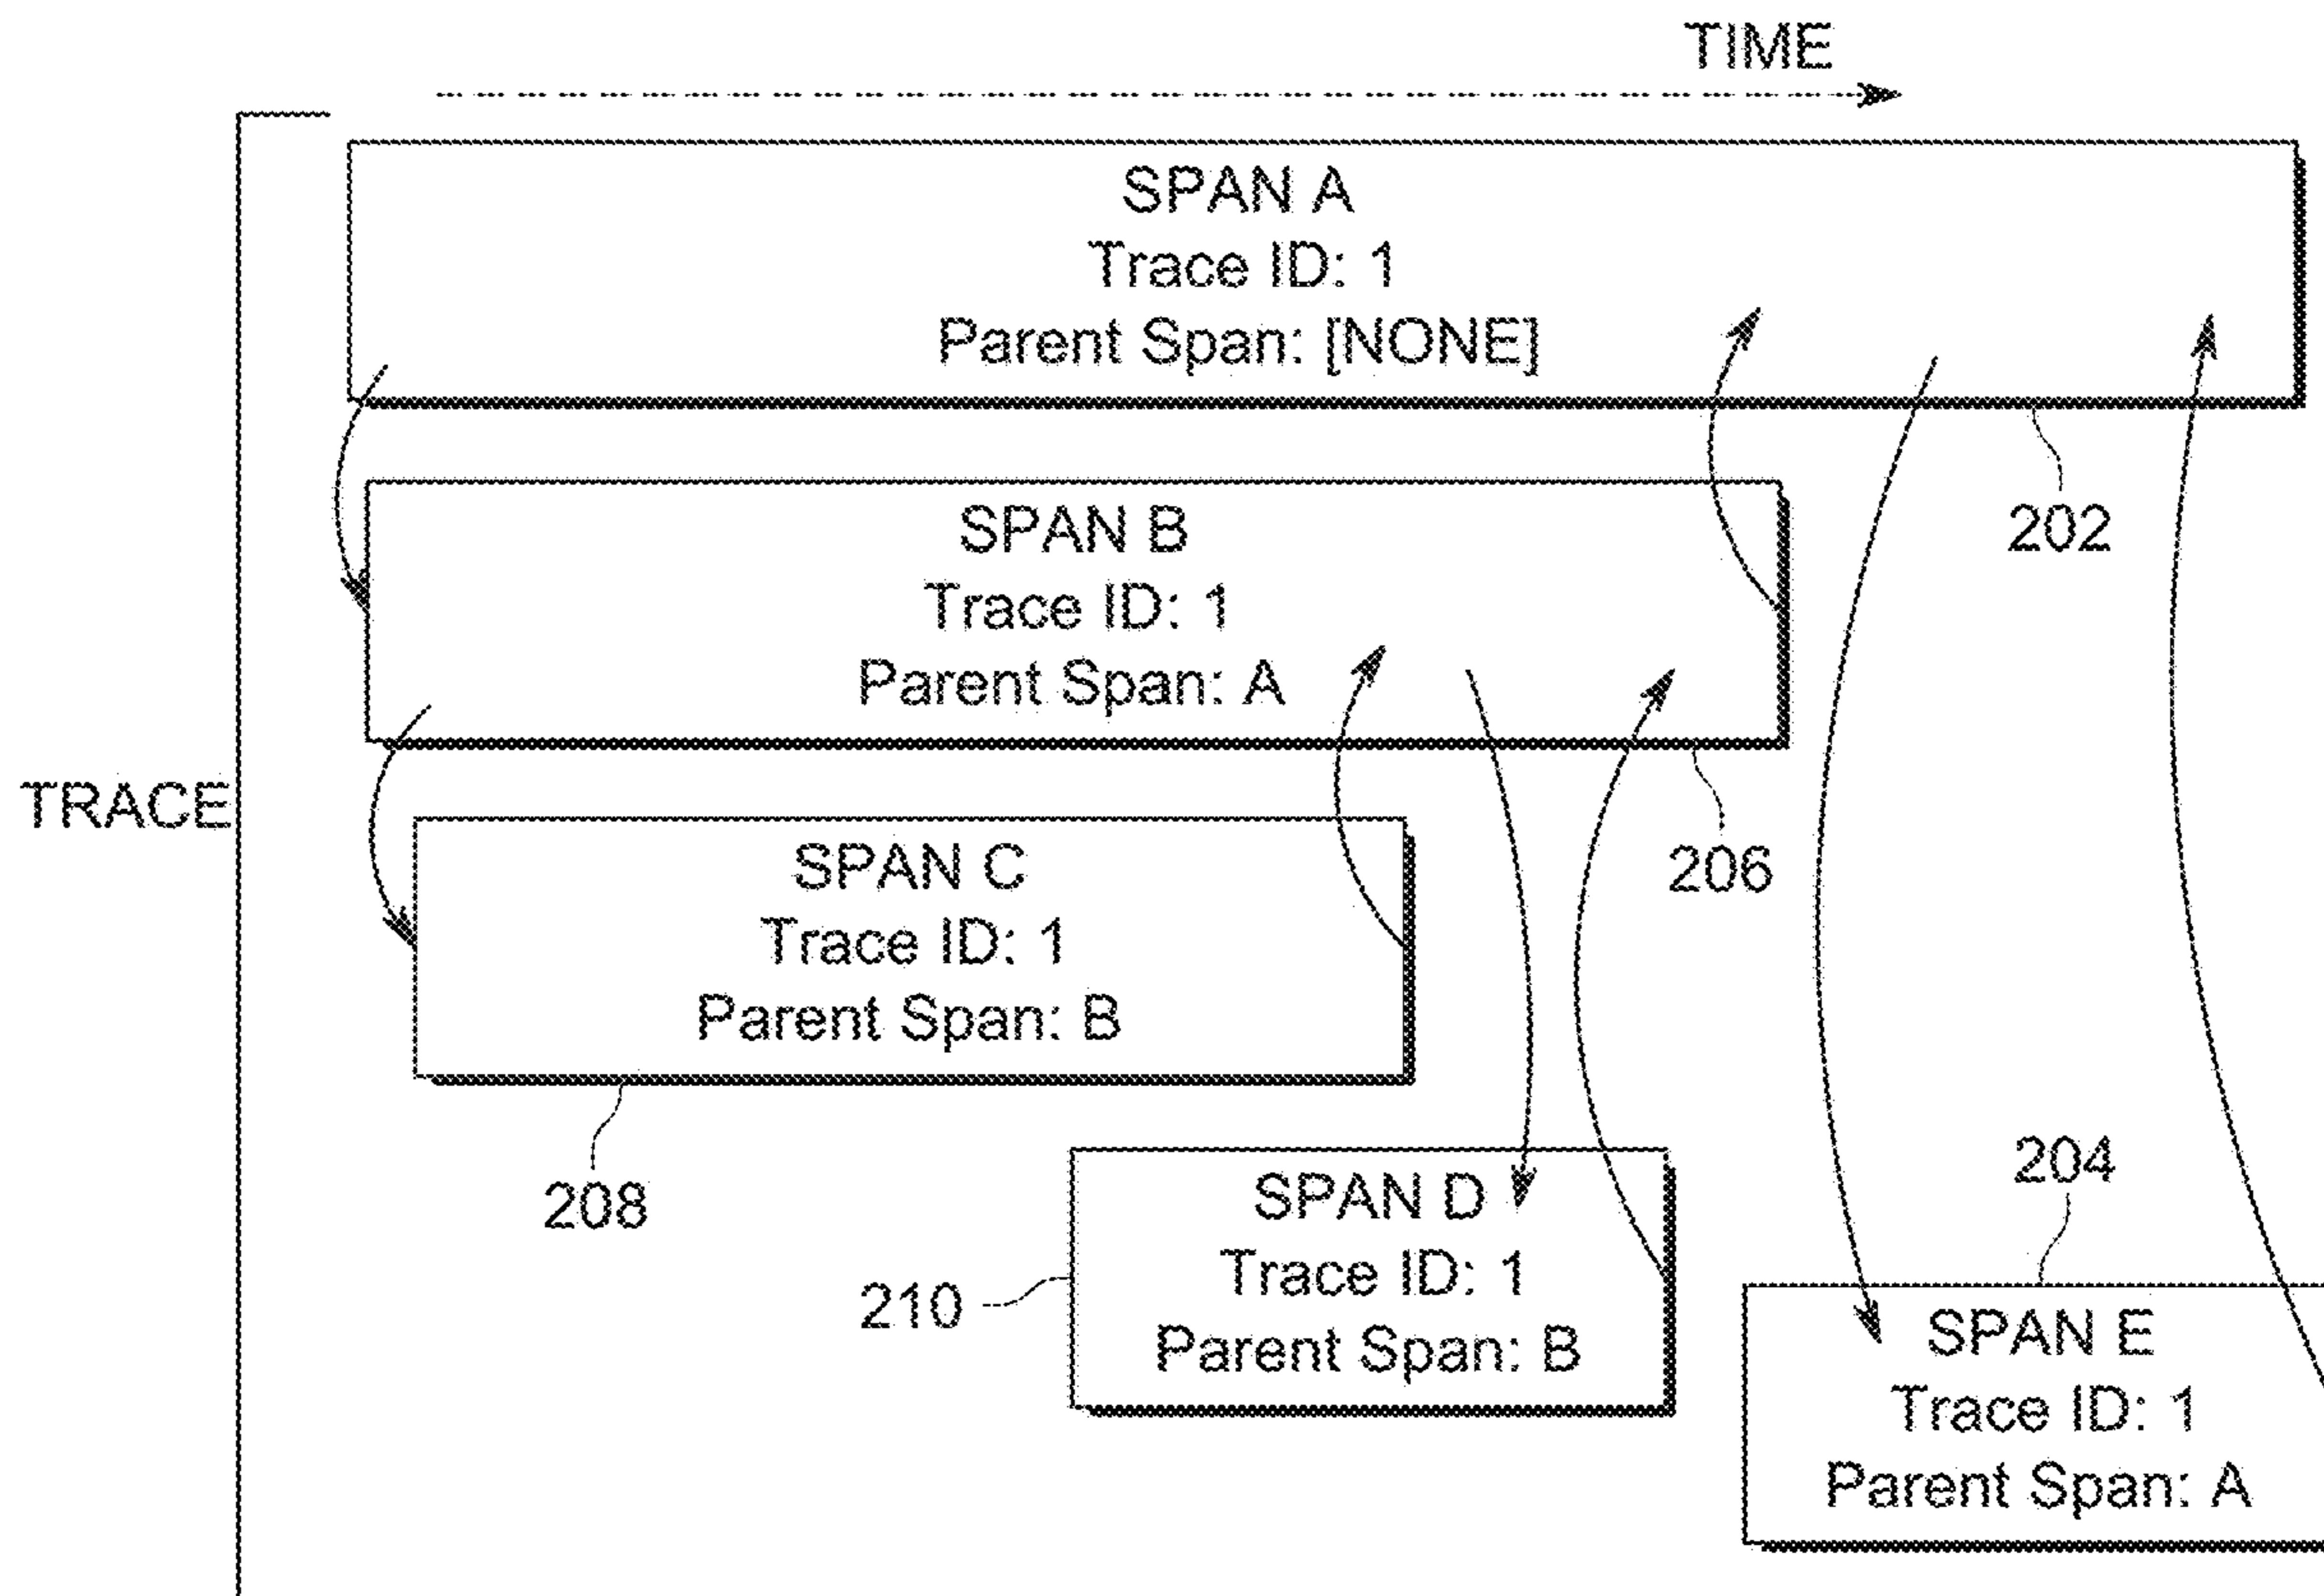


FIG. 2B

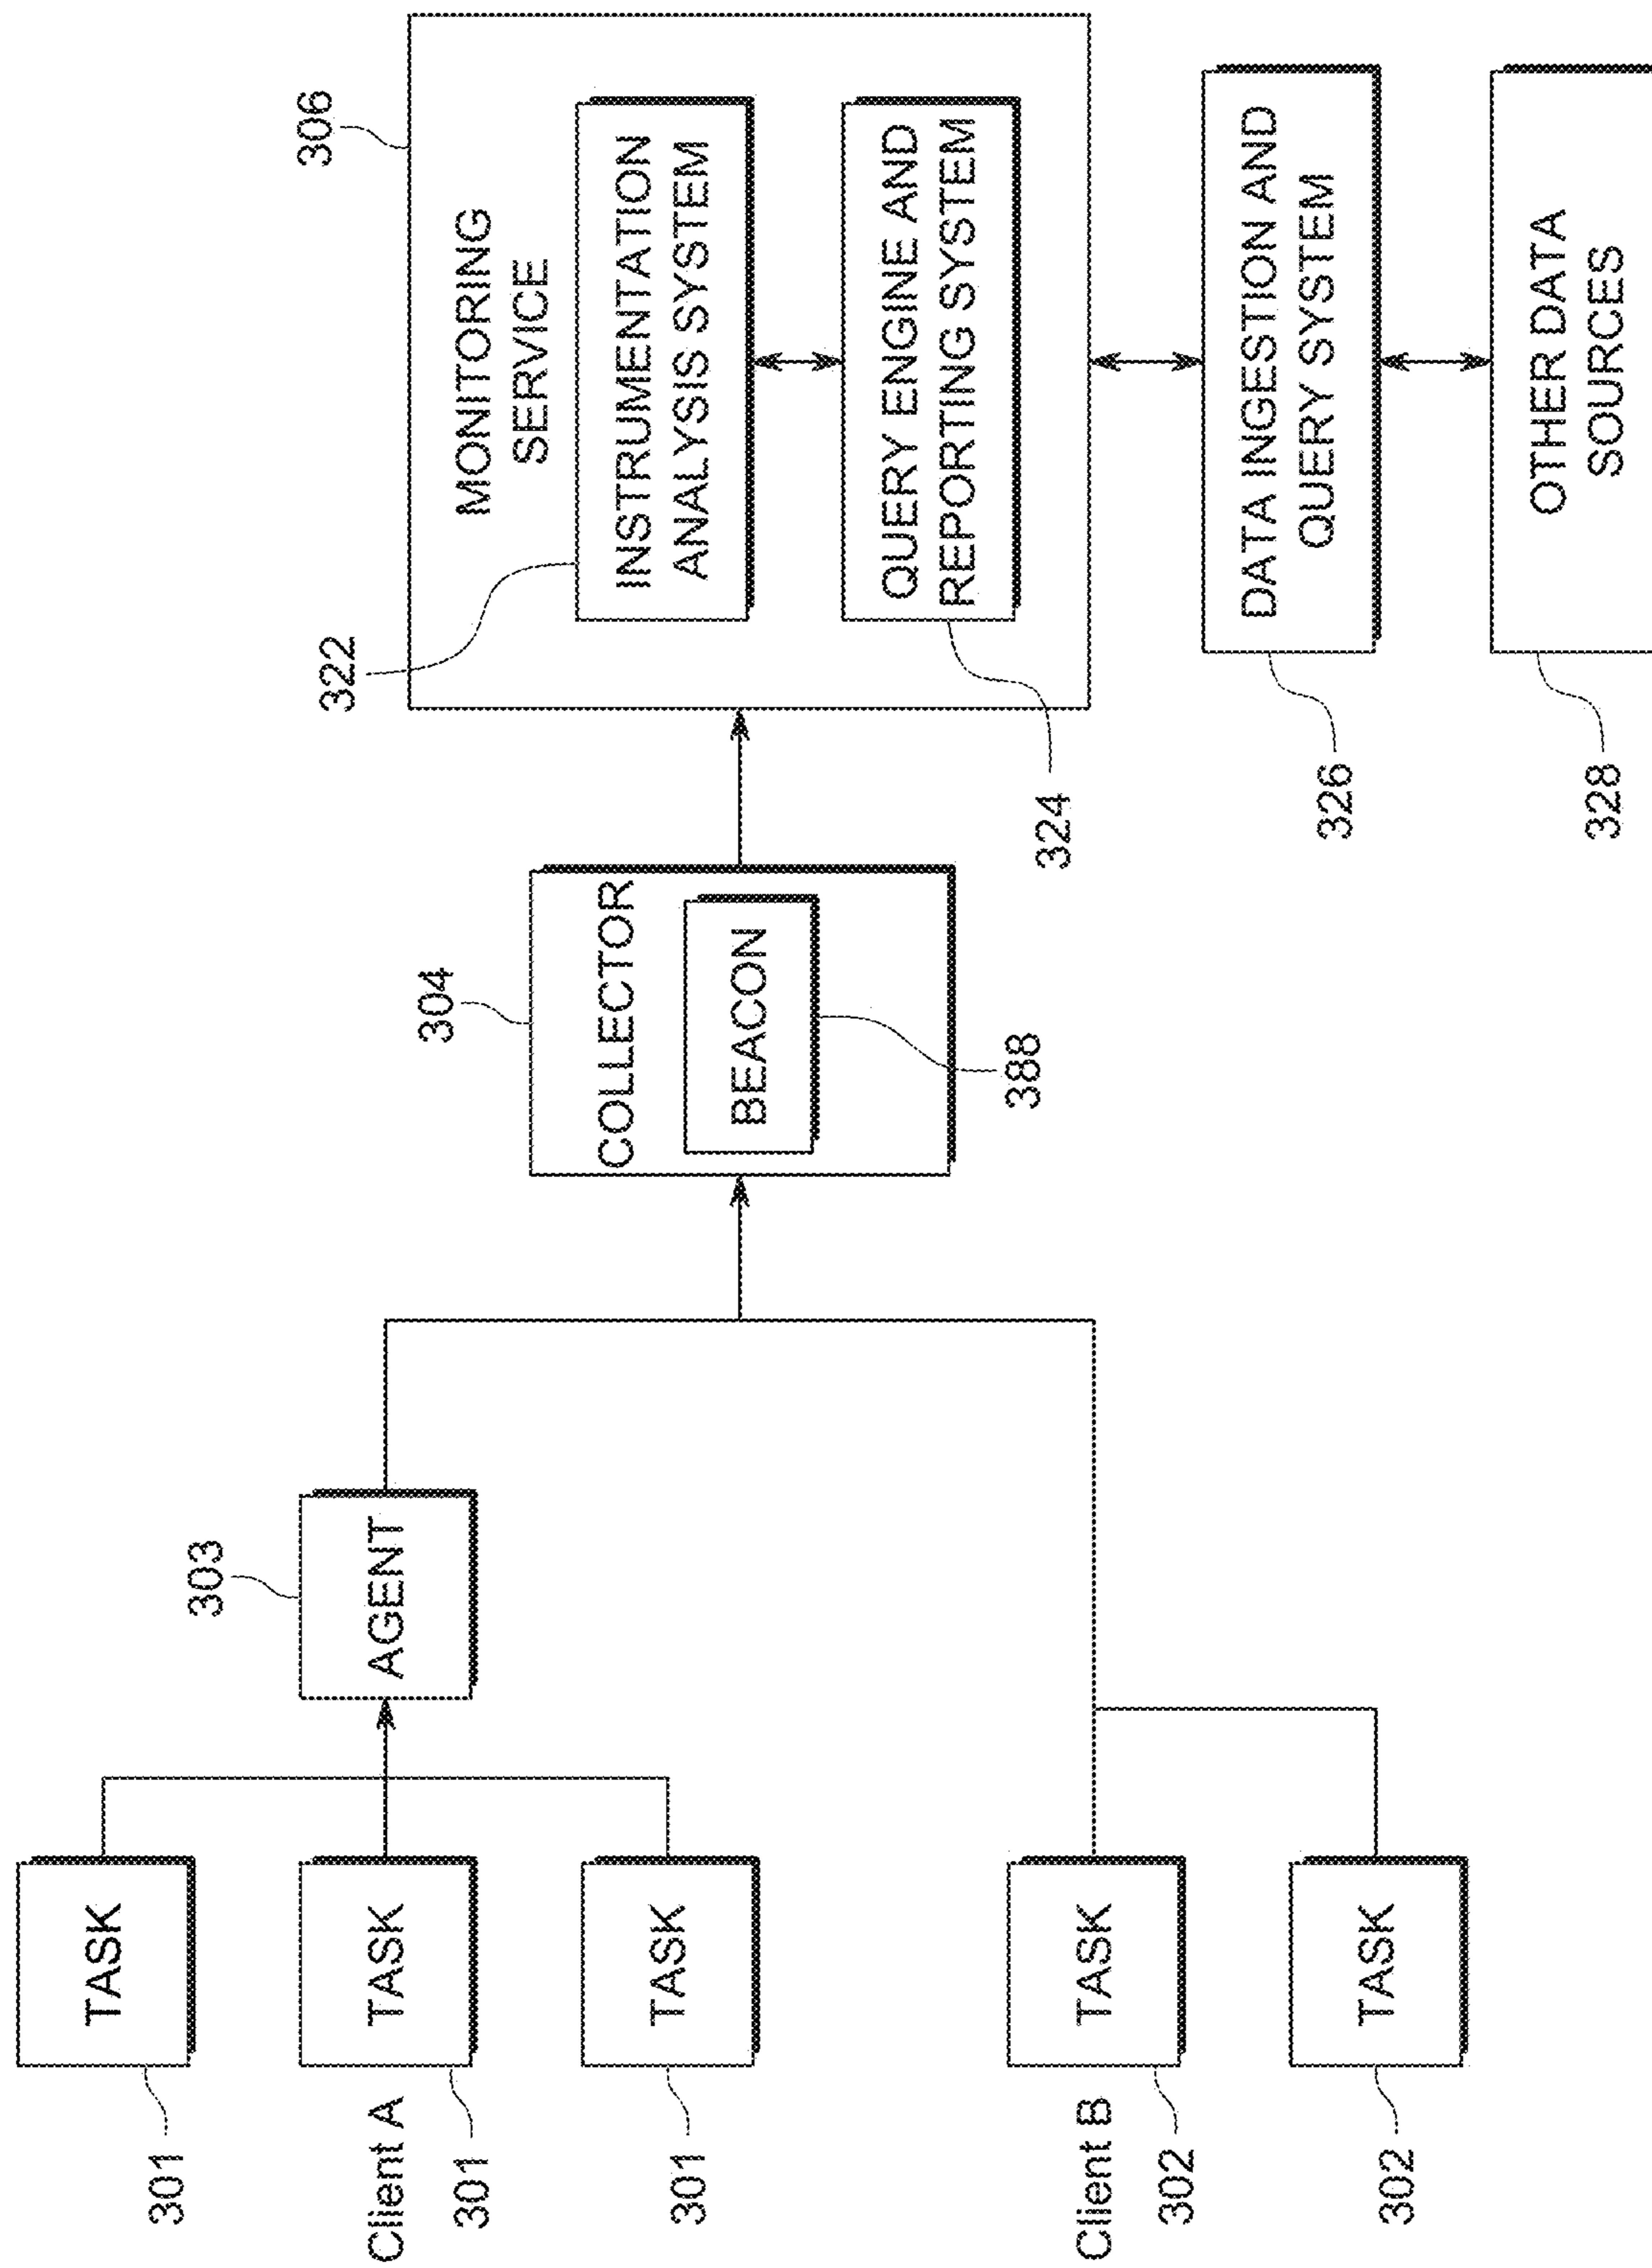


FIG. 3

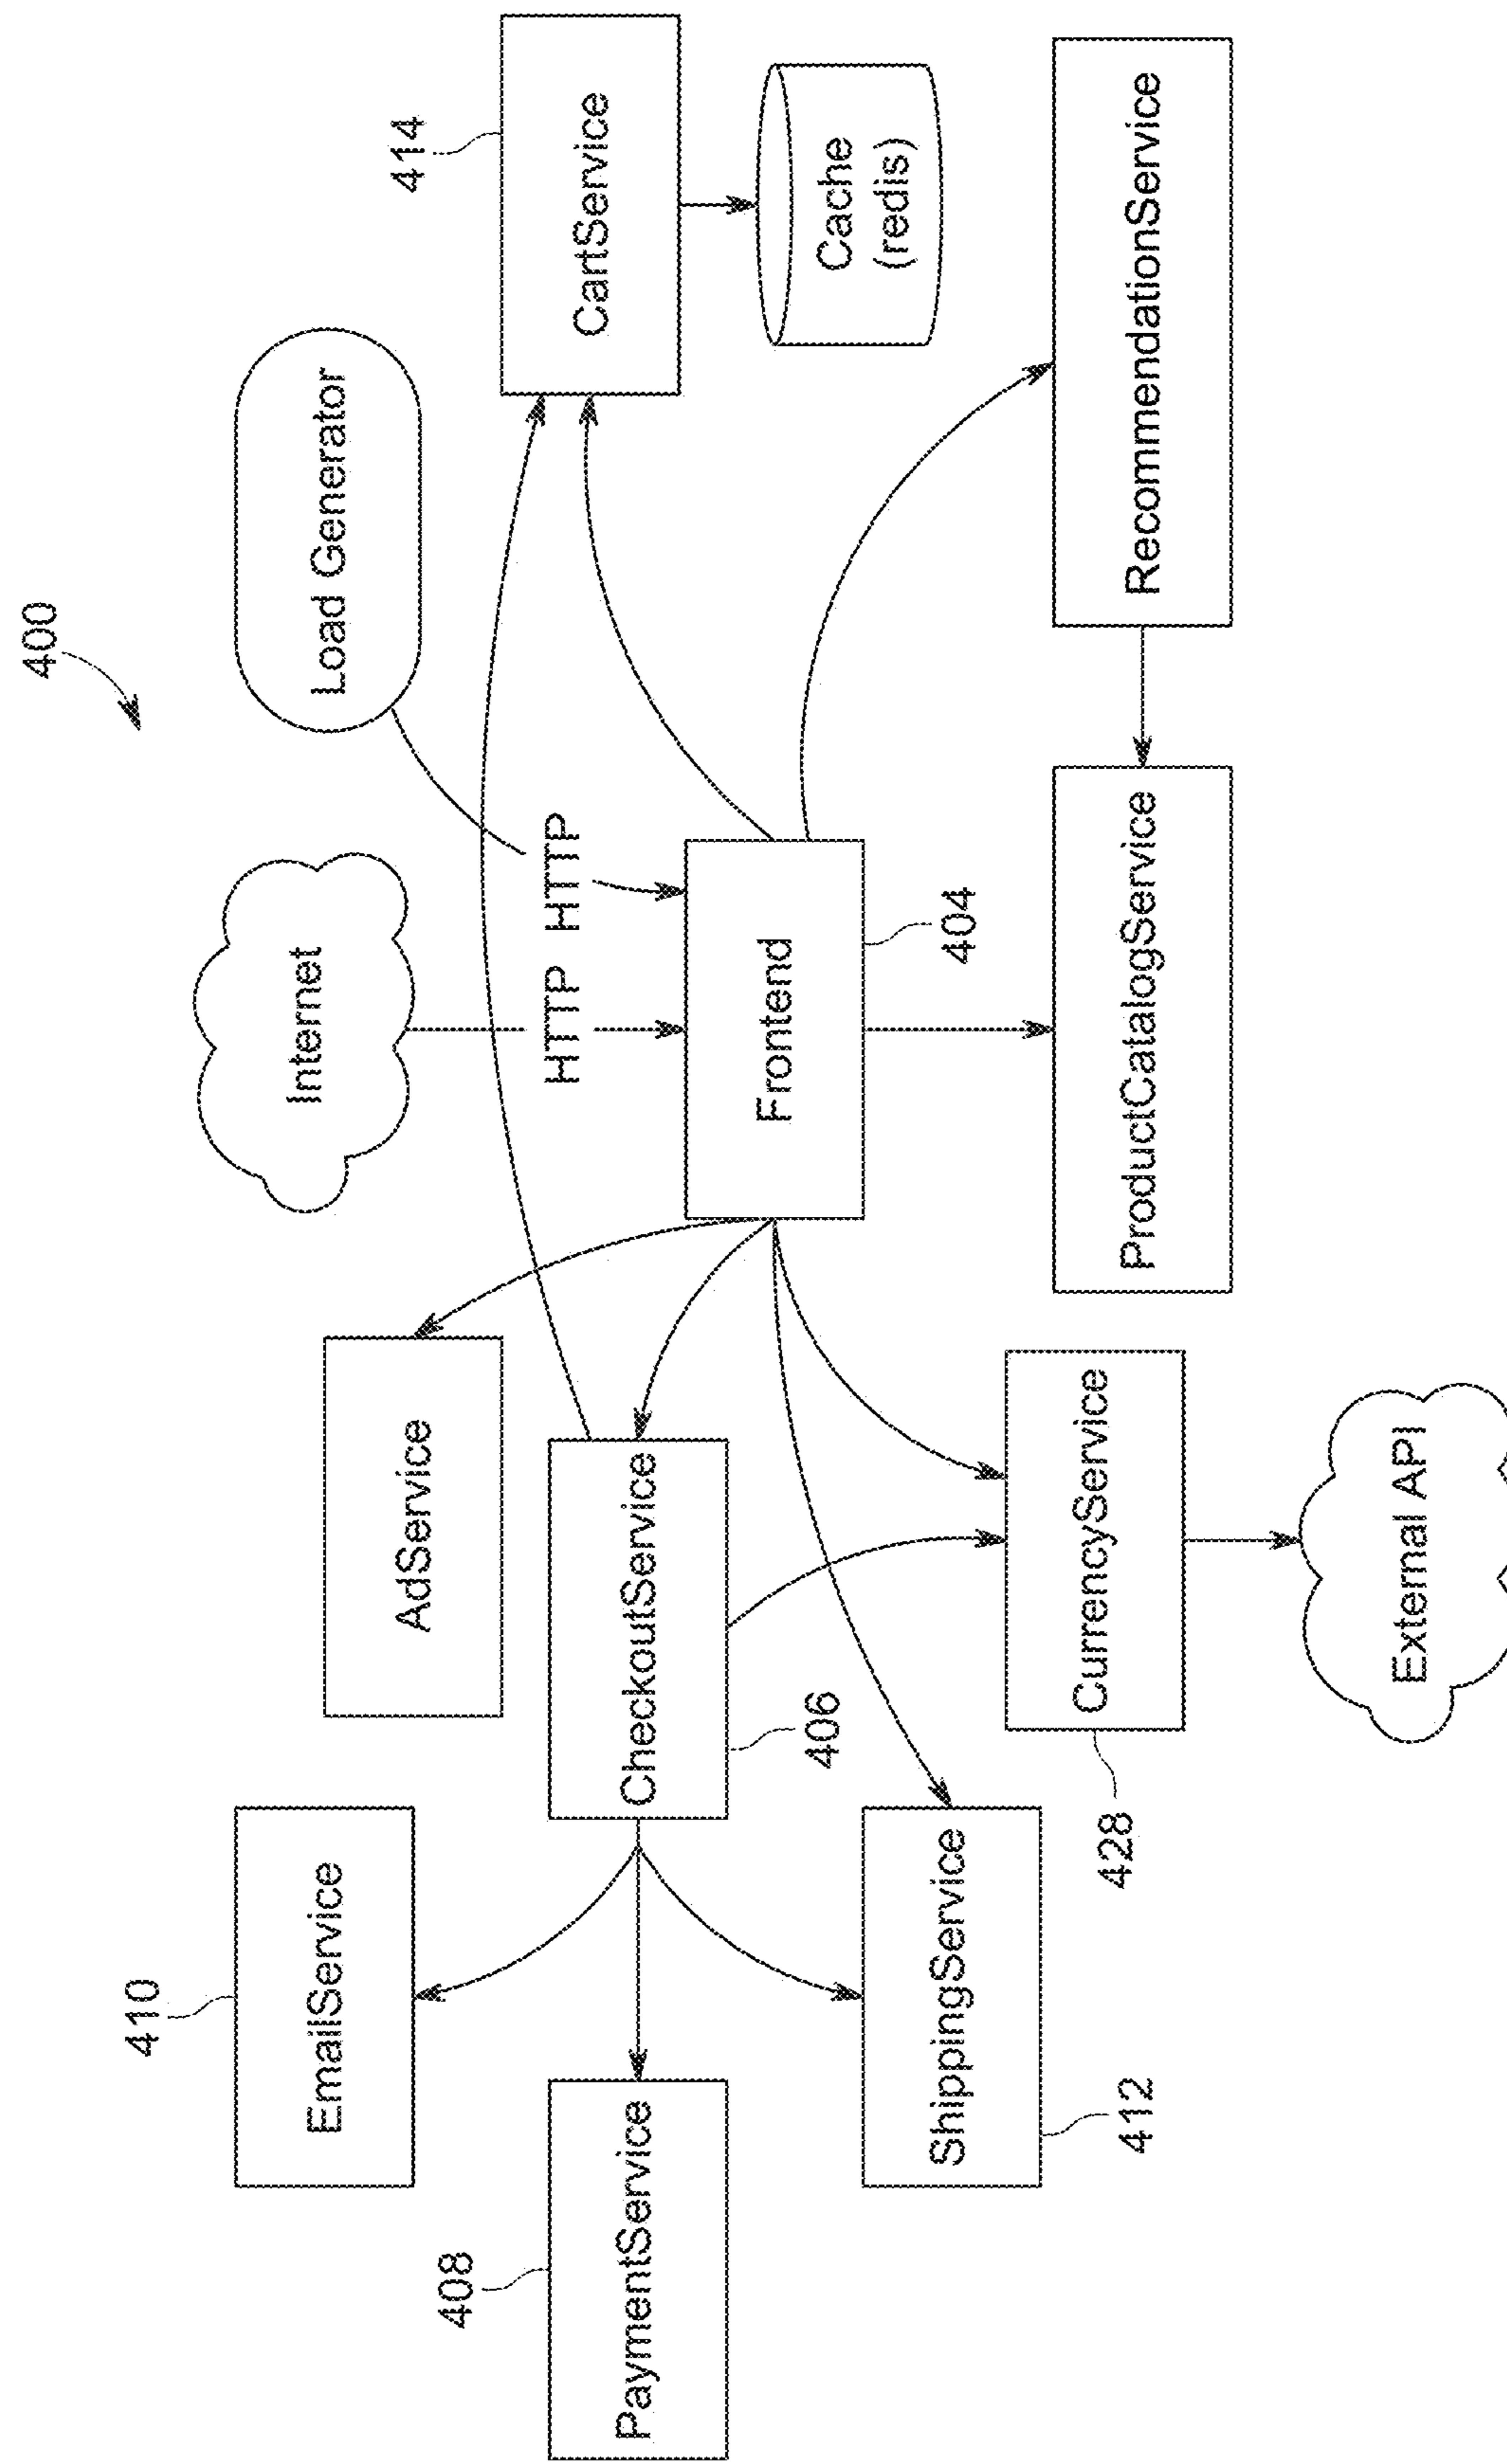


FIG. 4

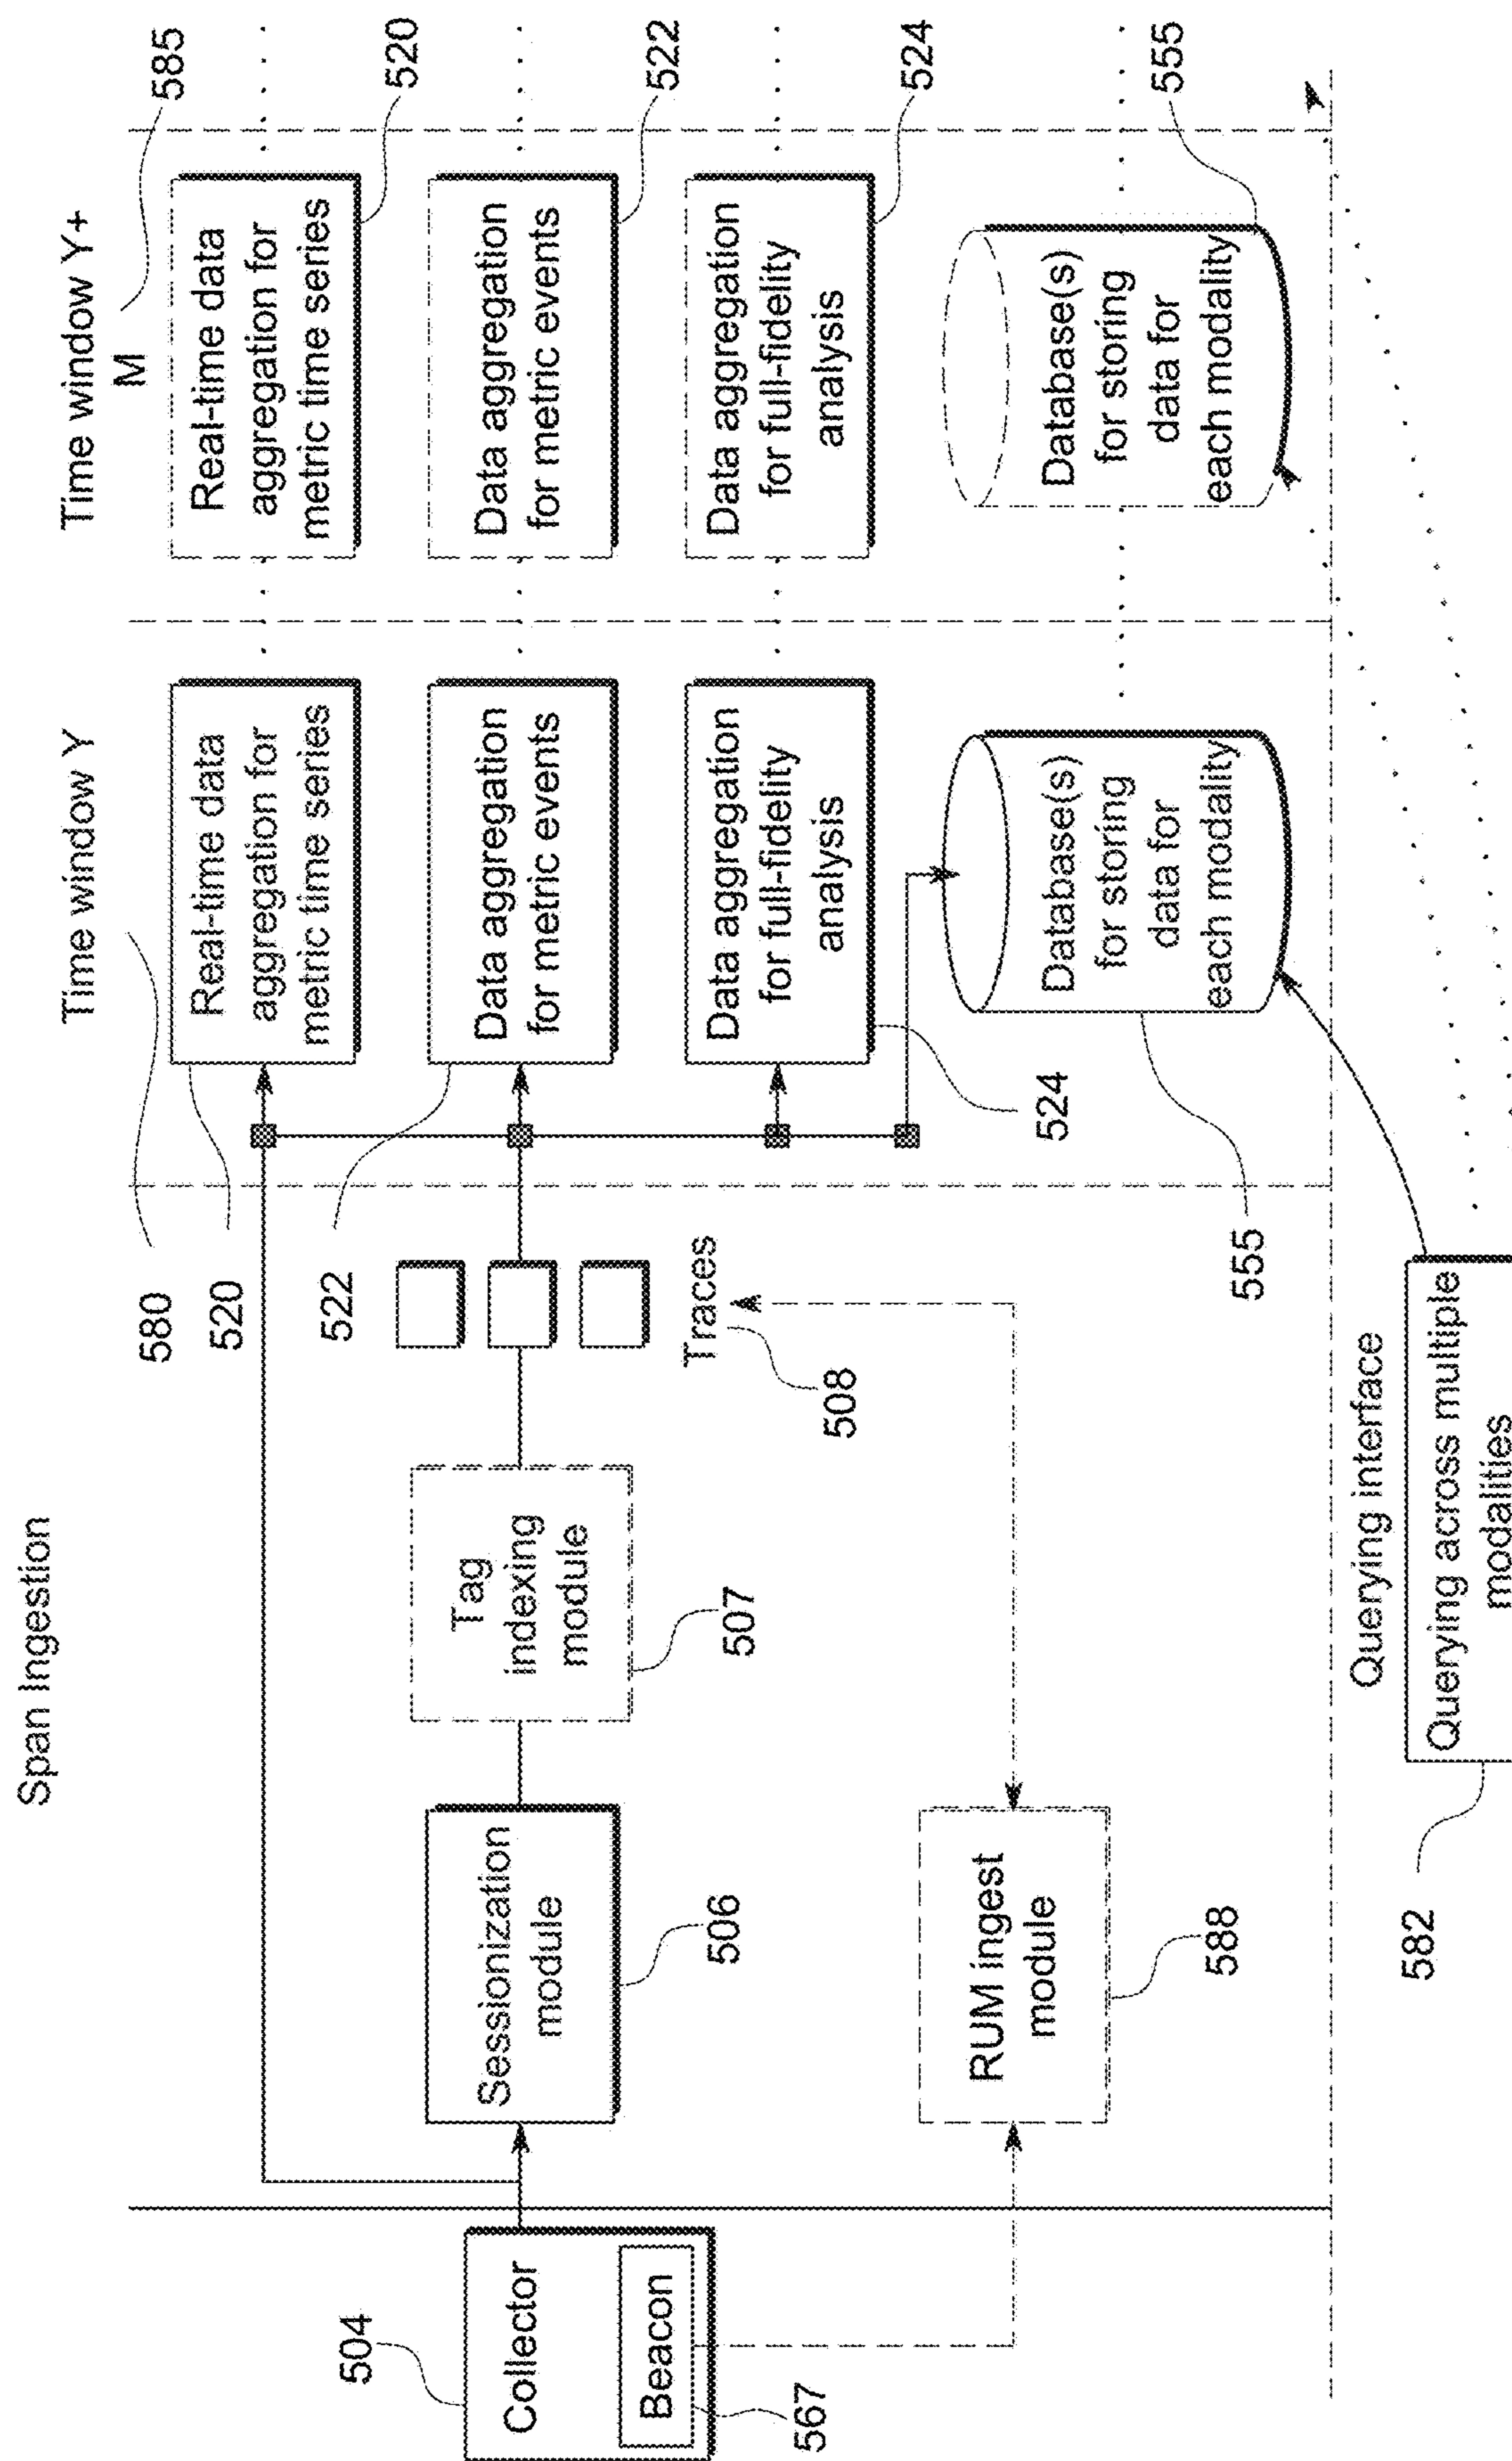


FIG. 5

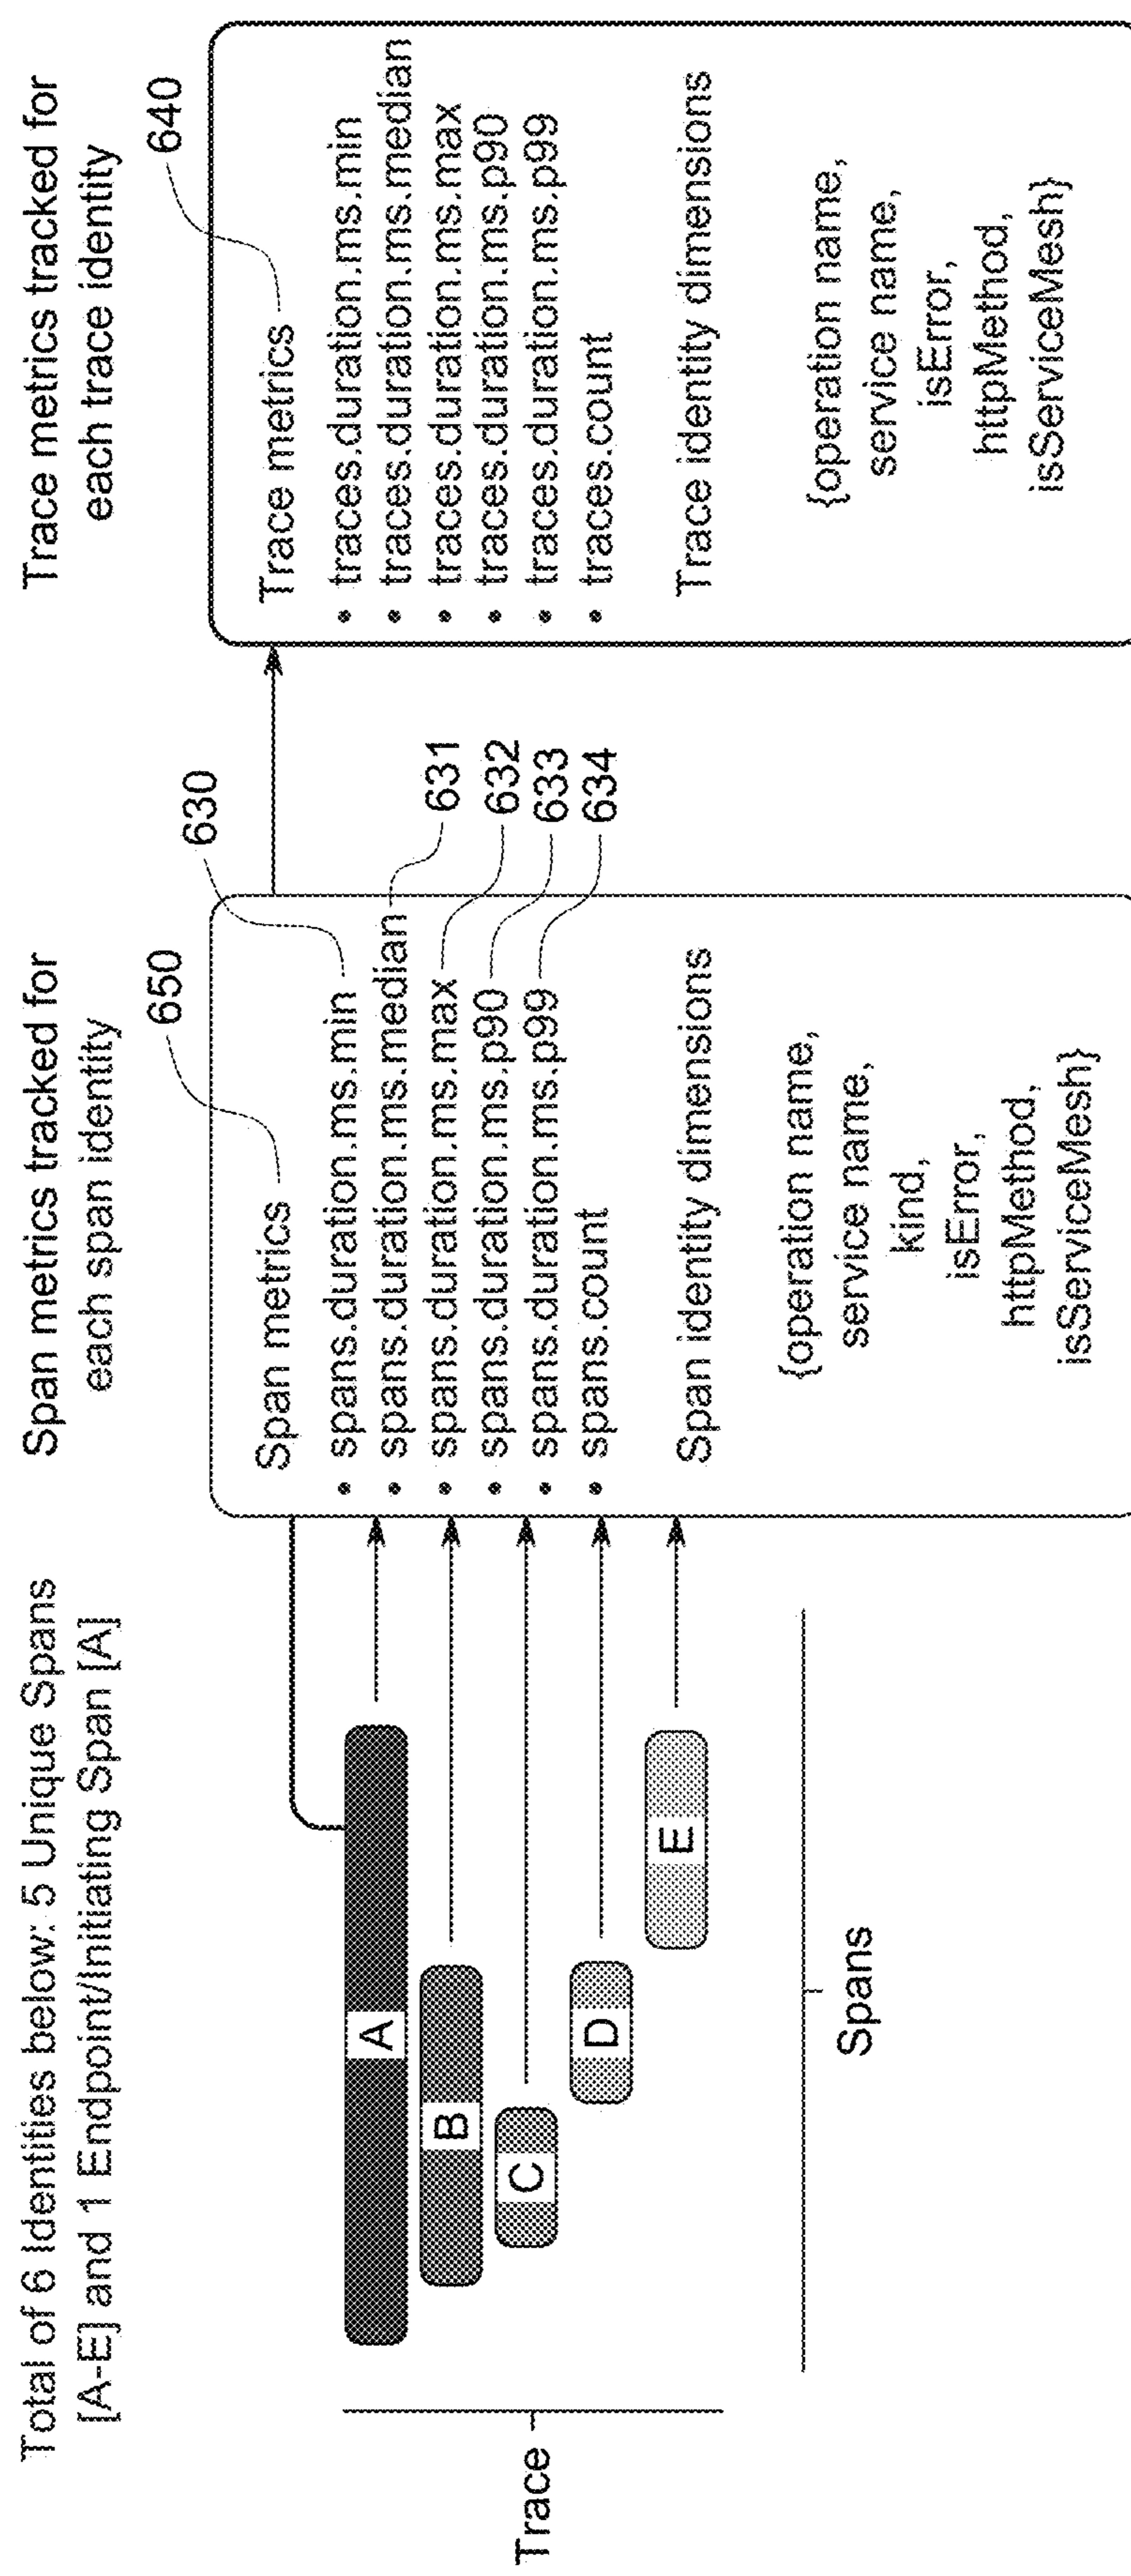


FIG. 6

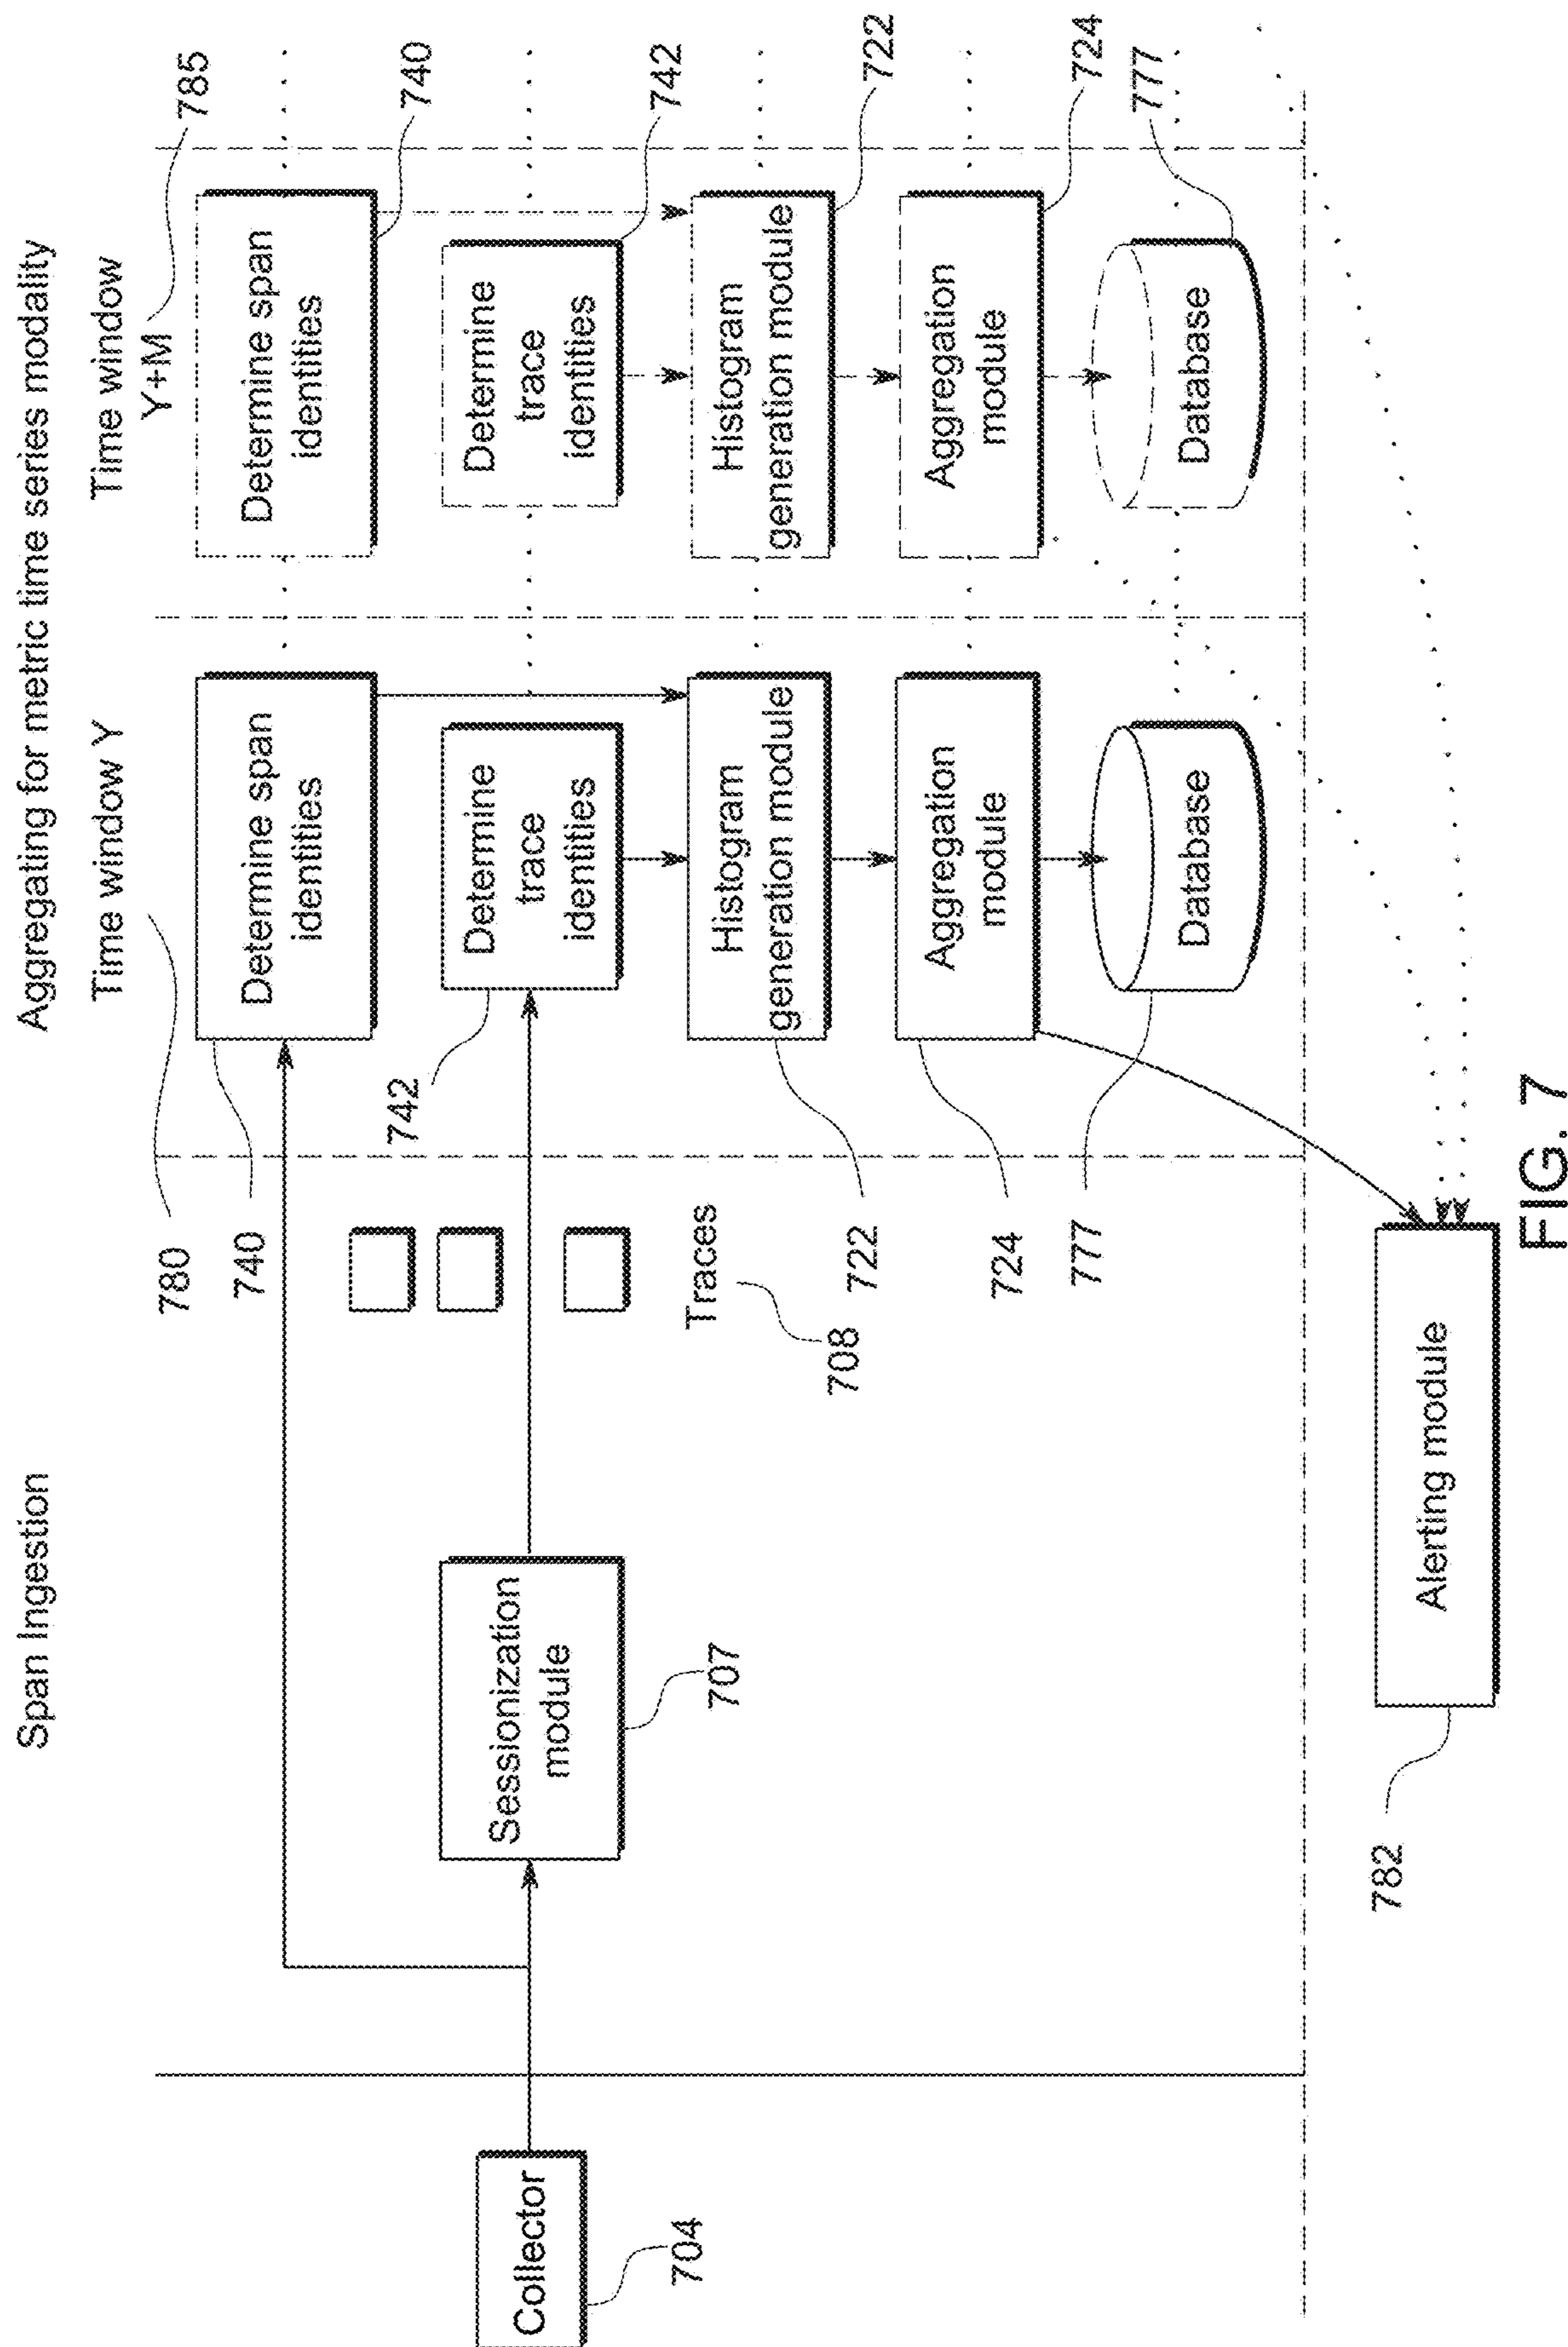


FIG. 7

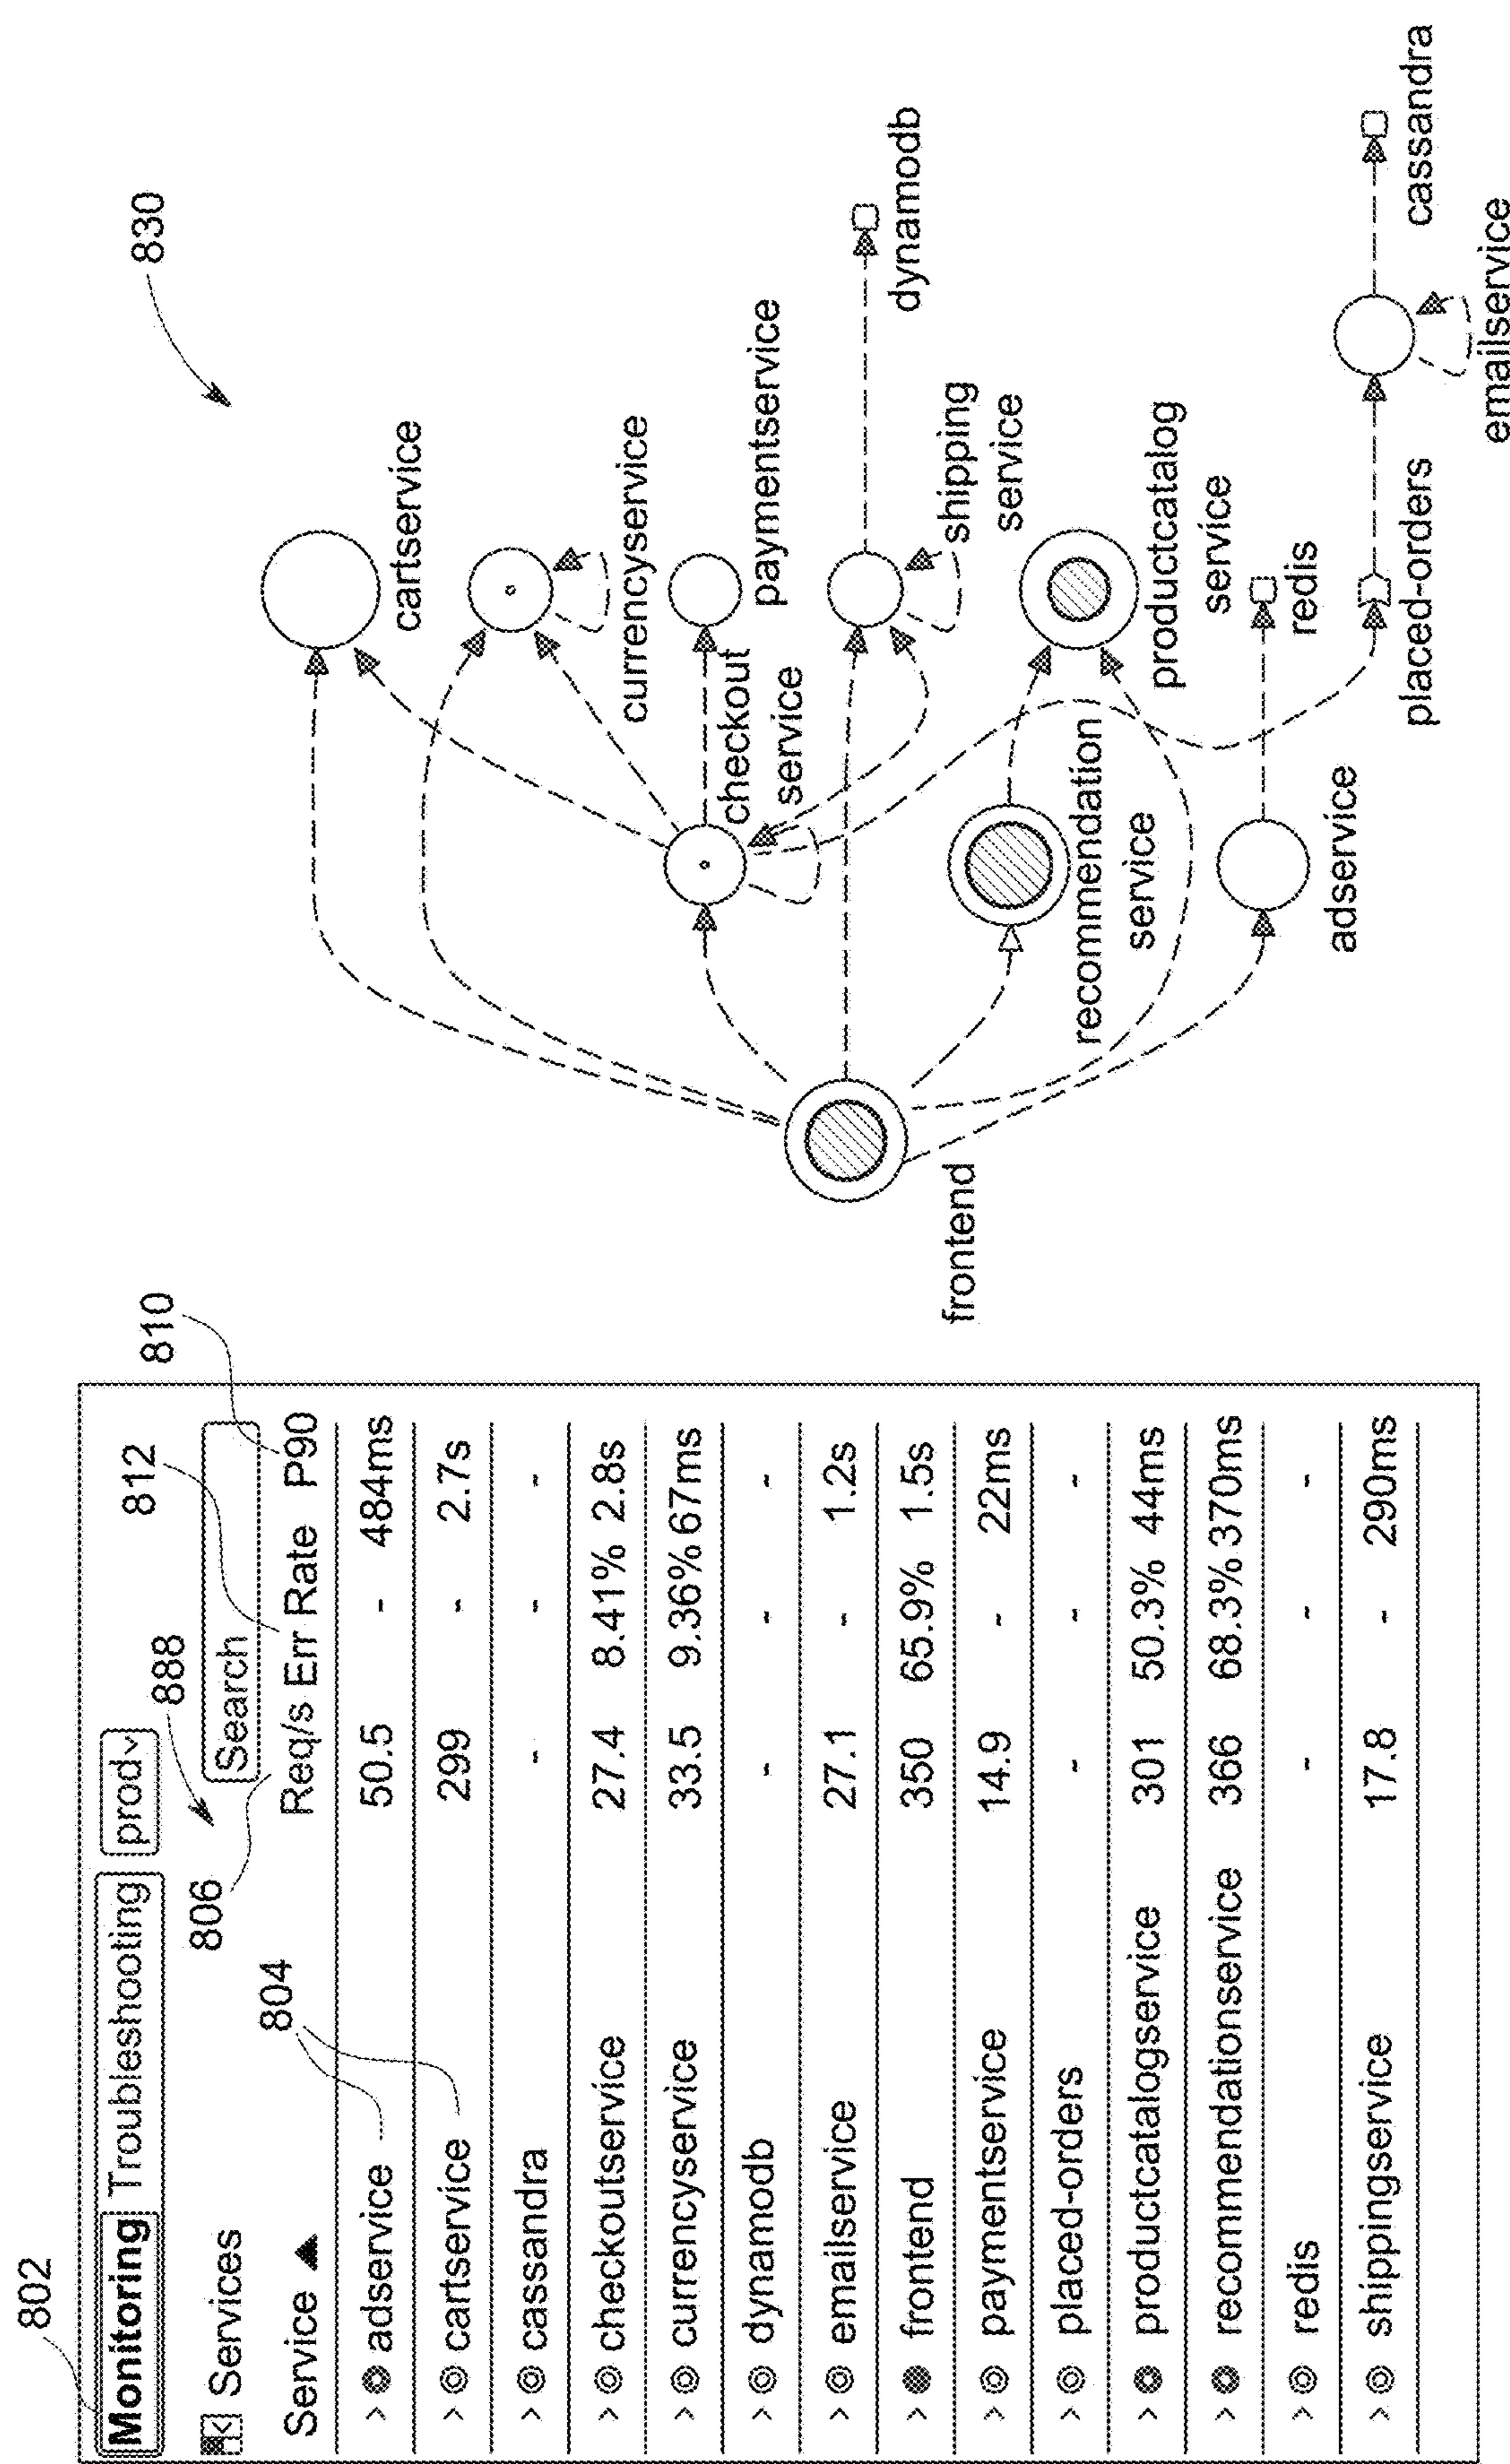


FIG. 8

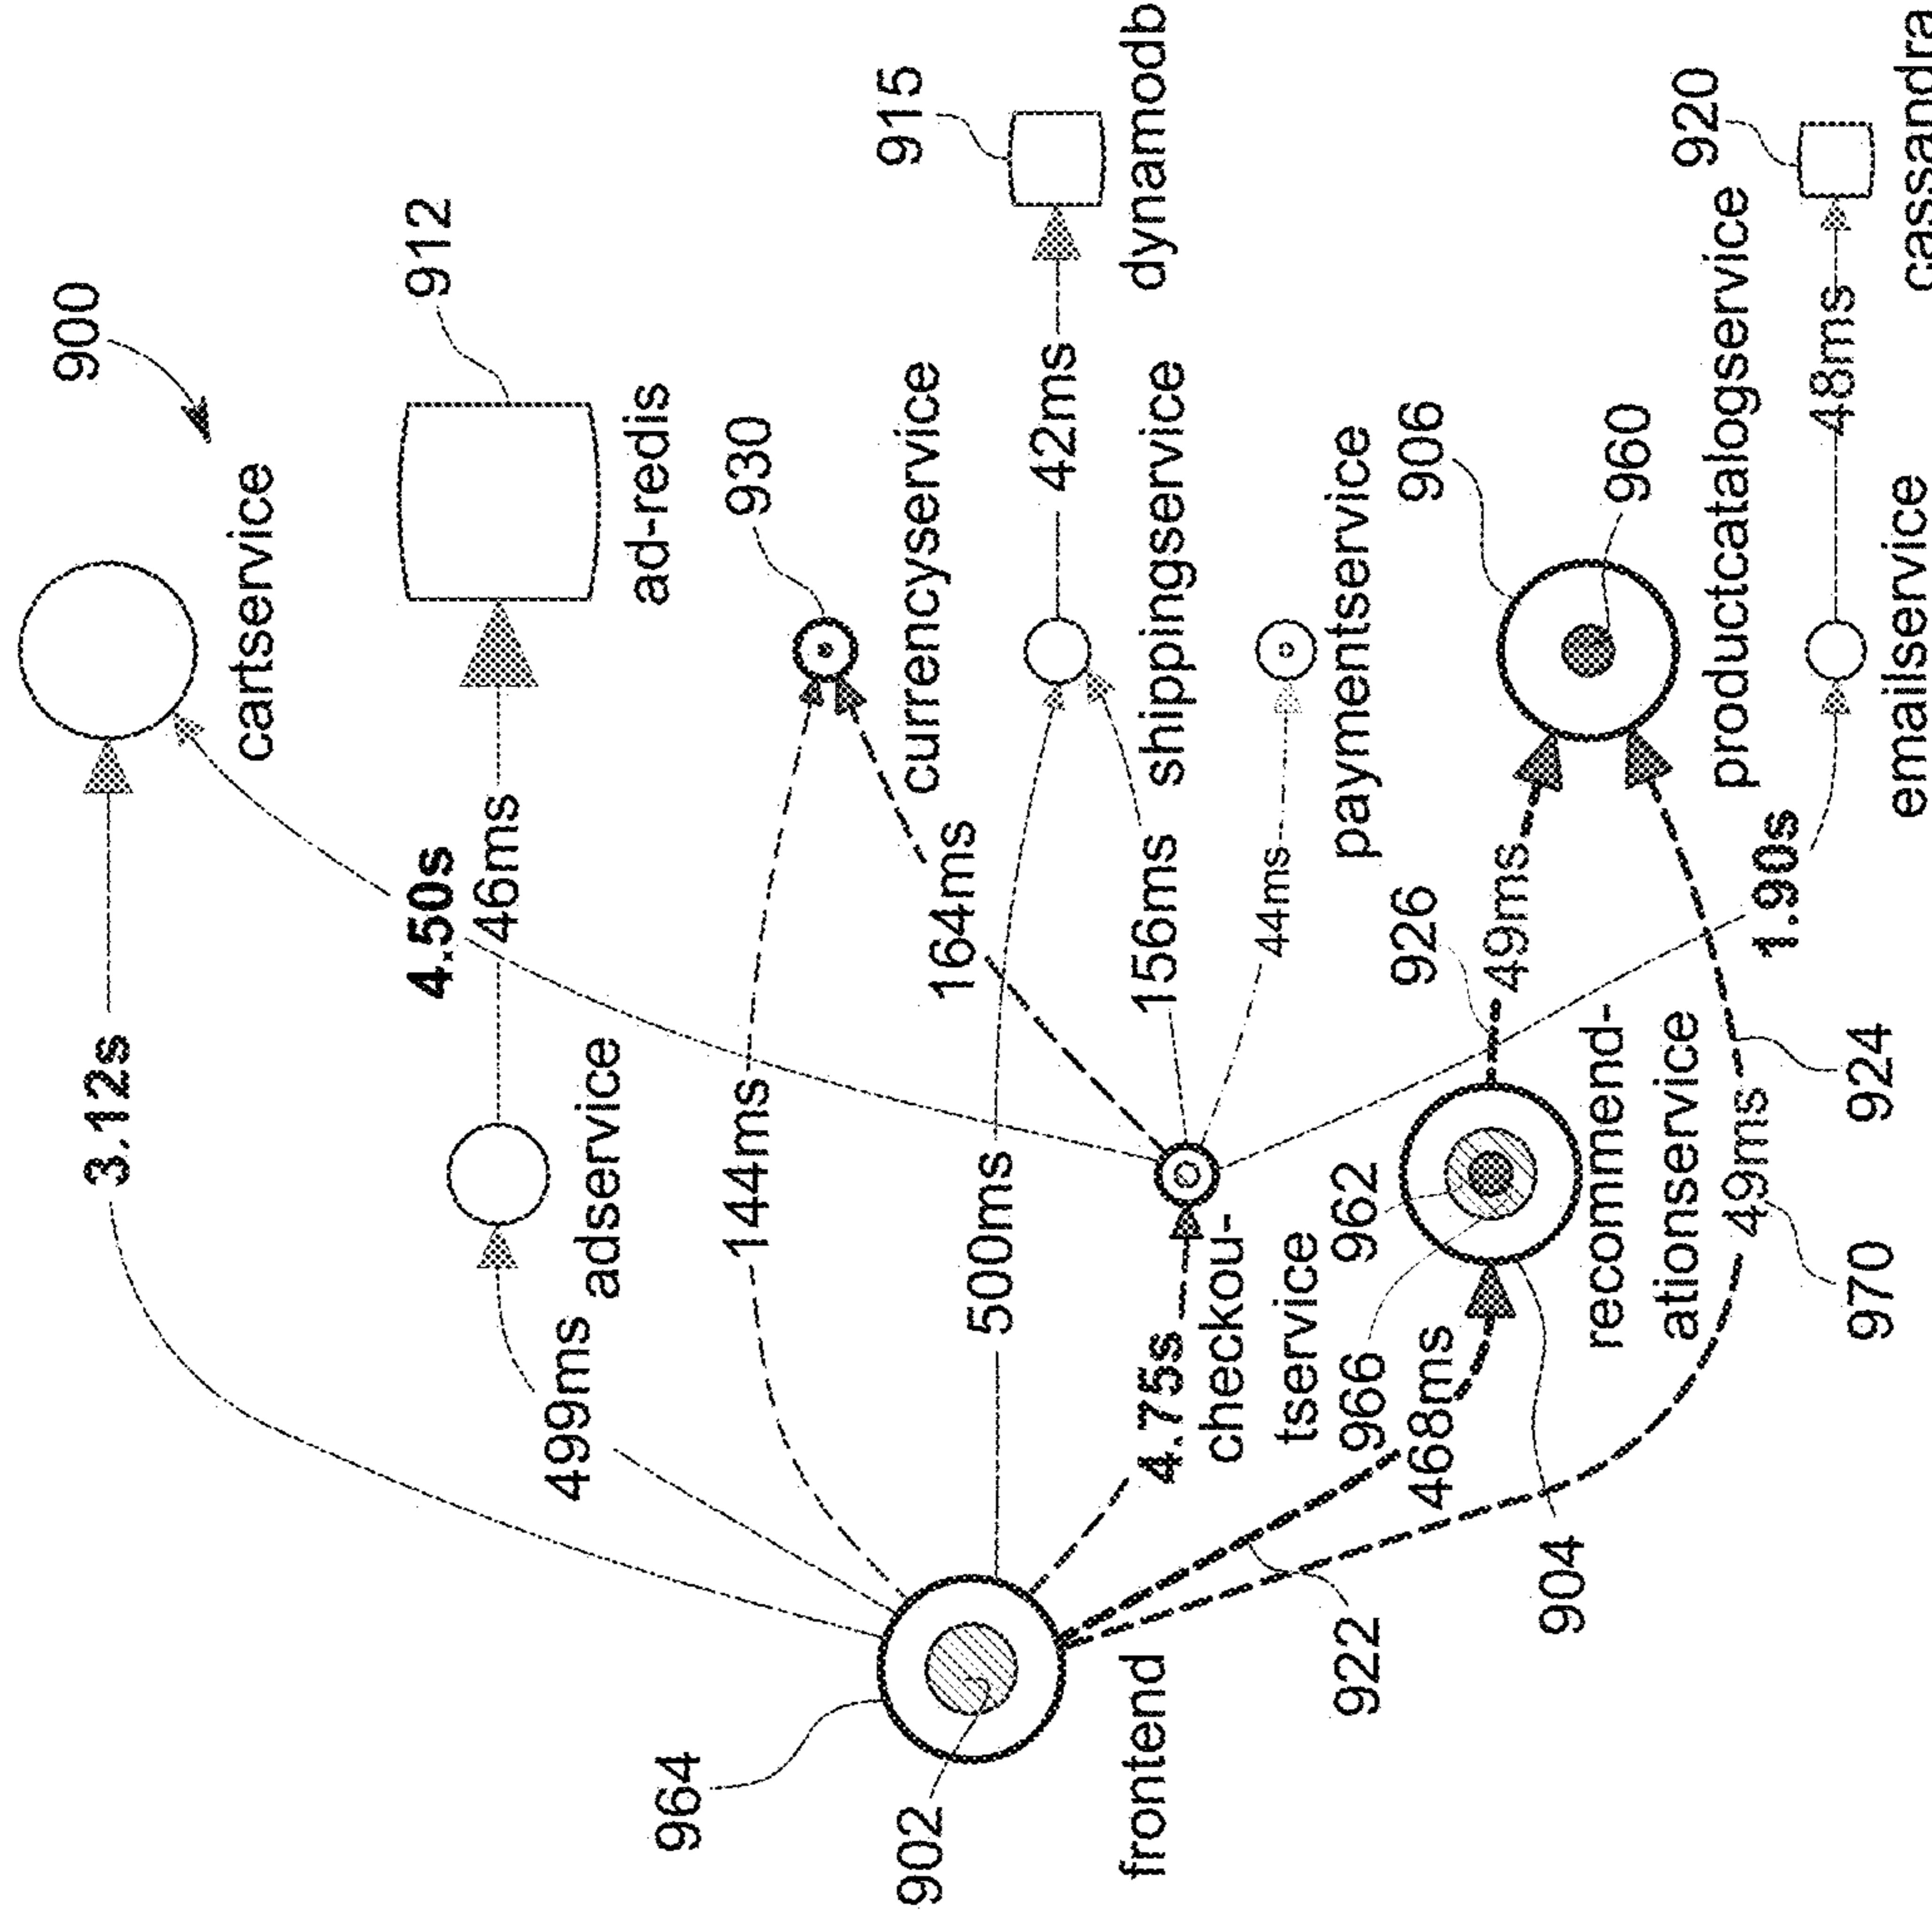
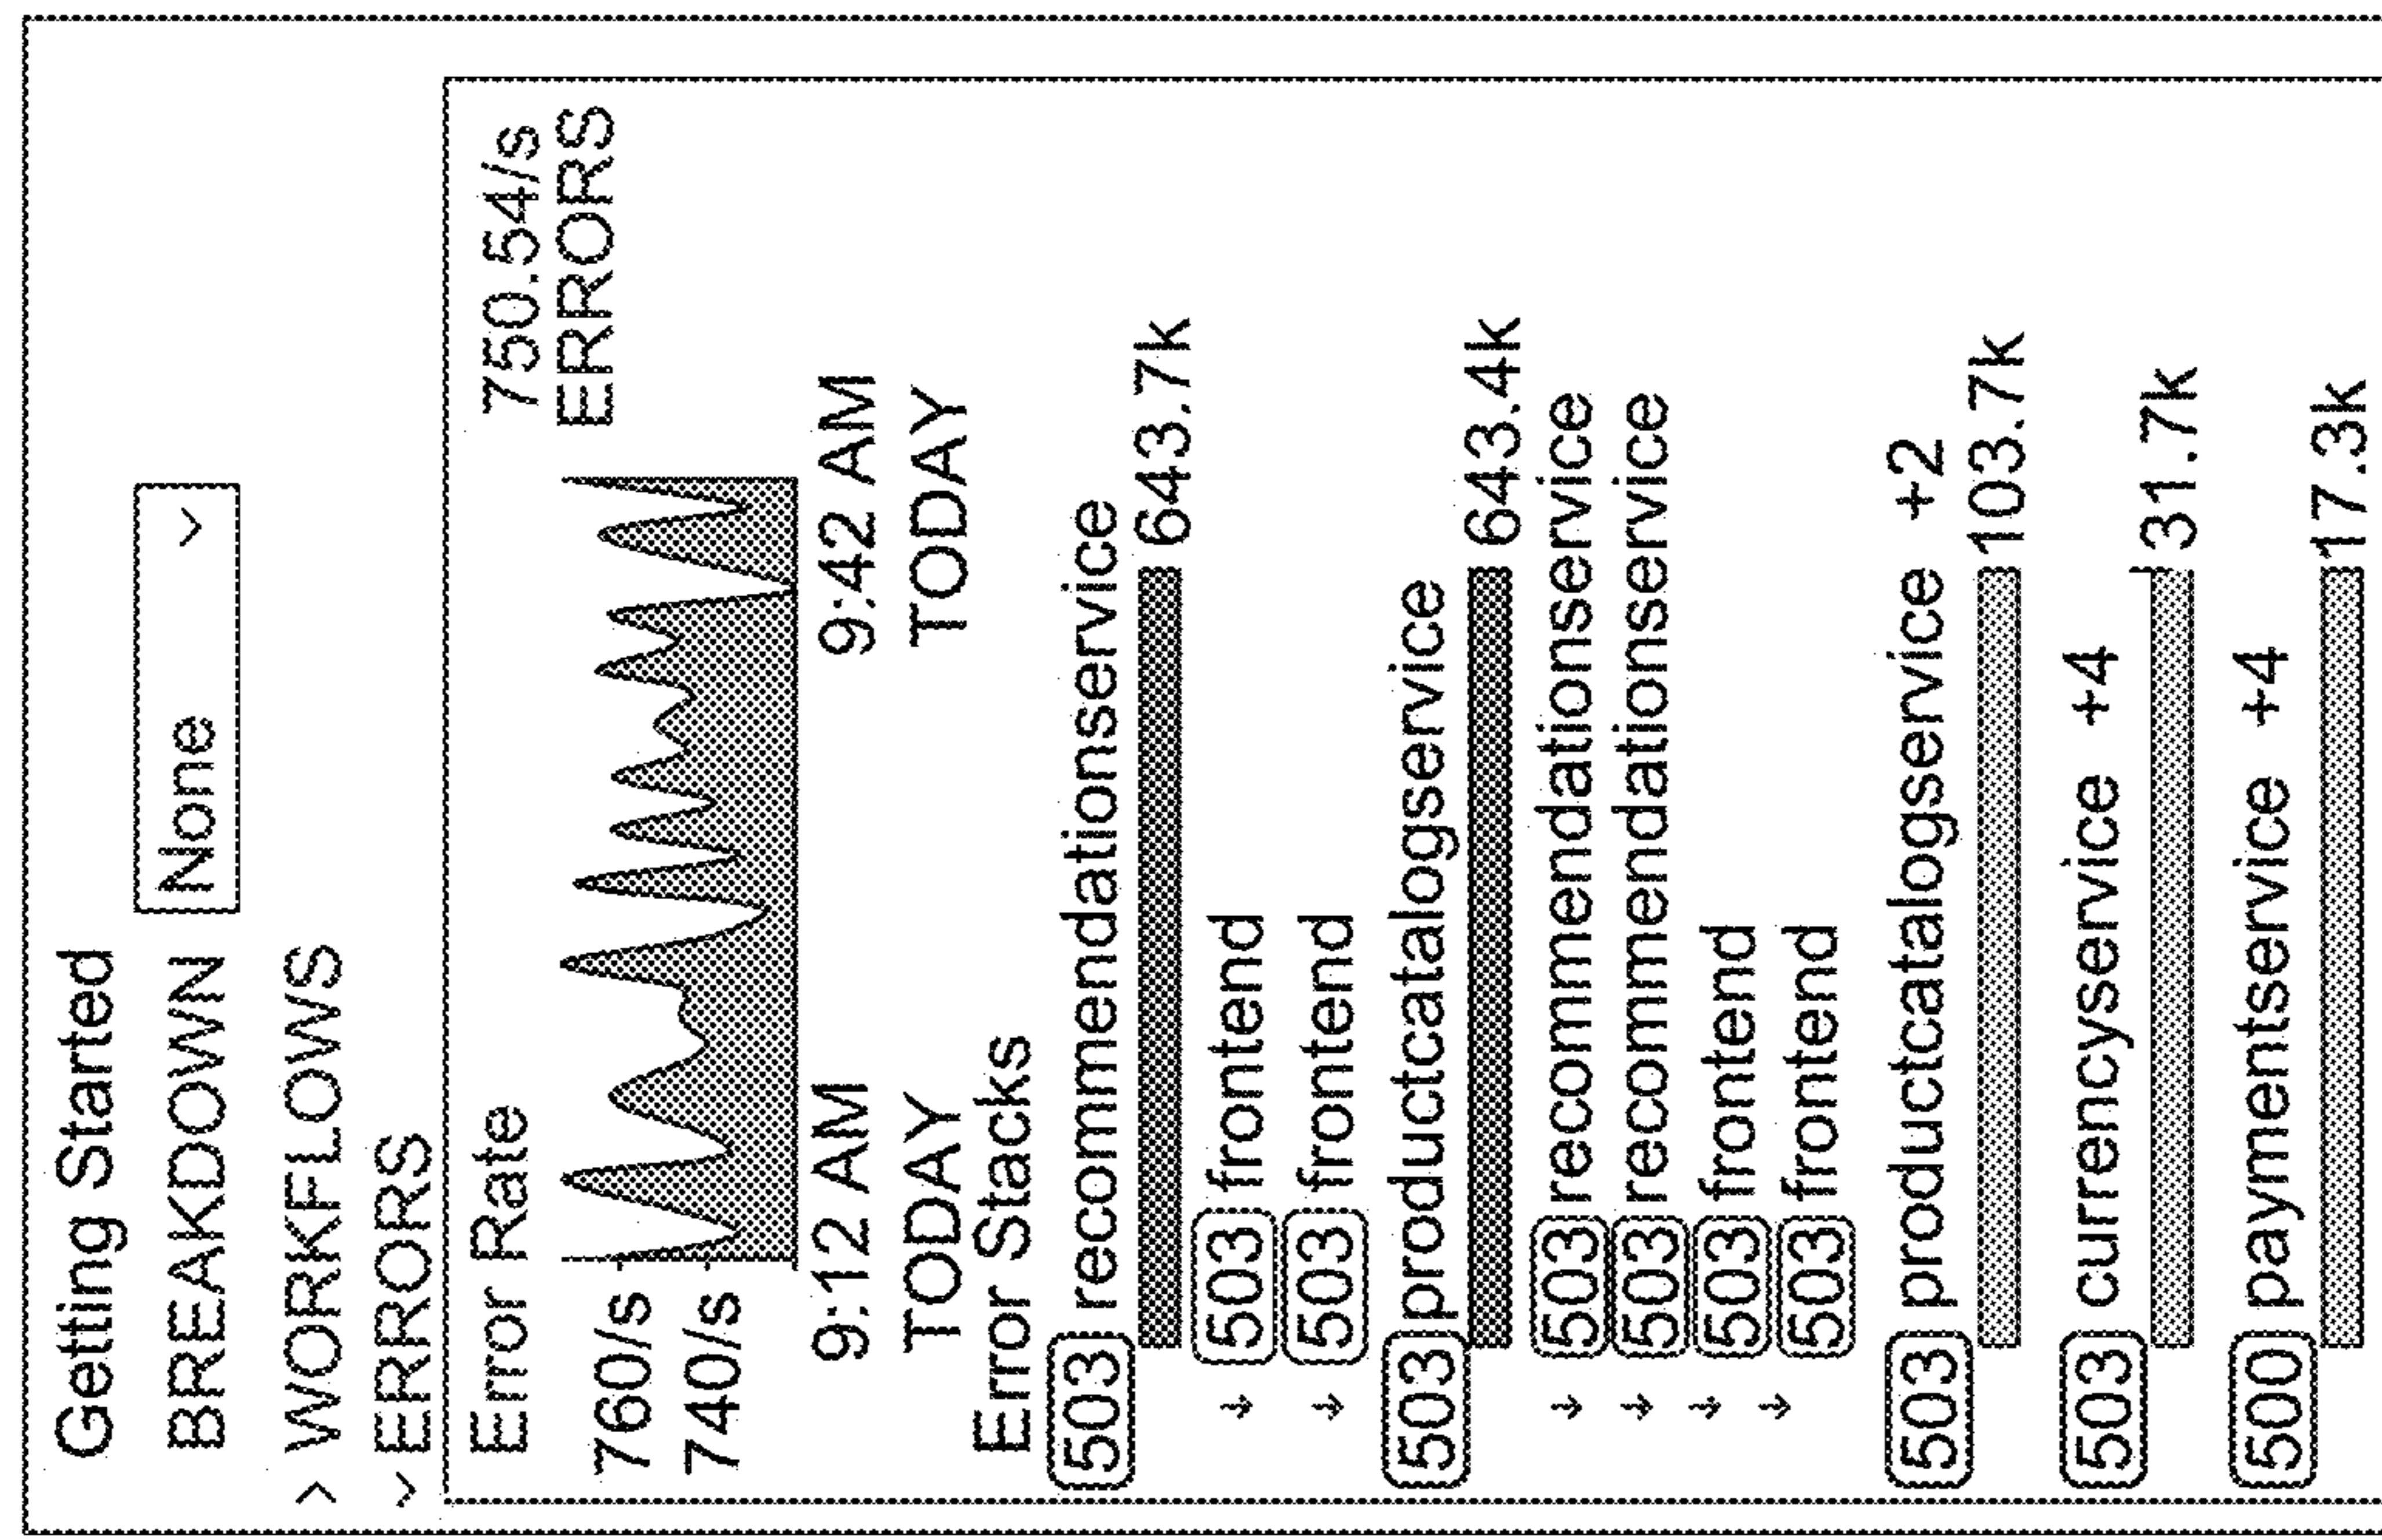
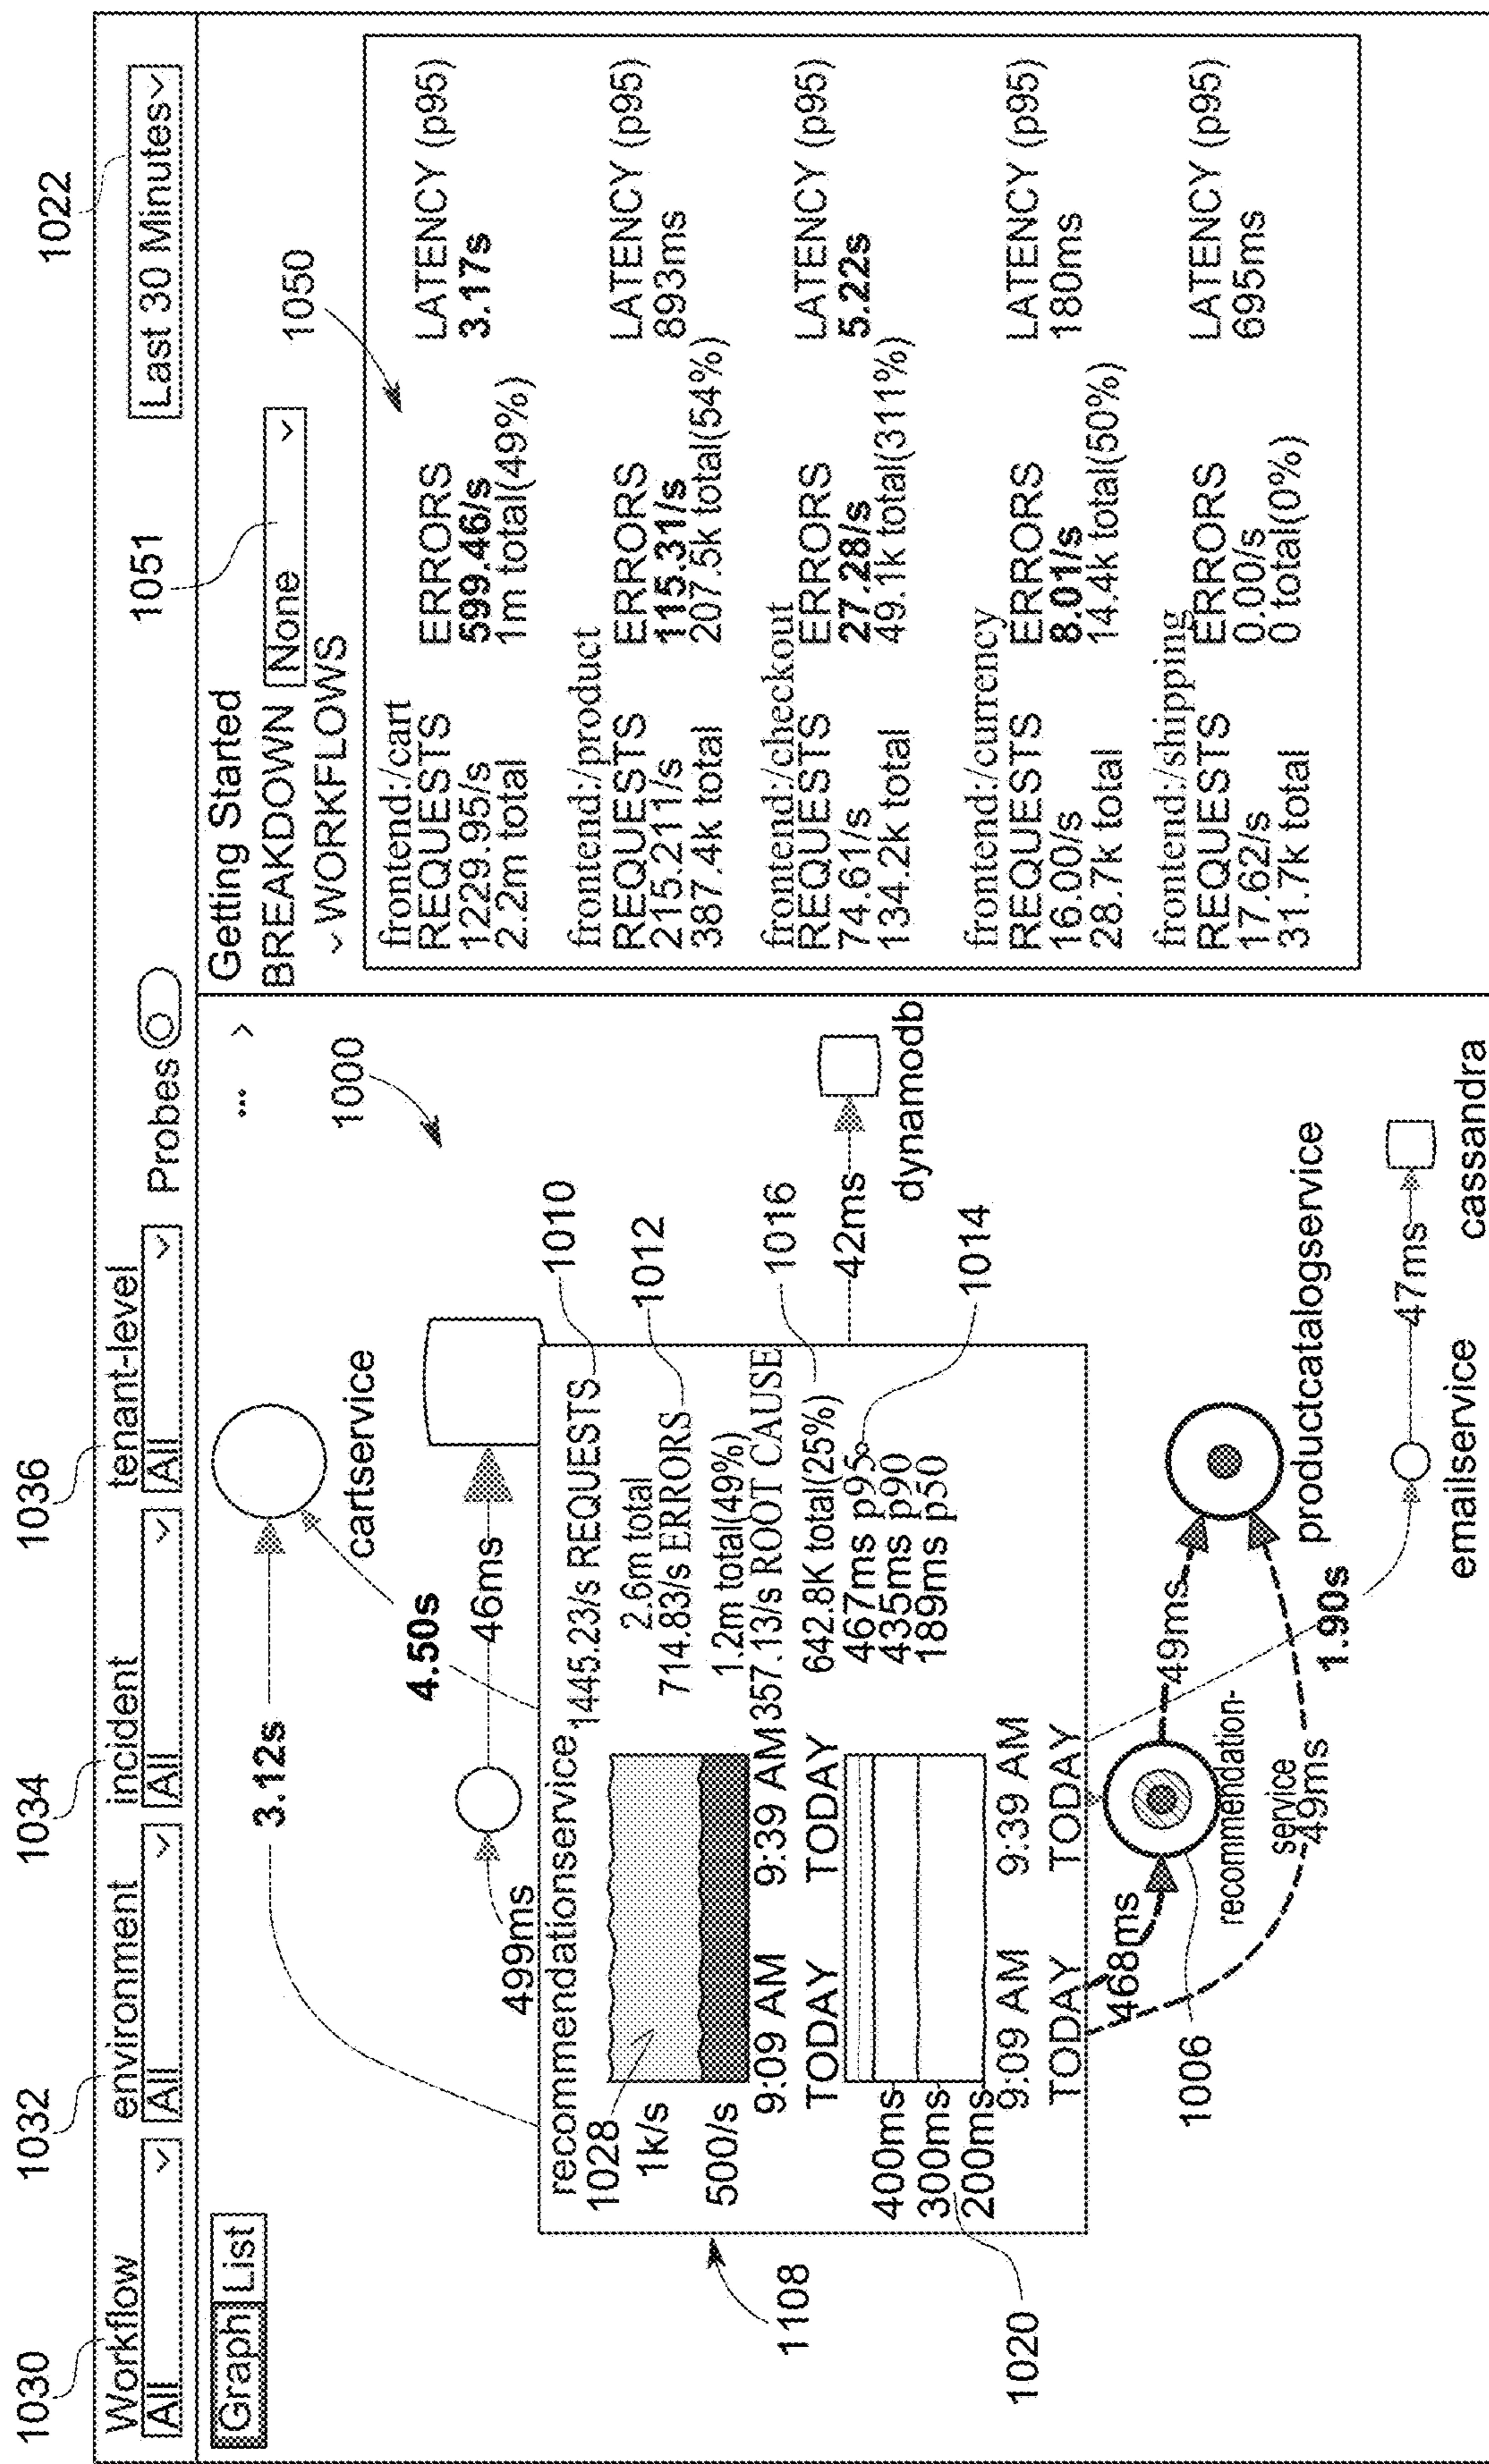


FIG. 9



卷之三

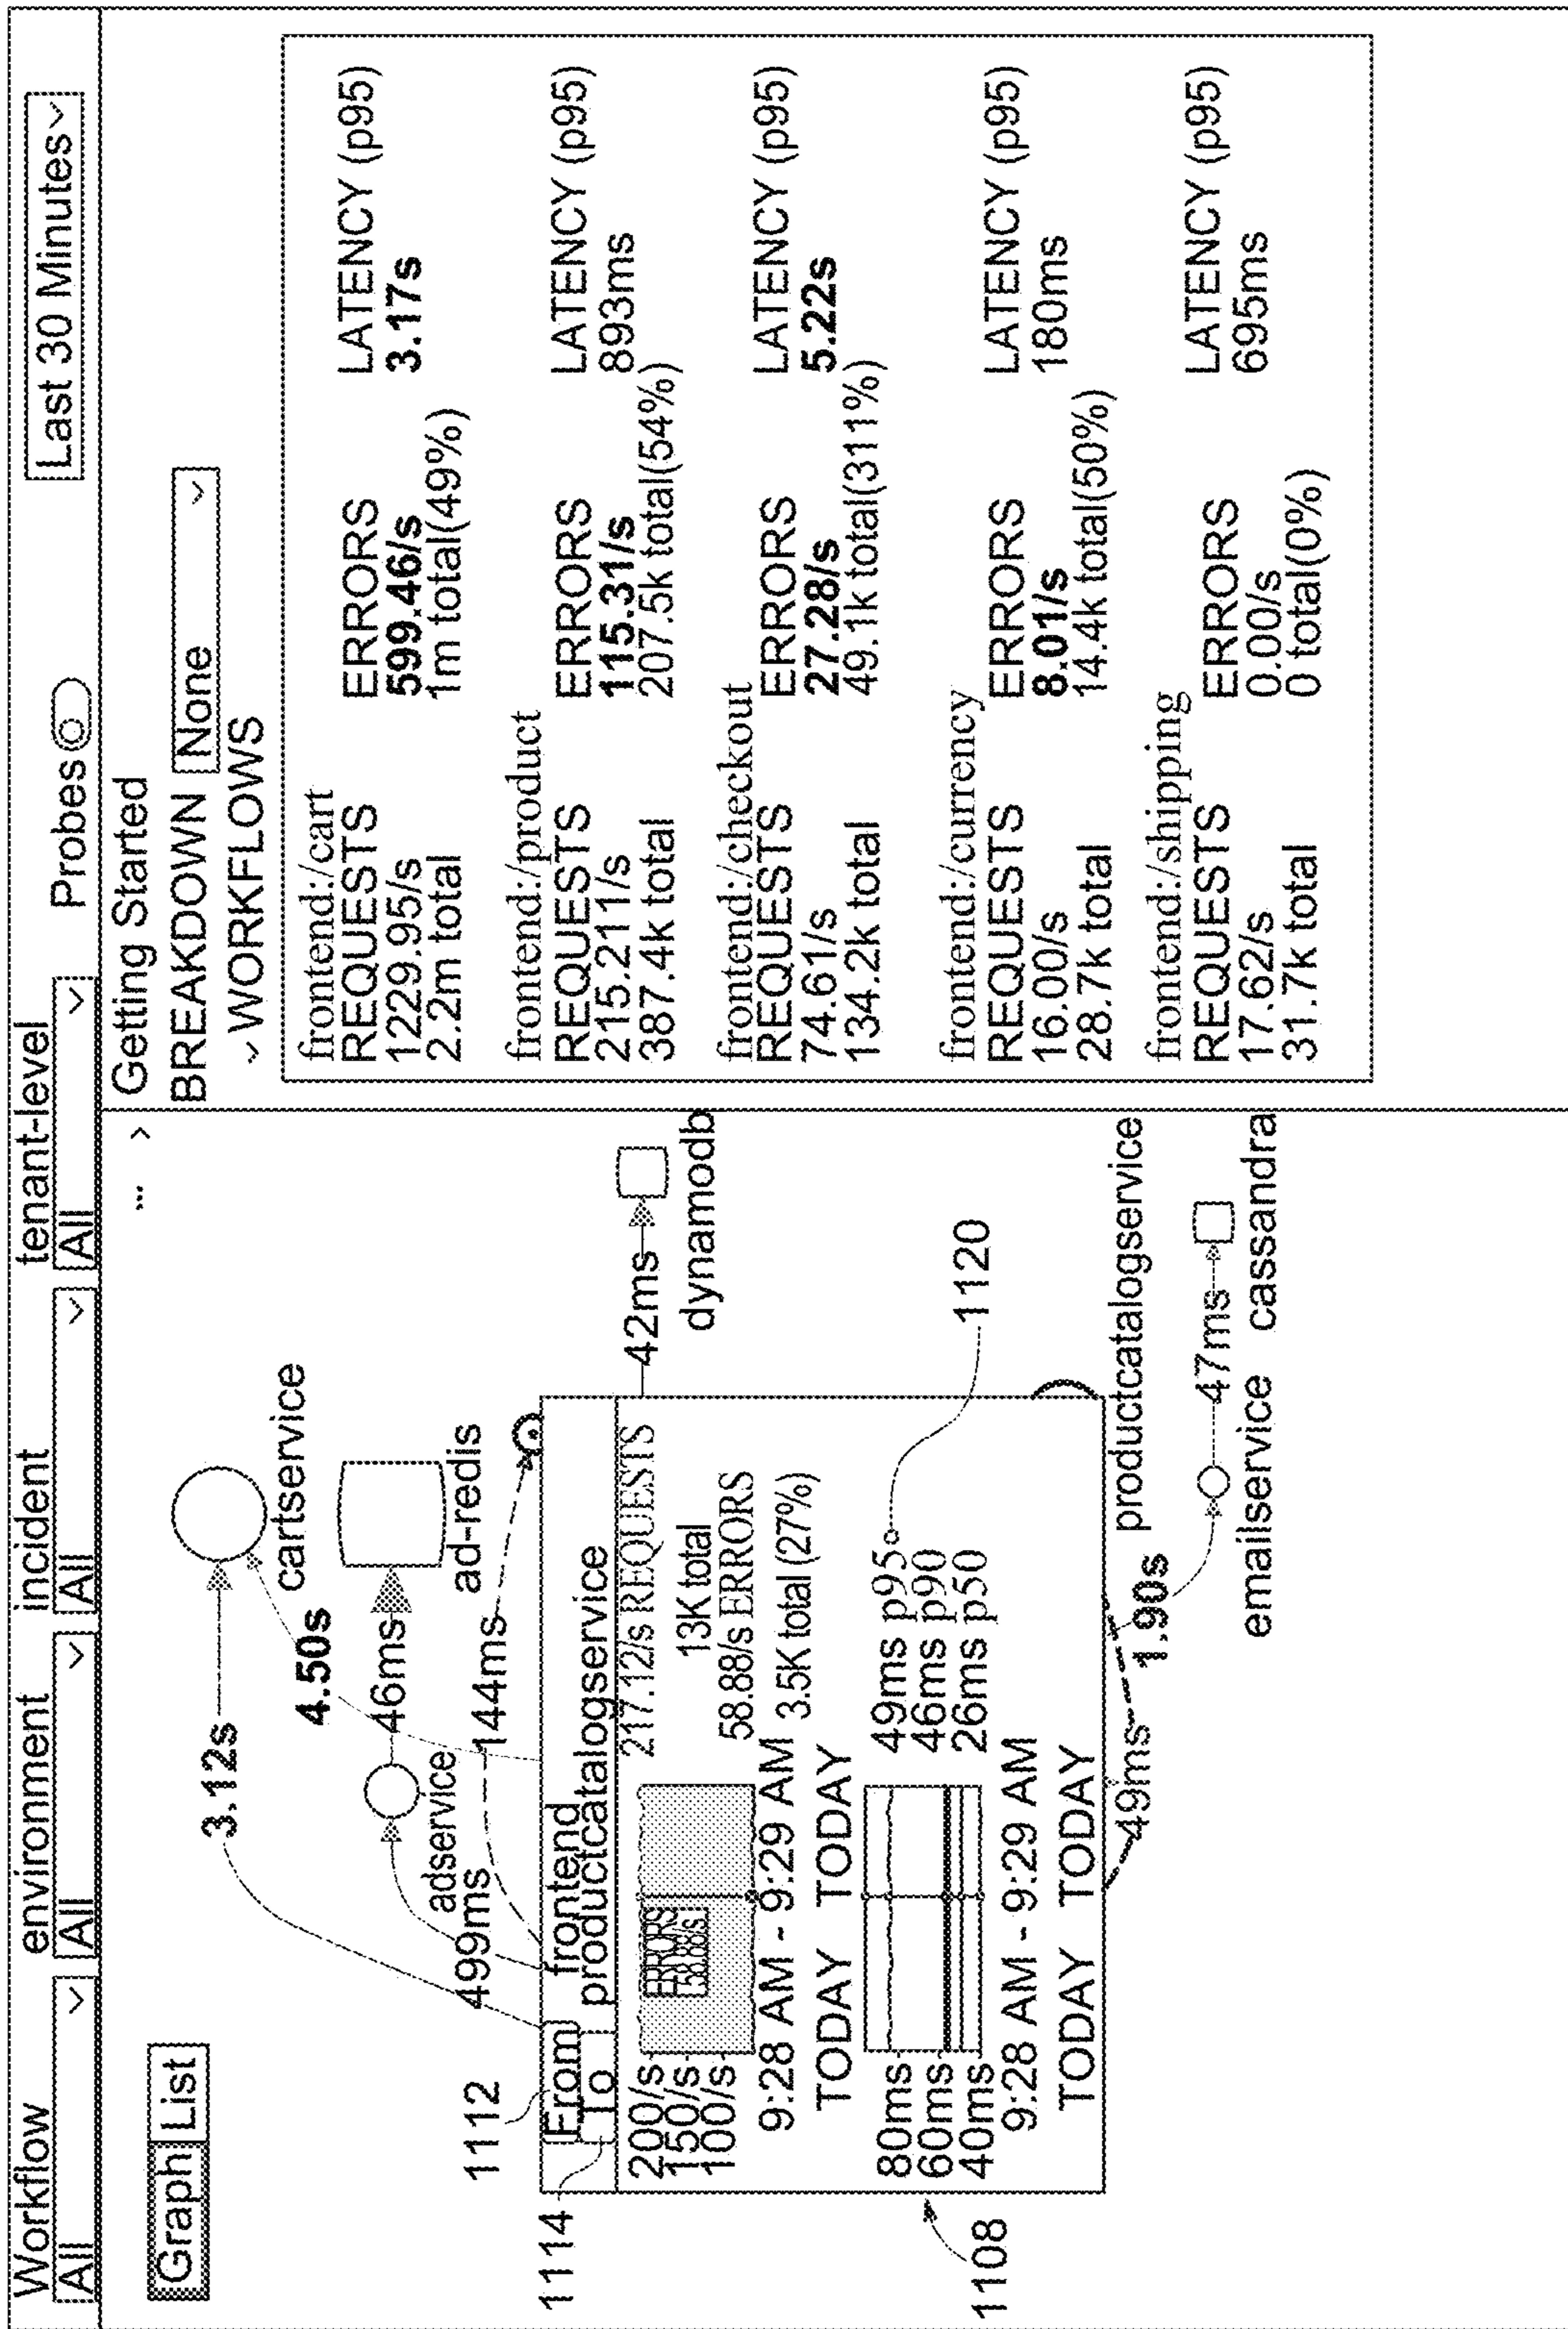


FIG. 11

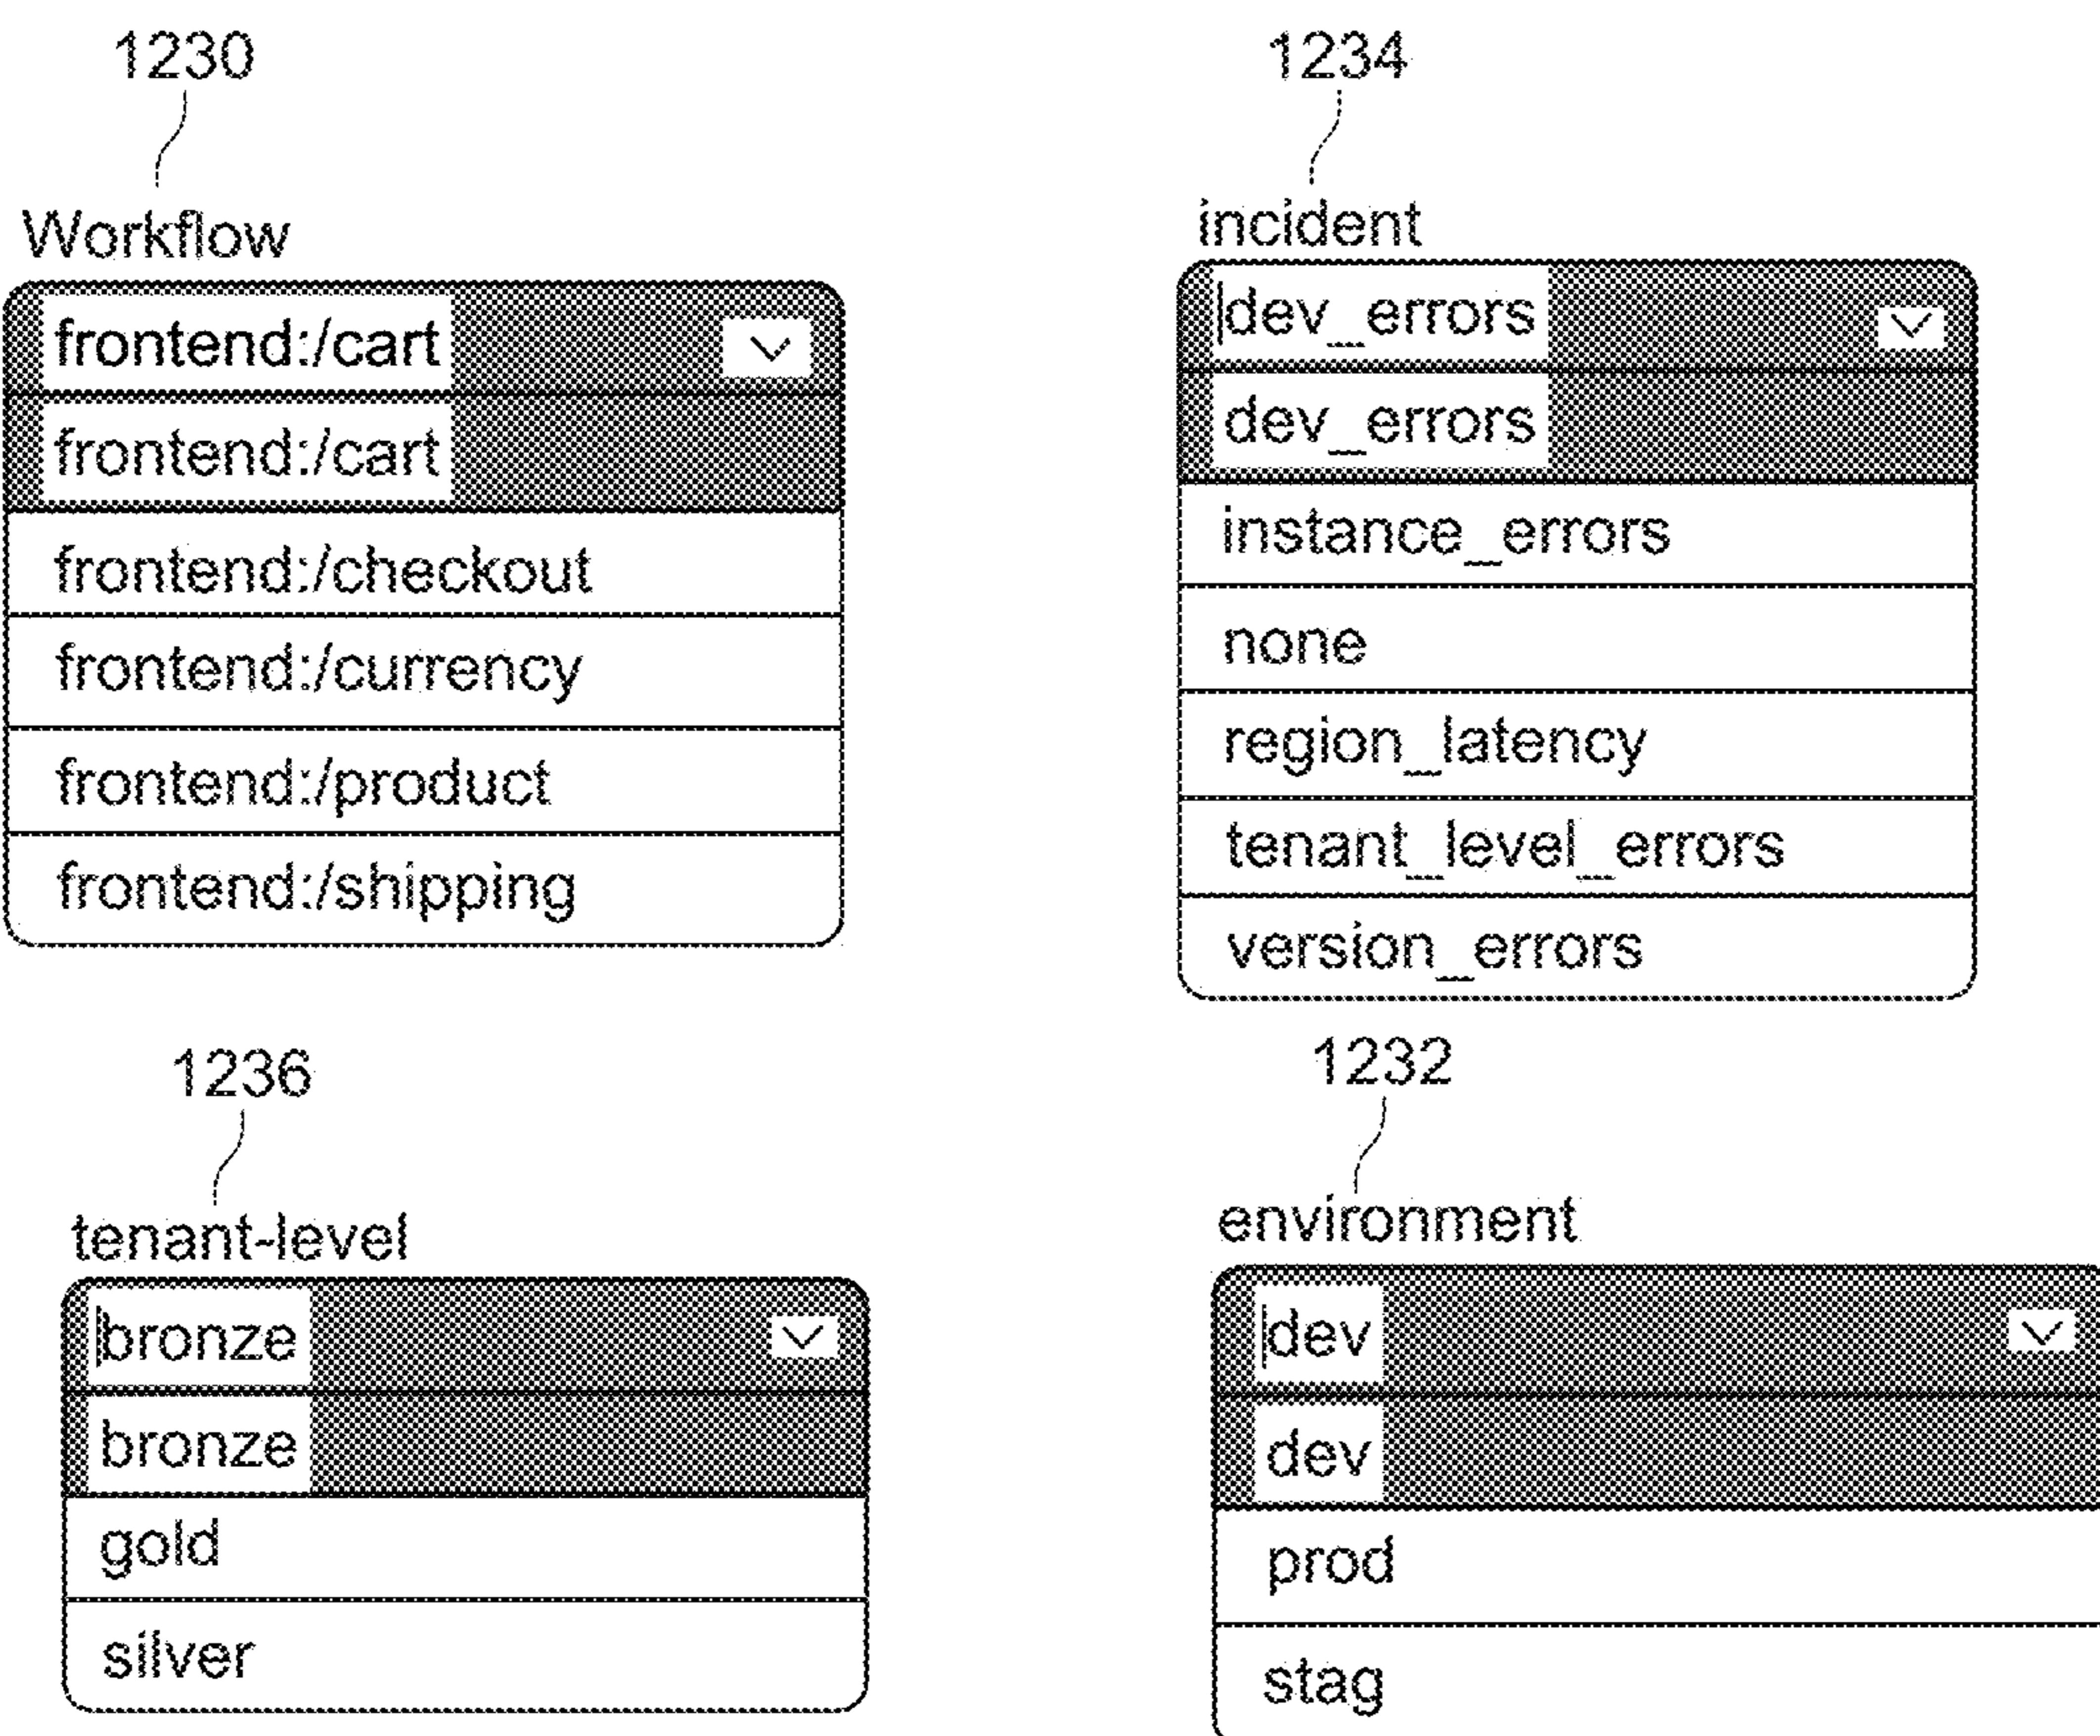


FIG. 12A

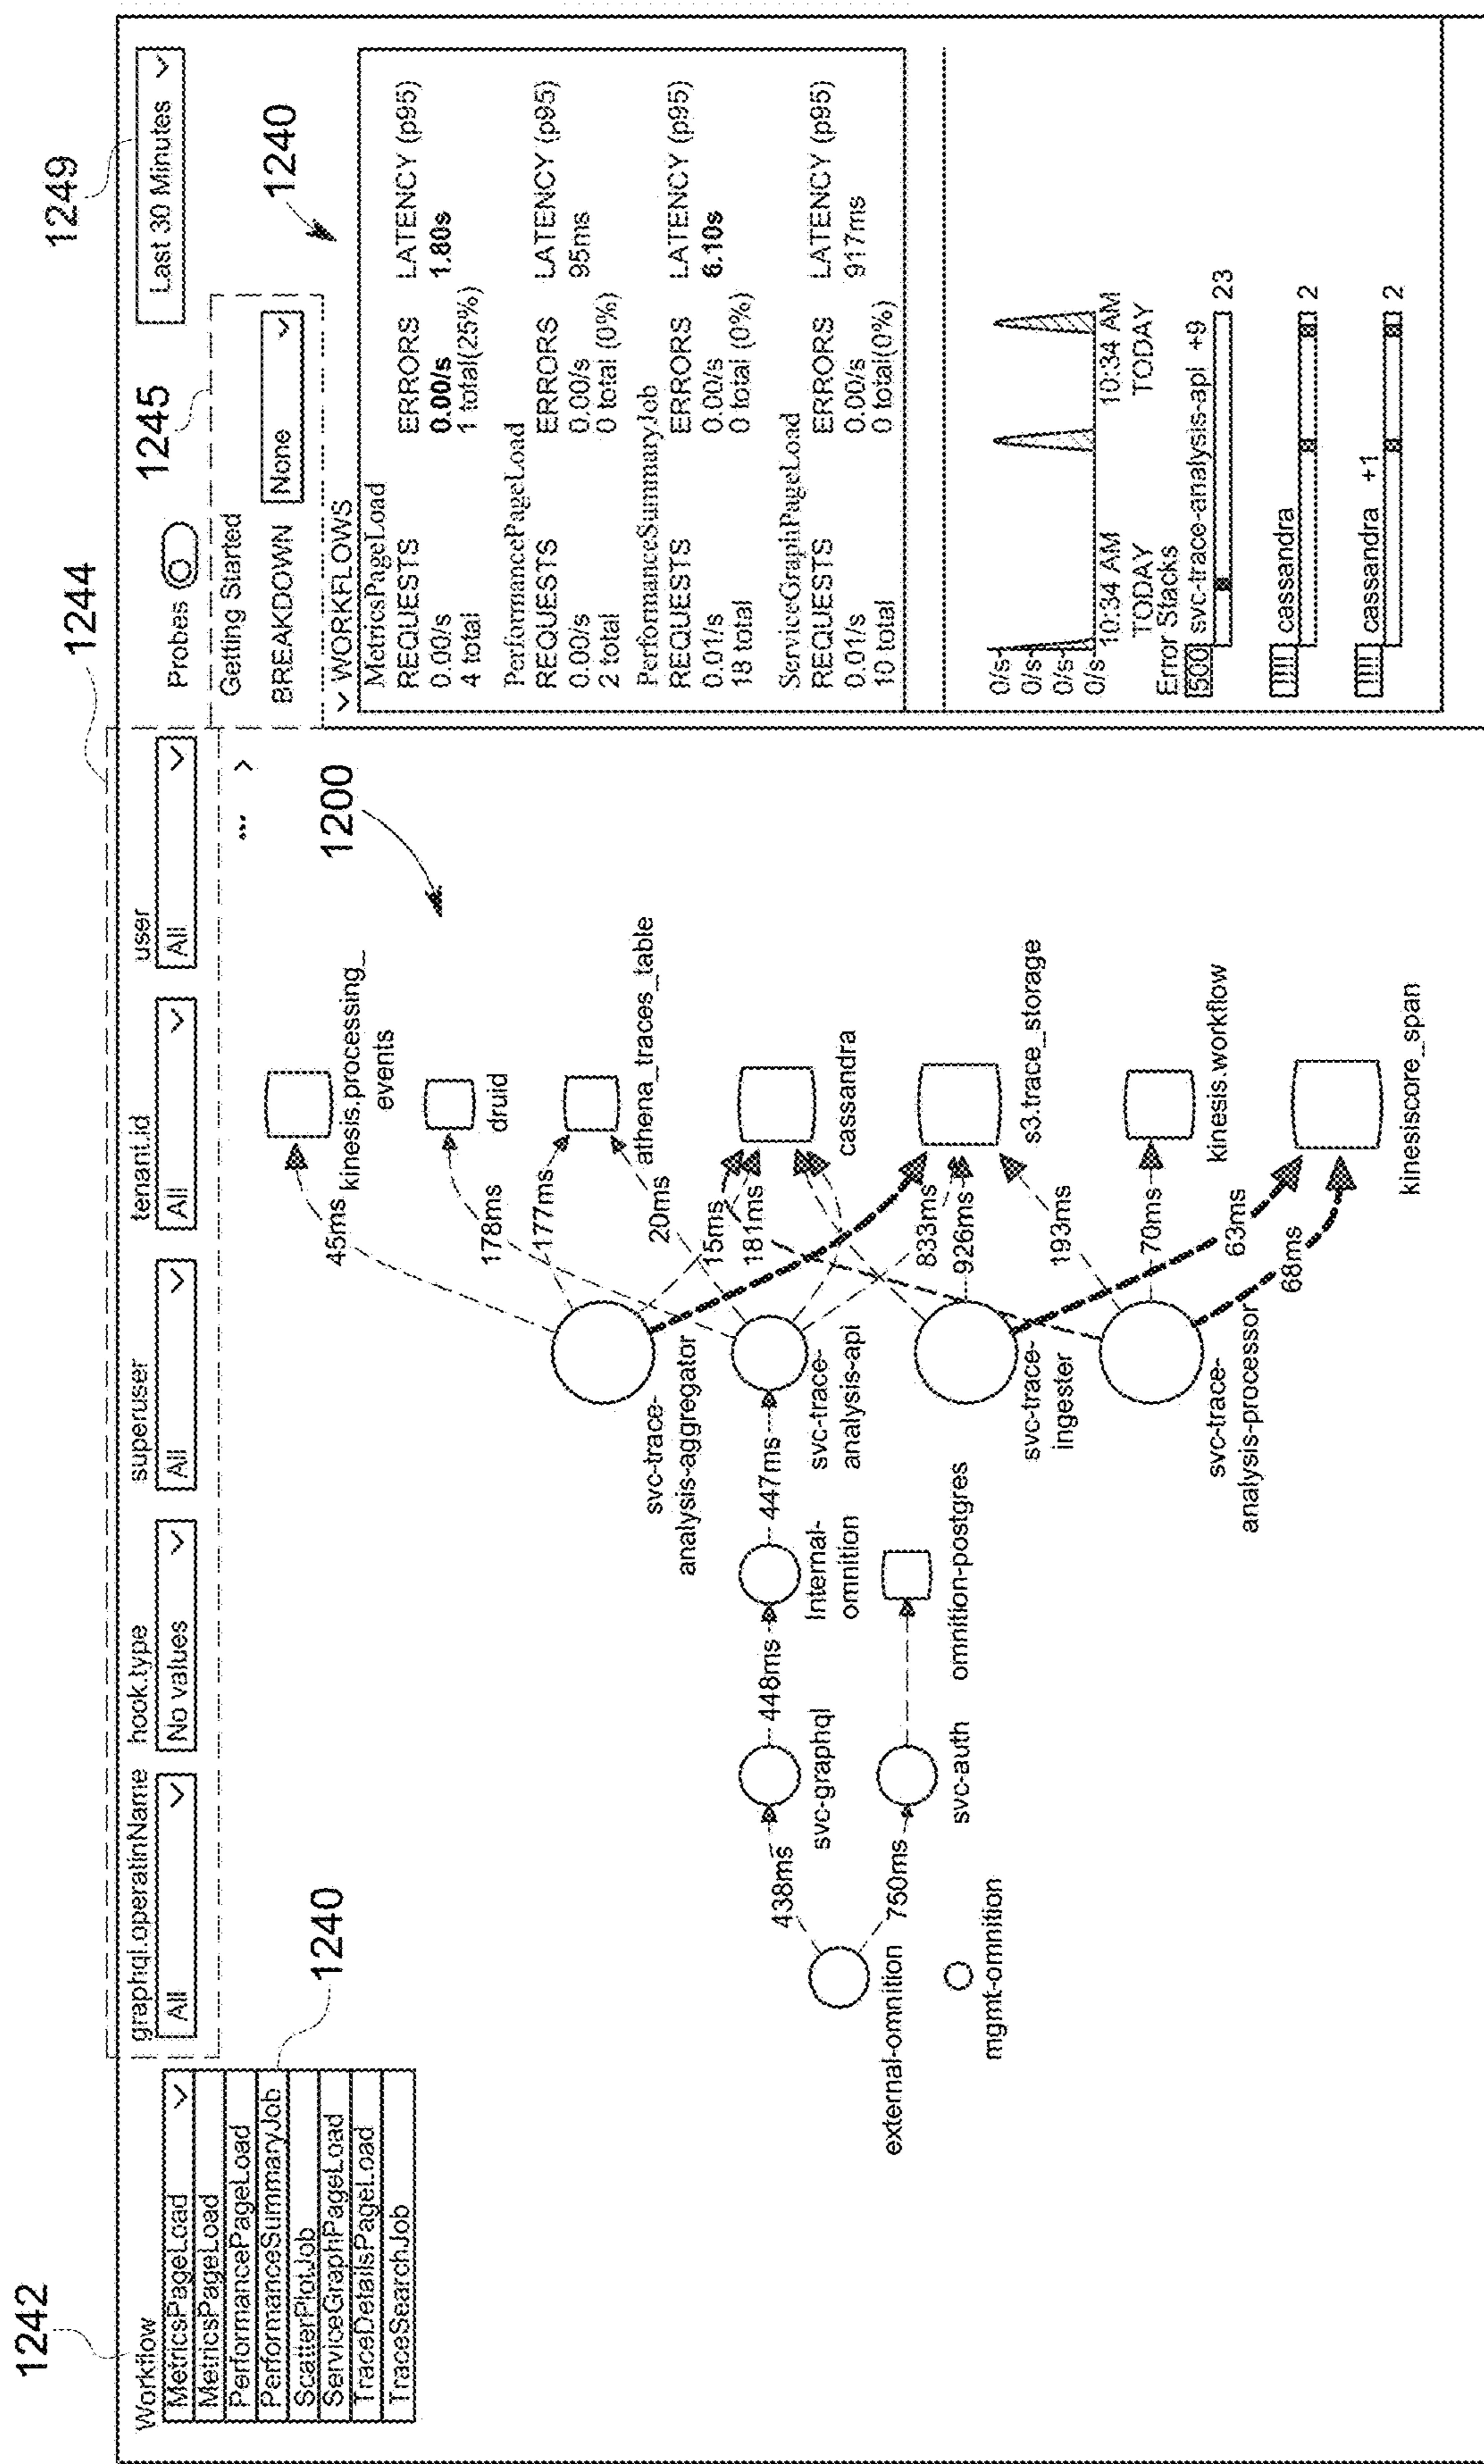
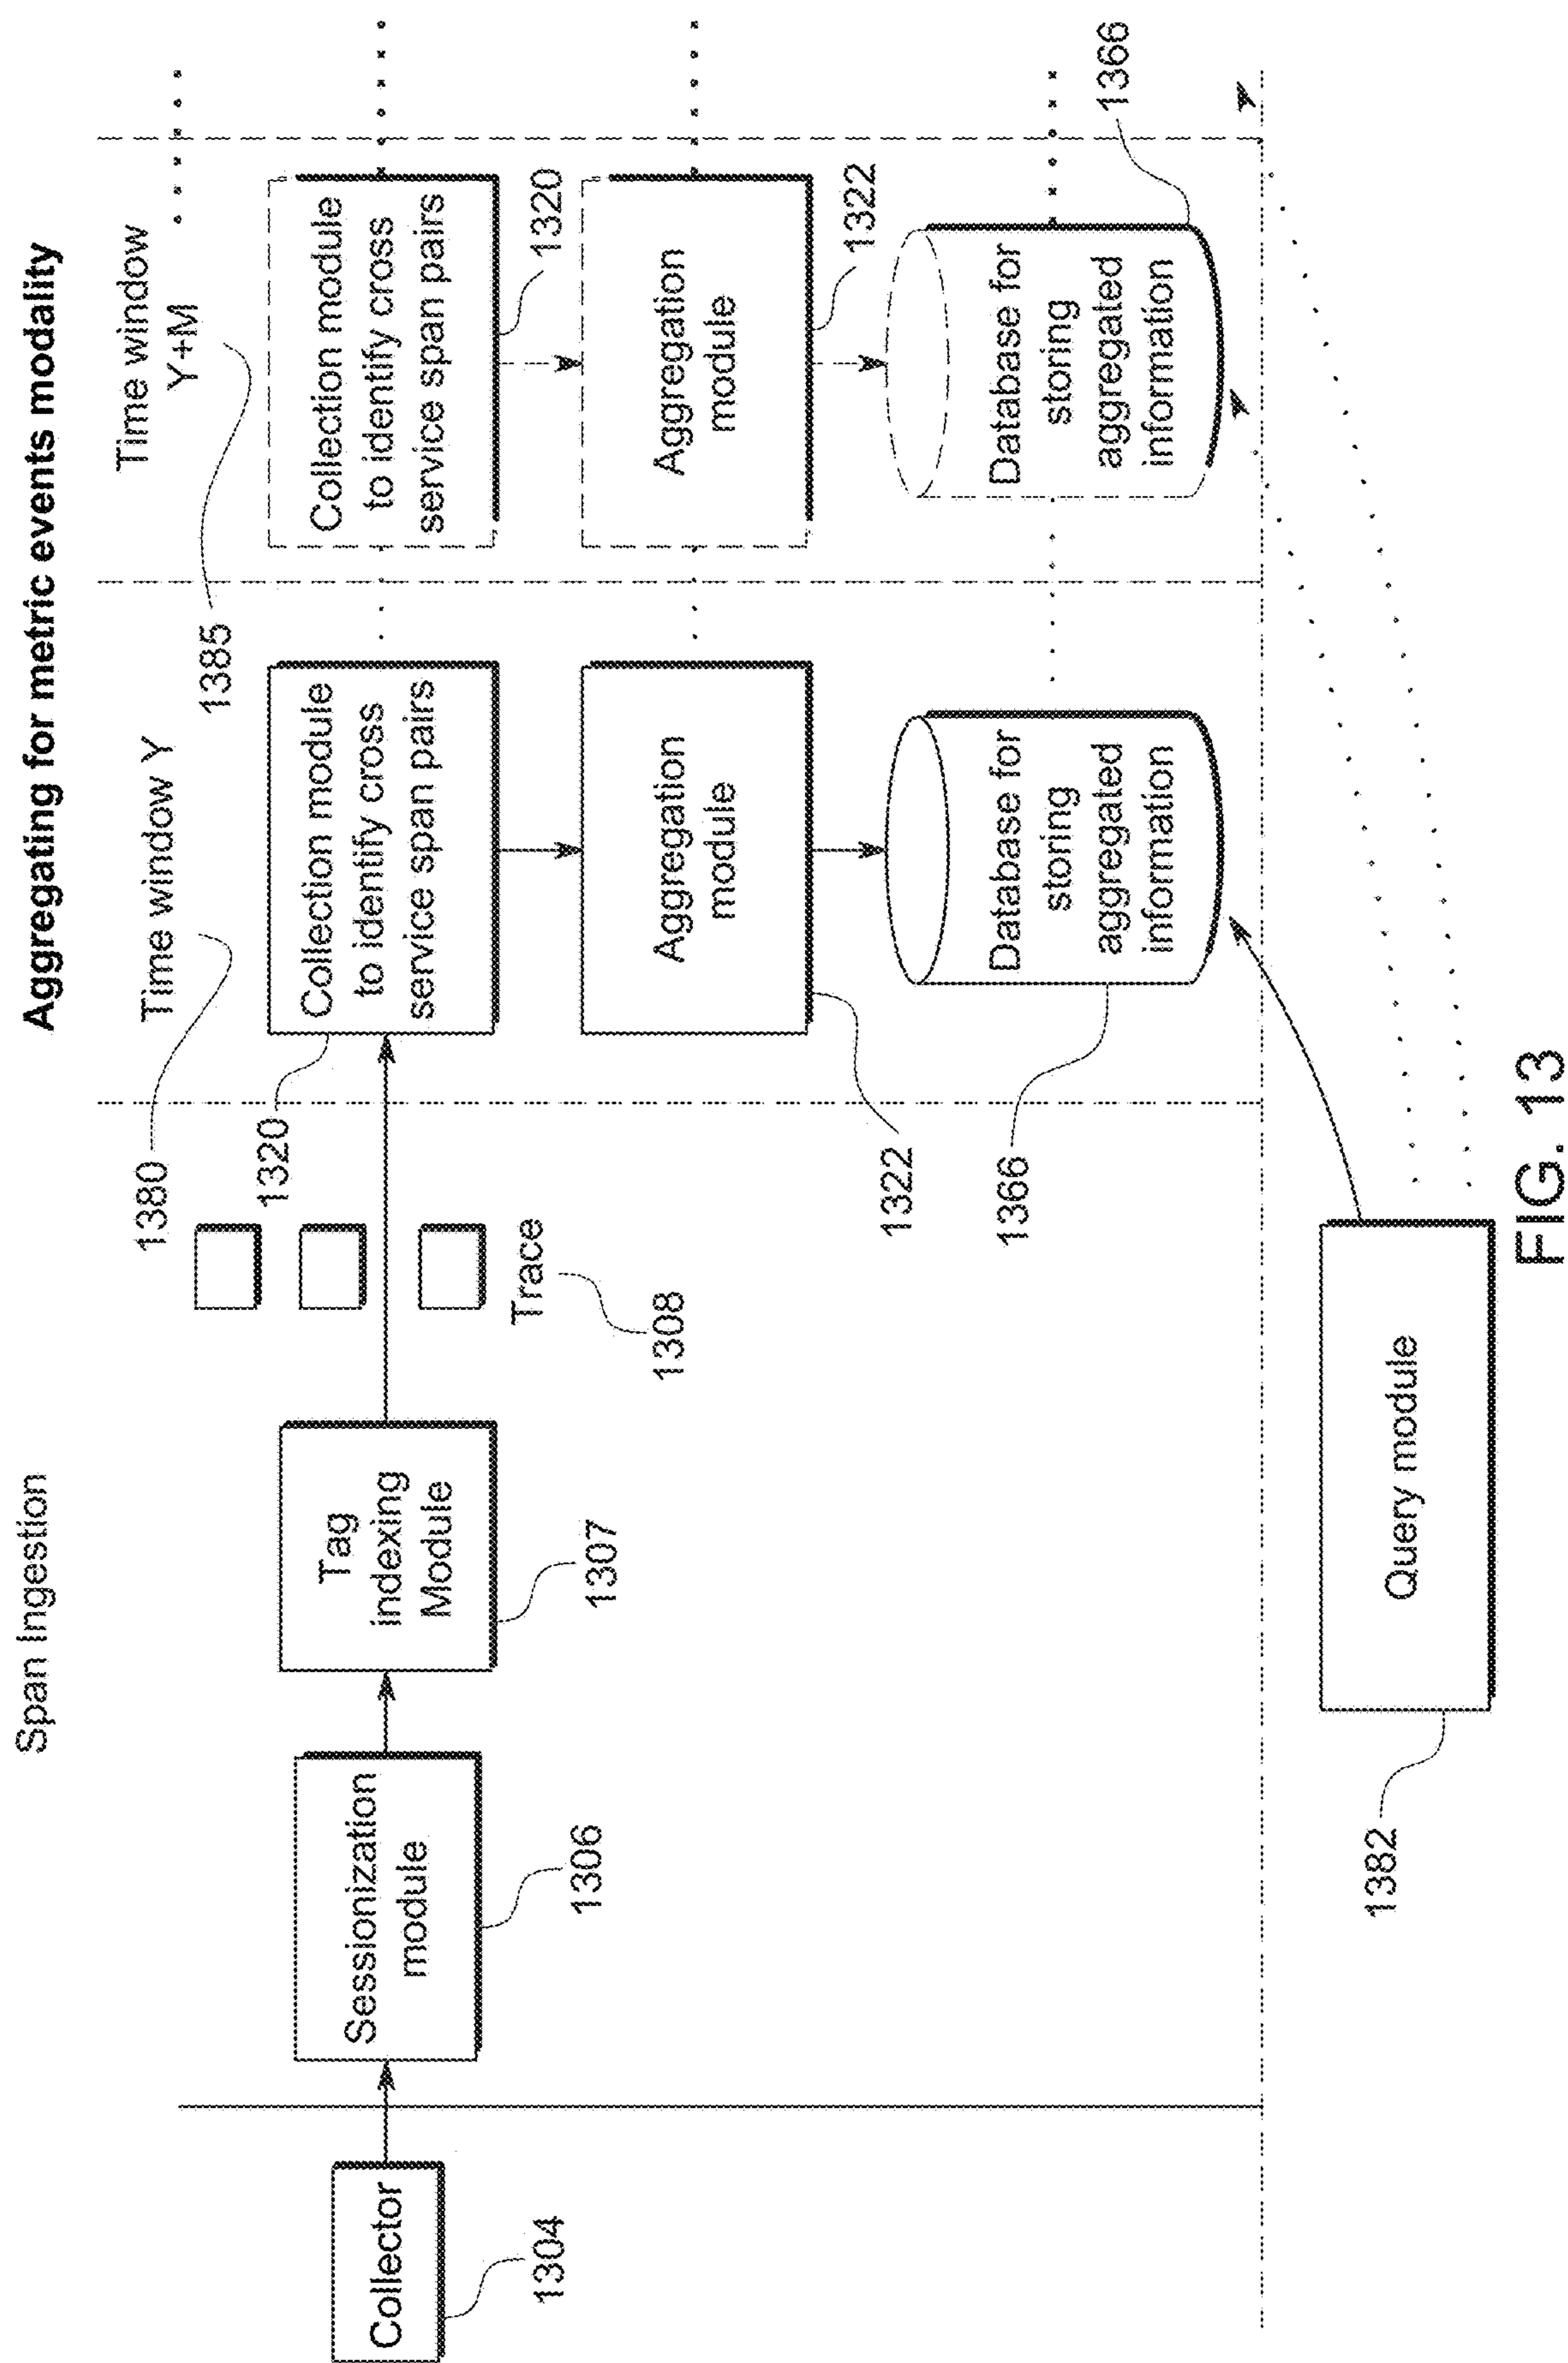


FIG. 12B



卷之三

from service name	to service name	from span.kind	to span.kind	from region	to region	Request count
Service A	Service B	client	server	us-west	us-east	2
Service A	Service B	client	server	us-west	us-west	1

1490 ↗
1492 ↗

FIG. 14

Service	Operation	1540	0ms	93ms	186ms	297ms	373ms
frontend	/product	1539	373ms				
	OPERATION						
	/product						
	TAGS						
	client		"ios"	1550			
	environment		"prod"				
	http.method		"GET"	1551			
	http.status_code		200				
	http.url		"http://frontend/product"				
	incident		"pubsub_errors"				
	ios.version		"v1.1.0"				
	k8s.io/cluster/name		"us-east-1-cluster"				
	k8s.io/container/name		"front-end"				
	k8s.io/namespace/name		"us-east-1-namespace-2"	1552			
	k8s.io/pod/name		"frontend-daliuytqq-dd58f"				
	region		"us-east-1"				
	sampler.param		true				
	sampler.type		"const"				
	status.code		0				
	tenant-level		"bronze"				
	version		"76c81d"				
	PROCESS						
	hostname		"util-svc-synthetic-load-generator-7c88bbf4fb-fts96"				
	ip		"100.96.12.172"				
	opencensus.exporterversion		"Jaeger-Java-0.33.1"				
frontend	parse-request	1546	15ms	■			
frontend	request/GetProducts	1547	30ms	■■■			
productcatalogservice	/GetProducts	1538	29ms	■■■■			
	OPERATION	1545					
	/GetProducts						
	TAGS						
	environment		"prod"	1560			
	http.method		"GET"	1561			
	http.status_code		200				
	http.url		"http://productcatalogservice/GetProducts"				
	k8s.io/cluster/name		"us-east-1-cluster"				
	k8s.io/container/name		"productcatalogservice"				
	k8s.io/namespace/name		"us-east-1-namespace-2"	1562			
	k8s.io/pod/name		"productcatalogservice-18iodsfh27-r234t"				
	products_returned		115				
	region		"us-east-1"				
	status.code		0				
	tenant-level		"bronze"				
	version		"deb52c"				
	PROCESS						
	hostname		"util-svc-synthetic-load-generator-7c88bbf4fb-fts96"				
	ip		"100.96.12.172"				
	opencensus.exporterversion		"Jaeger-Java-0.33.1"				

FIG. 15A

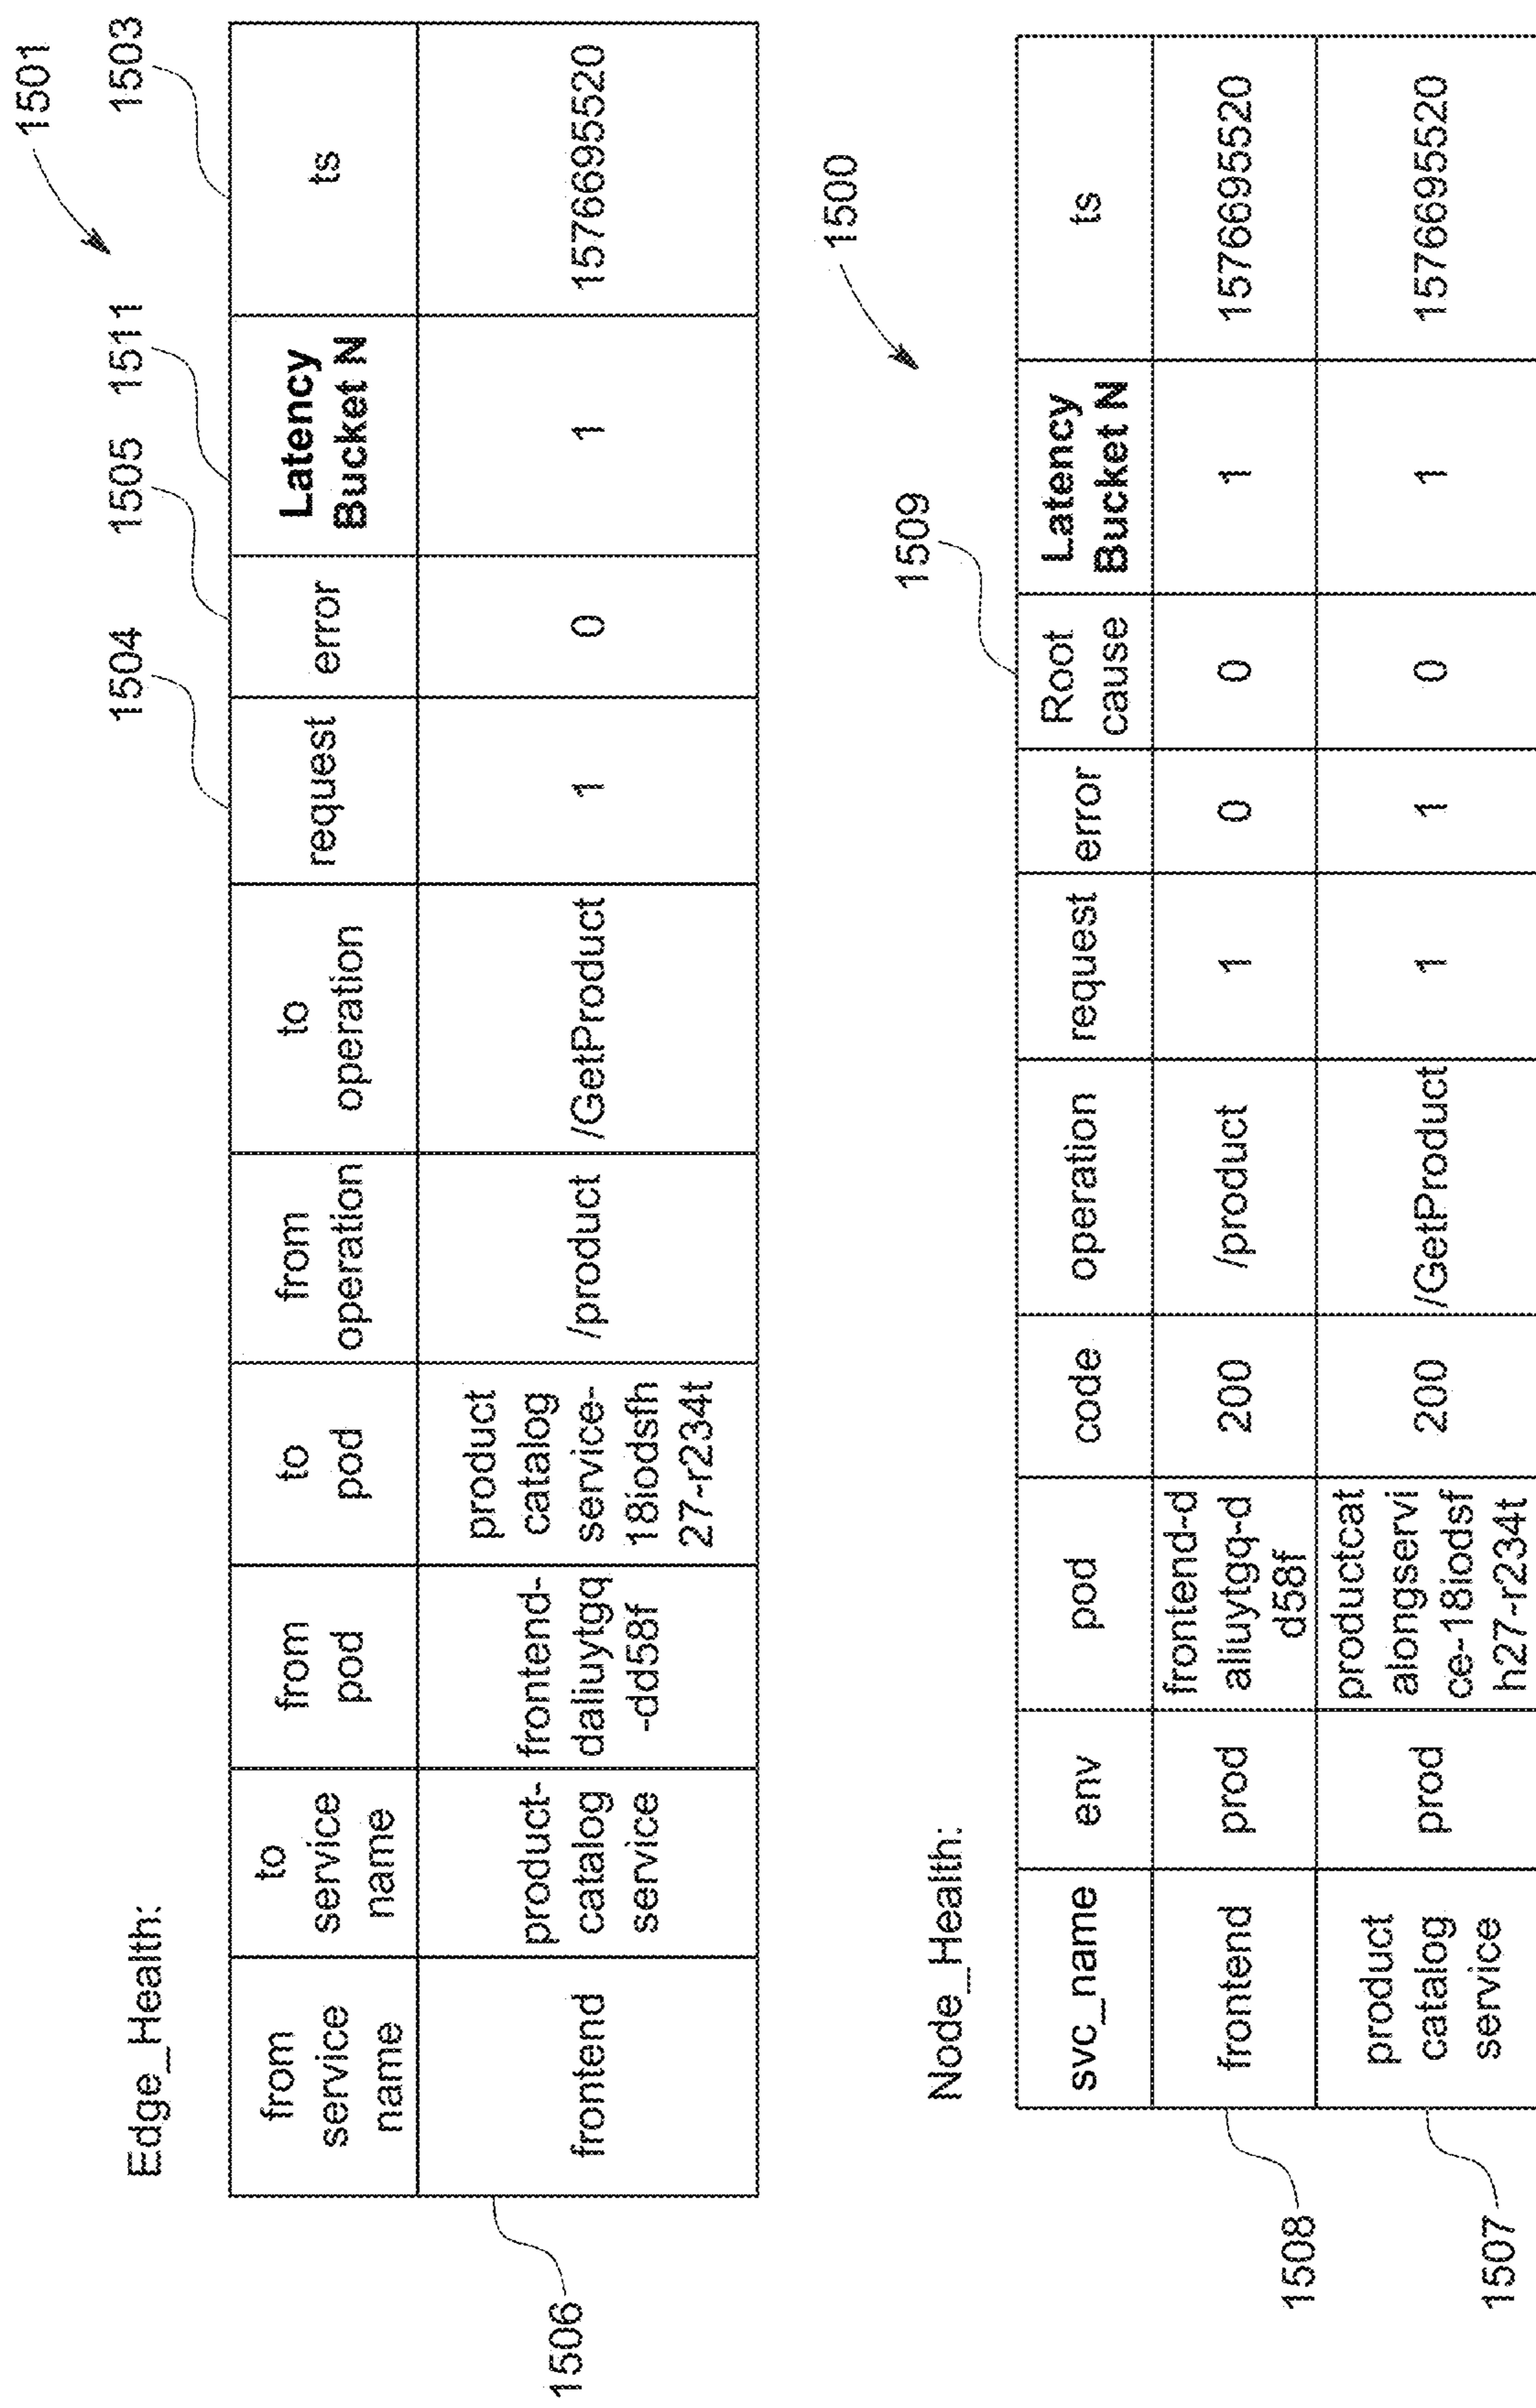


FIG. 15B

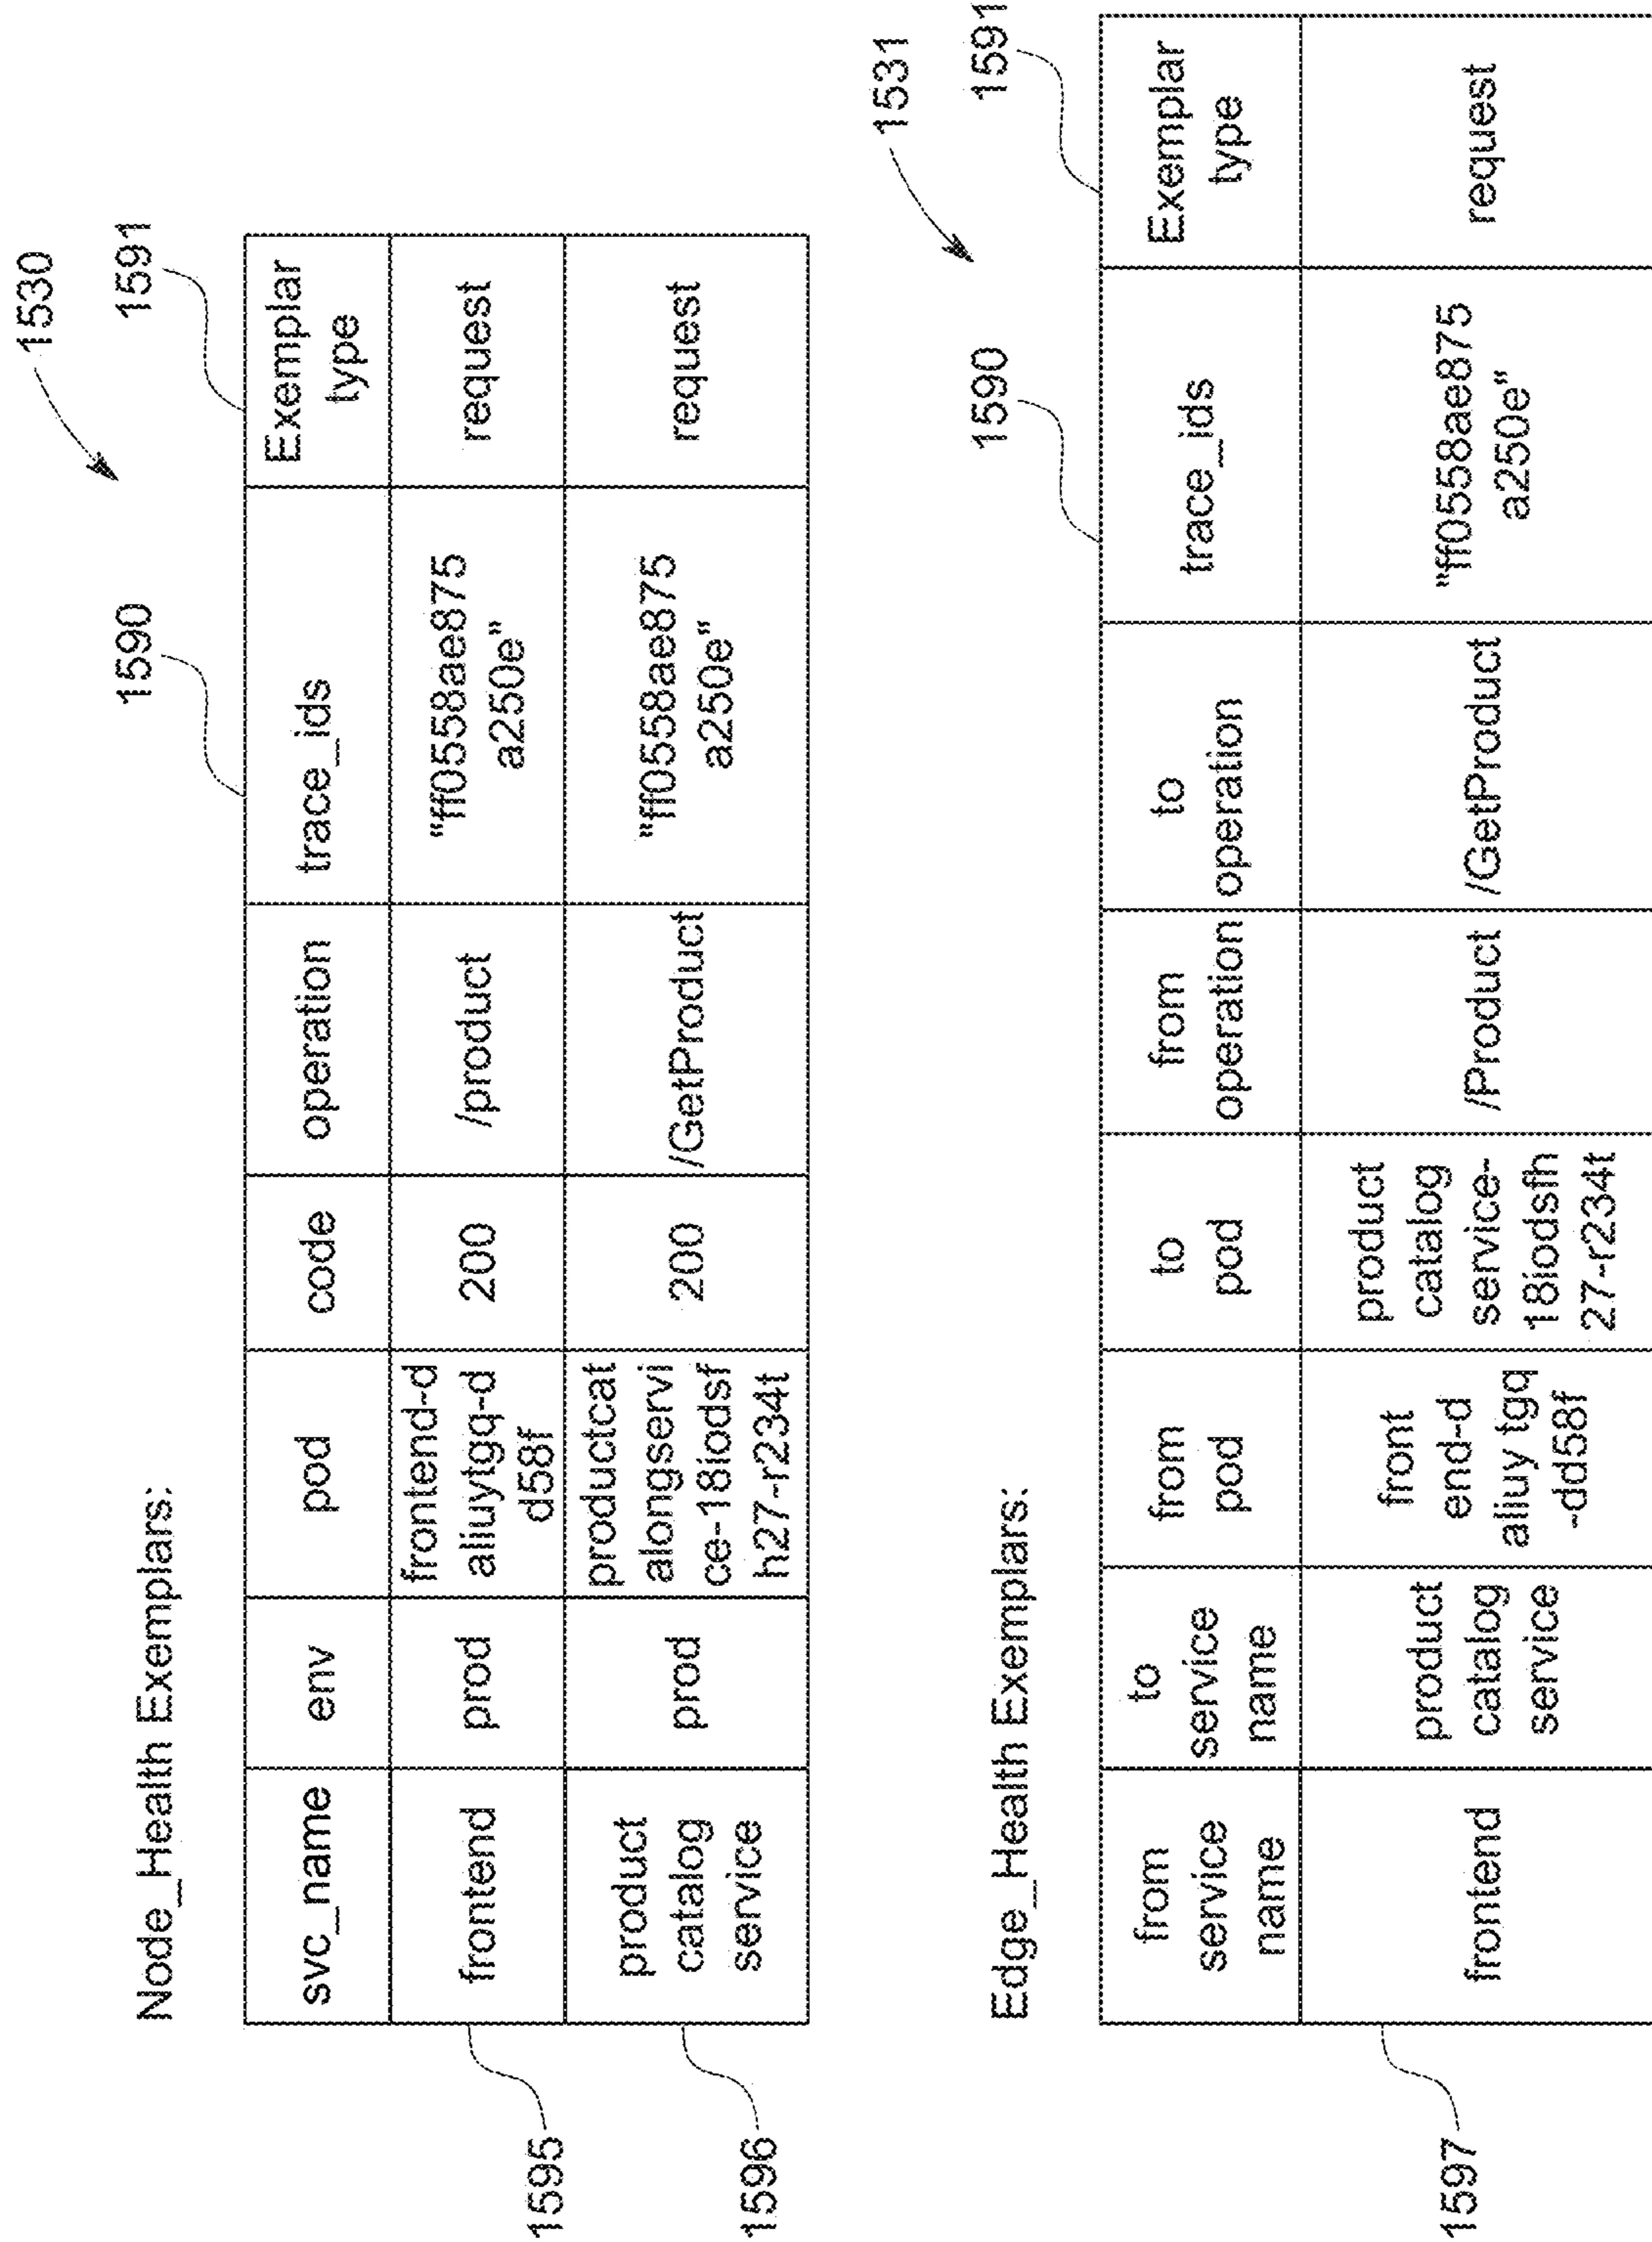


FIG. 15C

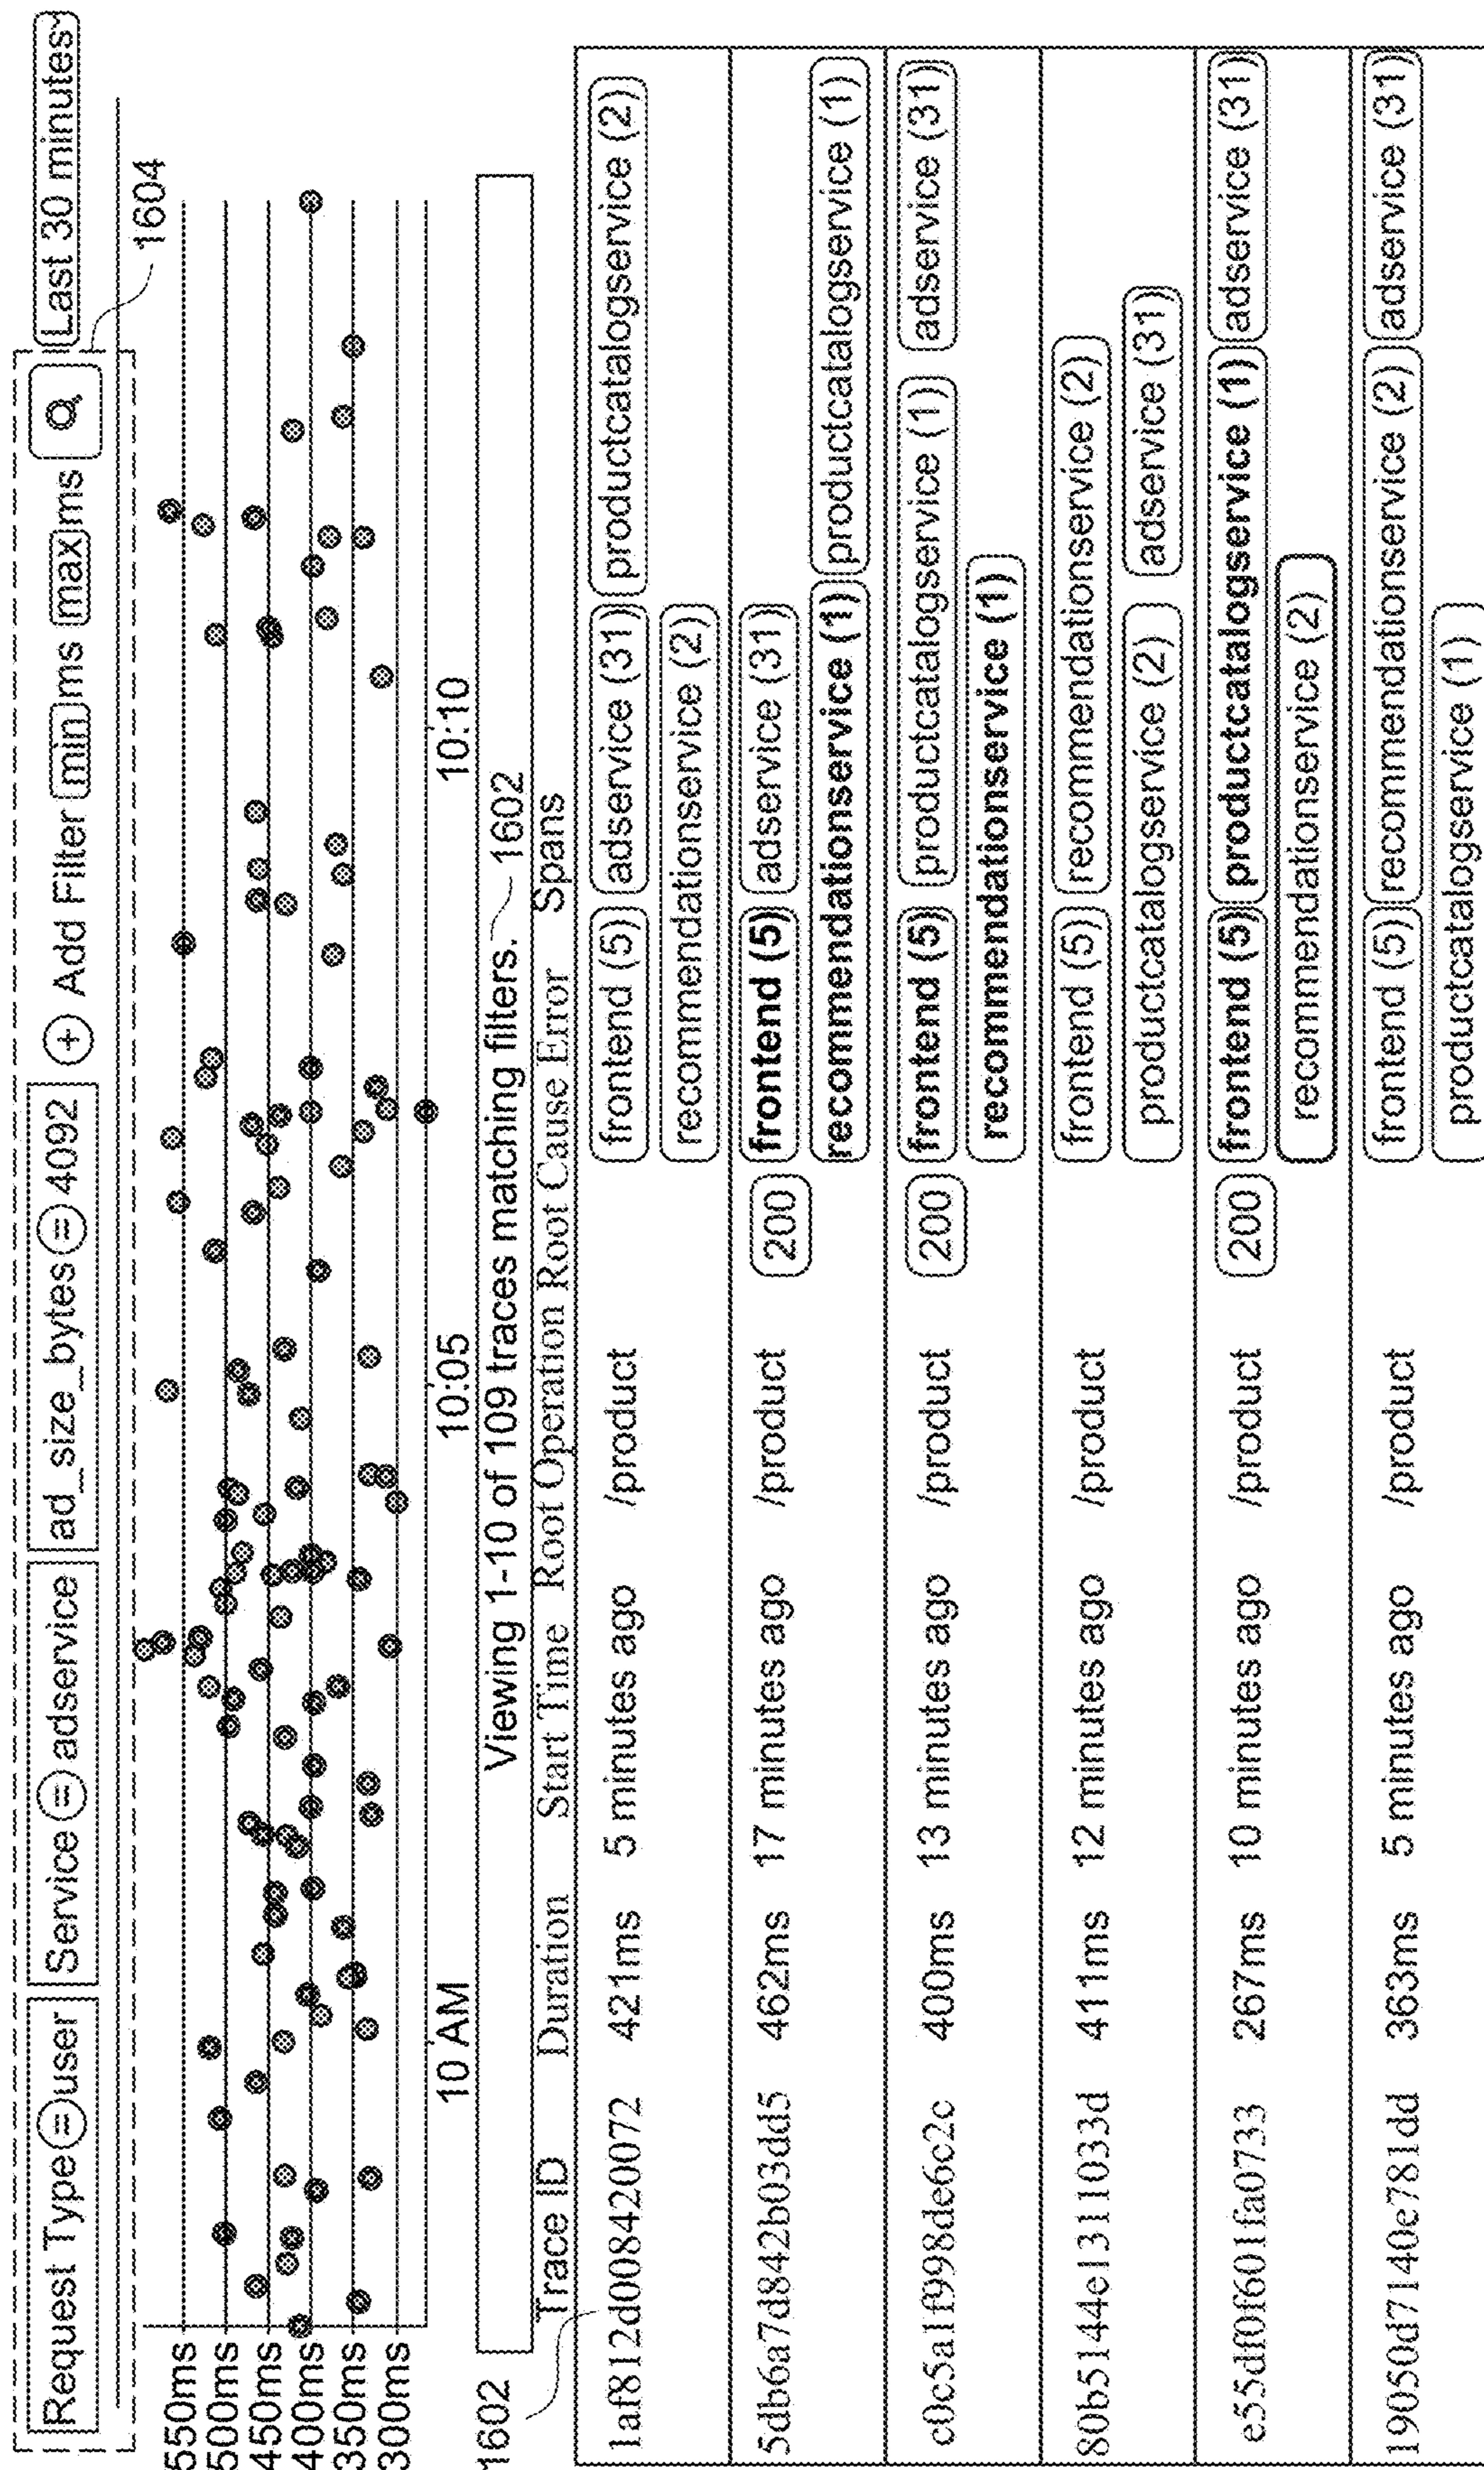


FIG. 16

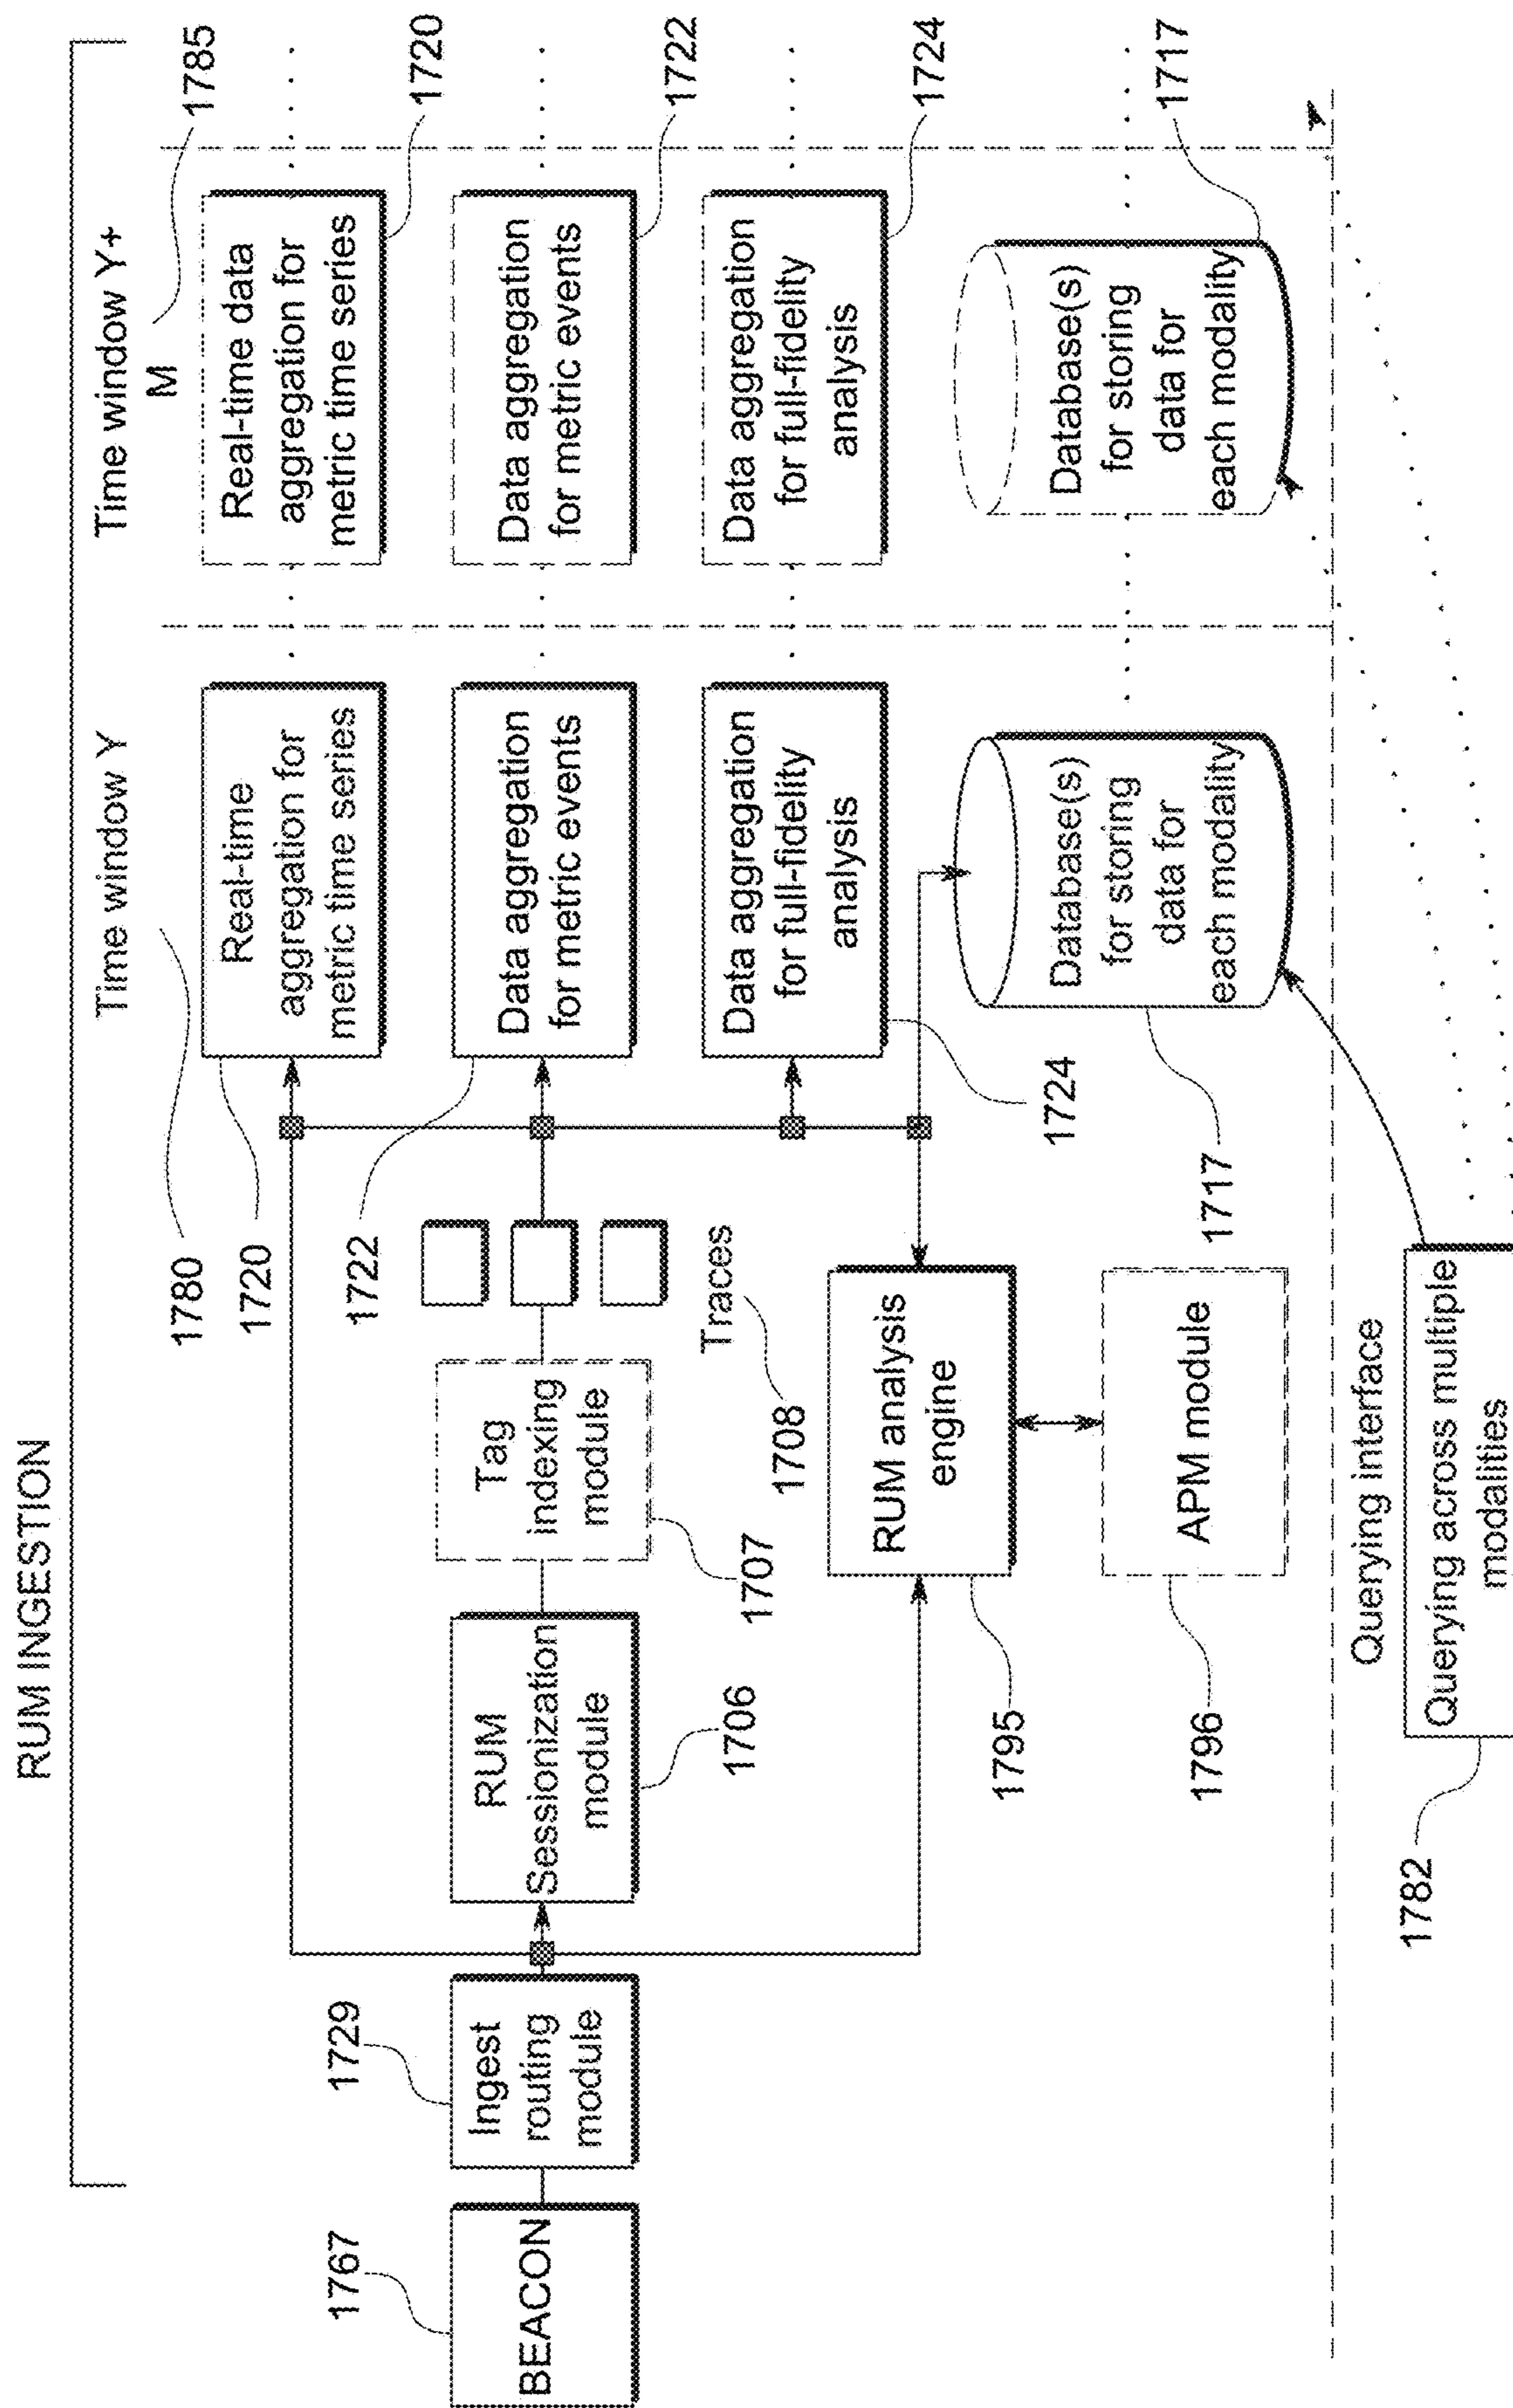


FIG. 17

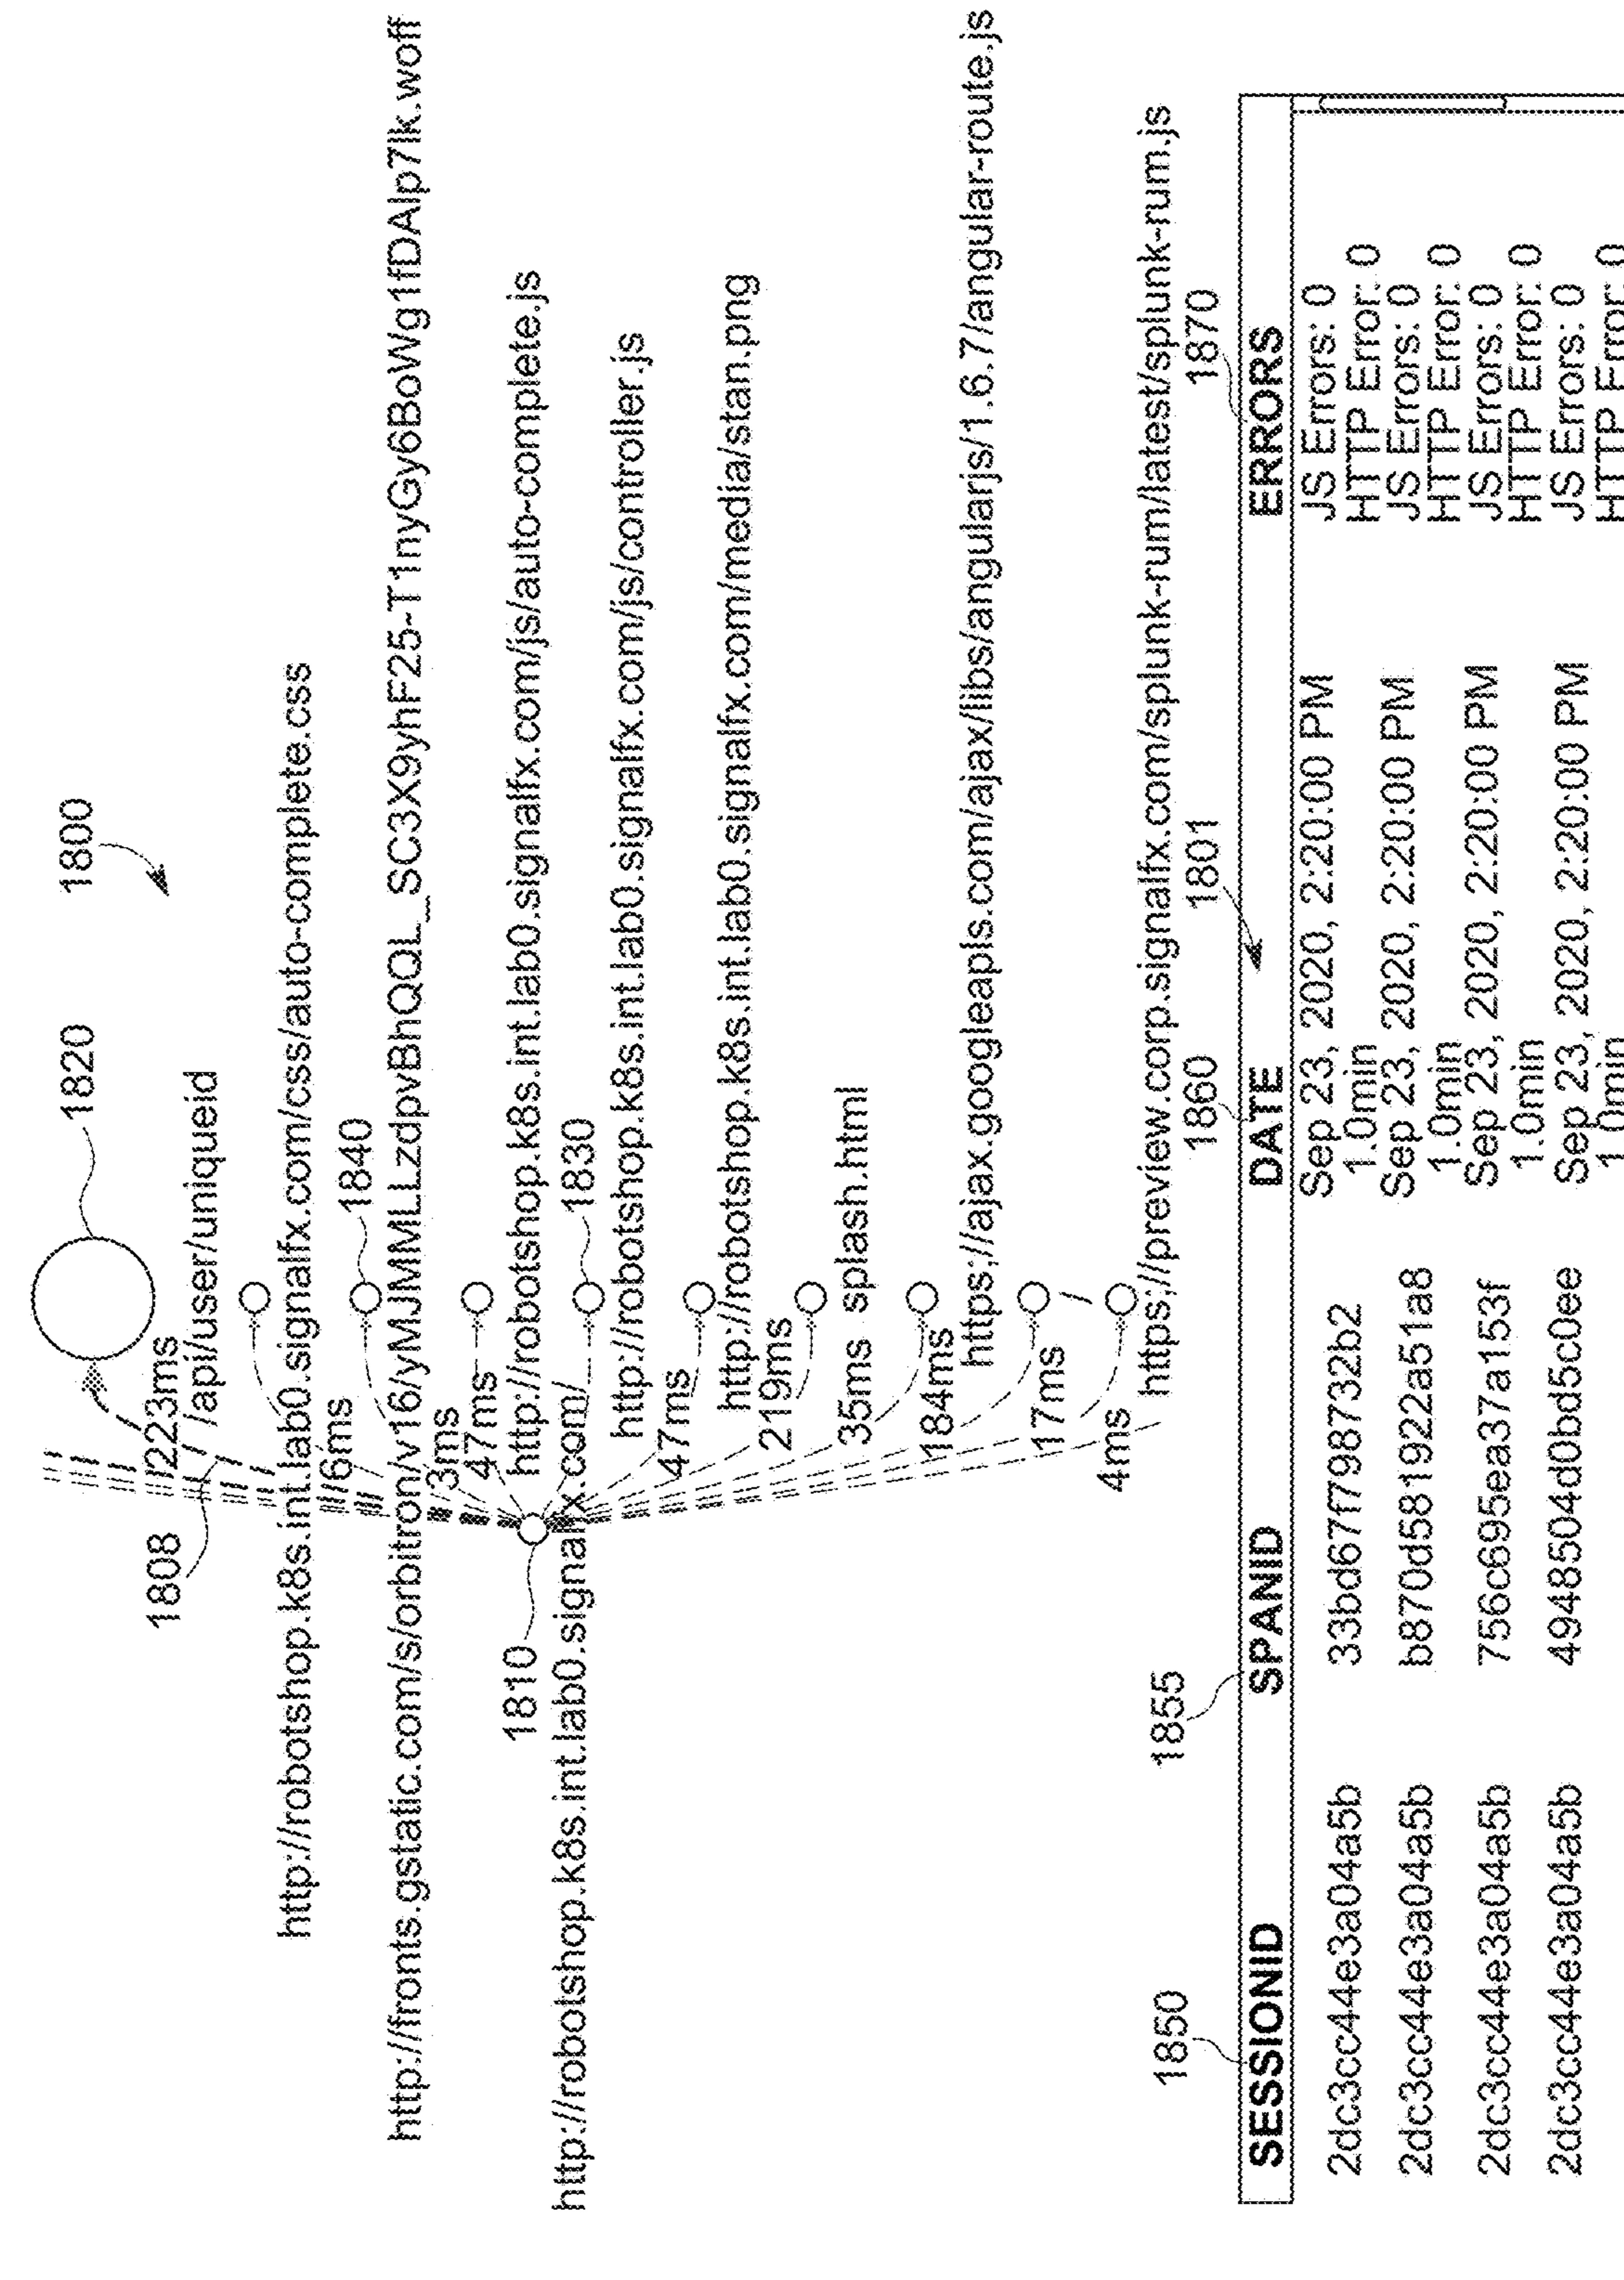


FIG. 18

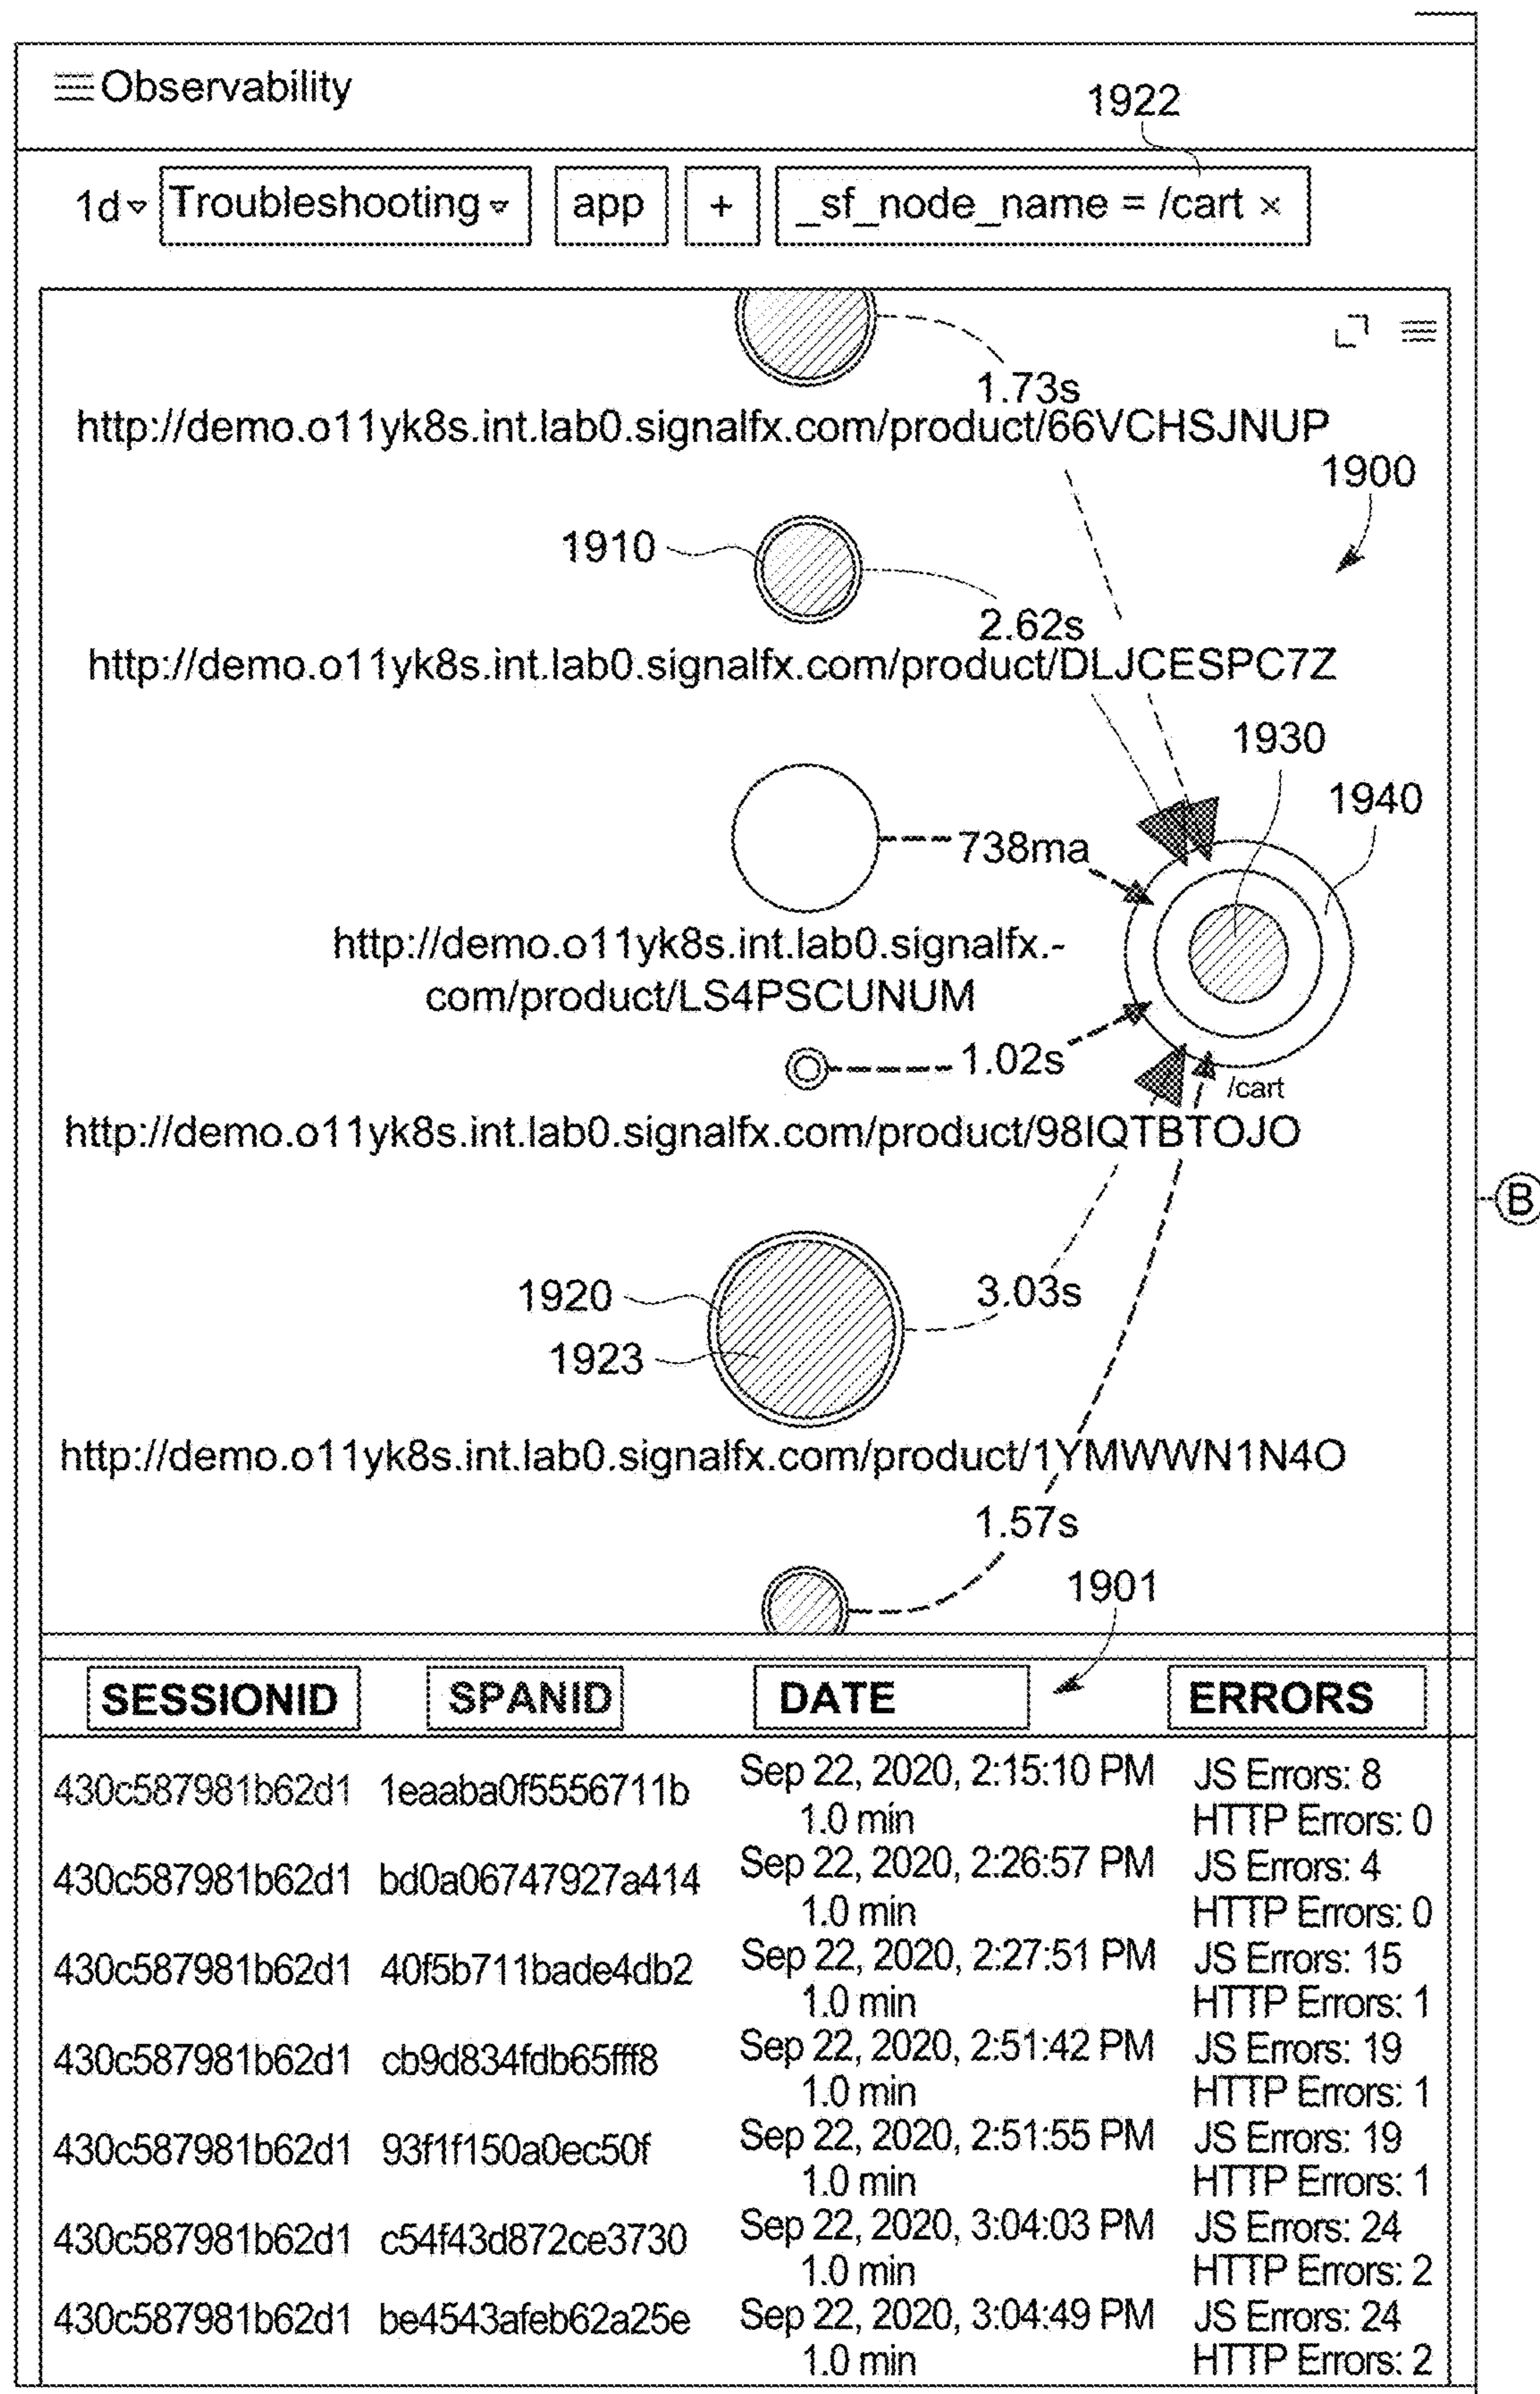


FIG. 19

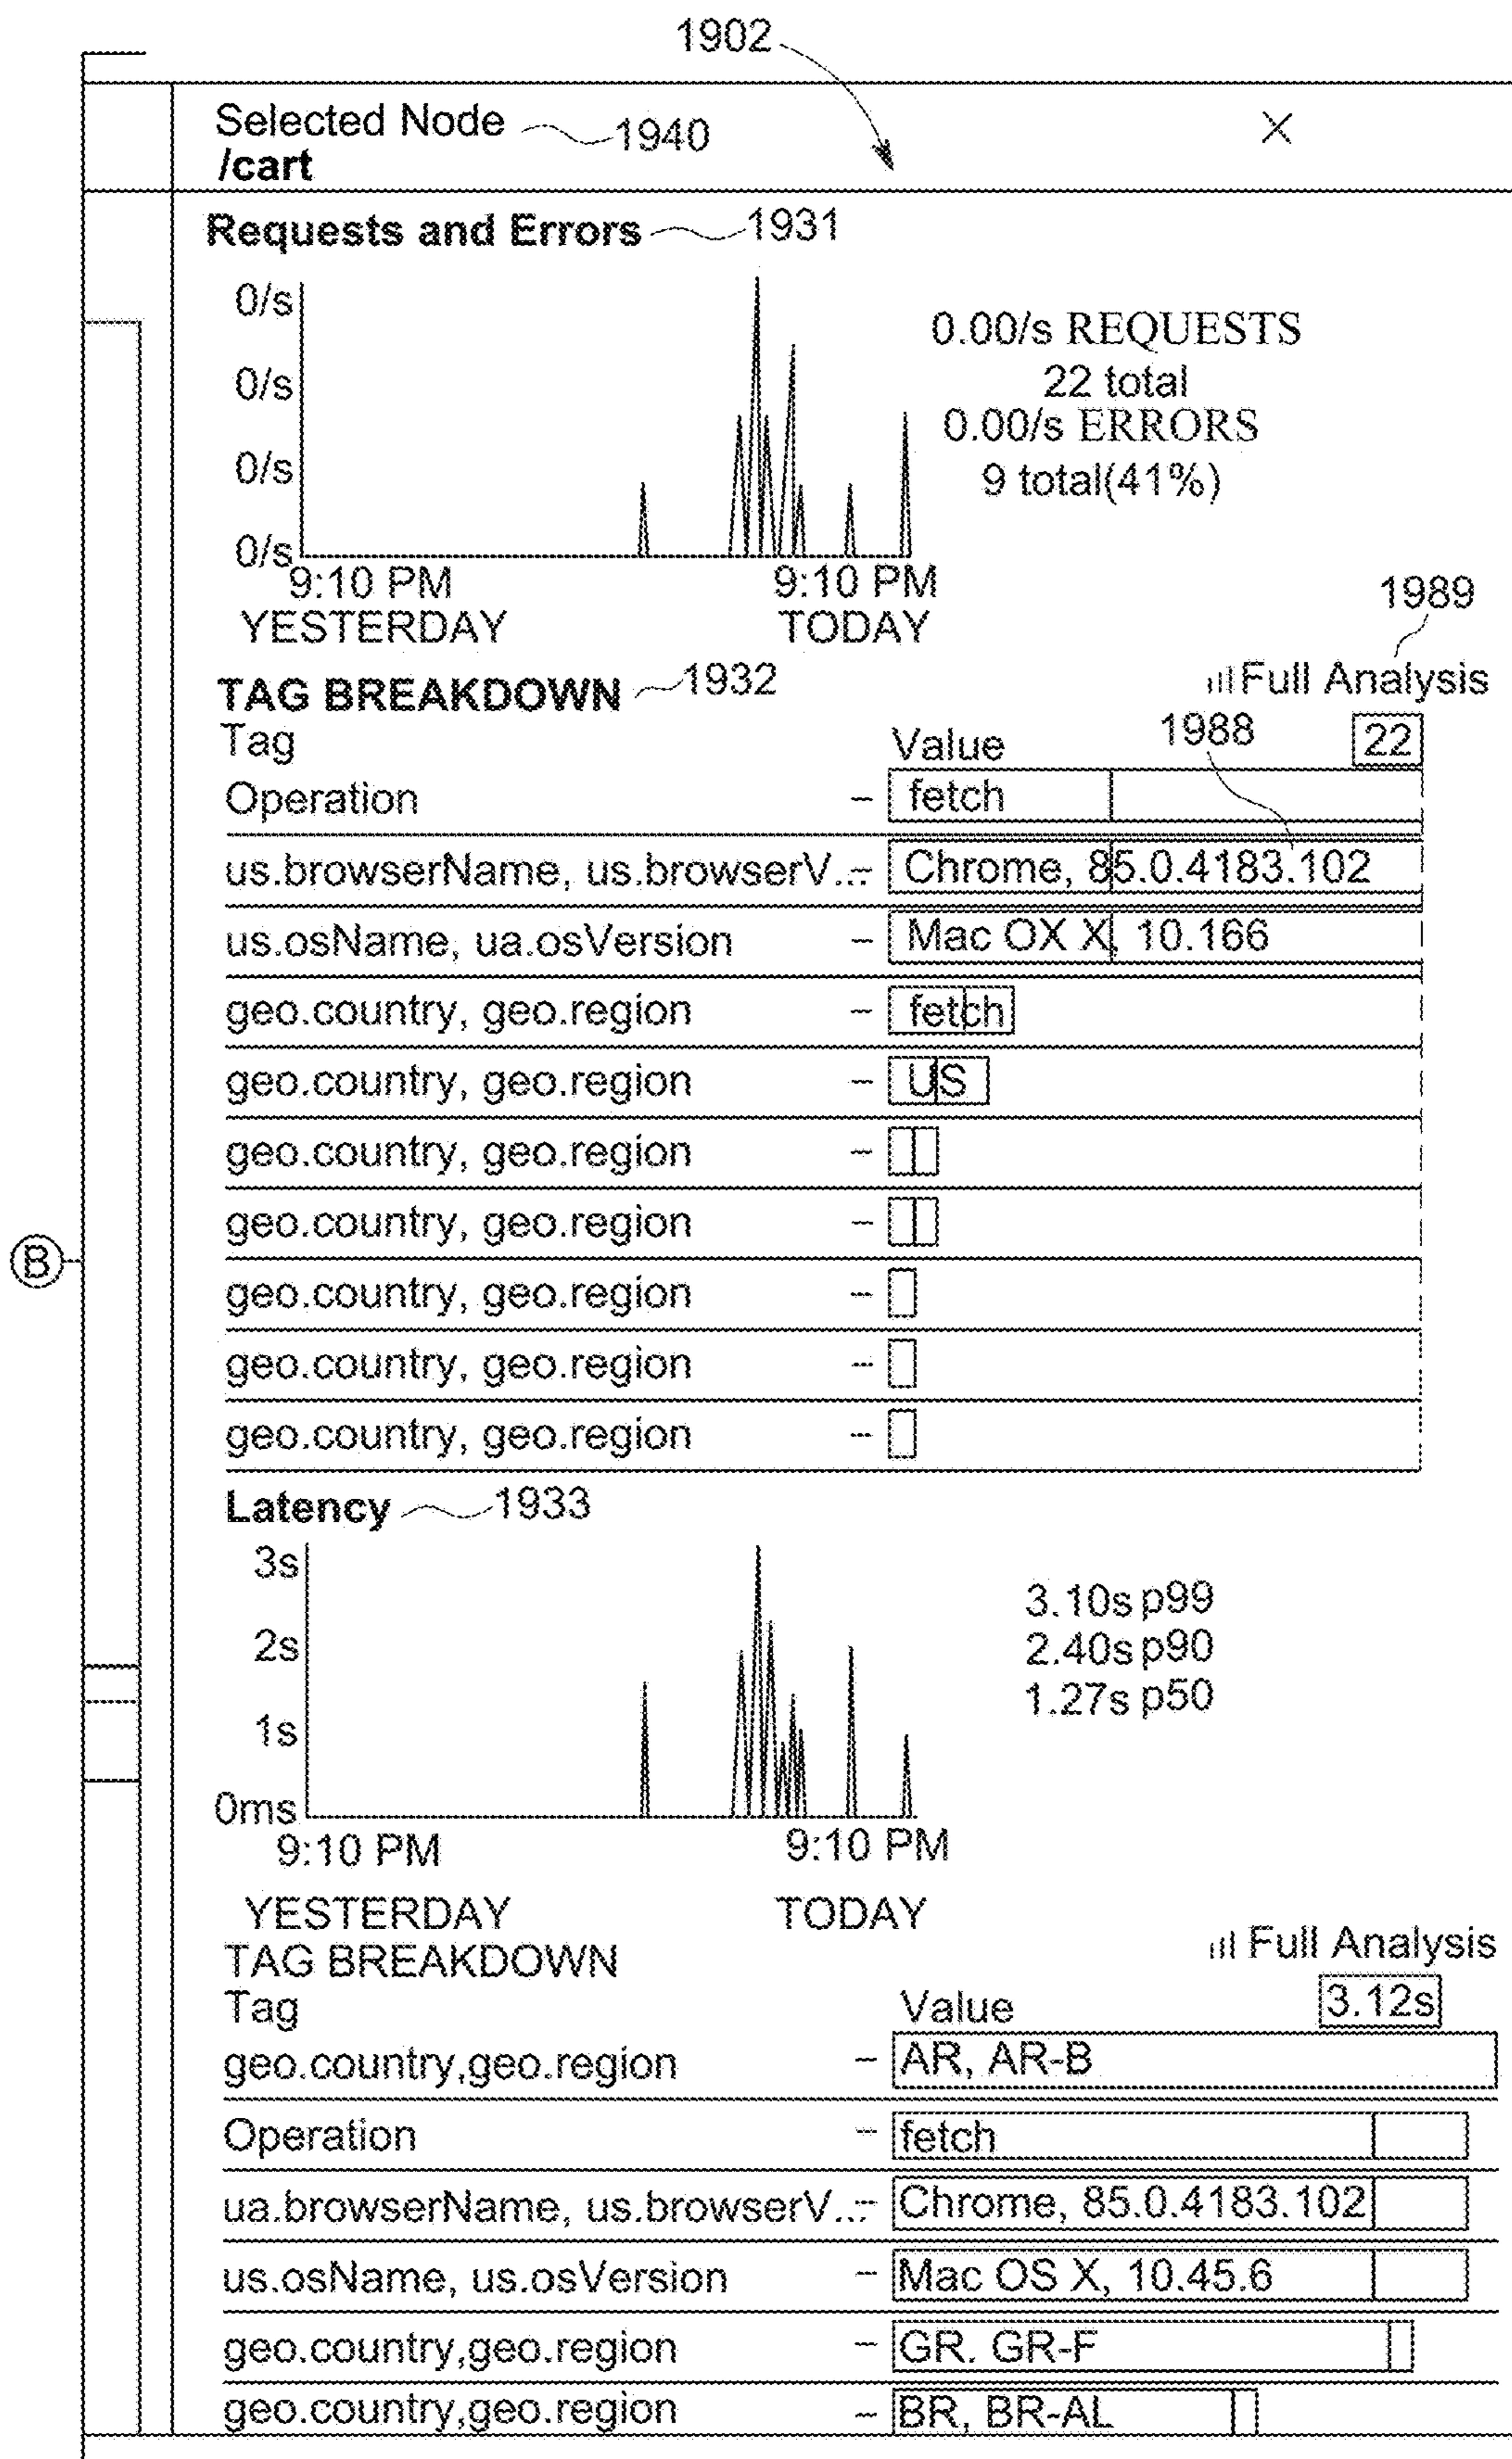


FIG. 19(Cont.)

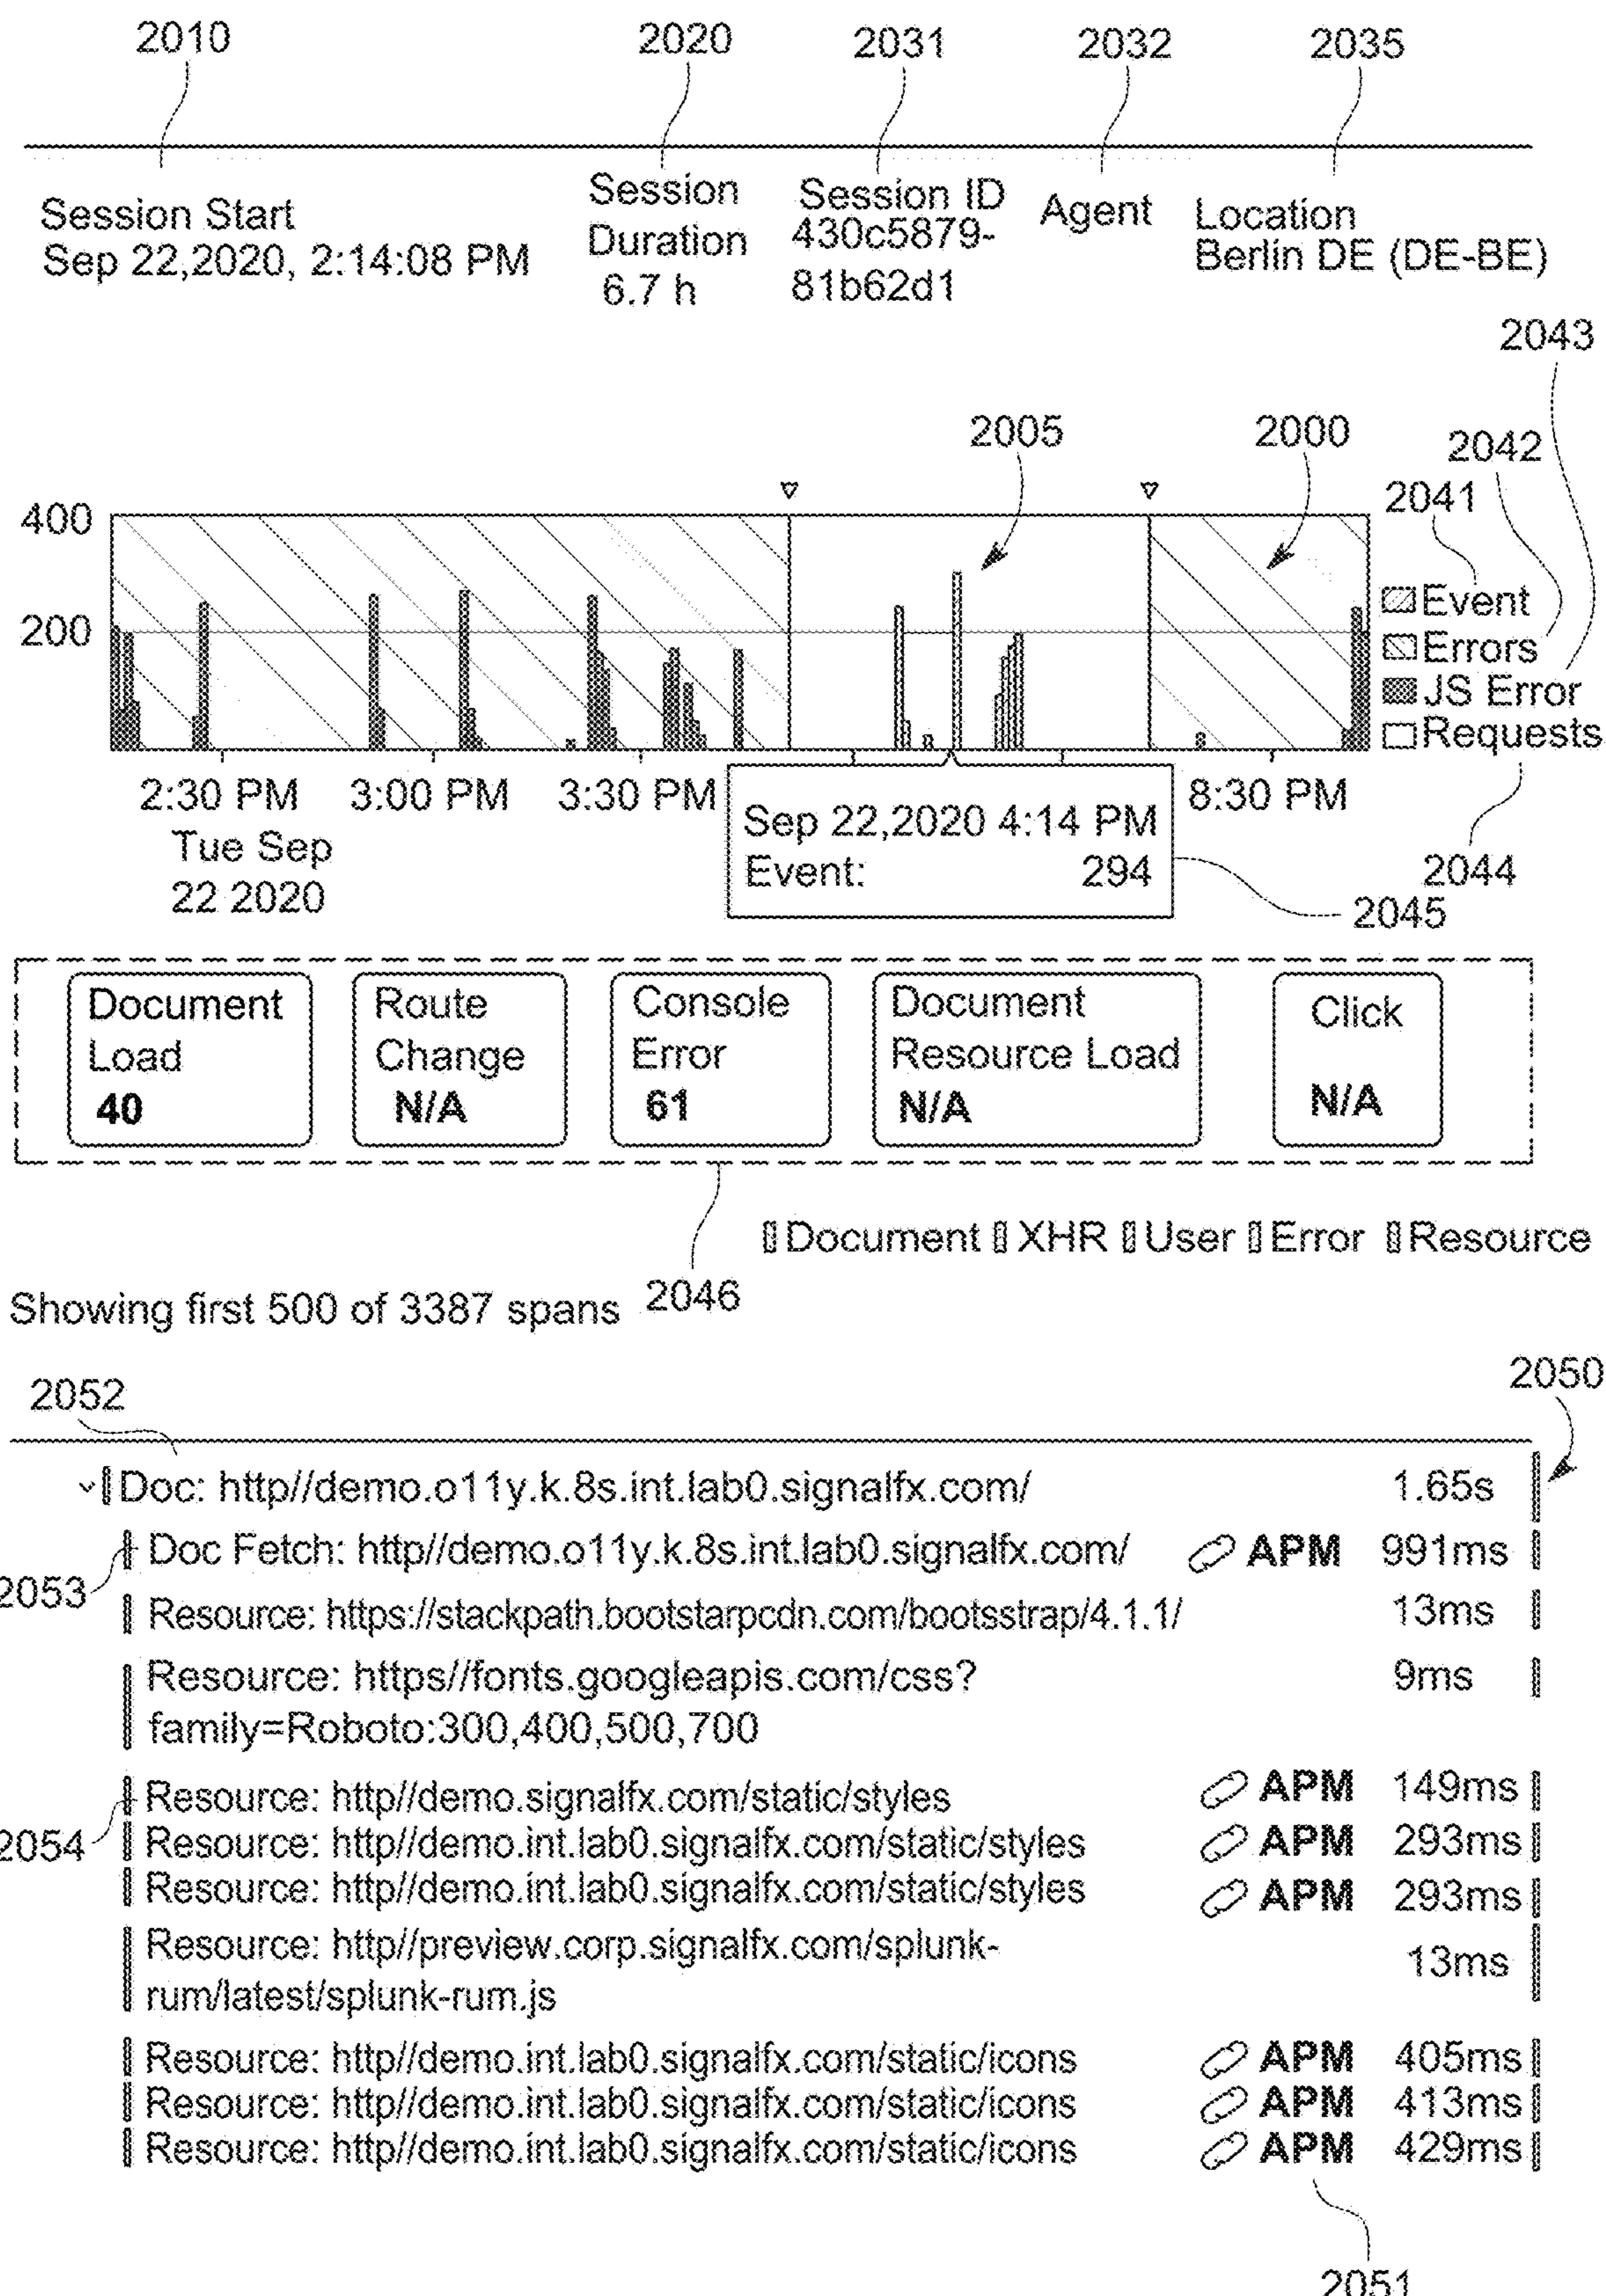
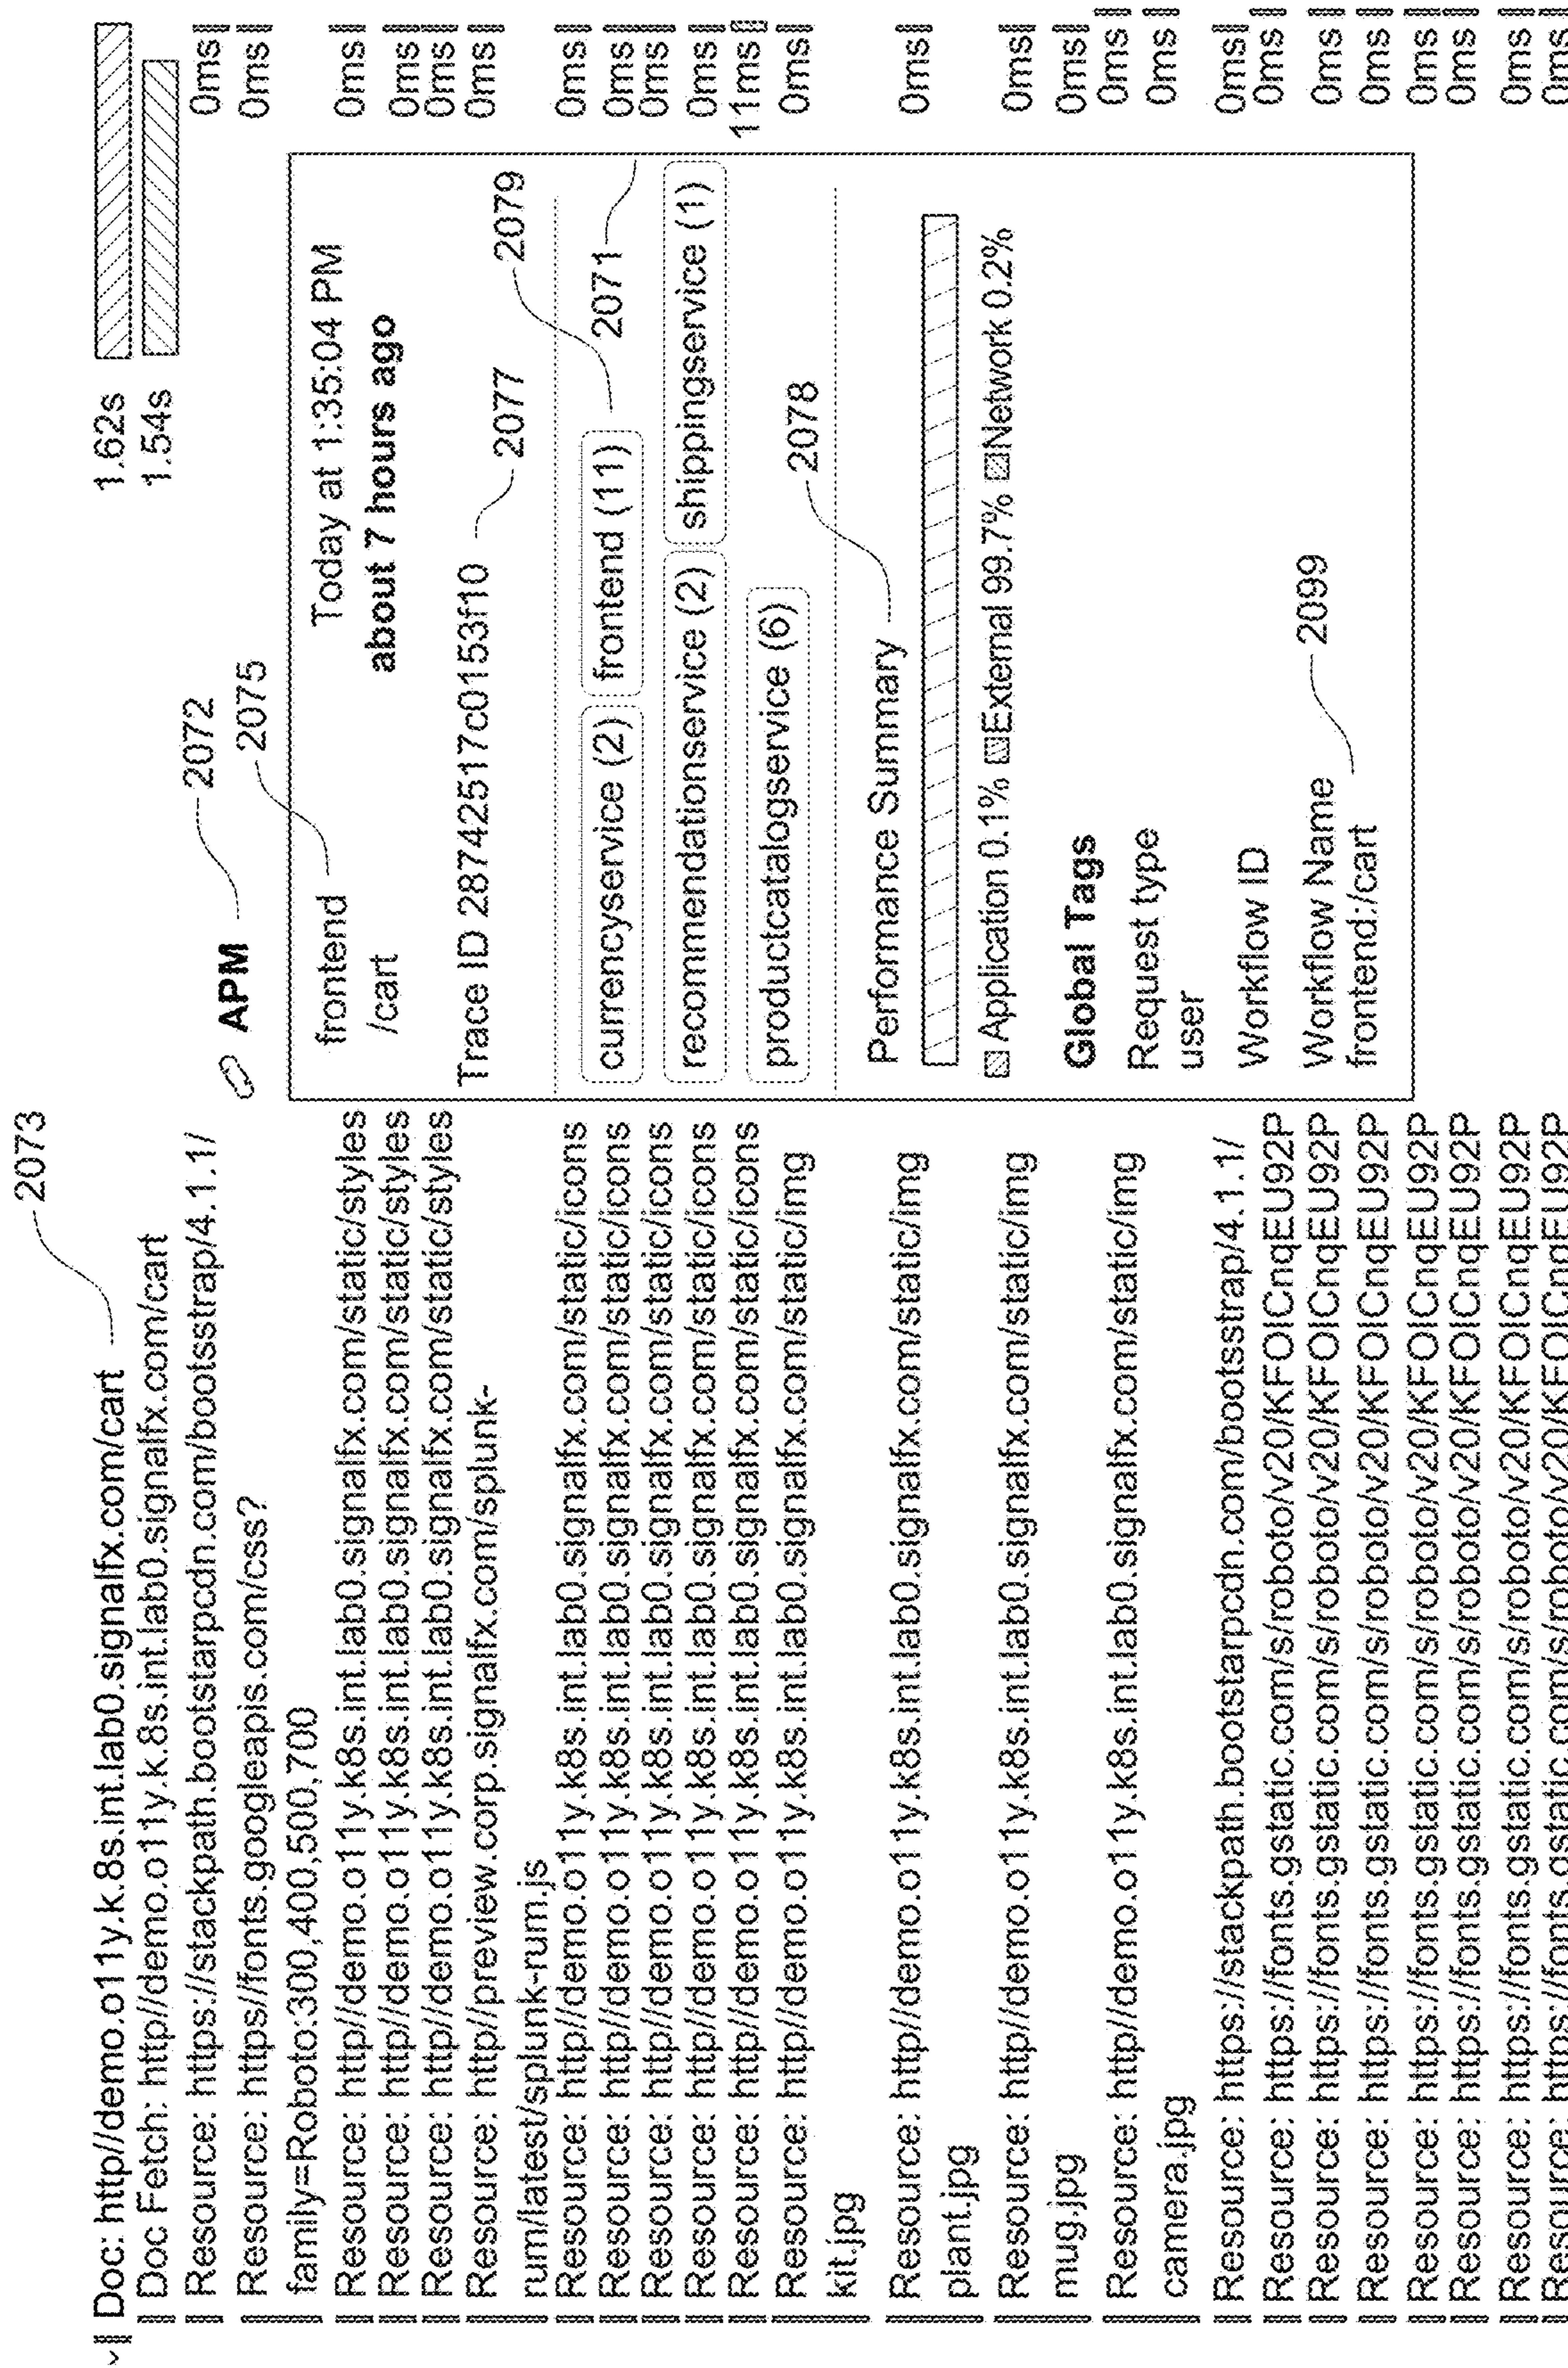


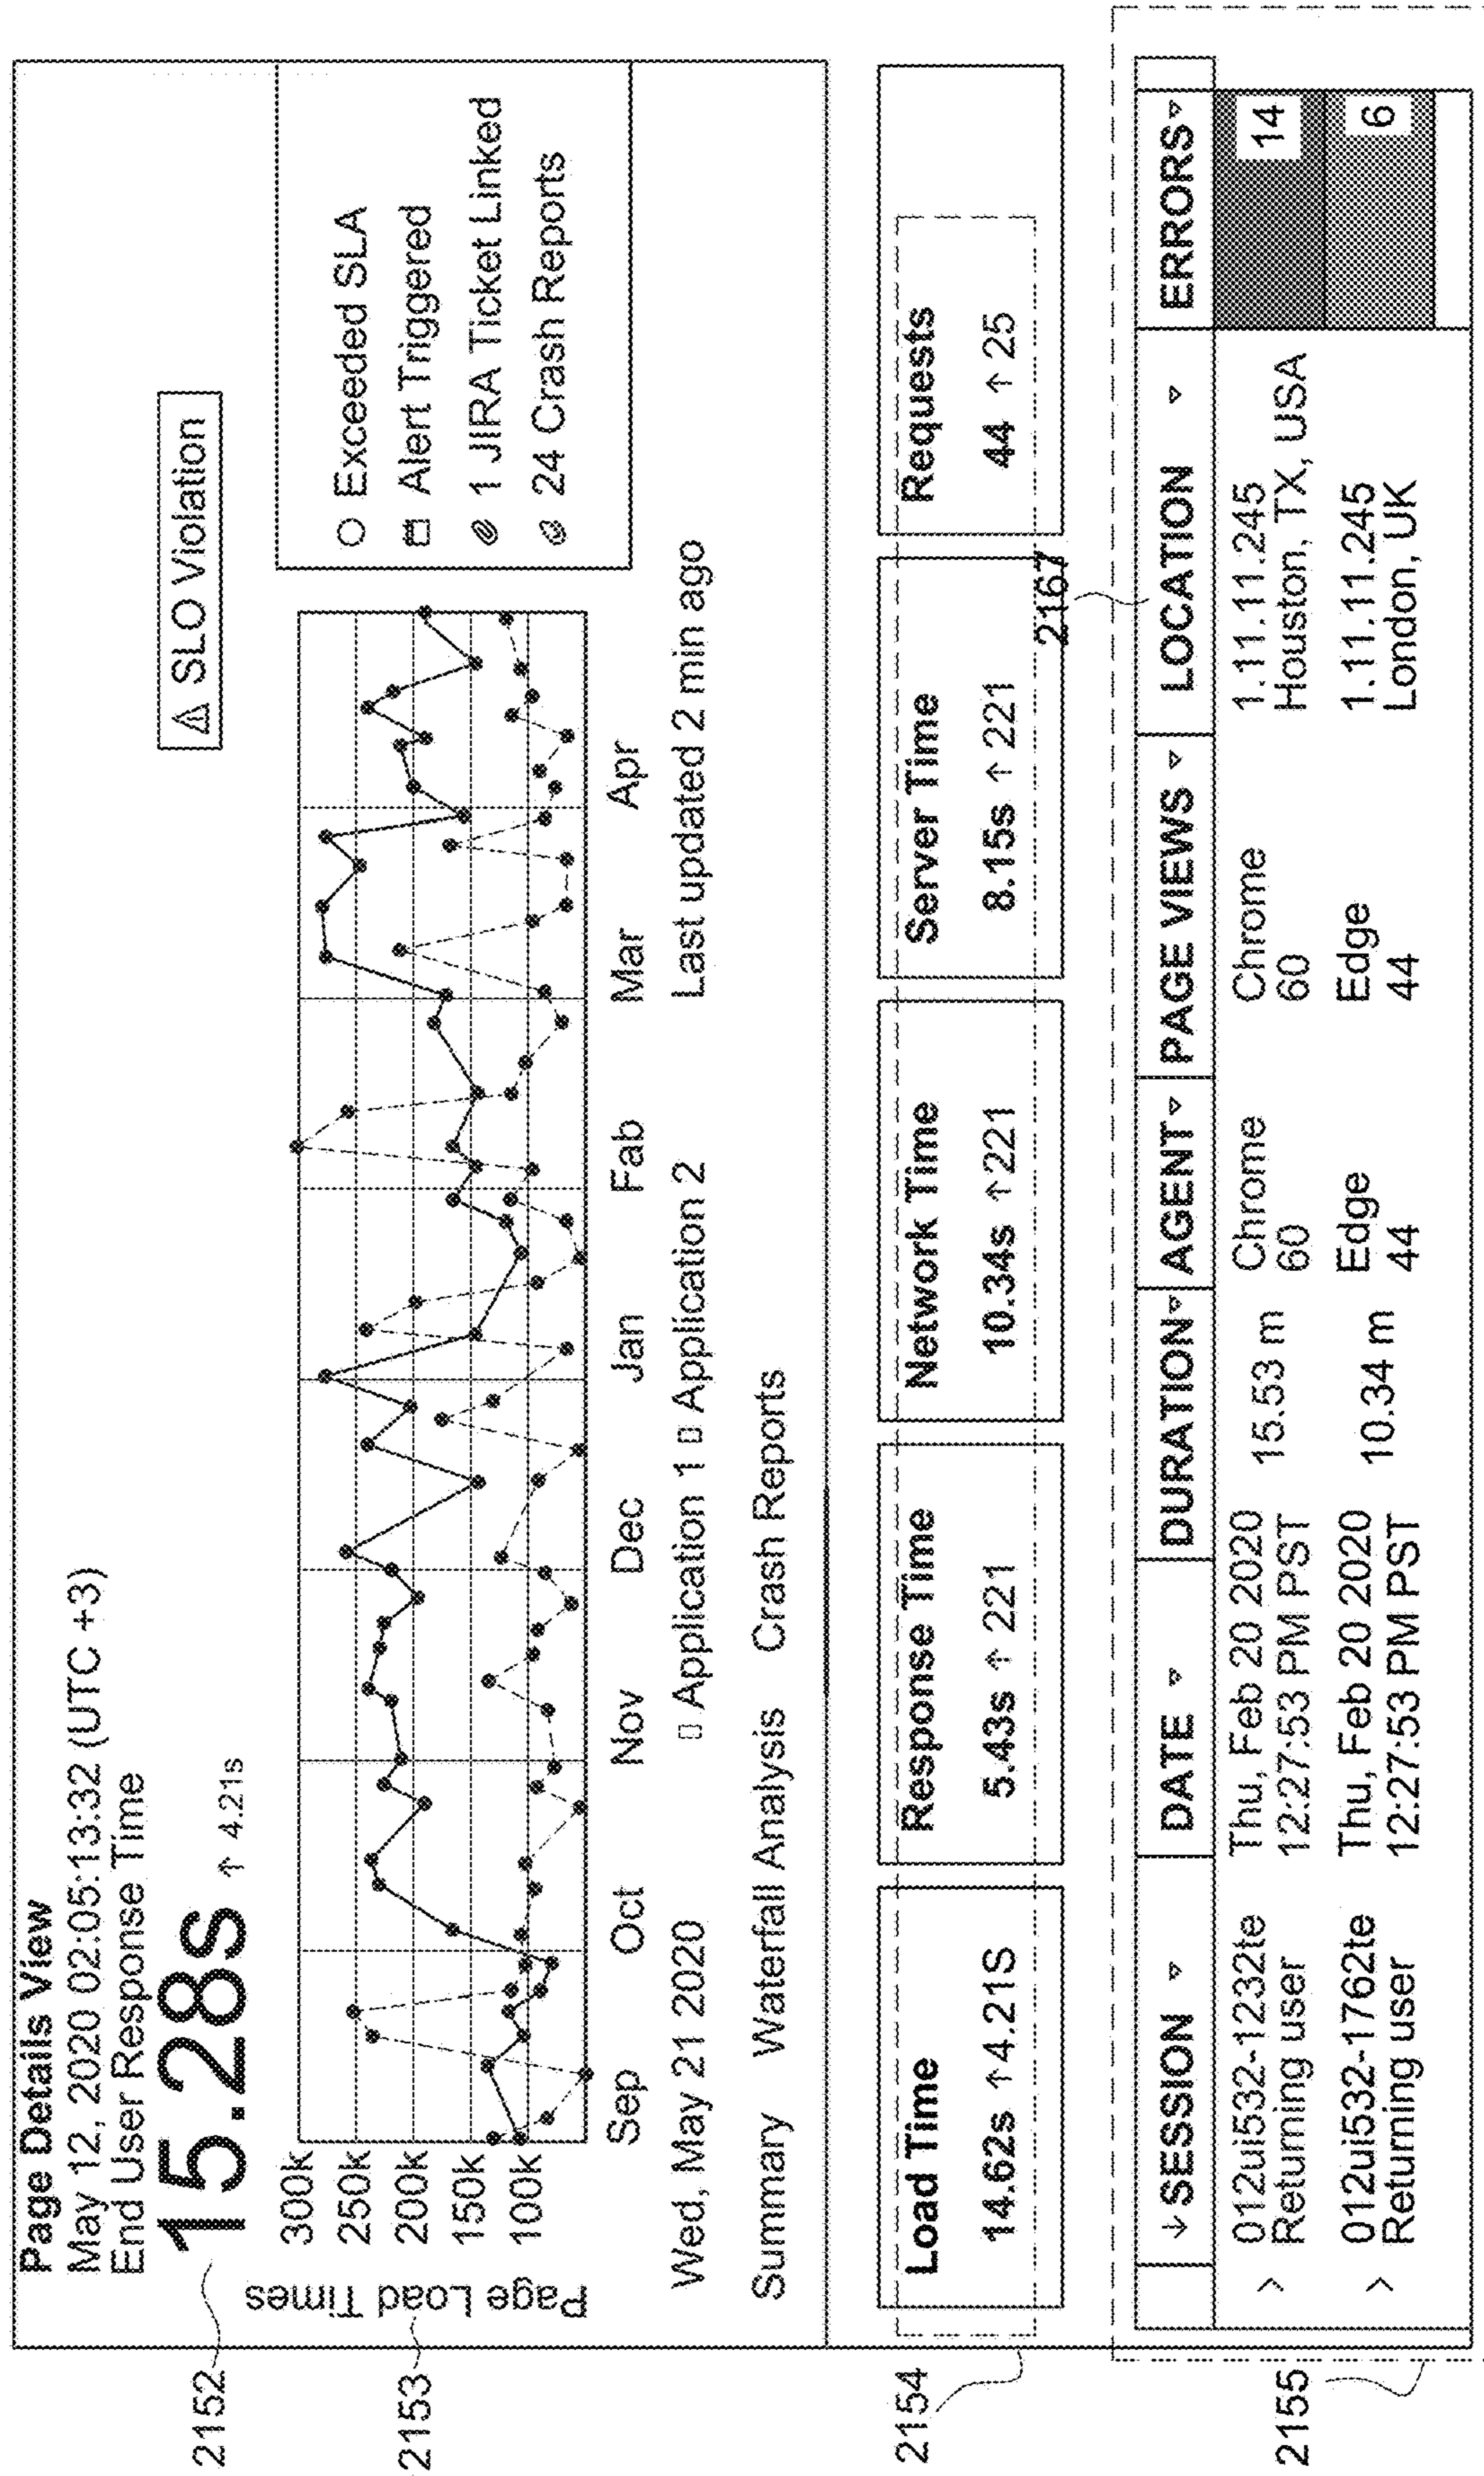
FIG. 20A

Resource: https://stackpath.bootstrapcdn.com/bootstrap/4.1.1/j	45ms	
Resource: https://demo.o11y.k8s.int.lab0.signalfx.com/static/image/2064	813ms	
SPAN 2063		2064
Span ID	619d8044d29b449	
Partner ID	d4cc7d9c3939r903	2062
OPERATION		
resourceFetch		
TAGS 2061		
app	hipster-shop-demo	
component	document-load	
geo.city	Miami	
geo.country	US	
geo.region	US-FL	
http.response_content_length	260426	
http.url	http://demo.images/HeroBanner2.png	2088
link.spanId	77caa92cb7ba4761	2067
link.traceId	00000000000000077caa92cb7ba4761	2068
location.href	http://demo.o11y-k8s.int.lab0.signalfx.com/	
norm.http.url	http://demo.images/HeroBannerImage2.png	
norm.location.href	http://demo.o11y-k8s.int.lab0.signalfx.com	
norm.sf_operation	resourceFetch	
orgId	EauzO4EAIAA	
ot.status_code	OK	
splunk.rumSessionId	c48c238baf98c353beb46519d3dcbae8	2089
splunk.rumVersion	0.0.9	
splunk.scriptInstanceId	896bd4fcee479d31	
ua.browserName		
ua.browserVersion		
ua.osName		
WORKFLOW_NAME	CHARTLOAD	2021 2022
LOGS		
691 ms		
message	fetchStart	
707 ms		
message	domainLookupStart	
707 ms		
message	domainLookupEnd	
707 ms		
message	connectStart	
746 ms		
message	connectEnd	
847 ms		
message	requestStart	
1.02s		
message	responseStart	
1.50s		

FIG. 20B



卷之三



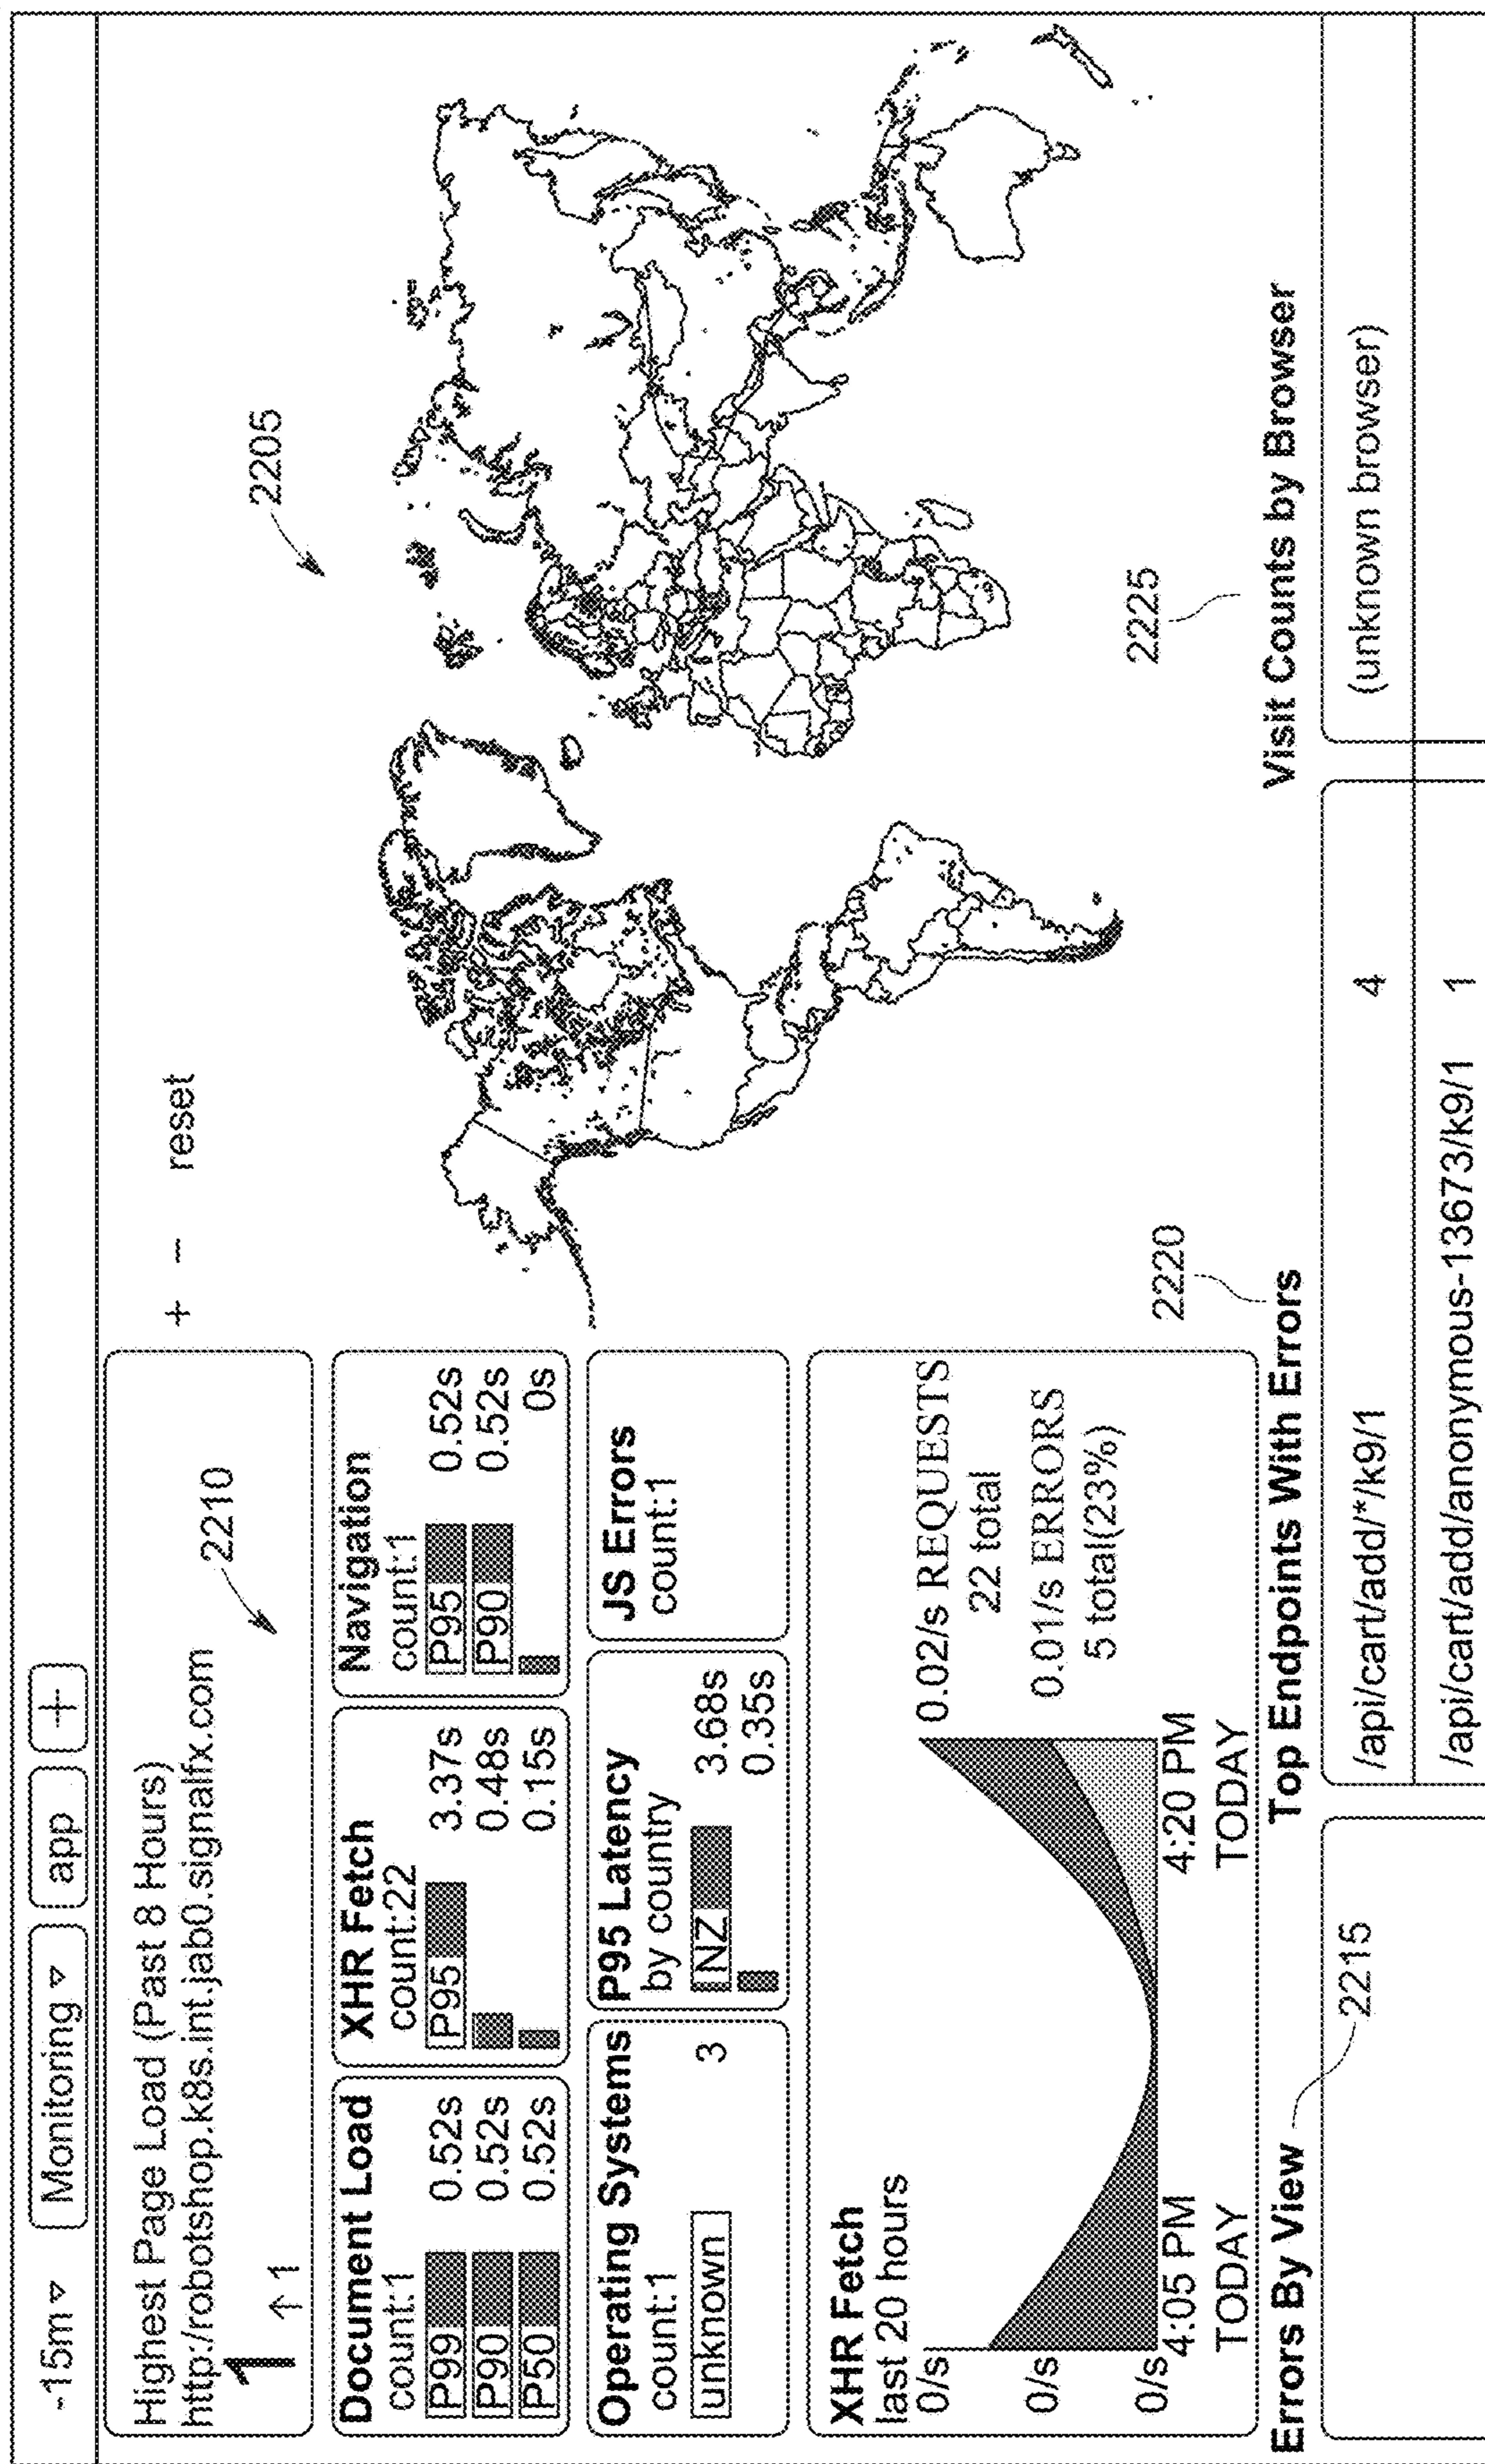


FIG. 22

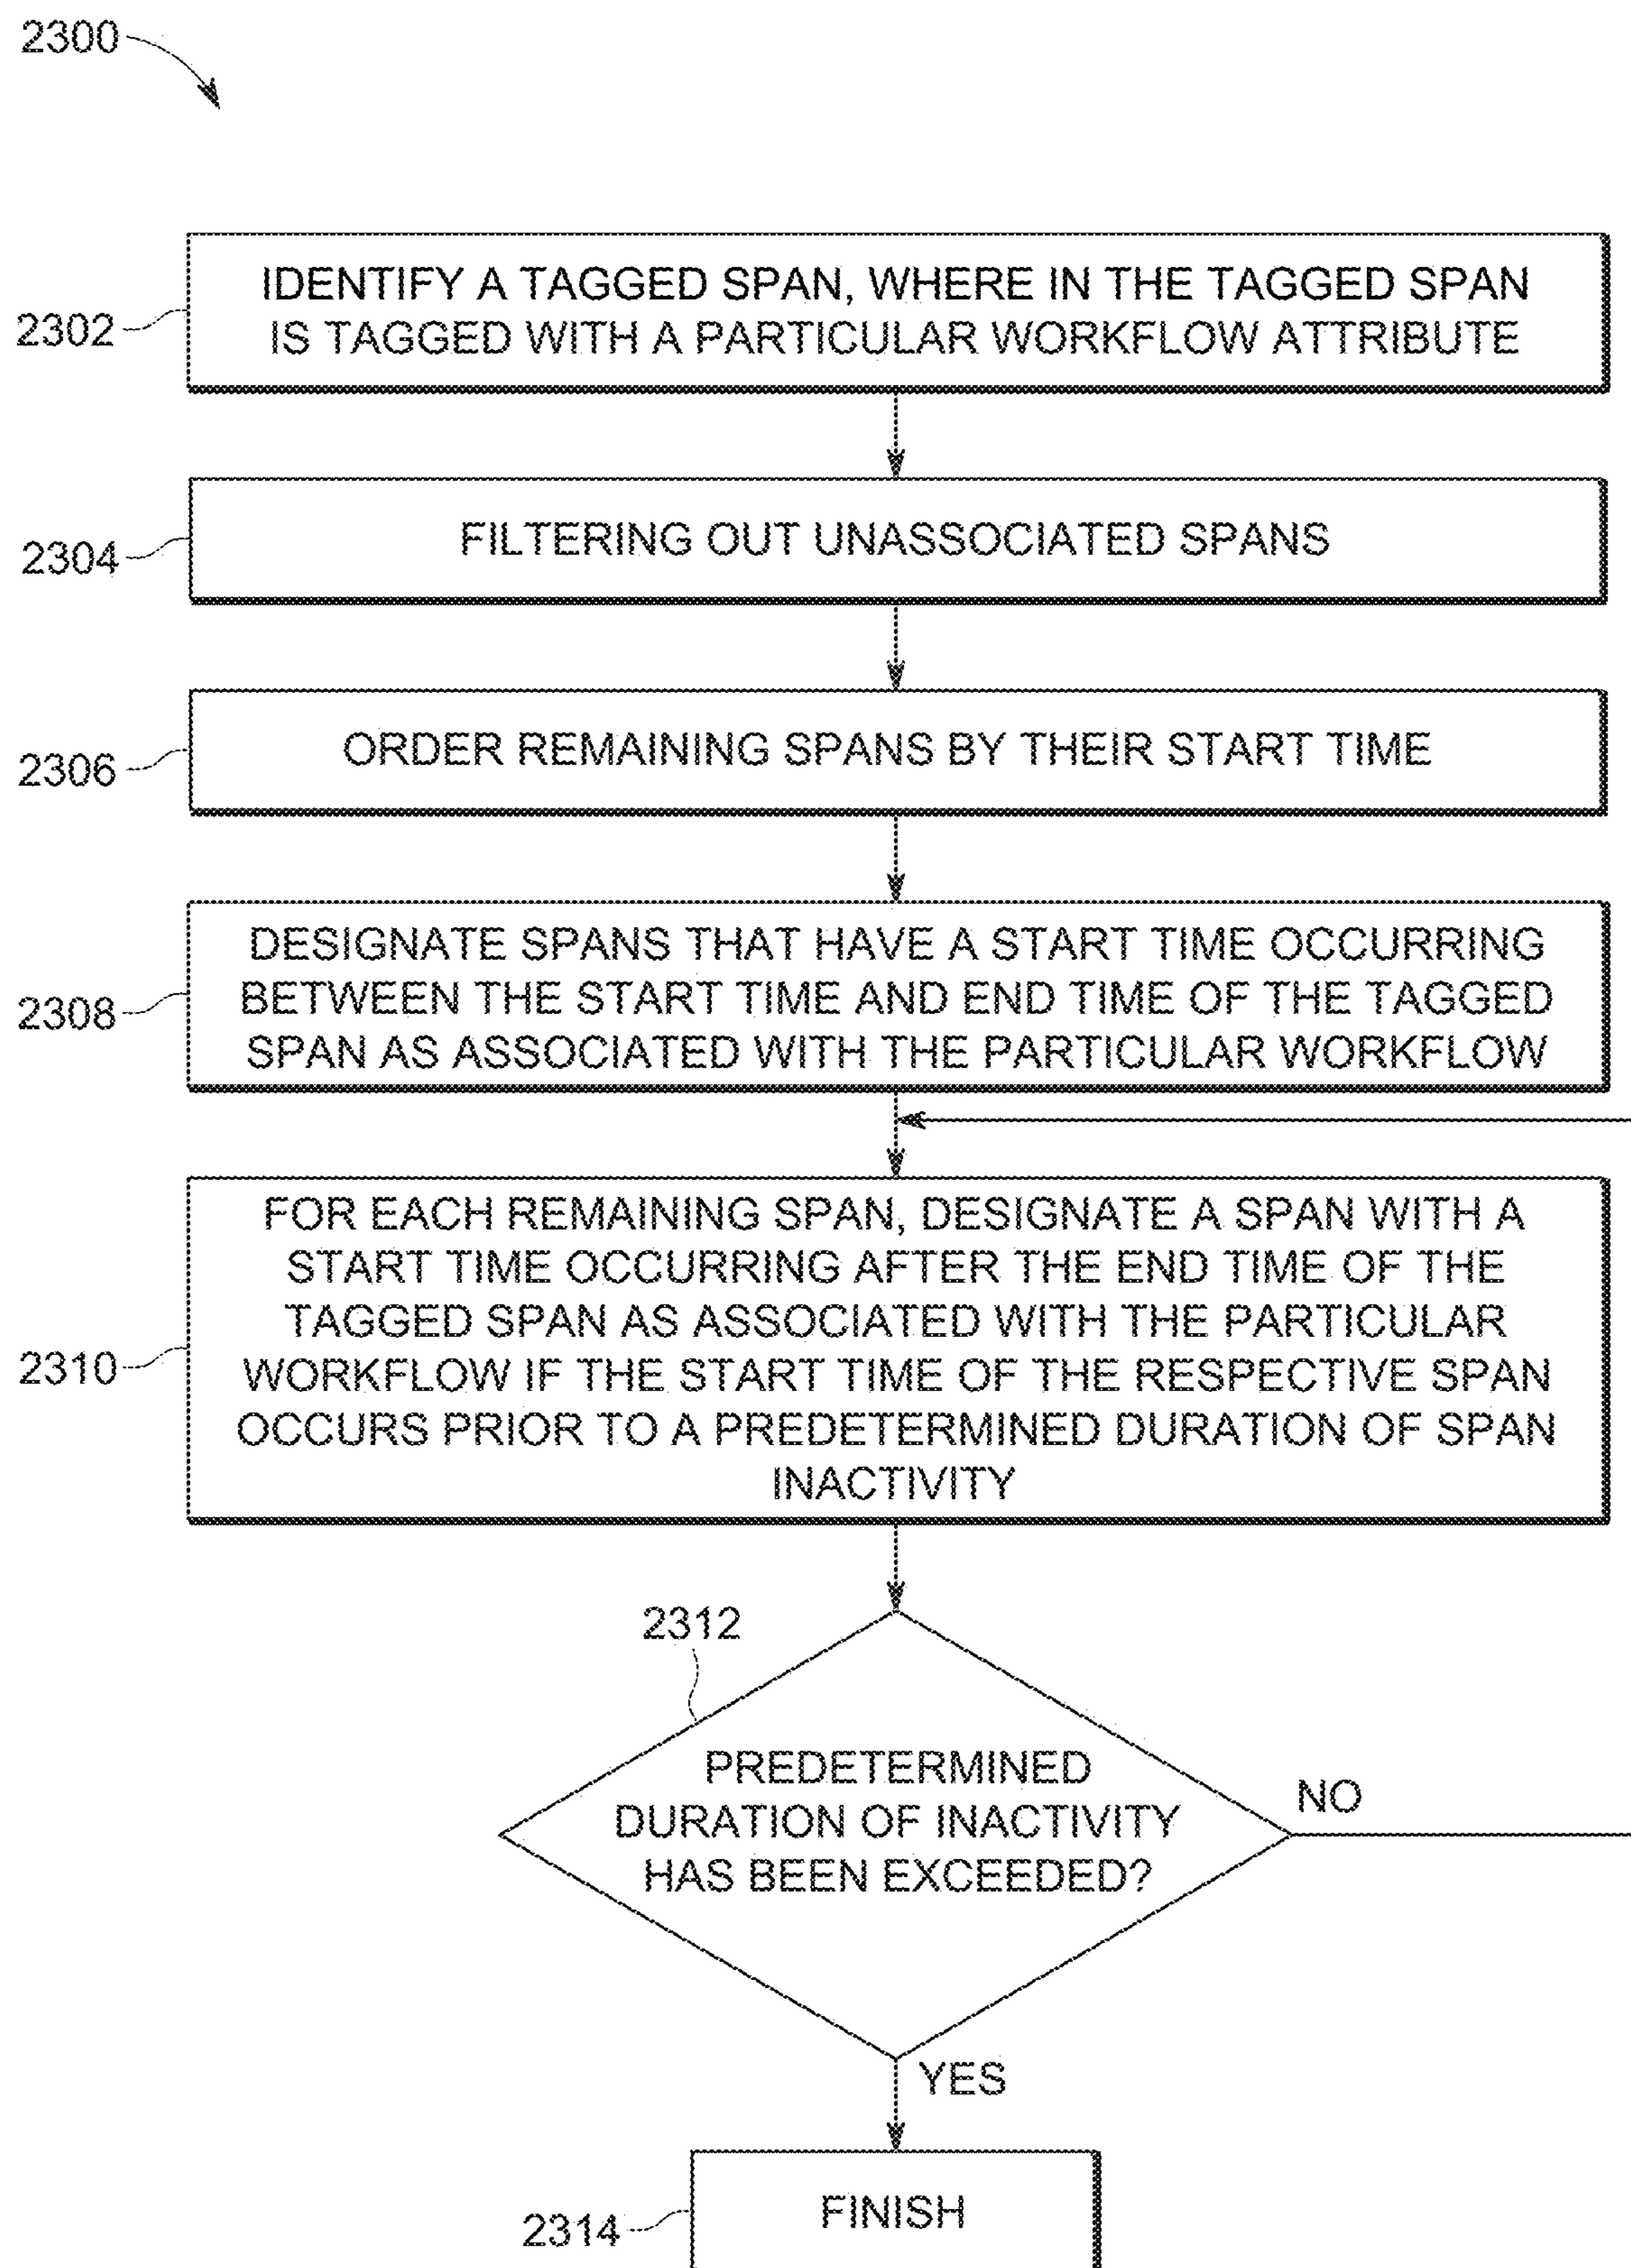


FIG. 23

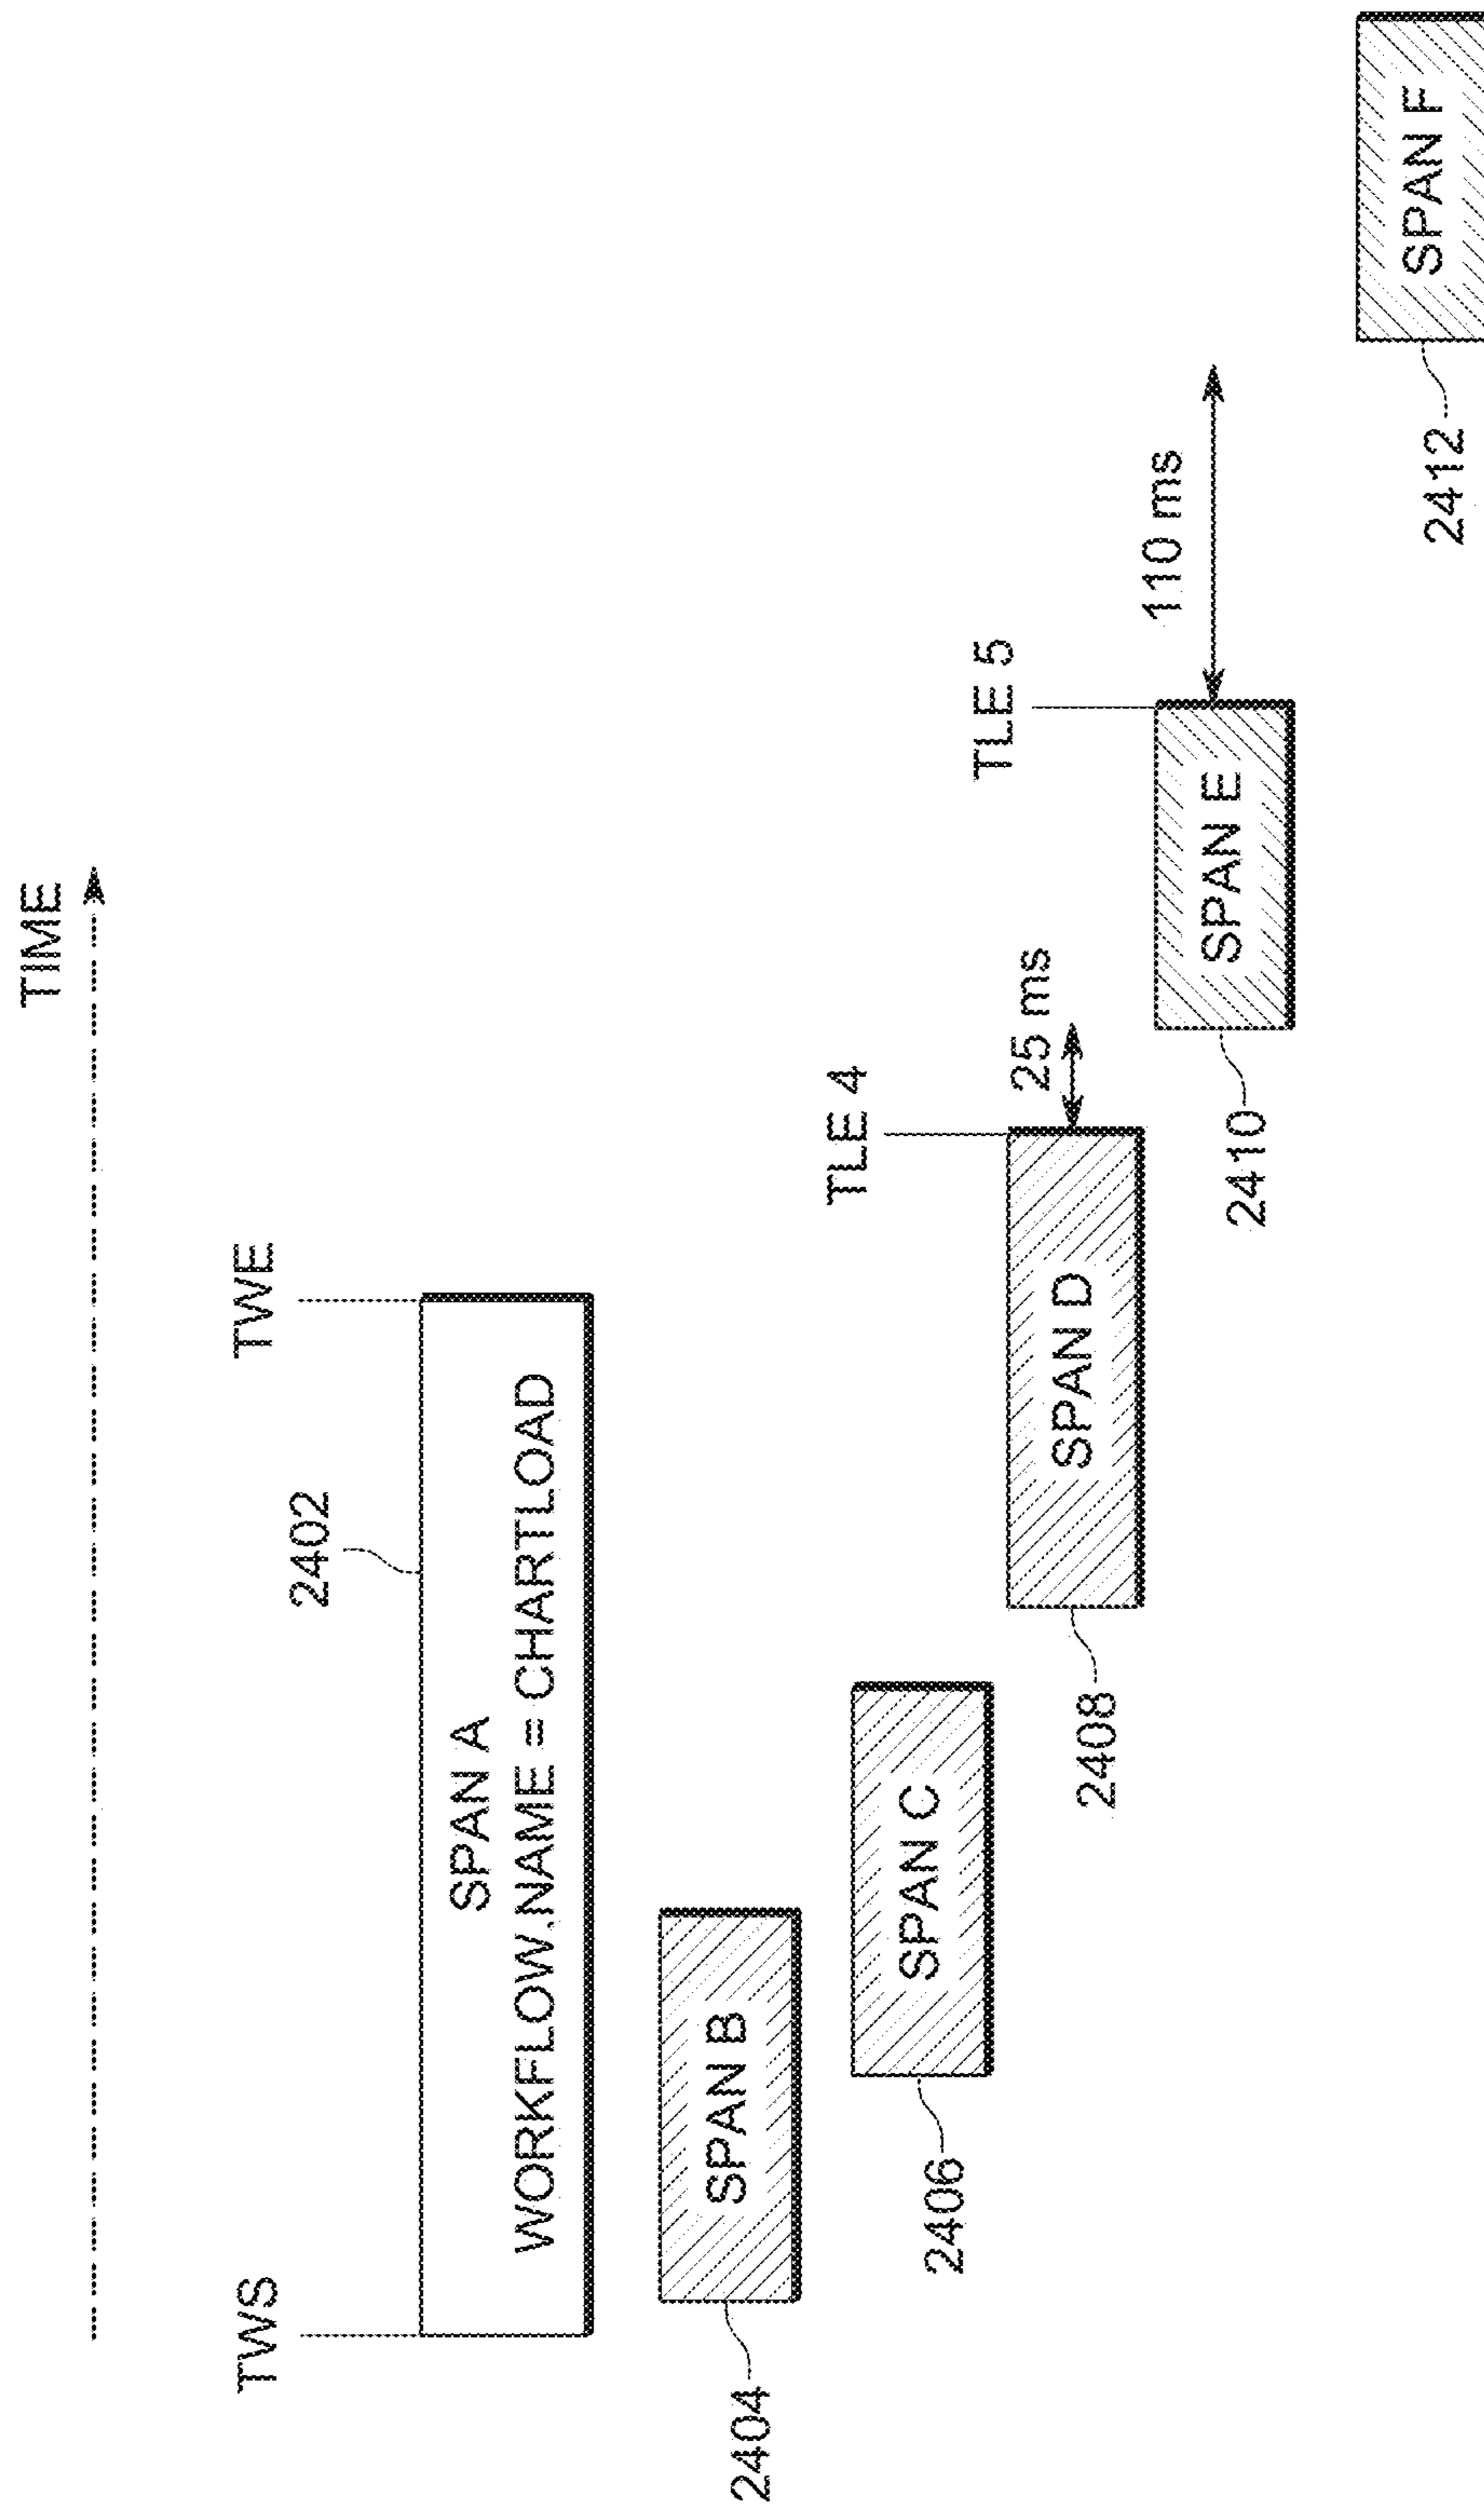


FIG. 24

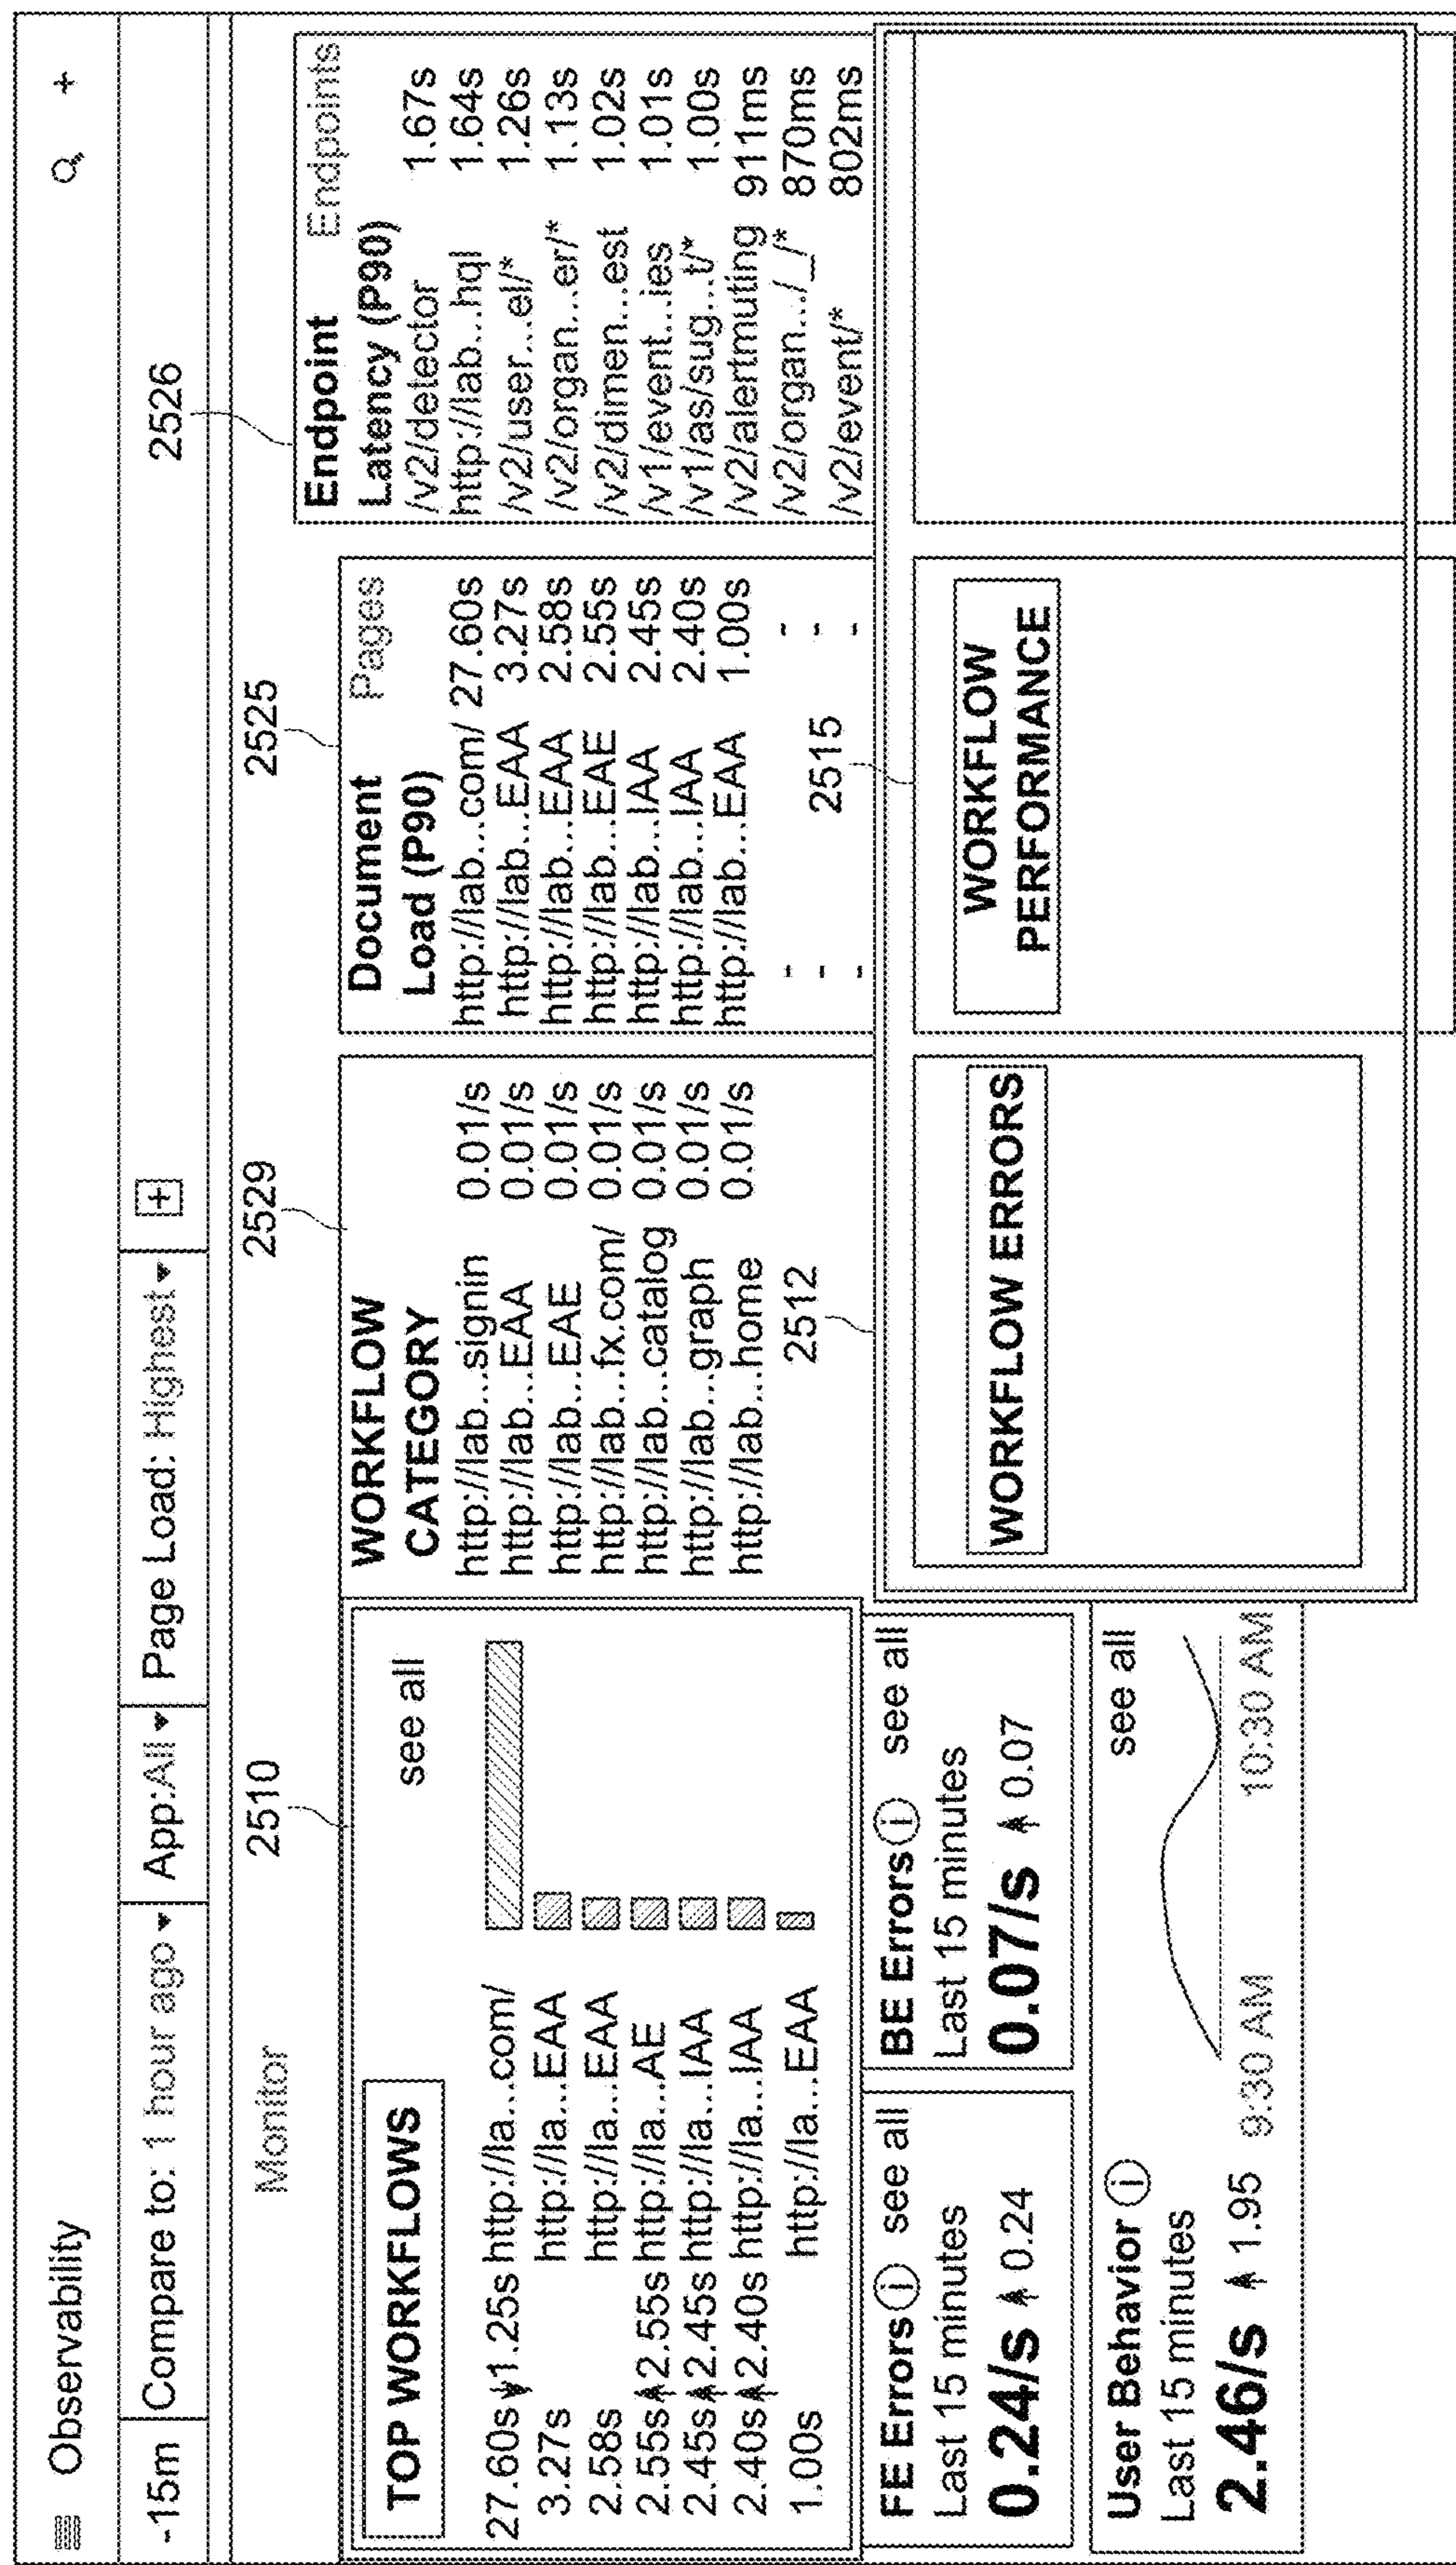


FIG. 25

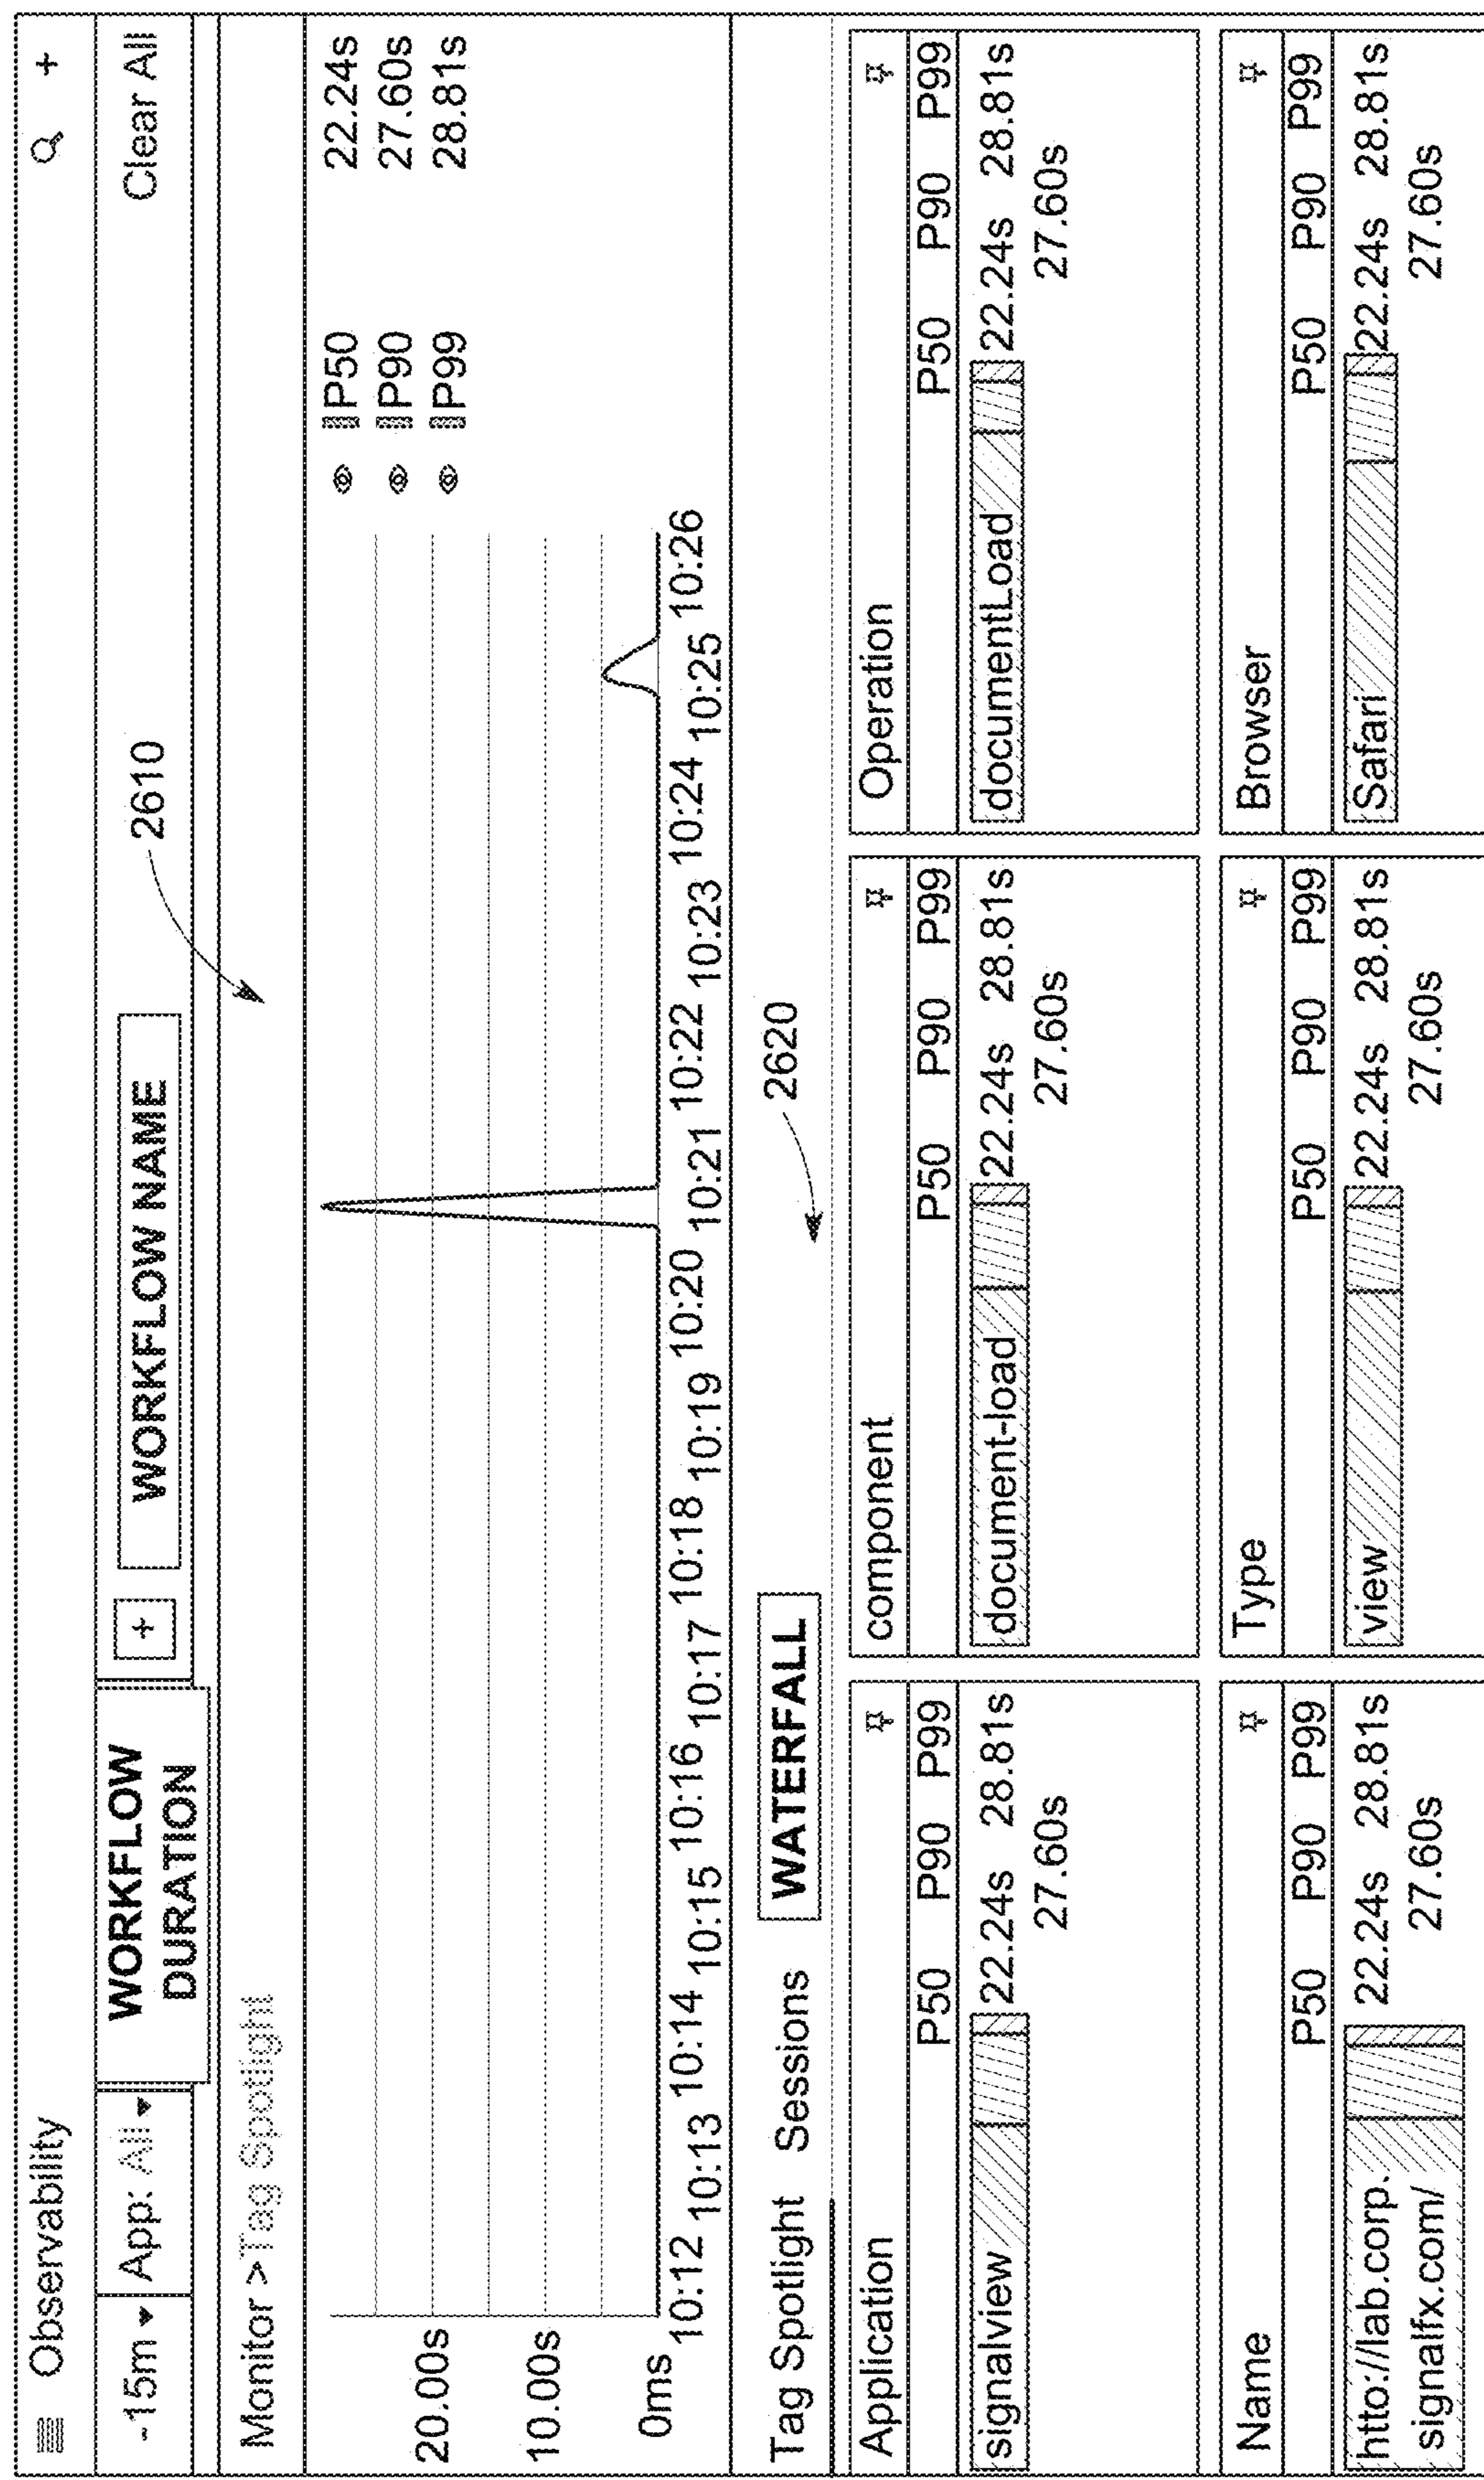


FIG. 26

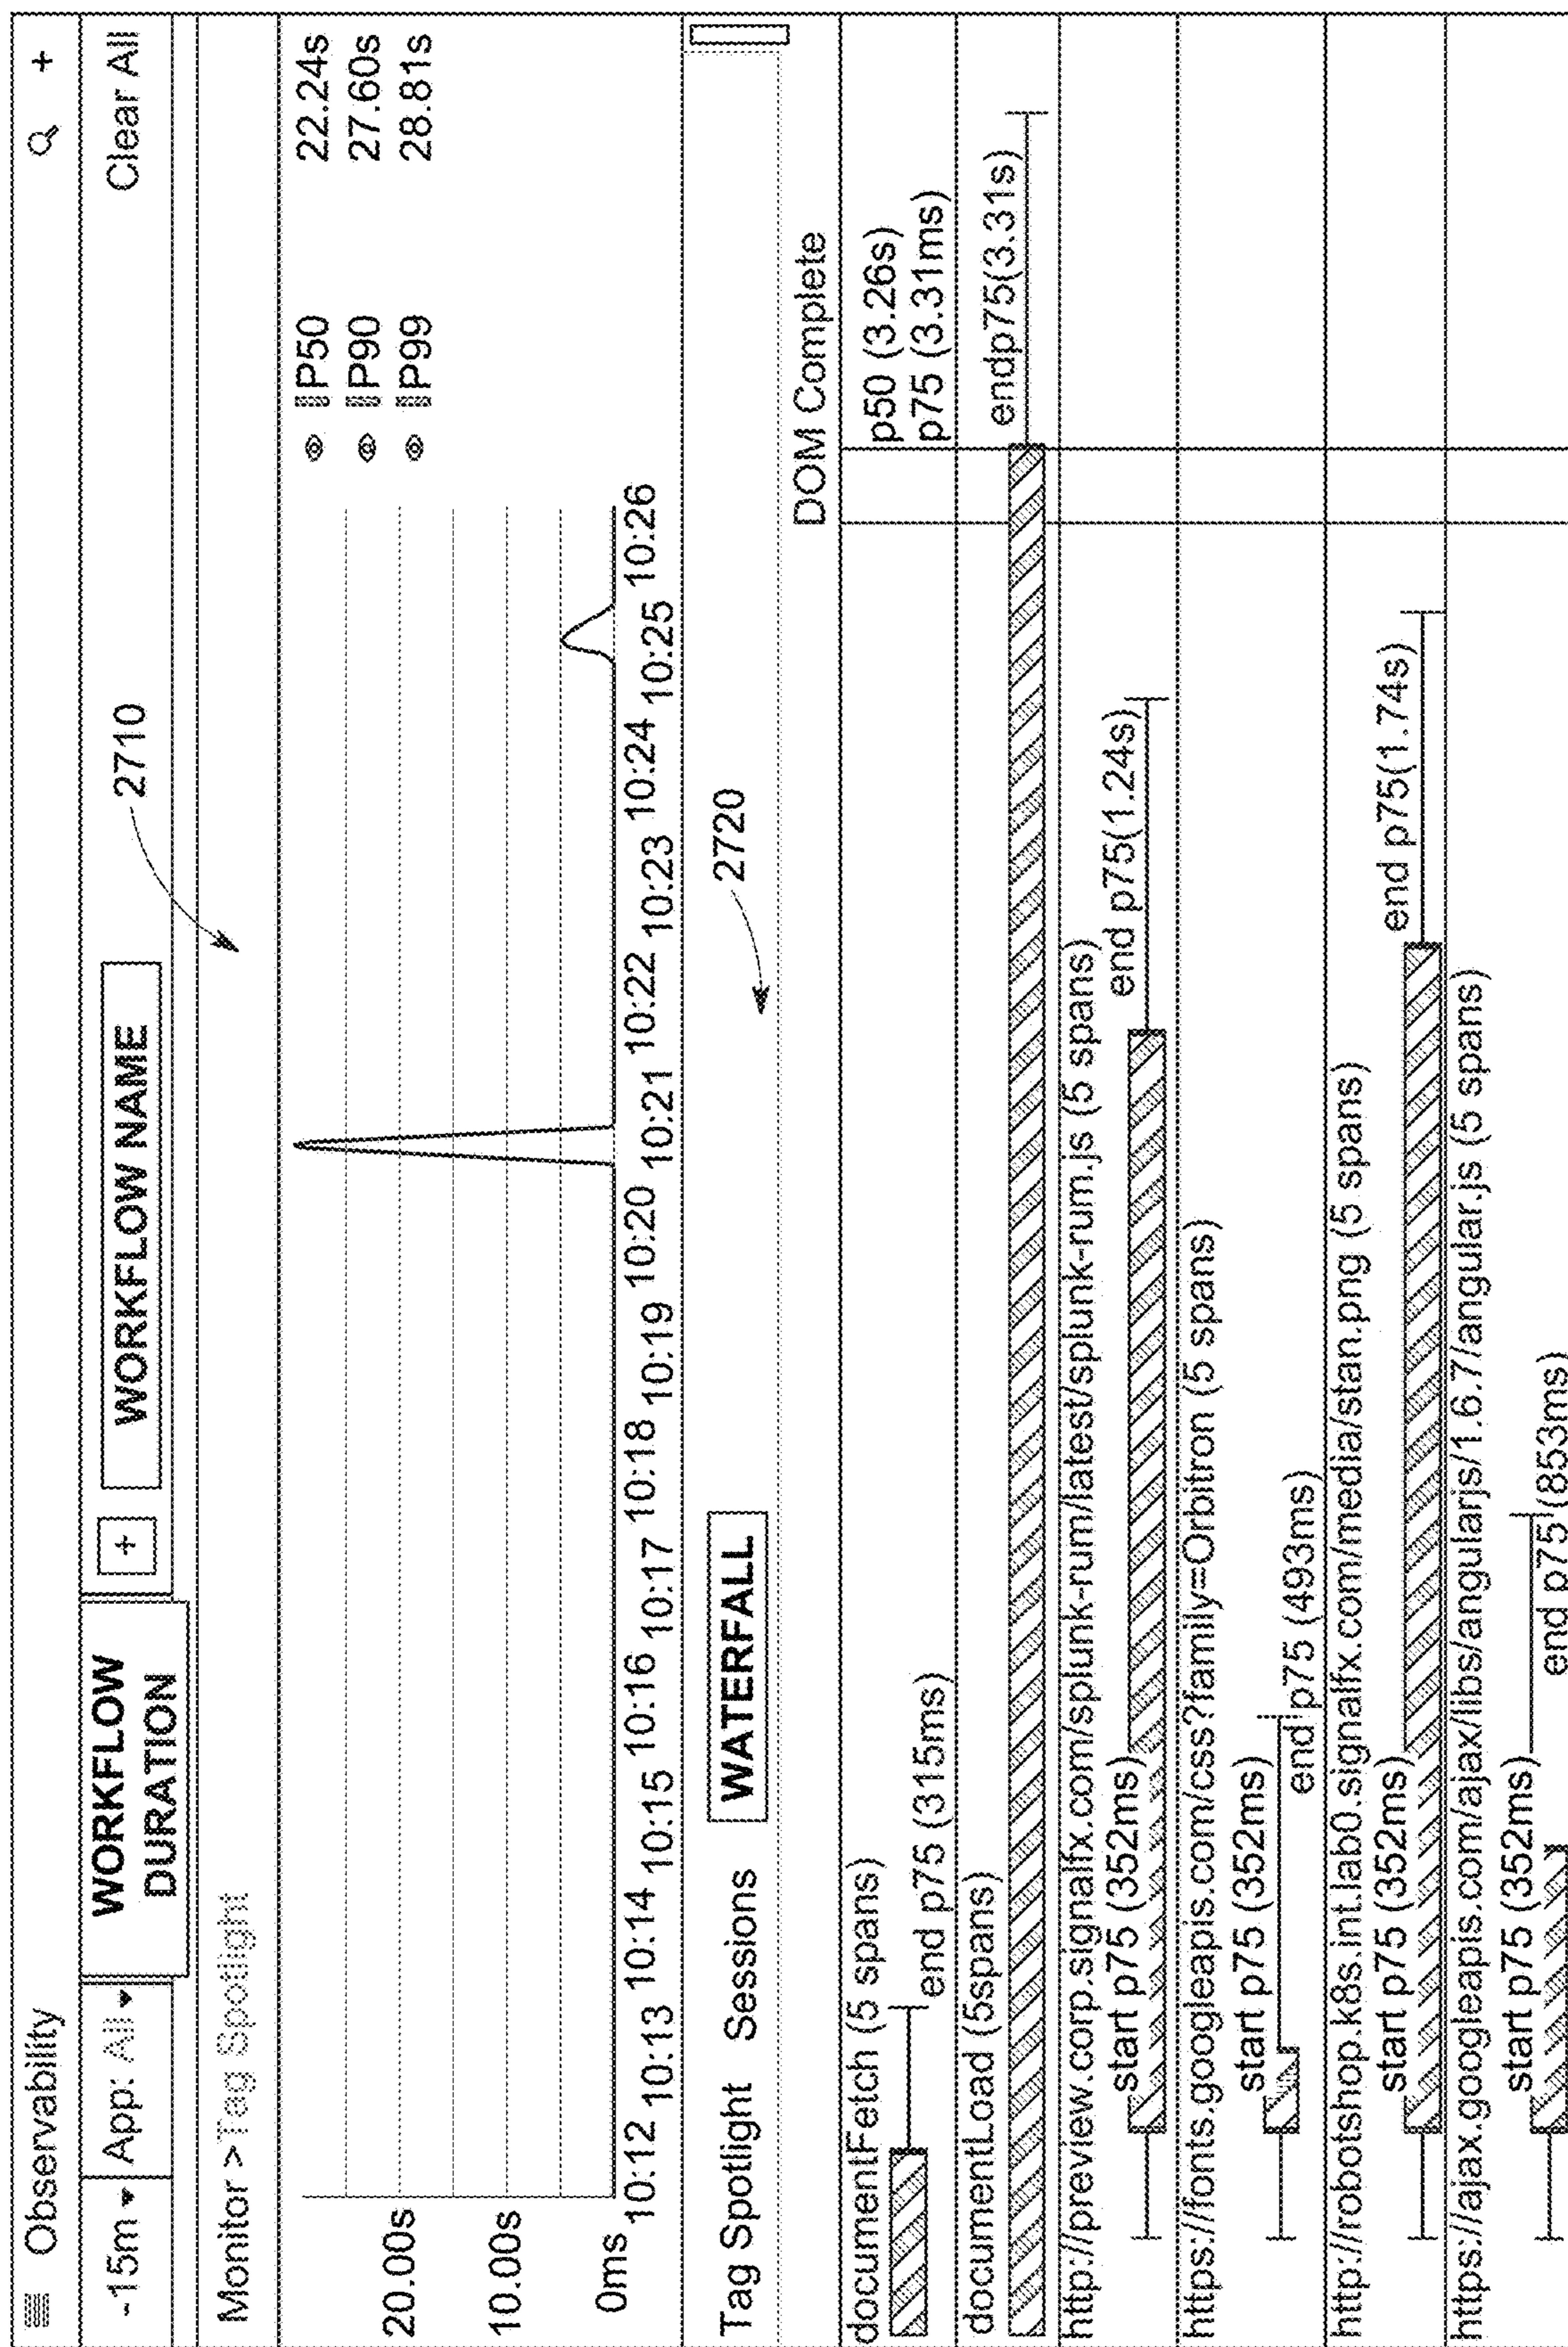


FIG. 27

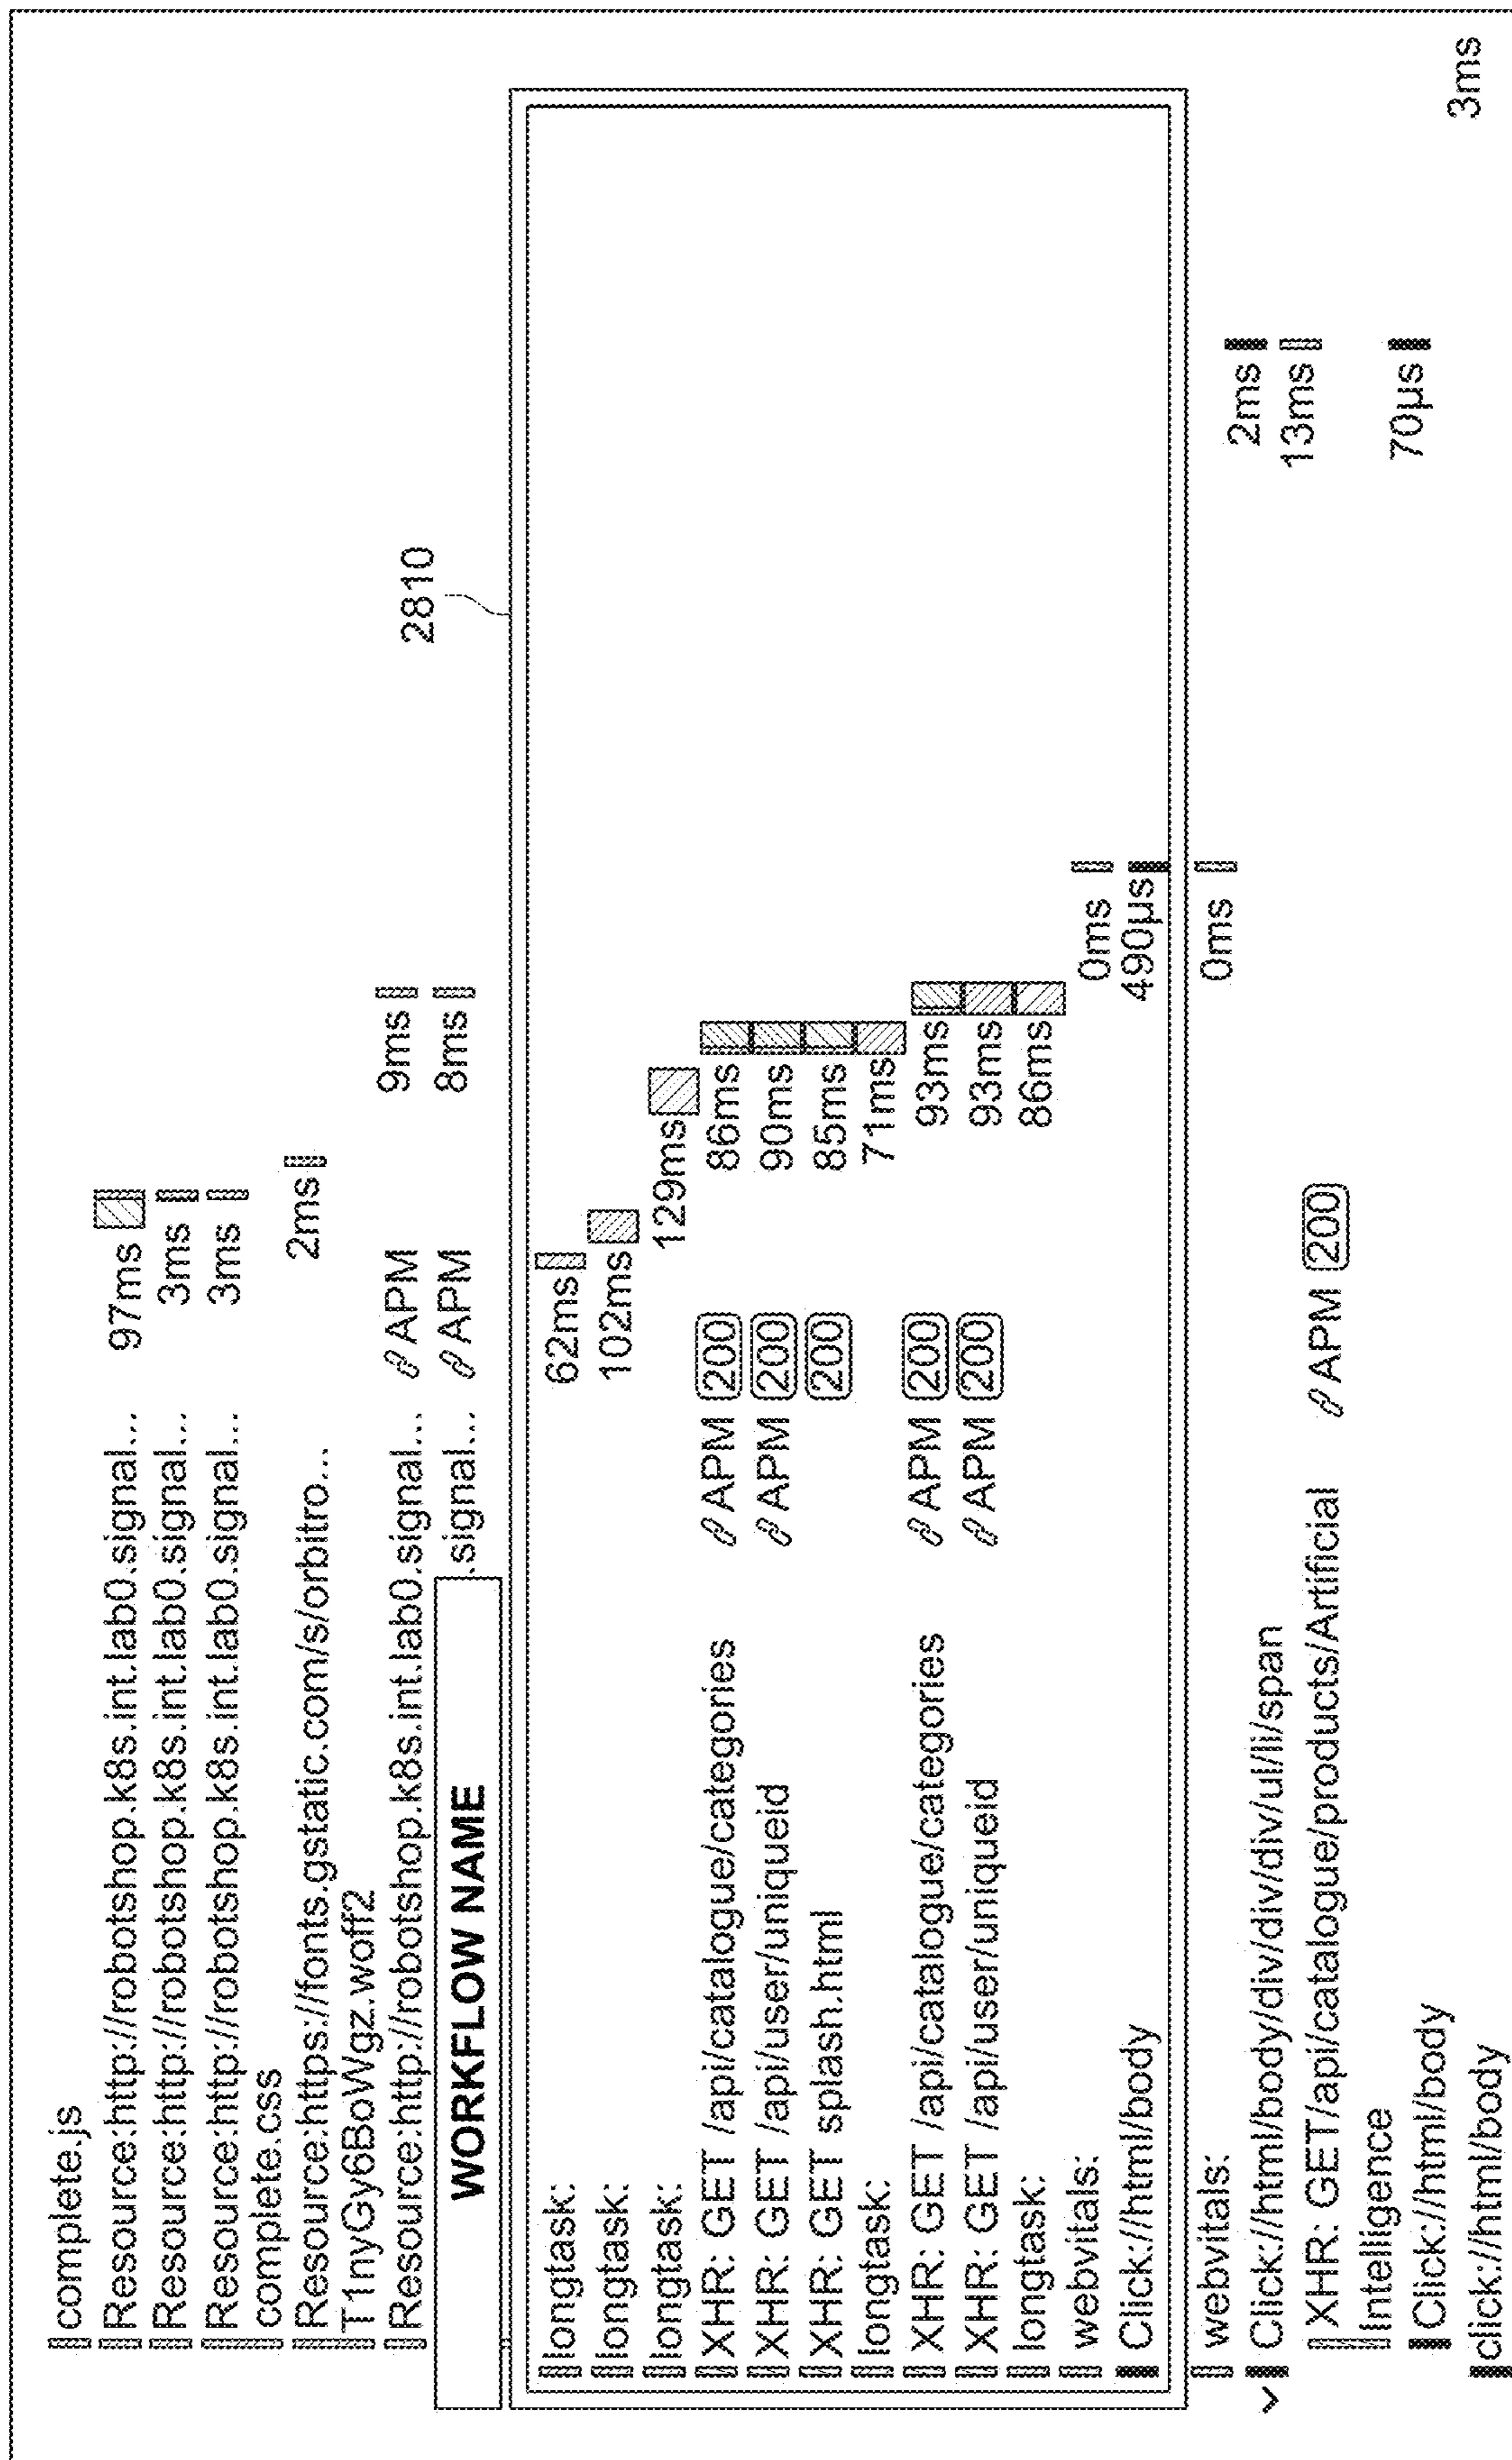


FIG. 28

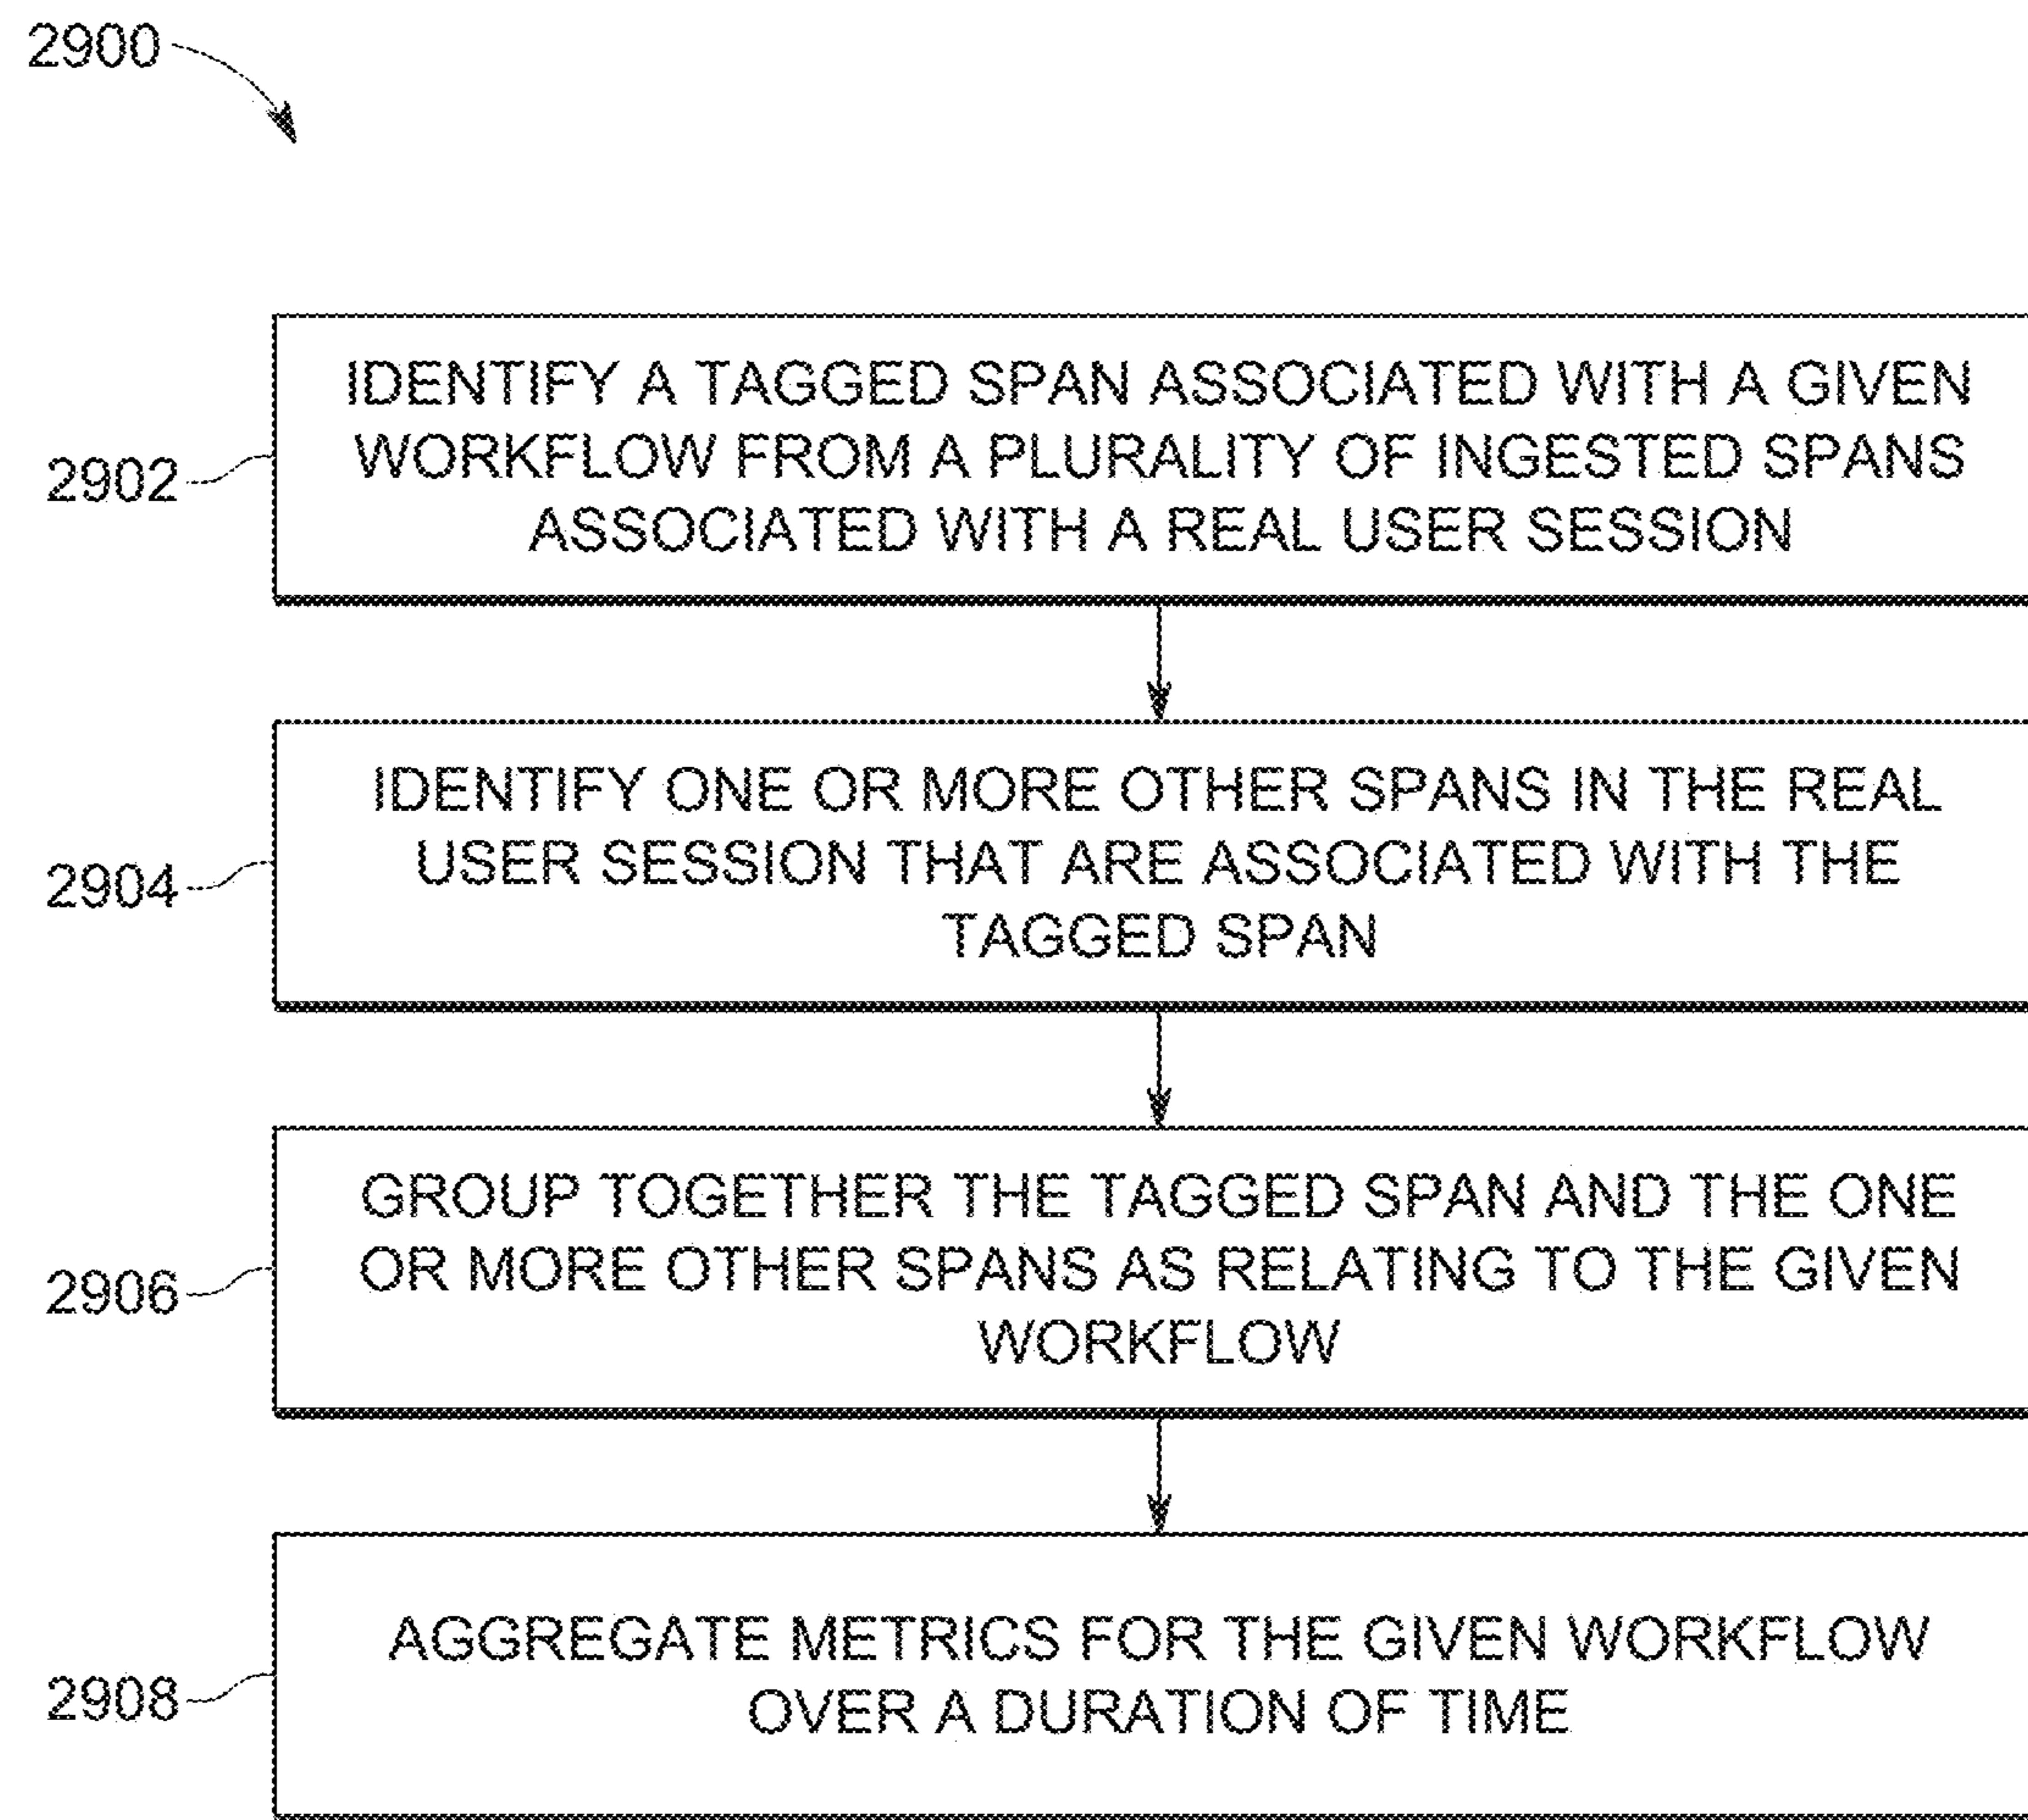


FIG. 29

1

**AGGREGATING METRICS FOR
WORKFLOWS ASSOCIATED WITH A REAL
USER SESSION**

BACKGROUND

Computer systems may run applications or services that are provided via a server or cloud computing environment. A client computer system may send a request to a server that retrieves application installation files in an underlying database. The applications or services may be developed and deployed as a single unit or as multiple units, such as a collection of microservices. Applications that are developed as a single unit may be monolithic applications that include a user interface and data access codes combined into a single program from a single platform. Monolithic applications are self-contained and independent from other computing applications. With the advent of cloud computing, however, these large centralized monolithic systems are being decoupled and distributed to address scalability needs and to allow companies to deliver value faster.

Microservices or a “microservices architecture” are used in a software development method wherein software applications are developed as a suite of independently deployable smaller cooperating services. The cooperating services run processes and communicate to serve a business goal to form an enterprise application. More specifically, in a microservices architecture, an application is developed as a collection of small services; each service implements business capabilities, runs in its own process and communicates via Application Program Interfaces (“APIs”), e.g., hypertext transfer protocol (HTTP) APIs, or messaging. Each micro-service may be deployed, upgraded, scaled and restarted independent of other services in the application, typically as part of an automated system, enabling frequent updates to live applications without impacting end customers.

With the rise of cloud native applications, e.g., Software as a Service (SaaS) applications, which include microservices, there has been a shift in the manner in which software is built and deployed, and also in the manner in which it is monitored and observed. Microservices-based applications operate within environments of dramatically increased complexity and many more layers of abstraction compared to previous generations of monolithic applications. Compared to monolithic applications, microservices architectures generally introduce complexity in network communication, feature short lifecycles and require resiliency in dynamic environments.

Diligent application performance monitoring (APM) is needed on the part of developers of microservices-based applications to ensure that their software delivers a steady baseline of performance. APM typically involves carefully managing the performance, availability and user experience of software applications, in particular micro-services based applications deployed in a cloud computing environment. Using APM-based tools, software developers for microservices-based applications monitor different aspects of the software they develop by instrumenting the software. These aspects include performance of the software, disk utilization of the software, CPU utilization of the software, errors encountered during execution of the software, significant events encountered during execution of the software, information describing which parts of code are being executed and which parts are not being executed, among others. After development, similar aspects of the software are also monitored during production, such as when software is being executed in a cloud architecture.

2

Conducting APM effectively involves not only monitoring the services deployed on the backend of a cloud computing environment but also monitoring service quality delivered by, for example, the frontend or user interface of an application or website (e.g., for a SaaS application). A critical aspect of APM involves gauging user experience, including key metrics such as the load time of a particular website or visualization. Real User Monitoring (RUM) (also referred to as real user measurement or end-user experiencing monitoring) is a type of passive performance monitoring that captures and analyzes each transaction by users of a website or an application (e.g., a cloud-based microservices-based application). Monitoring actual user interaction with a website or an application is important to operators (e.g., site reliability engineering teams or developer teams of a website or a cloud-based application) to determine if users are being served quickly and without errors and, if not, which part of a process is failing. SaaS and application service providers use RUM to monitor and manage service quality delivered to their clients and to detect errors or slowdowns on web sites. The data may also be used to determine if changes that are propagated to websites or cloud-based applications have the intended effect or cause errors.

As companies increasingly rely on cloud-native applications and microservices architectures, they run into operational complexity and struggle to efficiently monitor their environments. Performance monitoring for cloud-based applications is complicated because it involves monitoring many micro-transactions both on the frontend and the backend, where the transactions are handled by a variety of hosts, containers and infrastructure platforms.

One of the challenges associated with performance monitoring, for example, is providing a developer of a website, or cloud-based application, end-to-end visibility into the performance of the website or application for troubleshooting purposes. Providing a unified view that allows for end-to-end performance visibility is challenging, in part, because of the significant amounts of data (e.g., span data, trace data, etc.) generated by both the frontend (e.g., the website) or the backend (e.g., a database on the backend) in an application owner’s architecture, and using the data efficiently to provide a developer with insights into the performance of a website or application or to detect problematic conditions associated with, for example, browser performance, network performance, erroneous processes, failing services, etc. Further, providing a unified monitoring view for diagnosing problems and troubleshooting can be complex because it requires establishing a link between the manner in which backend problems surface (or present themselves) at the frontend, and vice versa.

Another challenge with respect to performance monitoring, for example, is efficiently ingesting and aggregating significant amounts of span and trace data generated by various frontend websites or applications and also backend services in an application owner’s architecture. Traditional monitoring and troubleshooting tools, designed as symptom-based solutions with single purpose capabilities, are simply unable to keep up with tracking the performance of dynamic cloud-native applications and analyzing the significant amounts of span and trace data they generate.

Conventional monitoring tools also are unable to group and track span data associated with discrete client functions or client flows. For example, conventional monitoring tools are typically unable to track and group spans associated with real user sessions that relate exclusively to a particular user-interaction or a client process associated with a website. Thus, systems that can efficiently and accurately conduct

performance monitoring for websites and cloud-based applications, and further aggregate the collected information in a meaningful way, are the subject of considerable innovation.

BRIEF DESCRIPTION OF THE DRAWINGS

In the drawings:

FIG. 1A illustrates an exemplary monolithic multi-layer architecture.

FIG. 1B illustrates an exemplary microservices architecture.

FIG. 2A illustrates an exemplary trace tree.

FIG. 2B illustrates an alternate view of the trace from FIG. 2A adjusted for timeline.

FIG. 3 is a flow diagram that illustrates the manner in which trace data may be collected and ingested for further analysis within a computer system, in accordance with implementations of the monitoring service disclosed herein.

FIG. 4 illustrates the backend components of an exemplary microservice application for an online retailer that are monitored using APM.

FIG. 5 is a flow diagram that illustrates an exemplary method of ingesting and aggregating span information to support multiple modalities of analysis for APM, in accordance with implementations of the monitoring service disclosed herein.

FIG. 6 illustrates the manner in which span metrics and trace metrics are generated, in accordance with implementations of the monitoring service disclosed herein.

FIG. 7 is a flow diagram that illustrates an exemplary method of generating metric time series from ingested spans, in accordance with implementations of the monitoring service disclosed herein.

FIG. 8 illustrates an exemplary on-screen GUI for APM illustrating a monitoring mode for an application displaying metric values aggregated from metric time series data, in accordance with implementations of the monitoring service disclosed herein.

FIG. 9 illustrates an exemplary on-screen GUI comprising an interactive topology graph for an application created from the aggregated metric events data, in accordance with implementations of the monitoring service disclosed herein.

FIG. 10 illustrates an exemplary on-screen displayed GUI showing the manner in which a client may access SLIs pertaining to a service within an interactive topology graph, in accordance with implementations of the monitoring service disclosed herein.

FIG. 11 illustrates an exemplary on-screen GUI showing the manner in which a client may access SLIs pertaining to an edge within an interactive topology graph, in accordance with implementations of the monitoring service disclosed herein.

FIG. 12A illustrates on-screen displays that represent exemplary categories of dimensions across which SLIs may be computed, in accordance with implementations of the monitoring service disclosed herein.

FIG. 12B illustrates an exemplary on-screen GUI showing the manner in which a client may filter a service graph for an application and associated metrics by the workflow dimension, in accordance with implementations of the monitoring service disclosed herein.

FIG. 13 is a flow diagram that illustrates an exemplary method of aggregating metrics data from ingested traces for the metric events modality, in implementations according to the present disclosure.

FIG. 14 is a table illustrating an exemplary manner in which selected tags for each service in a cross-service span

pair may be mapped to tag attributes and stored as part of a memory-resident data object associated with an edge in the service graph, in implementations according to the present disclosure.

FIG. 15A illustrates an exemplary on-screen GUI showing a visual representation of a portion of an exemplary trace illustrating a cross-service call, in implementations according to the present disclosure.

FIG. 15B illustrates the manner in which data in the metric events modality is persisted using an edge data object comprising a memory-resident table of tag attributes with associated counts and using a node data object comprising a memory-resident table of tags with associated counts, in implementations according to the present disclosure.

FIG. 15C illustrates the manner in which data in the metric events modality is persisted using an edge data object comprising a memory-resident table of extracted indexed tag attributes with associated Trace IDs and using a node data object comprising a memory-resident table of extracted tags with associated Trace IDs, in implementations according to the present disclosure.

FIG. 16 is an exemplary on-screen GUI showing the manner in which a client may submit a query to be executed against the full set of traces stored in connection with the full-fidelity modality, in accordance with implementations of the monitoring service disclosed herein.

FIG. 17 is a flow diagram that illustrates an exemplary computer-implemented method of ingesting and aggregating span information to support multiple modalities of analysis for RUM, in accordance with implementations of the monitoring service disclosed herein.

FIG. 18 illustrates an exemplary on-screen GUI comprising a service graph illustrating an access of one or more endpoints by a page with which a user is interacting, in accordance with implementations of the monitoring service disclosed herein.

FIG. 19 illustrates an exemplary on-screen GUI comprising a service graph illustrating an access of an endpoint by multiple pages, in accordance with implementations of the monitoring service disclosed herein.

FIG. 20A illustrates an exemplary on-screen GUI comprising an end-to-end view of a user session, in accordance with implementations of the monitoring service disclosed herein.

FIG. 20B illustrates an exemplary on-screen GUI illustrating the manner in which a span may be expanded within the waterfall view to obtain further information regarding the attributes of the span, in accordance with implementations of the monitoring service disclosed herein.

FIG. 20C illustrates an exemplary on-screen GUI illustrating the manner in which hovering a cursor over a backend trace link provided for a span in the waterfall view conveys further information regarding a backend trace in a pop-up window, in accordance with implementations of the monitoring service disclosed herein.

FIG. 21 illustrates an exemplary on-screen GUI displaying aggregate metrics for a specific page, in accordance with implementations of the monitoring service disclosed herein.

FIG. 22 illustrates an exemplary on-screen GUI displaying a geo-map associated with a particular website or application, in accordance with implementations of the monitoring service disclosed herein.

FIG. 23 presents a flowchart 2300 illustrating a computer-implemented process for identifying the entire set of RUM-related spans that are associated with a particular workflow, in accordance with implementations of the present monitoring service.

FIG. 24 illustrates an exemplary manner in which spans related to a particular workflow may be identified and grouped, in accordance with implementations of the monitoring service disclosed herein.

FIG. 25 illustrates an exemplary on-screen RUM monitoring GUI showing an aggregate view of instrumented workflows, in accordance with implementations of the monitoring service disclosed herein.

FIG. 26 illustrates an exemplary on-screen GUI showing a detailed view of a selected workflow, in accordance with implementations of the monitoring service disclosed herein.

FIG. 27 illustrates an exemplary on-screen GUI showing an aggregated waterfall view of an instrumented workflow, in accordance with implementations of the monitoring service disclosed herein.

FIG. 28 illustrates an exemplary on-screen GUI the manner in which associated spans from an individual session may be grouped into workflows, in accordance with implementations of the monitoring service disclosed herein.

FIG. 29 presents a flowchart 2900 illustrating a computer-implemented process for aggregating metrics associated with a user interaction during a real user session, in accordance with implementations of the present monitoring service.

DETAILED DESCRIPTION

Implementations are described herein according to the following outline:

- 1.0 Terms
- 2.0 General Overview
- 3.0 Data Collection
 - 3.1 Logs, Traces and Metrics
- 4.0 Multiple Modalities for Performing Application Performance Monitoring (APM)
 - 4.1 Metric Time Series
 - 4.1.1 Generating Metric Data Streams Using Span Identities
 - 4.1.2 Real-Time Monitoring Using Metric Time Series Data
 - 4.2 Metric Events
 - 4.2.1 Metric Events Data Generation and Persistence
 - 4.3 High-Fidelity Data
- 5.0 Multiple Modalities for Performing Real User Monitoring (RUM)
 - 5.1 End-to-End Visibility of a Real User Session
 - 5.1.1 Aggregating Metrics for Workflows Associated with a Real User Session

1.0 Terms

The term “trace” as used herein generally refers to a record of the manner in which a single user request, also referred to as a transaction, propagates from one microservice (hereinafter interchangeably referred to as “service”) to the next in a distributed application. A transaction is generally described as an end-to-end request-response flow, from the making of the user’s initial request to receiving the final response. A transaction often involves the interaction of multiple services. A trace is a record of a transaction and each trace may be identified using a unique trace identifier (“Trace ID”). The trace follows the course of a request or transaction from its source to its ultimate destination in a distributed system. In one implementation, a trace may be conceptualized as a highly dimensional structured log that captures the full graph of user-generated and background

request execution within an application, and includes valuable information about interactions as well as causality.

The term “span” as used herein generally refers to the primary building block of a trace, representing an individual unit of work done in a distributed system. A trace is composed of one or more spans where a span represents a call within the request. It is appreciated that a call may be to a separate microservice or a function within a microservice. The trace represents the work done by each microservice which is captured as a collection of linked spans sharing the same unique Trace ID. Each component of the distributed system may contribute a span—a named, timed operation representing a piece of the workflow. A span may also include a unique span ID, a service name (e.g., “analytics”), an operation name (e.g., “start”), duration (latency), start and end timestamps and additional annotations and attributes (e.g., tags such as key:value pairs). The annotations and attributes can describe and contextualize the work being done under a span. For example, each span may be annotated with one or more tags that provide context about the execution, such as the client instrumenting the software, a document involved in the request, an infrastructure element used in servicing a request, etc.

The term “tags” as used herein generally refers to key:value pairs that provide further context regarding the execution environment and enable client-defined annotation of spans in order to query, filter and comprehend trace data. Tag information is typically included with each span and there may be different levels of tag information included in a span.

Tag information (including the ‘key’ and corresponding ‘value’) is typically included with each span and there may be different levels of tag information included in a span.

“Global tags” generally represent properties of a user-request (e.g., tenant name, tenant level, user location, environment type, etc.) and may be extracted from any span of the trace based on configured rules. A global tag for a particular span in a trace may be attributed to the other spans in a trace, because each span within a single trace may comprise the same global attributes. For example, if one span within a trace comprises a tag relating it to a request from a “gold” level “tenant,” it may be inferred that other spans in the same trace are associated with the same request and, accordingly, from the same “gold” level “tenant.” Consequently, the “tenant:gold” key-value pair or tag may be attributed to the other spans in the same trace.

“Span-level tags” comprise attributes that are specific to a particular span.

The term “root span” as used herein generally refers to the first span in a trace. A span without a parent is called a root span.

The term “child span” as used herein generally refers to a span that follows a root span, including a child of a child.

The term “parent span” as used herein generally refers to a span that executes a call (to a different service or a function within the same service) that generates another span, wherein the span executing the call is the “parent span” and the span generated in response to the call is the “child span.” Each span may typically comprise information identifying its parent span, which along with the Trace ID, may be used to consolidate spans associated with the same user-request into a trace.

A “metric” as used herein generally refers to a single quantifiable measurement at a specific point in time. Combining the measurement with a timestamp and one or more dimensions results in a metric data point. A single metric data point may include multiple measurements and multiple dimensions. Metrics are used to track and assess the status

of one or more processes. A metric typically comprises a numeric value that is stored as a timeseries. A timeseries is a series of numeric data points of some particular metric over time. Each time series comprises a metric plus one or more tags associated with the metric. A metric is any particular piece of data that a client wishes to track over time.

2.0 General Overview

One of the fundamental shifts in modern day computing has been the shift from monolithic applications to microservices-based architectures. As previously mentioned, this is the shift from an application being hosted together (e.g., on a single system) to each piece of an application being hosted separately (e.g., distributed). FIG. 1A illustrates an exemplary monolithic multi-layer architecture. A monolithic application is traditionally built as a single unit. The monolithic application consists of a single self-contained unit in which code exists in a single codebase **100** and in which modules are interconnected. At deployment time, the entire codebase is deployed and scaling is achieved by adding additional nodes.

FIG. 1B illustrates an exemplary microservices architecture. A microservices architecture involves the building of modules (e.g., modules **104**, **106** and **108**) that address a specific task or business objective. As a result, these modules tend to exhibit low coupling and high cohesion. A microservices architecture is often achieved by decoupling a monolithic application into independent modules that each include the components necessary to execute a single business function. These services typically communicate with each other using language agnostic Application Programming Interfaces (“APIs”) such as Representational State Transfer (REST).

Microservices were created in order to overcome the issues and constraints of monolithic applications. Monolithic applications have a tendency to grow in size over time. As applications become larger and larger, the tight coupling between components results in slower and more challenging deployments. Because of the tight coupling, the potential for a failure of the entire application due to a recently deployed feature is high. In some cases, deployments may take several months to a year, greatly reducing the number of features that may be rolled out to users. This tight coupling also makes it difficult to reuse and replace components because of the effect they may have on other components throughout the application.

Microservices address these issues by being small in scope and modular in design. Modular design results in components being loosely coupled, which offers enormous benefits from the standpoint of being both fault tolerant and independently deployable. This results in functionality that may be frequently deployed and continuously delivered. The attribute of loosely coupled modules without a central orchestrator in a microservices architecture, however, leads to considerable challenges in terms of monitoring, troubleshooting and tracking errors.

These challenges have led to the rise of observability, a new generation of monitoring, the foundation for which is built, in part, on distributed tracing. Distributed tracing, also called distributed request tracing, is an application performance monitoring (APM) method used to profile and monitor applications, especially those built using a microservices architecture. Distributed tracing helps pinpoint where failures occur and what causes poor performance. Distributed tracing, as the name implies, involves tracing user requests

through applications that are distributed. A trace represents a single user request, also referred to as a transaction, and represents the entire lifecycle of a request as it traverses across the various services or components of a distributed system.

While distinct from the methodologies employed for APM, real user monitoring (RUM) is considered one of the critical strategies employed for performance monitoring by focusing on the manner in which end users’ experiences might inform application optimization strategies. RUM surfaces meaningful diagnostic information on frontend performance so developers can optimize frontend code and deliver the best possible user experience. APM meanwhile typically monitors the performance of server-side code and offers detailed insight on improving it to reduce infrastructure costs and creating faster applications for users.

RUM utilizes data related to the end users’ experiences to help developers track and improve a website or application’s performance. RUM focuses on measuring the experience of real users of a website or an application. It does this by tracking and reporting on several metrics including time-to-first-byte, full page load time, load time of specific elements, DNS timing, transaction paths, JavaScript errors, etc. With RUM, real user data can be tracked across browser versions, operating systems and end-user configurations. Tracking real users allows RUM to provide critical real-world measurements and helps developers identify whether certain user engagements or activities are triggering a lag in performance or causing errors. RUM, therefore, contributes to successful performance monitoring by analyzing how the end users’ experiences might inform application-optimization strategies.

RUM-based and APM-based methods together monitor the speed at which both frontend and backend transactions are performed both by end-users and by the systems and network infrastructure that support a software application, providing an overview of potential bottlenecks and service interruptions. This typically involves the use of a suite of software tools—or a single integrated SaaS or on-premises tool—to view and diagnose an application’s speed, reliability, and other performance metrics to maintain an optimal level of service.

Computing operations of instrumented software may be described by spans and traces. The spans and traces are produced by various instrumented services in an architecture and are communicated to an analysis system that analyzes the traces and spans to enable a software developer to monitor and troubleshoot the services within their software.

FIG. 2A illustrates an exemplary trace tree. The first span in the trace tree, Span A **202**, is known as the root span. A trace tree typically comprises a root span, which is a span that does not have a parent. It may be followed by one or more child spans. Child spans may also be nested as deep as the call stack goes. Span B **206** and Span E **204** are child spans of the parent span, Span A. Further, Span C **208** and Span D **210** are child spans of the parent Span B **208**.

FIG. 2B illustrates an alternate view of the trace from FIG. 2A adjusted for timeline. The trace starts with the Span A **202**, the root span, where the request starts. When the trace starts, a Trace ID is generated (e.g., Trace ID:1 as shown in FIG. 2B), which follows the request as it propagates through the distributed system. A new span is generated for each logical chunk of work in the request, where the new span includes the same Trace ID, a new Span ID and a Parent Span ID, which points to the span ID of the new span’s logical parent. The Parent Span ID creates a parent-child relationship between spans.

A given request typically comprises one span (e.g., the root Span A 202) for the overall request and a child span for each outbound call made to another service, database, or a function within the same microservice etc. as part of that request. For example, in the example of FIG. 2B, the Span A 202 is the root span for the overall request and generates several child spans to service the request. The Span A 202 makes a call to the Span B 206, which in turn makes a call to the Span C 208, which is a child span of the Span B 206. The Span B 206 also makes a call to the Span D 210, which is also a child span of the Span B 206. The Span A 202 subsequently calls the Span E 204, which is a child span of the Span A 202. Note, that the spans in a given trace comprise the same Trace ID. The Trace ID along with the Parent Span ID may be used to consolidate the spans together into a trace.

3.0 Data Collection

Distributed tracing data is generated through the instrumentation of browsers, microservices-based applications, libraries and frameworks. Software may be instrumented to emit spans and traces. The spans and traces may be generated according to an industry standard, such as the Open-Tracing standard. Other common open source instrumentation specifications include OPENTELEMETRY and OpenCensus. Each span may be annotated with one or more tags that provide context about the execution, such as the client instrumenting the software, a document involved in the request, an infrastructure element used in servicing a request, etc.

The instrumentation handles the creating of unique session IDs, trace and span IDs, tracking duration, adding metadata and handling context data. Handling context data, also known as context propagation is critical and is responsible for passing context (e.g. Trace ID) between function/microservice calls, thereby, enabling an observer to view the entire transaction at each step along the way. Context propagation may, for example, be based on REST. REST is header-based and requires a transaction to pass headers between service-to-service calls. In order to work properly, services within a request use the same context propagation format. Once the code has been instrumented and context propagation has been implemented using a standard format, the trace data generated by the services may be collected and analyzed to monitor and troubleshoot the microservices-based applications generating the trace data.

FIG. 3 is a flow diagram that illustrates the manner in which trace data may be collected and ingested for further analysis within a computer system, in accordance with an implementation of the monitoring service disclosed herein. Tasks 301 represent client applications that execute within a client data center for Client A. Similarly, tasks 302 represent client applications that execute within a client data center for Client B. The tasks 301 or 302 may comprise services or applications within a client's on-premises ("on-prem") software. Alternatively, they may comprise services or applications running in the cloud computing environment, e.g., in an AMAZON WEB SERVICES (AWS) Virtual Private Cloud (VPC).

The tasks 301 and 302 may be instrumented using open source or common commercial tracing libraries, from tracing applications (e.g., Jaeger or Zipkin), in-house formats, or auto-instrumentation. Each task may be configured to generate spans that describe the processing of a portion of a request as the request traverses through the various tasks (or services) on the client-side.

It should be noted that while the tasks 301 and 302 may comprise instrumented application software, the techniques disclosed herein are not limited to application software but are applicable to other kinds of software, for example, server software, software executing on customer devices, websites and so on. Furthermore, a client device (e.g., a device at a data center for Client A or Client B) may include any computing system that is configured to execute instrumented software, whether or not it is used for development of improved software. For example, the client device may be a computing system used for testing purposes, staging purposes, or any production system executing in an enterprise.

An agent 303 is typically configured at the client-side host or service for receiving spans collected from the various tasks on the client-side and transmitting the spans to a collector 304. An agent may receive generated spans locally using, for example, User Datagram Protocol (UDP). The tasks 302 may comprise instrumented tasks that are not using an agent and may be configured to span directly to the collector 304. The tasks may include various front-end tasks such as those performed by a web browser running on a client's computer. While spans may be collected from the client-side tasks without configuring an agent (e.g., in the case of Client B), using an agent may provide benefits including batching, buffering and updating trace libraries.

Batches of span data collected by the agent 303 are periodically received at the collector 304. The collector may be implemented within a client's on-prem software or in the cloud computing environment (e.g., in an AWS VPC). The collector 304 may also, for example, be implemented in a cloud computing environment by the same entity as the one implementing monitoring service 306. Traces often generate duplicative data that is not relevant for monitoring or troubleshooting. The collector 304 may avoid redundancies by sampling the data before processing and storing it. The collector 304 runs the span data through a processing pipeline and may store it in a specified storage or analytics backend such a monitoring service 306. It should be noted that the collector 304 may interact with the monitoring service 306 through a network (not shown).

In an implementation, the collector 304 may consolidate data from several client devices and combine the data to send to the monitoring service 306 (e.g., without sampling). For example, the collector 304 may comprise a server that receives data streams internally from different client devices and, periodically, sends the combined data (in batch form) to the monitoring service 306. The data streams may comprise trace related or metrics information. This improves efficiency of external communications from the enterprise.

In one implementation, the collector 304 may comprise a beacon module 388 configured to collect all data associated with RUM sessions, e.g., users' browsing sessions, users' interactions with an application or data generated by users' web browsers, etc. The beacon module 388 may, for example, be configured to collect all the spans generated by browser instrumentation configured on a client's device or a client's web-browser. The beacon may, among other functions, enrich the spans generated at the frontend (e.g., by a browser) with additional information (e.g., with HTTP client's IP address) before forwarding the information to be ingested by the monitoring service 306. Note that the beacon module 388 may not necessarily be a component within the collector 304 but may also be implemented as a standalone module. Further note that similar to the collector 304, the beacon module 388 may be implemented within a client's

on-prem software or in the cloud computing environment (e.g., in the same environment in which monitoring service 306 is implemented.).

In an implementation, the monitoring service 306 receives and analyzes the span data for monitoring and troubleshooting purposes. It should be noted that, in addition to monitoring service 306, span and tracing data might also be simultaneously transmitted to other types of storage and monitoring back-end services, e.g., a data ingestion and query system 326.

In one implementation, the monitoring service 306 may be a Software as a Service (SaaS) based service offering. Alternatively, in another implementation, it may also be implemented as an on-prem application. The monitoring service 306 receives the observability data collected by the collector 304 and provides critical insights into the collected trace data to a client of the monitoring service, who may be an application owner or developer. In an implementation, the monitoring service 306 may be hosted on a computing system that includes one or more processors, memory, secondary storage and input/output controller. The computing system used for hosting the monitoring service 306 is typically a server class system that uses powerful processors, large memory resources and fast input/output systems.

The monitoring service 306 may comprise an instrumentation analysis system 322 (also referred to herein as an “analytics engine”) and a query engine and reporting system 324. The instrumentation analysis system 322 receives data comprising, for example, trace information, span information and/or values of metrics sent by different clients. As noted previously, a task or software program may be instrumented to generate spans with a common field in their data structures to designate spans that are part of a common trace. For example, the spans may include a trace identifier such that spans with the same trace identifier are a part of the same trace.

The tasks (or software) executing on the client device are configured to send information generated as a result of instrumenting the software to the instrumentation analysis system 322 of the monitoring service 306. For example, the tasks may send span information collected from the various services at the client end to the instrumentation analysis system 322. Alternatively, traces may be sampled to generate metric values, and the tasks may send values corresponding to various metrics as they are generated to the instrumentation analysis system 322. The tasks may send group values of metrics periodically to the instrumentation analysis system 322. Different tasks may send the same metric or different metrics at different rates. The same task may send different metrics at different rates.

In an implementation, the tasks (e.g., tasks 301 and 302) and the collector 304 may send data to the monitoring service 306 by invoking an API supported by the monitoring service 306 and the instrumentation analysis system 322. In one implementation, a customer name may be specified for the instrumented software. The instrumented software includes the customer name when it identifies a data stream associated with that particular customer. The ability to associate a data stream with a customer allows the instrumentation analysis system 322 to perform customer specific analysis, for example, report on usages of systems for each customer, identify customers reporting more than a threshold number of errors and so on.

In one implementation, an application owner or developer may submit queries to the query engine and reporting system 324 to gain further insight into the spans and traces (or metrics) received and analyzed by the instrumentation

analysis system 322. For example, the query engine and reporting system 324 within the monitoring service 306 may be configured to generate reports, render graphical user interfaces (GUIs) and/or other graphical visualizations to represent the trace and span information received from the various clients. The query engine and reporting system 324 may, for example, interact with the instrumentation analysis system 322 to generate a visualization, e.g., a histogram or an application topology graph (referred to interchangeably 10 as a “service graph” herein) to represent information regarding the traces and spans received from a client. Alternatively, the query engine and reporting system 324 may be configured to respond to specific statistical queries submitted by a developer regarding one or more services within a client’s 15 application.

3.1 Logs, Traces and Metrics

As mentioned above, the shift from monolithic applications 20 to microservices-based architectures has increased the usefulness of analyzing traces in a distributed system. In one or more implementations, the tracing data may be coupled with log data and/or metrics data, in order to provide clients with a more complete picture of the system. For example, 25 the trace data may be coupled with log or other data from the data ingestion and query system 326. In one implementation the data ingestion and query system 326 may be comprised within the monitoring service 306.

One example of a data ingestion and query system 326 is 30 the event-based data intake and query SPLUNK® ENTERPRISE system developed by Splunk Inc. of San Francisco, California. The SPLUNK® ENTERPRISE system is the leading platform for providing real-time operational intelligence that enables organizations to collect, index and search machine-generated data from various data sources 328, for 35 example, websites, applications, servers, networks and mobile devices that power their businesses. In one implementation the other data sources 328 may be associated with the same clients (e.g., Client A and Client B) that generate 40 the trace data received by the monitoring service 306.

The SPLUNK® ENTERPRISE system is particularly useful for analyzing data which is commonly found in system log files, network data and other data input sources. In another example, the data ingestion and query system 326 45 may be an on-premises application or based on a distributed or cloud-based service.

In one implementation, the trace data may be ingested into the data ingestion and query system 326, or may be coupled with outputs from the data ingestion and query system 326 50 e.g., from searches that may be based on trace data and run on the data ingestion and query system 326. In some implementations, the data ingestion and query system 326 described above may be integrated with or into the monitoring service 306 that analyzes trace data, e.g., the monitoring service 306. The monitoring service 306 may, accordingly, 55 comprise a full suite of services including, for example, analyzing spans generated by users’ browsing sessions and other frontend activities, analyzing trace data, generating metrics data from the trace data, ingesting and analyzing log data, ingesting metrics data and providing insights generated from the metrics data, including aggregating and/or correlating trace data, log data and metrics data, in order to gain insights into a computing platform.

As described above, the span, trace and other data 60 received from the collector 304 may be sent to systems configured to ingest and search data, such as the data ingestion and query systems 326 described above. In some

implementations data ingestion and query system 326 may be configured to generate metrics data from the trace data received from the collector 304. Additionally, other implementations may use a stream processor that may perform transformations and other operations on incoming data prior to, concurrently with, and/or as an alternative to, ingestion of the data. In some implementations, the system may also be configured to ingest metrics data and may be optimized to ingest, query and generate insights from metrics data.

In other implementations, metrics may be generated by instrumentation (e.g., from instrumenting client software and tasks, e.g., tasks 301, 302 etc. as described above) and sent to a SaaS-based processing system, e.g., the monitoring service 306. For example, software may be instrumented to send metrics to a gateway or to a instrumentation analysis engine, where metrics may be aggregated, queried and alerted.

As above, the trace data may be paired with data from the data ingestion and query system 326, metrics generated by instrumentation, and other data sources, and correlated in various ways to provide insights. For example, as a broad-based correlation example, the metrics data may be used in a thresholding comparison to determine that there is an issue that needs attention, the trace data may be used to determine which component or microservice requires attention, and log data from the data ingestion and query system 326 may be used to determine exactly why the component or microservice needs attention. Other correlations and uses for the combination of metrics data, log data and event data are also contemplated herein. As noted above, the various features and services may be provided within an integrated monitoring platform (e.g., the monitoring service 306), wherein the platform comprises, among other things, an instrumentation analysis system (e.g., the instrumentation analysis system 322), a query engine and reporting system (e.g., the query engine and reporting system 324) and a data ingestion and query system (e.g., the data ingestion and query system 326).

4.0 Multiple Modalities for Performing Application Performance Monitoring (APM)

As noted previously, APM methods such as distributed tracing are used to profile and monitor applications, especially those built using a microservices architecture, at the backend of a website or application. Historically, there have been several challenges associated with implementing an analytics tool such as the monitoring service 306 within a heterogeneous distributed system. One of the challenges associated with APM for example, is efficiently ingesting and aggregating significant amounts of span and trace data generated by various services in an architecture. Conventional tracing and monitoring systems are typically unable to ingest the vast amounts of span and tracing data generated by clients' application and have to resort to sampling the data intelligently to reduce the volume of stored trace data. Using sampling exclusively, however, results in loss of data and, as a result, conventional monitoring tools do not allow clients access to all the traces generated by their application. Furthermore, conventional monitoring tools may calculate metrics (e.g., requests, errors, latency, etc.) based on the sampled set of data and, accordingly, the calculations may be approximate at best and inaccurate at worst.

Advantageously, implementations of the monitoring service (e.g. monitoring service 306) disclosed herein allow clients of the monitoring service the ability to ingest up to 100% of the spans and create streams of metric data using

the ingested spans prior to consolidating the spans into traces (through a sessionization process). The metric time series provide valuable real-time information pertaining to services or endpoints within an application and also allow alerts to be configured to manage anomalous behavior on the endpoints.

Implementations of the monitoring service disclosed herein also sessionize and store up to 100% of the spans received from the client in real time. Implementations of the monitoring service disclosed herein comprise an ingestion streaming pipeline that is able to ingest and consolidate the incoming spans into traces, and is further able to use advanced compression methods to store the traces. Additionally, because incoming trace and span information may be efficiently ingested and aggregated in real time, a monitoring platform is able to advantageously convey meaningful and accurate information regarding throughput, latency and error rate (without the need for sampling) for the services on the backend in the microservices-based application. High-cardinality metrics pertaining to throughput, latency and error rate may be calculated with a high degree of accuracy because all incoming data is accounted for and there is no data loss as a result of sampling. It should be noted that that monitoring service disclosed herein is able to ingest and store up to 100% of the spans for both APM and RUM.

Implementations of the monitoring service disclosed herein further allow a client to store and analyze the span and trace data using multiple modalities of analysis for both APM and RUM. The manner in which multiple modalities of analysis are supported for RUM will be discussed further below in connection with FIG. 17. In one implementation, a first modality comprises converting incoming spans from one or more clients into a plurality of metric data streams (also referred to as metric time series) prior to sessionizing the spans. Each metric time series is associated with a single span identity, where a base span identity comprises a tuple of information corresponding to an associated type of span. Each metric time series in this modality (referred to herein as "metric time series modality") represents a plurality of tuples, each tuple representing a data point. Key performance metrics (KPIs) can be extracted directly from the metric time series in real-time and reported to a client. Because the metric time series are created without paying a time penalty associated with sessionization, they can be used to perform real-time monitoring with sub-second resolution and generate alerts within two to three seconds if a condition is violated.

In one or more implementations, a second modality of analysis sessionizes the incoming spans and supports deriving higher-cardinality metrics (as compared with metric time series data) for a selected set of indexed tags, e.g., user-selected tags, global tags of the trace, etc. over selected time durations (referred to herein as the "metric events modality"). This modality is particularly useful for clients that need accurate SLI information for a larger set of high-value indexed tags. The metric events modality enables developers to aggregate metrics that have been pre-generated using the sessionized trace data to efficiently respond to queries submitted by a client. The aggregated metrics provide a client visibility into the performance of services within a microservices-based application. The metric events modality may deprioritize speed as compared to the metric time series to provide a client resolution into a larger set of indexed tags. As such, responses provided by the metric events modality are typically slightly slower as compared with the sub-second response rates of the metric time series.

In one or more implementations, the metric events modality may also keep track of exemplary traces associated with a pre-configured set of indexed tags. The tags to be indexed may be pre-selected by the client or the monitoring platform. The Trace IDs may be used to retrieve the associated traces and analysis on the actual traces may be performed to generate more particularized information, e.g., span duration, span count, span workload percentage, etc. for each span in a given trace. In one implementation, once the traces are retrieved, an analysis may be run on an arbitrary set of tags (in addition to the pre-configured indexed tags).

Additionally, in one or more implementations, a third modality of analysis may comprise a “full-fidelity” modality where a full-fidelity analysis may be conducted on any dimension or attribute of data to gauge the performance of services in the microservices-based application. The full-fidelity modality allows clients to search most or all of the incoming trace data (including all the tag data) that was ingested by the monitoring platform without relying on sampling. The full-fidelity modality may sacrifice speed for accuracy, and may be used by clients that need a more thorough analysis of the services across every dimension or attribute.

In an implementation, the three modalities may be supported by the monitoring platform simultaneously by storing ingested trace data using three different formats, wherein each format corresponds to one of the three available modalities of analysis. Note that implementations of the monitoring service disclosed herein are not restricted to three discrete data sets. The data sets for the different modalities may overlap or may be saved as part of a single data set. When a client submits a query, the monitoring platform may determine which of the data sets is most suitable for addressing the query. Thereafter, the monitoring platform executes the query against the selected data set to deliver results to the client. By comparison, conventional monitoring systems typically focus on a single modality and do not provide clients the ability to seamlessly navigate between different modalities. Conventional monitoring systems also do not provide the ability to automatically select the most appropriate modality based on the content, structure, syntax or other specifics pertaining to an incoming query.

FIG. 4 illustrates the backend components of an exemplary microservice application for an online retailer that are monitored using APM. A user needing to conduct a transaction may visit the website of the online retailer which would initiate a call to the retailer’s Front-end service **404** on a server. The call to the Front-end service **404** may subsequently trigger a chain of calls on the retailer’s back-end that would not be transparent to the client. For example, if the user proceeds to complete the transaction by checking out, several calls may be made to the back-end to services such as a CheckOutService **406**, a PaymentService **408**, an EmailService **410**, a ShippingService **412**, a CurrencyService **428** and a CartService **414** that may be involved in processing and completing the user’s transactions. Note, that a given request submitted by a user to the website would involve a subset of the services available and, typically, a single request would not result in a call to each of the services illustrated in FIG. 4.

As mentioned above, a request that the user initiates would generate an associated trace at the backend. It is appreciated that each user request will be assigned its own Trace ID, which will then propagate to the various spans that are generated during the servicing of that request. Each service may process a portion of the request and generate

one or more spans depending on the manner in which instrumentation is configured for a respective service. The Trace ID may then be used by the server to group the spans together into a trace with that Trace ID. So, for example, the user’s checkout transaction may generate a call at the Front-end service **404**, which may in turn generate calls to various microservices including the CheckoutService **406**. The CheckoutService **406** may, in turn, generate calls to other services such as the PaymentService **408**, the EmailService **410** and the ShippingService **412**. Each of these calls passes the Trace ID to the respective service being called, wherein each service in the call path could potentially generate several child spans.

It should be noted that a service does not necessarily need to make calls to other services—for instance, a service may also generate calls to itself (or, more specifically, to different operations and sub-functions within the same service), which would also generate spans with the same Trace ID. Through context propagation then, each of the spans generated (either by a service making a call to another service or a service making a call to various operations and sub-functions within itself) is passed the Trace ID associated with the request. Eventually, the spans generated from a single user request would be consolidated (e.g., by the collector **304** or the monitoring service **306** of FIG. 3) together using the Trace ID (and the Parent Span IDs) to form a single trace associated with the request.

As noted above, conventional distributed tracing tools are not equipped to ingest the significant amounts of span and tracing data generated by clients’ application and have to resort to sampling the data intelligently to reduce the volume of stored trace data. Further, conventional distributed tracing tools do not provide application owners multiple modalities of storing and querying trace data with the flexibility of switching between the different modalities depending on the level of detail required to respond to a client’s query.

Referencing FIG. 4 again, an owner of the application **400** may, for example, need varying degrees of detail regarding the services in the application. For example, the application owner may need to monitor certain metrics (e.g., RED metrics associated with Request, Errors, Durations) in real-time associated with a particular service, e.g., CheckoutService **406**. Assuming there are errors generated by a call made from the Frontend service **404** to the CheckoutService **406**, the owner may require further information pertaining to additional tags (indexed or non-indexed) associated with CheckoutService **406**. The application owner may also need to access each of the spans or full trace(s) associated with the request from the Frontend service **404** to the CheckoutService **406** to perform a more detailed analysis. Each of the requests requires a different degree of detail extracted from the span and trace information.

In one implementation, the metric time series modality allows the client to monitor RED metrics associated with a given service, e.g., CheckoutService **406** in the online retailer’s application in real-time. In one implementation, the metric time series modality can also be configured to deliver real-time alerts to a client based on each of the RED metrics, e.g., anomalies related to the request rate, error rate, or latency (duration).

If the client needs Service Level Indicators (SLIs) pertaining to certain indexed tags related to the call between Frontend service **404** and CheckoutService **406** for a given time duration, the metric event modality may enable the client to perform aggregations of metrics data computed from the indexed tags associated with the spans generated by the call between the Frontend service **404** and the Check-

outService **406**. The metrics aggregation may be a numeric summation, for example, and may be performed relatively quickly.

The metric event modality, in accordance with implementations of the monitoring service disclosed herein, associates the selected tags indexed from the incoming span data (e.g., the same indexed tags used for performing metrics extraction) with Trace IDs (or Span IDs in the case of RUM data as will be discussed in connection with FIG. 17 below) for exemplary traces. The Trace IDs (or Span IDs in the case of RUM data) may be used to retrieve the exemplary traces (or spans) associated with indexed tags. Thereafter, the monitoring platform may analyze the exemplary spans or traces to generate more particularized information, e.g., span duration, span count, span workload percentage, etc. for each span in a given trace. For the example of FIG. 4, if the client requires a performance summary for the spans generated by the call made from the Frontend service **404** to the CheckoutService **406**, the associated query submitted by the client may access the data set associated with the metric event modality. Using the Trace IDs (or Span IDs for RUM) corresponding to the indexed tags, the monitoring platform may then perform the computations necessary on the corresponding exemplary traces (or spans) to provide the client further information regarding the span performances. In an implementation, the client may also be able to extract meaningful information from the unindexed tags associated with the spans generated by the call using the exemplary traces.

If the client wants to search all the incoming trace data associated with the call between Frontend service **404** to the CheckoutService **406**, implementations of the monitoring service disclosed herein provide a third modality of analysis. In the full-fidelity modality, a full-fidelity analysis may be conducted on any dimension, tag or attribute of the span or trace data. For example, the client may be able to search previously indexed or unindexed tags across each of the traces associated with the call between the Frontend service **404** and the CheckoutService **406**. The full-fidelity modality allows an analysis to be performed across any relevant span or trace. Conventional tracing systems are unable to provide that level of flexibility and detail for developers or application owners needing to investigate performance issues with their applications. Note that this modality of analysis may be more time-consuming because trace data may be detailed and require significant storage space.

Implementations of the monitoring service disclosed herein ingest and aggregate the span information from the online retailer's application (for APM) and from a user's interactions with a browser or other interface on the frontend (for RUM). Further, implementations of the monitoring service disclosed herein extract information from the incoming span data and store the information using multiple formats to support multiple modalities of data analysis for a client. Each modality is configured to allow the clients access to a different format in which incoming trace information may be represented and stored, where each format conveys a different degree of resolution regarding the ingested traces to a client and, accordingly, may occupy a different amount of storage space.

FIG. 5 is a flow diagram that illustrates an exemplary method of ingesting and aggregating span information to support multiple modalities of analysis for APM, in accordance with implementations of the monitoring service disclosed herein. As mentioned in connection with FIG. 3, span

information is received at the monitoring service **306** from the collector (e.g., the collector **504** in FIG. 5).

As noted previously, in one implementation, incoming spans from one or more clients are converted into a plurality of metric data streams prior to consolidating the spans into traces through a sessionization process. The incoming spans are received and the metric data streams are generated by module **520** prior to the spans being sessionized. Because the metric time series are created without paying a time penalty associated with sessionization, they can be used to perform real-time monitoring and alerting.

The incoming spans for APM (e.g., monitoring microservices at the backend of an application) are also sessionized where the span information is combined into traces in a process called sessionization. The APM sessionization module **506** is responsible for stitching together or combining the traces **508** using, among other things, the Trace IDs associated with each user-request (and typically also the Parent Span IDs of each span). Note that, in one implementation, the sessionized traces may also be inputted to the module **520** to create metric time series to track traces (separately from the time series created to track spans).

The spans associated with RUM (ingested, for example, from the beacon **567**), are, in one implementation, ingested and analyzed separately from the spans associated with APM. In one implementation, RUM-related spans may need to be treated differently from APM-related spans. For example, the spans related to RUM may need to be ingested and sharded by a session identifier (session ID) (and, optionally, an organization identifier) instead of using the Trace ID. A session ID is an identifier that connects a series of traces. RUM data is typically organized into page views (which show details of a page visit) and sessions (which group all the page views by a user in a single visit). A session ID is typically used to filter for all the views in a specific session.

For RUM, a developer is typically more interested in the behavior of a user over the course of a session, e.g., a user session interacting with a particular website or application. Accordingly, spans associated with RUM are usually sharded and tracked using a session identifier (or session ID). Spans associated with RUM that are received from the collector **504** are, therefore, ingested using a separate RUM ingestion module **588** (details of which will be covered in FIG. 17).

In an implementation, information extracted from the traces **508** may also be transmitted to the RUM ingest module **588** in order to facilitate a connection between the frontend RUM traces and the backend APM traces. In this implementation, a RUM span on the RUM frontend may comprise the associated Trace ID/Span ID of an APM span, so the RUM frontend would initiate the retrieval of the connection information from the APM backend.

In addition to a Trace ID, each trace also comprises a time-stamp; using the time-stamps and the Trace IDs, the APM sessionization module **506**, which is associated with APM-related spans, creates traces **508** from the incoming spans in real time and sessionizes them into discrete time windows. For example, the sessionization process may consolidate traces (from spans) within a first time window (associated with time window Y **580**) before transmitting the traces to modules **520**, **522** or **524**. Each of the modules **520**, **522** and **524** support a different modality of analysis for APM. Thereafter, the sessionization process may consolidate traces within the subsequent time window (associated with time window "Y+M" **585**) before transmitting those traces to the modules **520**, **522**, or **524**. It should be noted that the time windows associated with each of the modules

520, **522**, and **524** may be different. In other words, the metric time series data may be collected over short time windows of 10 seconds each. By comparison, traces for the metric events modality (associated with the module **522**) may be collected over 10 minute time windows.

In some implementations of the monitoring service disclosed herein, the sessionization module is able to ingest, process and store all or most of the spans received from the collector **504** in real time. By comparison, conventional monitoring systems do not accept all of the incoming spans or traces; instead, they sample incoming spans (or traces) to calculate SLIs at the root level of a trace before discarding the spans. Implementations of the monitoring service disclosed herein, by comparison, comprise an ingestion streaming pipeline that is able to ingest and consolidate all the incoming spans into traces in real time, and is further able to use advanced compression methods to store the traces. Further, implementations of the monitoring service disclosed herein are able to generate metric time series from the span data (prior to sessionizing the spans) to provide real-time monitoring and alerting of certain KPIs.

As noted above, the APM sessionization module **506** has the ability to collect all the traces within a first time window Y **580** using the time-stamps for the traces. Subsequently, the sessionized traces are fed to the modules **522** and **524**, for the respective modalities (metric events and full-fidelity) for extraction and persistence.

In one implementation, subsequent to consolidation, the trace data is indexed by an optional tag indexing module **507**, which indexes one or more tags in the trace data. The tags may be client-selected tags or tags that the monitoring platform is configured to index by default. In a different implementation, tag indexing may be performed as part of data aggregation, e.g., by module **522** associated with metric events.

In an implementation, data sets associated with each of the modalities may be persisted in one or more databases **555**. As noted previously, the data sets for the respective modalities may be separate data sets, overlapping data sets or a single data set that supports all the modalities. Note that the databases **555** may be a single database that stores data sets corresponding to all three modalities. Alternatively, the databases **555** may represent different respective databases for each of the three modalities. Furthermore, the databases **555** may also represent distributed databases across which relevant information for each of the three modalities is stored.

In one implementation, data associated with each of the three modalities is generated at the time of ingestion and stored separately from each other. The structure, content, type or syntax of query submitted by a client will typically dictate which of the three modalities and corresponding data set will be selected. In one implementation, an interface through which the query is submitted may also determine which of the three modalities and corresponding data set is selected. In an implementation, there may be some commonality in the data for the three modalities in which case the storage for the data may overlap. An alternative implementation may also comprise one or two of the three modalities (instead of all three) described above.

A client may send in a request to retrieve information pertaining to an application through query interface **582**. The underlying querying engine (e.g., the query engine and reporting system **324** from FIG. 3) will analyze the structure, content, type and/or syntax of the query, and also the interface through which the query is submitted, to determine which of the three modalities and respective data set to

access to service the query. In an implementation, the three data sets corresponding to the three modalities are structured in a way that allows the querying engine to navigate between them fluidly. For example, a client may submit a query through the query interface **582**, which may potentially result in the query engine accessing and returning data associated with the metric events modality. Thereafter, if the client requires more in-depth information, the querying engine may seamlessly navigate to data associated with a different modality (e.g., full-fidelity) to provide the client with further details. Conventional monitoring systems, by comparison, do not provide more than a single modality or the ability to navigate between multiple modalities of data analysis.

4.1 Metric Time Series

Implementations of the monitoring service disclosed herein allow trace data associated with APM to be stored and analyzed using multiple modalities of analysis. In one implementation, incoming spans from one or more clients are converted into a plurality of metric data streams (also referred to as metric time series) and transmitted to the analytics engine (e.g., the instrumentation analysis system **322**) for further analysis. Most of the metric data streams are created directly from the incoming spans prior to the sessionization process to generate metric time series related to spans. Each metric time series is associated with a single “span identity,” where a base span identity comprises a tuple of information corresponding to an associated type of span. Each metric time series in the metric time series modality represents a plurality of tuples with each tuple representing a data point. KPIs can be extracted in real-time directly from the metric time series and reported to a client. Because the metric time series are created without paying a time penalty associated with sessionization, they can be used to perform real-time monitoring with sub-second resolution and generate alerts within two to three seconds if some condition is violated.

4.1.1 Generating Metric Data Streams Using Span Identities

A client application associated with, for example, an online retailer’s website may potentially generate millions of spans from which a monitoring platform may need to extract meaningful and structured information. To organize the significant amounts of incoming span data, in an implementation, incoming spans may be automatically grouped by mapping each span to a base “span identity,” wherein a base span identity comprises some key attributes that summarize a type of span. An exemplary span identity may be represented as the following exemplary tuple: {operation, service, kind, isError, httpMethod, isServiceMesh}, where the operation field represents the name of the specific operation within a service that made the call, the service field represents the logical name of the service on which the operation took place, the kind field details relationships between spans and may either be a “server” or “client,” the isError field is a “TRUE/FALSE” flag that indicates whether a span is an error span, the httpMethod field relates to the HTTP method of the request for the associated span and the isServiceMesh field is a flag that indicates whether the span is part of a service mesh. A service mesh is a dedicated infrastructure layer that controls service-to-service communication over a network. Typically, if software has been instrumented to send data from a service mesh, the trace data transmitted therefrom may generate duplicative spans that may need to be filtered out during monitoring. Accordingly, the ‘isSer-

viceMesh' flag allows the analytics engine to filter out any duplicative spans to ensure the accuracy of the metrics computations.

In some implementations, the tuple used to represent the span identity may include other identifying dimensions as well. For example, if a client needs visibility into metadata tags from the spans in addition to the dimensions extracted for a base span identity by default (e.g., service, operation, kind, etc.), an extended identity may be created. An extended identity supports custom dimensionalization by a client, where dimensionalization refers to the ability to extract information pertaining to additional tags or metadata in a span. An extended identity provides a customer the ability to dimensionalize the span using pre-selected dimensions. Conventional methods of monitoring by comparison did not offer customers the flexibility to add custom dimensions to streams of metric data. An extended identity comprises the span's base identity and additionally a map of the span's tag key:value pairs that matched a client's configuration settings. An exemplary extended identity may be represented as the following exemplary tuple: {operation, service, kind, isError, httpMethod, isServiceMesh, keyValueMap . . . }, where the keyValueMap field represents one or more additional tags or dimensions configured by the client to be extracted as part of the span's identity, e.g., 25 customer name, member ID, etc.

By extracting information related to additional tags, higher cardinality metrics may be computed using the metric time series modality. Further, a client is able to configure alerts on the custom dimensions as well, wherein the alerts inform a client if a particular dimension has crossed some critical threshold. In alternate implementations of the monitoring service disclosed herein, the tuple used to represent a span's base or extended identity may contain fewer elements.

If the tuple of information of an incoming span happens to be the same as another span, both spans relate to the same identity. In an implementation, spans with the same base identity may be grouped together. A fixed size bin histogram is generated for each span identity to track metrics associated with the span identity. In this way the same type of spans are organized together and the client can track one or more metrics associated with each group of spans sharing a common identity. In an implementation, a fixed size bin histogram is generated for each unique span identity. The fixed size bin histogram may be a data structure, for example, that is preserved in memory.

As noted above, each span identity may be tracked with a respective histogram. The histograms associated with the corresponding span identities, in one implementation, are generated and updated in fixed time duration windows. For example, histogram data may be generated for the incoming spans in memory every 10 seconds. At the end of each fixed duration, metrics associated with the histograms are emitted and the histogram is reset for the next time window. By emitting metrics for each time duration, data streams of metrics may be generated from the histogram data. The streams of metric data associated with each span identity, in one implementation, may be aggregated by a monitoring platform to provide a client of the monitoring platform meaningful information regarding the application being monitored.

FIG. 6 illustrates the manner in which span metrics and trace metrics are automatically generated, in accordance with implementations of the monitoring service disclosed herein. FIG. 6 illustrates 5 unique spans (A-E) including a root span (an initiating span) A. In an implementation, each

group of spans identified by the same span identity is associated with one or more span metrics 650. For example, a minimum span duration 630, a median span duration 631, a maximum span duration 632, a p90 latency value 633, a 5 p99 latency value 634 and a span count (how many times a particular identity was counted) may be tracked for each span identity. A histogram corresponding to the span identity may track these metrics over fixed sized durations, e.g., 10 10 seconds. For example, over a 10 second window, the histogram may comprise fixed size bins that track a minimum span duration, a median span duration, a maximum span duration, a p90 value, a p99 value and a count of all spans received corresponding to a given identity. At the end of each duration, the metrics are emitted and the histogram is 15 reset. The emitted metrics are used to generate streams of metrics data corresponding to each span identity. Each data point on a metric data stream comprises the span identity dimensions or the extended identity dimensions if the client has configured additional metadata to be extracted from the 20 spans.

As shown in FIG. 6, in an implementation, the initiating span A comprises a trace identity that is used to emit trace metrics 640. The initiating span A helps define an identity for a trace which allows the monitoring platform to logically group together all traces that represent the same flow through an endpoint of the application. The duration of the identity for a trace is calculated as the end time of the latest span in the trace minus the start time of its initiating span. An exemplary trace identity may be represented as the 30 following exemplary tuple: {operation, service, isError, httpMethod, isServiceMesh}, where the operation field represents the name of the specific operation within a service that made the call, the service field represents the logical name of the service on which the operation took place, the isError field is a "TRUE/FALSE" flag that indicates whether the trace is associated with an error, the httpMethod field relates to the HTTP method of the request for the associated trace and the isServiceMesh field is a flag that indicates whether the trace is part of a service mesh. The trace metrics 35 40 45 50 55 60 65 70 75 80 85 90 95 100 105 110 115 120 125 130 135 140 145 150 155 160 165 170 175 180 185 190 195 200 205 210 215 220 225 230 235 240 245 250 255 260 265 270 275 280 285 290 295 300 305 310 315 320 325 330 335 340 345 350 355 360 365 370 375 380 385 390 395 400 405 410 415 420 425 430 435 440 445 450 455 460 465 470 475 480 485 490 495 500 505 510 515 520 525 530 535 540 545 550 555 560 565 570 575 580 585 590 595 600 605 610 615 620 625 630 635 640 645 650 655 660 665 670 675 680 685 690 695 700 705 710 715 720 725 730 735 740 745 750 755 760 765 770 775 780 785 790 795 800 805 810 815 820 825 830 835 840 845 850 855 860 865 870 875 880 885 890 895 900 905 910 915 920 925 930 935 940 945 950 955 960 965 970 975 980 985 990 995 1000 1005 1010 1015 1020 1025 1030 1035 1040 1045 1050 1055 1060 1065 1070 1075 1080 1085 1090 1095 1100 1105 1110 1115 1120 1125 1130 1135 1140 1145 1150 1155 1160 1165 1170 1175 1180 1185 1190 1195 1200 1205 1210 1215 1220 1225 1230 1235 1240 1245 1250 1255 1260 1265 1270 1275 1280 1285 1290 1295 1300 1305 1310 1315 1320 1325 1330 1335 1340 1345 1350 1355 1360 1365 1370 1375 1380 1385 1390 1395 1400 1405 1410 1415 1420 1425 1430 1435 1440 1445 1450 1455 1460 1465 1470 1475 1480 1485 1490 1495 1500 1505 1510 1515 1520 1525 1530 1535 1540 1545 1550 1555 1560 1565 1570 1575 1580 1585 1590 1595 1600 1605 1610 1615 1620 1625 1630 1635 1640 1645 1650 1655 1660 1665 1670 1675 1680 1685 1690 1695 1700 1705 1710 1715 1720 1725 1730 1735 1740 1745 1750 1755 1760 1765 1770 1775 1780 1785 1790 1795 1800 1805 1810 1815 1820 1825 1830 1835 1840 1845 1850 1855 1860 1865 1870 1875 1880 1885 1890 1895 1900 1905 1910 1915 1920 1925 1930 1935 1940 1945 1950 1955 1960 1965 1970 1975 1980 1985 1990 1995 2000 2005 2010 2015 2020 2025 2030 2035 2040 2045 2050 2055 2060 2065 2070 2075 2080 2085 2090 2095 2100 2105 2110 2115 2120 2125 2130 2135 2140 2145 2150 2155 2160 2165 2170 2175 2180 2185 2190 2195 2200 2205 2210 2215 2220 2225 2230 2235 2240 2245 2250 2255 2260 2265 2270 2275 2280 2285 2290 2295 2300 2305 2310 2315 2320 2325 2330 2335 2340 2345 2350 2355 2360 2365 2370 2375 2380 2385 2390 2395 2400 2405 2410 2415 2420 2425 2430 2435 2440 2445 2450 2455 2460 2465 2470 2475 2480 2485 2490 2495 2500 2505 2510 2515 2520 2525 2530 2535 2540 2545 2550 2555 2560 2565 2570 2575 2580 2585 2590 2595 2600 2605 2610 2615 2620 2625 2630 2635 2640 2645 2650 2655 2660 2665 2670 2675 2680 2685 2690 2695 2700 2705 2710 2715 2720 2725 2730 2735 2740 2745 2750 2755 2760 2765 2770 2775 2780 2785 2790 2795 2800 2805 2810 2815 2820 2825 2830 2835 2840 2845 2850 2855 2860 2865 2870 2875 2880 2885 2890 2895 2900 2905 2910 2915 2920 2925 2930 2935 2940 2945 2950 2955 2960 2965 2970 2975 2980 2985 2990 2995 3000 3005 3010 3015 3020 3025 3030 3035 3040 3045 3050 3055 3060 3065 3070 3075 3080 3085 3090 3095 3100 3105 3110 3115 3120 3125 3130 3135 3140 3145 3150 3155 3160 3165 3170 3175 3180 3185 3190 3195 3200 3205 3210 3215 3220 3225 3230 3235 3240 3245 3250 3255 3260 3265 3270 3275 3280 3285 3290 3295 3300 3305 3310 3315 3320 3325 3330 3335 3340 3345 3350 3355 3360 3365 3370 3375 3380 3385 3390 3395 3400 3405 3410 3415 3420 3425 3430 3435 3440 3445 3450 3455 3460 3465 3470 3475 3480 3485 3490 3495 3500 3505 3510 3515 3520 3525 3530 3535 3540 3545 3550 3555 3560 3565 3570 3575 3580 3585 3590 3595 3600 3605 3610 3615 3620 3625 3630 3635 3640 3645 3650 3655 3660 3665 3670 3675 3680 3685 3690 3695 3700 3705 3710 3715 3720 3725 3730 3735 3740 3745 3750 3755 3760 3765 3770 3775 3780 3785 3790 3795 3800 3805 3810 3815 3820 3825 3830 3835 3840 3845 3850 3855 3860 3865 3870 3875 3880 3885 3890 3895 3900 3905 3910 3915 3920 3925 3930 3935 3940 3945 3950 3955 3960 3965 3970 3975 3980 3985 3990 3995 4000 4005 4010 4015 4020 4025 4030 4035 4040 4045 4050 4055 4060 4065 4070 4075 4080 4085 4090 4095 4100 4105 4110 4115 4120 4125 4130 4135 4140 4145 4150 4155 4160 4165 4170 4175 4180 4185 4190 4195 4200 4205 4210 4215 4220 4225 4230 4235 4240 4245 4250 4255 4260 4265 4270 4275 4280 4285 4290 4295 4300 4305 4310 4315 4320 4325 4330 4335 4340 4345 4350 4355 4360 4365 4370 4375 4380 4385 4390 4395 4400 4405 4410 4415 4420 4425 4430 4435 4440 4445 4450 4455 4460 4465 4470 4475 4480 4485 4490 4495 4500 4505 4510 4515 4520 4525 4530 4535 4540 4545 4550 4555 4560 4565 4570 4575 4580 4585 4590 4595 4600 4605 4610 4615 4620 4625 4630 4635 4640 4645 4650 4655 4660 4665 4670 4675 4680 4685 4690 4695 4700 4705 4710 4715 4720 4725 4730 4735 4740 4745 4750 4755 4760 4765 4770 4775 4780 4785 4790 4795 4800 4805 4810 4815 4820 4825 4830 4835 4840 4845 4850 4855 4860 4865 4870 4875 4880 4885 4890 4895 4900 4905 4910 4915 4920 4925 4930 4935 4940 4945 4950 4955 4960 4965 4970 4975 4980 4985 4990 4995 5000 5005 5010 5015 5020 5025 5030 5035 5040 5045 5050 5055 5060 5065 5070 5075 5080 5085 5090 5095 5100 5105 5110 5115 5120 5125 5130 5135 5140 5145 5150 5155 5160 5165 5170 5175 5180 5185 5190 5195 5200 5205 5210 5215 5220 5225 5230 5235 5240 5245 5250 5255 5260 5265 5270 5275 5280 5285 5290 5295 5300 5305 5310 5315 5320 5325 5330 5335 5340 5345 5350 5355 5360 5365 5370 5375 5380 5385 5390 5395 5400 5405 5410 5415 5420 5425 5430 5435 5440 5445 5450 5455 5460 5465 5470 5475 5480 5485 5490 5495 5500 5505 5510 5515 5520 5525 5530 5535 5540 5545 5550 5555 5560 5565 5570 5575 5580 5585 5590 5595 5600 5605 5610 5615 5620 5625 5630 5635 5640 5645 5650 5655 5660 5665 5670 5675 5680 5685 5690 5695 5700 5705 5710 5715 5720 5725 5730 5735 5740 5745 5750 5755 5760 5765 5770 5775 5780 5785 5790 5795 5800 5805 5810 5815 5820 5825 5830 5835 5840 5845 5850 5855 5860 5865 5870 5875 5880 5885 5890 5895 5900 5905 5910 5915 5920 5925 5930 5935 5940 5945 5950 5955 5960 5965 5970 5975 5980 5985 5990 5995 6000 6005 6010 6015 6020 6025 6030 6035 6040 6045 6050 6055 6060 6065 6070 6075 6080 6085 6090 6095 6100 6105 6110 6115 6120 6125 6130 6135 6140 6145 6150 6155 6160 6165 6170 6175 6180 6185 6190 6195 6200 6205 6210 6215 6220 6225 6230 6235 6240 6245 6250 6255 6260 6265 6270 6275 6280 6285 6290 6295 6300 6305 6310 6315 6320 6325 6330 6335 6340 6345 6350 6355 6360 6365 6370 6375 6380 6385 6390 6395 6400 6405 6410 6415 6420 6425 6430 6435 6440 6445 6450 6455 6460 6465 6470 6475

seen in FIG. 6), including client-configured custom metrics, corresponding to each type of span to the analytics engine.

In one implementation, the span information is also combined into traces **708** using a sessionization module **707** as discussed in connection with FIG. 5. The sessionization process may consolidate traces (from spans) within a first minute window (associated with time window Y **780**). Thereafter, the sessionization process may consolidate traces within the subsequent window (associated with time window "Y+M" **785**). Trace identities are determined for the sessionized traces **708** using module **742** after which the trace metrics (as discussed in connection with FIG. 6) are determined using the histogram generation module **722** in a process similar to the manner in which span metrics are generated.

In an implementation, an aggregation module **724** may aggregate the periodic metric data from the histogram generation module **722** and create metric time series from the data for each span identity. In some implementations, the aggregation module **724** may generate quantized data streams from the metric data received from the histogram generation module **722**. The quantized data stream has data values occurring periodically at fixed time intervals.

In one implementation, the aggregation module **724** may identify a function for aggregating the metric for which values are provided by one or more input data streams. The aggregation module **724** generates the quantized data streams by determining an aggregate value for each input data stream for each fixed time interval by applying the identified function over data values of the input data stream received within the fixed time interval. The aggregation module **724** may further receive a request to evaluate an expression based on the data values from the input data streams. The system periodically evaluates the expression using the data values of the quantized data streams.

In one implementation, the aggregation module **724** may, for example, perform aggregations on the various metric time series to provide real-time monitoring of certain higher priority endpoints in the application. For example, aggregations may be performed to determine request, error and latency metrics for certain designated services. In order to do that, the aggregation module **724** may, for example, aggregate values across all span identities that are associated with the designated service.

Further, in some implementations, alerting module **782** may monitor one or more metric time series from the aggregation module **724** and may be configured to generate alerts if certain metrics being monitored exhibit anomalous behavior. For example, if a maximum span duration associated with a given span identity crosses over a certain threshold, an alert configured using the alerting module **782** may be triggered. The alert may, for example, be responsive to a metric time series associated with span metric **632** from FIG. 6, wherein the alert is triggered if the maximum span duration exceeds a given threshold.

In one implementation, the histograms generated by the histogram generation module **722** may be stored in database **777**. In an implementation, the histogram data may be stored as parquet-formatted files.

4.1.2 Real-Time Monitoring Using Metric Time Series Data

FIG. 8 illustrates an exemplary on-screen GUI for APM illustrating a monitoring mode for an application displaying metric values aggregated from metric time series data, in accordance with implementations of the monitoring service disclosed herein. In one implementation, the GUI of FIG. 8 displays a monitoring mode indication when a correspond-

ing monitoring mode option **802** is selected. The monitoring mode displays a panel **888** listing services **804** comprised within the application being monitored. Each service is displayed alongside metrics pertaining to requests/second **806**, error rate **812** and P90 latency values **810**. The metrics data displayed in the panel **888** is computed in real-time and is aggregated using the metric time series data. In an implementation, an aggregation module similar to the aggregation module **724** discussed in connection with FIG. 7 performs the necessary aggregations from the various metric time series to display metrics associated with each of the services. The service level KPIs may be computed through the real-time aggregation pipeline discussed in connection with FIG. 7 before the histogram metadata is stored in the backend of the analytics engine.

The monitoring mode also comprises an application topology graph **830**. An application topology graph (or service graph) typically decomposes an application into all its component services and draws the observed dependencies between the services so a client can identify potential bottlenecks and get a better understanding of the manner in which data flows through the software architecture. The service graph **830** also facilitates visualizing cross-service relationships between services comprised within the application and external to the application (as will be discussed further in connection with the metric events modality). In an implementation, the service graph may be created using information gleaned from the metric time series data aggregated by the aggregation module **724** discussed in connection with FIG. 7.

By ingesting up to 100% of the incoming spans from the client software and implementing monitoring service **306** as a Software as a Service (SaaS) based service offering, implementations of the monitoring service disclosed herein advantageously retain valuable information pertaining to the spans that is further analyzed in the SaaS backend. Span identities and histogram information (e.g., various counts and metrics data) associated with the incoming spans that are stored may be used to conduct further analysis. For example, metadata may be analyzed to identify certain offending services or operations, and data regarding those services or operations may be surfaced for further analysis.

Conventional monitoring systems typically expunged the span data after extracting the relevant metrics from them. By comparison, implementations of the monitoring service disclosed herein retain high-fidelity information related to all the incoming spans for deeper analysis. The metadata retained provides a client the ability to filter based on certain dimensions and services that would not have been possible using conventional monitoring systems. Further, the metadata retained may be used in conjunction with data sets for other modalities such as metric events and full-fidelity to allow a client to provide a thorough investigation of an alert.

In one implementation, using, for example, the "service," "operation," and "kind" fields in the tuple, the aggregation module **724** (from FIG. 7) may be able to determine span identities associated with cross-service calls. Spans associated with inter-service calls are of interest to a client because they provide client information regarding the manner in which two services within an application are interacting. Implementations of the monitoring service disclosed herein are able to advantageously use the metadata saved for the metric time series to perform post-processing and determine services associated with inter-services calls. For example, the value of the "kind" field related to a span identity may be either "client" or "server" where the analytics engine may

be able to use that information in post-processing to determine if the span is related to a cross-service call.

If it is determined that a particular span is related to a cross-service call, those spans could be processed through the analytics engine to discover further information regarding the dependencies. For example, in one implementation, if a client identifies a span identity associated with a cross-service call or a span identity associated with a high value operation, the client may create an extended identity for the corresponding span identities and supplement those identities with additional custom dimensions to be monitored. For example, the client may want to monitor a customer name association with such spans. The client may simply reconfigure the analytics engine to extract the additional customer name dimension as part of the spans' extended identity.

Retaining span information associated with incoming spans provides a client additional metadata to perform intelligent processing. In an implementation, the client may only collect data pertaining to select operations. In other words, the client may filter out data pertaining to select operations that are of less interest to a client.

The number of unique span identities may typically roughly correlate with the number of unique operation names present on the span. In an implementation, the client is able to turn off or filter out span identities associated with certain operations if they are not particularly useful. In other words, the monitoring platform can be configured to turn off metric generation related to selected span identities. This advantageously reduces loads on the metrics analytics engine because it does not need to track and store metric time series for spans that are of little interest to a client. For example, spans associated with calls that a service makes to operations internal to the service do not convey information and can be filtered. Accordingly, additional resources can be directed towards processing spans associated with services and operations that are of greater interest to a client. Conventional monitoring systems by comparison would not have the flexibility to selectively focus on spans associated with high value services or operations by filtering out the less valuable spans.

4.2 Metric Event Modality

The metric event modality generates and stores aggregated rows of metrics values for selected indexed tags from the incoming trace data for given time durations. The selected tags may, for example, be indexed from the incoming spans when the spans are ingested. Metrics data may, for example, comprise, but is not limited to, number of requests (e.g., between two services), number of errors and latency. The aggregated rows of metrics data are stored efficiently for fast aggregation. The metric events data may be rapidly vectorized and aggregated in response to queries from a client.

Implementations of the monitoring service disclosed herein use the aggregated rows of metrics data created in association with the metric events modality to generate a full-context application topology graph using the metric events data (e.g., by module 522 in FIG. 5). As noted above, an application topology graph (or service graph) typically decomposes an application into all its component services and draws the observed dependencies between the services so a client can identify potential bottlenecks and get a better understanding of the manner in which data flows through the software architecture. FIG. 9 illustrates an exemplary on-screen GUI comprising an interactive topology graph for an application created from the aggregated metric events data, in accordance with implementations of the monitoring service disclosed herein. The service graph facilitates visual-

izing cross-service relationships between services comprised within the application and external to the application. The exemplary GUI of FIG. 9 also enables customers to track the causal chain of operations resulting in an error.

It should be noted that the service graph may also be generated using the metric time series data as noted earlier, however, storage for the metric events data set may be significantly less because it does not need to store as much metadata as metric time series data. Accordingly, generating the service graph using metric events data is more efficient from a storage standpoint.

FIG. 9 illustrates an on-screen GUI comprising an interactive full-context service graph 900, which is constructed for an exemplary microservices-based application using the metrics data generated in connection with the metric events modality. Each circular node (e.g., nodes associated with services 902, 904 and 906 of FIG. 9) represents a single microservice. Alternatively, in an implementation, a circular node may also represent a group of multiple microservices, where the GUI for the monitoring platform (associated with, for example, the monitoring service 306) provides a client the ability to expand the node into its sub-components.

In an implementation, services that are part of the client's application may be represented differently from services that are external to the client's application. For example, circular nodes (e.g., nodes associated with services 902, 904 and 906) of the exemplary application represented by service graph 900 are associated with services comprised within the client's application. By contrast, squarish nodes (e.g., nodes associated with databases dynamodb 915, Cassandra 920, ad-redis 912) are associated with services or databases that are external to the client's application.

A user may submit a request at the front-end service 902; the user's request at the front-end service 902 may set off a chain of subsequent calls. For example, a request entered by the user at the front end of the platform may generate a call from the front-end service 902 to the recommendation service 904, which in turn may generate a further call to the product catalog service 906. As noted previously, a chain of calls to service a request may also comprise calls that a service makes to internal sub-functions or operations within the same service.

Each edge in the service graph 900 (e.g., the edges 922, 924 and 926) represents a cross-service dependency (or a cross-service call). The front-end service 902 depends on the recommendation service 904 because it calls the recommendation service 904. Similarly, the recommendation service 904 depends on the product catalog service 906 because it makes a call to the product catalog service 906. The directionality of the edge represents a dependency of a calling node on the node that is being called. Each of the calls passes the Trace ID for the request to the respective service being called. Further, each service called in the course of serving the request could potentially generate several spans (associated with calls to itself or other services). Each of the spans generated will then carry the Trace ID associated with the request, thereby, propagating the context for the trace. Spans with the same Trace ID are, thereafter, grouped together to compose a trace.

In some implementations, the GUI comprising service graph 900 may be configured so that the nodes themselves provide a visual indication regarding the number of errors that originated at a particular node versus errors that propagated through the particular node but originated elsewhere. In an implementation, the high-cardinality metrics data aggregated in association with the metric events modality

may be used to compute the number of errors that are used to render the nodes of the service graph.

For example, as shown in the service graph of FIG. 9, the front-end service 902 makes calls to the recommendation service 904. Errors may be generated at the recommendation service 904 not only in response to calls from the front-end service 902, but also in response to calls that the recommendation service 904 makes to itself (e.g., in response to sub-functions or operations that are part of recommendation service). For such errors, the recommendation service 904 would be considered the “originator” for the error. The recommendation service 904 also makes calls to the product catalog service 906 and these calls may result in their own set of errors for which the product catalog service 906 would be considered the error originator. The errors originating at the product catalog service 906 may propagate upstream to the front-end service 902 through the recommendation service 904; these errors would be observed at the recommendation service 904 even though the recommendation service 904 is not the originator of those errors.

It is appreciated that conventional monitoring technologies would not provide adequate means for a client to distinguish between errors that originated at the recommendation service 904 versus errors that propagated through the recommendation service 904 but originated elsewhere. By performing computations using the metrics data associated with the metric events modality, implementations of the monitoring service disclosed herein are able to render a service graph that visually indicates critical information regarding the services in an architecture, e.g., number of requests between services, the number of errors generated by a service, number of errors for which the service was the root cause, etc. The service graph 900 allows clients the ability to visually distinguish between errors that originated at the recommendation service 904 as compared with errors that simply propagated through the recommendation service 904. As shown in FIG. 9, the node associated the recommendation service 904 comprises a solid-filled circular region 966 and a partially-filled region 962, where the region 966 represents errors that originated at the recommendation service 904 while the region 962 represents errors that propagated through the recommendation service 904 but originated elsewhere (e.g., at the product catalog service 906).

Similarly, solid-filled region 960 within the node associated with the product catalog service 906 represents the errors that originated at the product catalog service. Note that the errors returned by the product catalog service 906 originated at the product catalog service. In other words, the product catalog service 906 does not have errors from another downstream service propagating through it because it does not make calls to another service that is further downstream in the execution pipeline. Conversely, the front-end service 902 comprises a partially-filled region 964 because the errors observed at the front-end service 902 propagated to it from other downstream services (e.g., the recommendation service 904, the currency service 930, the product catalog service 906, etc.) The front-end service 902 was not the originator of errors in the example shown in FIG. 9. Note that in other implementations solid-filled regions (e.g., region 966) and partially-filled regions (e.g., region 964) may be represented differently. For example, different shades, patterns, or colors may be used to distinguish these regions from each other.

Implementations of the monitoring service disclosed herein use the aggregated rows of metrics data created for the metric events modality to determine full-fidelity SLIs

associated with the services in an application (e.g., by the module 522 in FIG. 5). An SLI is a service level indicator—a defined quantitative measure of some aspect of the level of service that is provided. The SLIs are aggregated and extracted for the various services in a microservices architecture so that the behavior of applications may be understood. Most clients consider request latency—how long it takes to return a response to a request—as a key SLI. Other common SLIs include the error rate (often expressed as a fraction of all requests received) and system throughput, typically measured in requests per second. The measurements are often aggregated over a measurement window using the metrics data associated with the metric events modality and then turned into a rate, average, or percentile.

In one implementation, the GUI comprising service graph 900 is interactive, thereby, allowing a developer to access the SLIs associated with the various nodes and edges within the application by interacting with respective portions of the service graph. Referring to FIG. 9, in an implementation, a client may be able to hover their cursor over various regions of the on-screen displayed service graph 900, including but not limited to the nodes (e.g., the nodes associated with services 904, 906 etc.) and edges (e.g., the edges 922, 926, etc.), to receive SLI-related information for the associated microservices through a pop-up window or other interface.

FIG. 10 illustrates an exemplary on-screen displayed GUI showing the manner in which a client may access SLIs pertaining to a service within an interactive topology graph, in accordance with implementations of the monitoring service disclosed herein. As shown in FIG. 10, when a client hovers the cursor over the node associated with, for example, the recommendation service 1006, a pop-up window 1008 is overlaid on the service graph 1000 comprising SLIs pertaining to the recommendation service 1006. Specifically, SLIs pertaining to Requests 1010, Errors 1012 and Latency percentiles 1014 are provided. Furthermore, in an implementation, information pertaining to Root Cause 1016 is also provided to the client.

For example, the SLIs related to Requests 1010 comprise information regarding the rate of requests and number of requests serviced by the recommendation service 1006 during a specific time duration. The time duration over which the SLIs are calculated may be adjusted using drop-down menu 1022. The time duration over which SLIs are calculated may vary, for example, from 1 minute to 3 days. As indicated by the time axis on hover chart 1028, for this example, a time window of 30 minutes (from 9:09 to 9:39 a.m.) is selected.

In an implementation, the pop-up window 1008 also provides the client information pertaining to SLIs related to Errors 1012. In the example of FIG. 10, the pop-up window 1008 provides information regarding the error rate and the total number of errors that occurred during the specified time duration. The client is also provided information regarding what percentage of the total number of requests resulted in errors.

In an implementation, the pop-up window 1008 also provides the client information pertaining to Latency Percentiles 1014 and a graphical representation 1020 of the same. For example, SLI p95 indicates that for 95% of the users, the latency for servicing the requests was less than 467 ms. Latency-related SLIs also include information regarding p90 and p50 percentiles. The graphical representation 1020, in the example of FIG. 10, shows the latency information regarding the p95 percentile graphically.

In one implementation of the monitoring service disclosed herein, the pop-up window 1008 also displays information

pertaining to errors for which the selected service was the root-cause. The Root Cause information **1016** includes the number of errors for which the selected service (e.g., the recommendation service **1006** in the example of FIG. 10) was the originator, the associated error rate and the percentage of the total number of requests that represents. In this way, implementations of the monitoring service disclosed herein, in addition to providing clients visual cues for identifying root cause error originators, are also able to provide meaningful and accurate quantitative information to help clients distinguish between root cause-related errors and errors associated with downstream causes.

Note that the SLIs displayed in the pop-up window **1008** are computed accurately using the metrics data gathered for the metric events modality. Because implementations of the monitoring service disclosed herein are able to ingest up to 100% the incoming span data (without sampling), the SLIs are computed factoring in all the incoming data, which results in accurate measurements. For the example of FIG. 10, there were a total of 2.6 million requests served by the recommendation service **1006** at a rate of 1445.23 requests/second (“sec”). Of these, 1.2 million of the requests resulted in errors at a rate of 714.83/sec, which represents approximately 49% of the total number of requests. In this way, implementations of the monitoring service disclosed herein provide a modality of analysis that enables a client to gather critical SLIs pertaining to the recommendation service **1006** including an indication of how many of the errors originated at the recommendation service **1006**.

FIG. 11 illustrates an exemplary on-screen GUI showing the manner in which a client may access SLIs pertaining to an edge within an interactive topology graph, in accordance with implementations of the monitoring service disclosed herein. The SLIs pertaining to edges are also computed using the metrics data associated with the metric events modality. As shown in FIG. 11, if a client hovers over or selects a particular edge, e.g., the edge **924** (as shown in FIG. 9) (which represents the cross-service dependency of the front-end service **902** on the product catalog service **906**) a pop-up dialog box **1108** opens up on-screen that reports SLIs specific to the dependency. The “From” field **1112** represents the service that executes the call and the “To” field **1114** represents the service that is called (the service that the calling service depends on). As shown in the dialog box **1108**, SLIs pertaining to the number of requests (or calls) that were made, the number of those that returned in errors, and the latency associated with servicing the requests are provided. It should be noted that a latency value **1120** of 49 ms shown in FIG. 11 for this particular dependency may be annotated directly on the edge of the service graph. For example, as shown in service graph **900** of FIG. 9, edge **924** of the service graph **900** in FIG. 9 indicates the latency value **970** (e.g., 49 ms) directly on the edge in the service graph allowing a client to efficiently gather information regarding latency associated with the dependency.

In an implementation, as shown in FIG. 9, the edges within the application topology graph are annotated with their corresponding latency values. In this way, implementations of the monitoring service disclosed herein efficiently compute SLI data from the metrics information aggregated for this modality and advantageously enable developers to gather meaningful and accurate information regarding cross-service dependencies directly from the service graph **900**.

In one implementation, the metrics data associated with the metric events modality are used to compute accurate SLIs across multiple dimensions. Further, implementations of the monitoring service disclosed herein support high

dimensionality and high cardinality tags for the metric events modality. In one implementation, the GUI of FIG. 10 may display one or more attribute (or tag) categories that comprise dimensions that may be varied across the service graph **1000**. In other words, the metrics data and the service graph may both be scoped by one of the various dimensions and also a time-range, which is helpful for keeping track of an architecture that is evolving.

For example, attribute categories (e.g., Workflow **1030**, environment **1032**, incident **1034** and tenant-level **1036**) may be depicted within the GUI, each of which may correspond to attributes that may be varied to compute SLIs and error-related information for different combinations of the attributes. The categories of dimensions across which the SLIs may be computed, include, but are not limited to, workflow **1030**, environment **1032**, incident **1034** and tenant-level **1036**. Each of the categories comprises a drop-down menu with options for the different dimensions. Using the drop-downs to select a different scope may result in a re-drawing of the service graph or a re-calculation of the metrics data to correspond with the selected scope. The metrics events data allows clients to easily and rapidly compute measurements across various cross-combinations of tags or attributes.

In an implementation, the GUI may include a panel **1050** that may display SLIs across the various workflows. Further, the GUI allows clients the ability to break down the workflows across multiple different attributes using drop down menu **1051**. The computations for each of the break-downs may be efficiently determined using the metrics data aggregated for the metric events modality.

FIG. 12 illustrates on-screen displays that represent exemplary categories of dimensions across which SLIs may be computed, in accordance with implementations of the monitoring service disclosed herein. The exemplary category of dimensions corresponds to the categories associated with drop-down menus (e.g., **1030**, **1032**, **1034** and **1036**) discussed in connection with FIG. 10. The metrics data aggregated using the metric event modality allows clients to easily and rapidly compute measurements across various cross-combinations of attributes. As noted above, using the drop-downs to select a different scope may result in a re-drawing of the service graph or a re-calculation of the metrics data to correspond with the selected scope. Drop-down on-screen menu **1230**, for example, corresponding to workflow, illustrates different workflows specific to the application discussed in connection with FIG. 9. A “workflow” is a type of category of dimension of the request that was processed. A workflow may be conceptualized as a type of “global tag” that is attributed to each span in a given trace. A workflow may, for example, be associated with a type of client process, e.g., “checkout,” that is generated on the back-end in response to a request.

Similarly, drop down on-screen menus **1234**, **1236** and **1232**, relating to incident, tenant-level and environment respectively, provide further categories of dimensions across which SLIs may be computed. Each of the drop down on-screen menus **1230**, **1232**, **1234** and **1236** comprises various dimensions (associated with the respective categories) across which aggregations may be scoped. For example, the client may submit a query asking for the number of requests in a trace where “Workflow=frontend:/cart” and “incident=instance errors” and “tenant-level=gold.” By aggregating metrics data associated with the indexed tags, the metric events modality is able to respond to the client’s query rapidly and efficiently.

Note that SLIs may be computed for each attribute of the categories in FIG. 12 and also for each combination of attributes associated with the categories. In an implementation, for each combination of attributes selected using one or more of the drop-down menus, the client may be able to determine the computed SLIs (e.g., by hovering a cursor over the various nodes and edges of the graph after the dimensions have been selected using, for example, the drop-down menus shown in FIG. 10). In this way, implementations of the monitoring service disclosed herein enable a client to use the metric events modality to slice the application topology graph across several different attributes.

It should be noted that clients might have different attributes or dimensions that may be of interest for their respective application. In an implementation, the monitoring platform may be configured to provide insight into client-specific dimensions. Consequently, the specific attributes or dimensions available in each of the drop-down menus may vary by client.

For APM, in certain instances, a user may want to monitor the interactions between services in a microservices architecture on the backend that are related to a particular user-interaction or to a particular client process. An example of a client process may be a checkout process on a website of an online retailer. Conventional tracing and monitoring tools do not provide users the ability to effectively isolate and monitor a group of services on the backend of a microservices architecture that is associated with a particular client process, e.g., pertaining to a checkout process on an online retailer's website.

Implementations of the monitoring platform disclosed herein (e.g., monitoring service 306) allow users to monitor a path or sequence of events that occur on the backend of a distributed application in response to a particular user-interaction or a client process. For example, a user may need to monitor a chain of calls and associated services that are invoked on the backend of the application in response to a user electing to conduct a checkout transaction on a website for an online retailer. Based on metadata extracted from the tags of one or more spans ingested into the monitoring platform, implementations of the monitoring service disclosed herein extract a unique workflow dimension from spans and traces associated with a particular user-interaction or a client process.

Extracting a workflow name from traces associated with a particular user-interaction or a client process advantageously allows implementations of the monitoring service disclosed herein to track metrics associated with the respective user-interaction or process more efficiently. Further, implementations of the monitoring service disclosed herein are also able to construct a topology graph that facilitates visualizing cross-service relationships between services associated with the workflow. Implementations of the monitoring service disclosed herein also extract a unique workflow identifier for traces within a workflow that are associated with a discrete user-interaction over time. For user-interactions that span multiple traces, implementations of the monitoring service disclosed herein are able to use the workflow identifier to reconstruct a client-side view for a user of the monitoring platform based upon the shared workflow identifier between the traces.

As noted earlier, software developed by a client (e.g., clients A and B in FIG. 3) may be instrumented in order to monitor different aspects of the software. Instrumentation results in each span being annotated with one or more tags that provide context about the execution, e.g., a client

process related to the span. The workflow dimension is one of the attributes (or tags) that a span may be annotated with through instrumentation to provide contextual information regarding the client process to which the span relates. For example, if a user clicks the checkout button on the website of an online retailer, the user-interaction may generate spans that are tagged with the workflow attribute where the attribute comprises a value of "checkout." Any metrics extracted for the online retailer's application can then be filtered to compute metrics associated specifically with the checkout workflow (as also discussed in connection with FIGS. 10 and 12).

A workflow for APM is, therefore, a dimension that may be extracted from spans or traces associated with a user-interaction or a client process. In an implementation, the workflow dimension may be included within a set of tags (e.g., global tags) attributed to each span in a given trace. In an implementation, the workflow tag is included in each span within a trace associated with a user-interaction or process. A workflow may, for example, be associated with a type of client process, e.g., "checkout," "update cart," "add to cart," etc. that originates in response to a user request. Attributing the workflow dimension to traces (e.g., as part of the instrumentation process) allows the monitoring platform to create logical groupings of services involved in a particular client process. The workflow dimension therefore allows ingested traces to be grouped based on the respective value of the workflow dimension and allows for metrics to be calculated for each type of workflow (e.g., each value associated with the workflow dimension).

Typically, for APM, each trace is associated with a single workflow. As noted previously, a trace generally refers to a record of the manner in which a single user request, also referred to as a transaction, propagates from one service to the next in a distributed application. The spans resulting from one user request may be consolidated into a trace in the backend of the monitoring platform where the trace correlates with a single workflow.

In an implementation, the workflow dimension may be instrumented as a global tag, where a workflow dimension associated with one span in a trace is attributed to all the other spans in the trace. For example, if one span in a trace is associated with a value of "checkout" for the workflow attribute, other spans in the trace would also be assigned the same value.

In an implementation, for APM, the monitoring service 306 (referenced in FIG. 3) will extract the tags related to the workflow dimension and their respective values from the ingested spans and will analyze them to determine the paths associated with each respective workflow. The path for a workflow comprises the various services and cross-service dependencies in the backend of the application that are invoked by a particular user-interaction or client process.

In one or more implementations, the monitoring service 306 may be programmed to execute a complex set of rules to extract information pertaining to the workflows and their associated values from the tags of ingested spans. Further, pattern matching may also be employed to extract workflow information from ingested spans. In different implementations, the workflow names or values are explicitly written into the instrumentation by the clients and can be extracted from the spans after they are ingested into the monitoring service 306.

FIG. 12B illustrates an exemplary on-screen GUI showing the manner in which a client may filter a service graph for an application and associated metrics by the workflow dimension, in accordance with implementations of the moni-

toring service disclosed herein. As shown in FIG. 12B, the GUI may display a workflow dimension **1242** in addition to other dimensions **1244** that may be varied across the service graph **1200**. For example, a workflow may be selected from menu **1240** in order to display services and dependencies associated with a selected workflow. The workflow dimension **1242** and other dimensions **1244** may be varied to compute metrics, SLIs and error-related information for different combinations of the attributes over a selected time duration **1249**.

Further, as shown in FIG. 13, the GUI may include a side-panel **1240** that may display metrics across the various workflows (or any other selected dimension). For example, side-panel **1240** may display Request, Error, Latency (RED metrics) related metrics for each of the workflows aggregated over the selected time duration **1249**. Further, the GUI allows users the ability to break down the workflows across multiple different attributes using drop down menu **1245**. The computations for each of the break-downs may be efficiently determined using the metrics data aggregated for the metric events mode.

4.2.1 Metric Events Data Generation and Persistence

FIG. 13 is a flow diagram that illustrates an exemplary method of aggregating metrics data from ingested traces for the metric events modality, in implementations according to the present disclosure. As mentioned previously, span information is received at a monitoring service from a collector **1304**. The span information is then combined into traces **1308** in real time using module **1306** in a process called sessionization as discussed in connection with FIG. 5. The sessionization process may consolidate traces (from spans) within a first time window (associated with time window Y **1380**) before transmitting the traces to the collection module **1320**. Thereafter, the sessionization process may consolidate traces within the subsequent window (associated with time window “Y+M” **1385**).

Subsequent to consolidation, the trace data is indexed by tag indexing module **1307**, which indexes one or more tags in the trace data. The tags may be client-selected tags or tags that the monitoring platform is configured to index by default. In one implementation, the metric events modality indexes a subset of tags associated with the spans of a trace, but indexes that set of tags with perfect accuracy because the metrics calculated take into account all the ingested spans.

In one or more implementations, collection module **1320** receives one or more traces **1308** generated within a predetermined time window Y **1380**, and traverses the traces to identify and collect cross-service span pairs that represent cross-service calls. To collect the cross-service span pairs, the collection module **1320** identifies parent-child span pairs in a given trace where the service name for the parent and the child are different. Stated differently, the collection module **1320** will collect each pair of spans that has a parent-child relationship and where each of the two spans in the pair are associated with a different service. The service name of a span may be identified in a span-level tag included with each span. Alternatively, there may be other conventions for identifying a service name associated with a span, e.g., a special field within the span for the service name.

Identifying and collecting the cross-service span pairs from the incoming spans are advantageous because they enable the monitoring platform to track information that will be most relevant to a client, e.g., to render the service graph and display the SLIs associated with the various dependencies between services. Spans associated with calls to internal operations that a service might make may not be of interest

to an application owner and may, therefore, be ignored by the collection module **1320** when determining the cross-service span pairs.

It should be noted that, in one implementation, once the cross-service span pair is identified, indexed tags may be extracted for the cross-service span pair by determining a service tier for the respective parent and child spans of the span pair. A service tier is a subset of spans in a trace that logically identifies a single request to a service. Accordingly, both a parent span and a child span in the cross-service span pair are associated with a respective subset of related spans known as a service tier. Indexed tags are extracted by the collection module **1320** from service tiers associated with a cross-service span pair. In a different implementation, however, the tags may be extracted directly from the parent span and child span in a cross-service span pair rather than the respective service tier associated with the parent span or child span.

In one or more implementations, once the cross-service span pairs are collected and the indexed tags extracted from the respective service tiers, the collection module **1320** maps one or more selected tags for each service in the cross-service span pair to tag attributes, e.g., selected tags in a parent span (associated with the originating service) are mapped to a “FROM” tag attribute and selected tags in a child span (associated with the target service) are mapped to a “TO” tag attribute. This enables directionality information for the cross-service calls to be preserved. It will be appreciated that while the discussion herein focuses on “FROM” and “TO” tag attributes to indicate the direction of the dependency between services in a cross-service call, there may be several different ways to record dependency information between the two services.

In one implementation, the aggregation module **1366** of the monitoring platform aggregates across the cross-service span pairs by maintaining a count for each unique set of “FROM” tag attributes (and their corresponding values) to “TO” tag attributes (and their corresponding values) for a cross-service pair. It should be appreciated that in this implementation, counts are maintained at the tag level for the cross-service span pair (rather than at the service level). Accordingly, a separate count is maintained for each set of parent span tags (mapped to a “FROM” tag attribute) and child span tags (mapped to a “TO” tag attribute) for a cross-service pair. The count is increased each time the aggregation module encounters the same unique set of “FROM” tag attributes (associated with tags of a parent span) and “TO” tag attributes (associated with tags of a child span) for the same cross-service span pair in one or more traces. In a different implementation, the count may be maintained at the service level. Accordingly, the count may be increased each time the same cross-service span pair is encountered within the trace information ingested from the client.

The aggregation module **1322** advantageously prevents duplication by storing a single instance of each unique set of “FROM” tag attributes and “TO” tag attributes for a given cross-service span pair with an associated count in the storage module **1366**. The information in the storage module **1366** may be accessed by querying module **1382** where the querying module **1382** determines that the query is associated with the metric events modality. The querying module **1382** may, for example, be associated with the query engine and reporting system **324** discussed in FIG. 3.

The aggregated cross-service “FROM” and “TO” tag attribute sets and associated count values stored in the storage module **1366** may be used by the querying module

1382 to respond to queries in accordance with the metric events modality. Note that the collection and aggregation process is repeated for subsequent time windows (including window Y+M **1385**) after time window Y **1380**. In this way, the aggregation process is performed over time. This allows the metric events modality to deliver query results over varying time durations (as discussed, for example, in connection with the drop-down menu **1022** in FIG. 10).

FIG. 14 is a table illustrating an exemplary manner in which selected tags for each service in a cross-service span pair may be mapped to tag attributes and stored as part of a memory-resident data object associated with an edge in the service graph, in implementations according to the present disclosure. As noted above, in one or more implementations, once the cross-service span pairs are collected, the monitoring platform maps selected tags associated with each service in the cross-service span pair to tag attributes, e.g., selected tags in a parent span are mapped to a “FROM” tag attribute and selected tags in a child span are mapped to a “TO” tag attribute. The mapping is performed to allow directionality information for the cross-service calls to be preserved. For example, a data object for an “edge” (corresponding to an edge or dependency in the topology graph) may be created that comprises both the FROM-type of tag attributes and the TO-type of tag attributes. In one implementation, one or more edge data objects similar to the one shown in FIG. 14 is used to persist the data for the metric events modality (in addition to node data objects which will be discussed in connection with FIG. 15B).

The table of FIG. 14 illustrates an exemplary manner of storing a data object associated with an edge in the service graph. The table comprises two services, Service A and Service B, in an application. Both Service A and Service B comprise indexed tags “span.kind” and “region.” Tag “span.kind” may have two possible values, “client” and “server.” Similarly, tag “region” may have two possible values, “us-west” and “us-east.”

If all possible combinations exist in Service A, there may be 4 unique tag combinations associated with the “FROM” tag attribute, e.g., {(span.kind=client, region=us-west) (span.kind=client, region=us-east) (span.kind=server, region=us-west) (span.kind=client, region=us-east)}. Similarly, if all possible combinations exist in Service B, there may also be 4 unique tag combinations associated with the “TO” tag attribute. Assuming there is a complete interaction between Service A and Service B, there may be 16 (4×4) different edges between the two services based on the unique set of “FROM” and “TO” type tag attributes.

Note that the example in FIG. 14 illustrates information for two unique sets of “FROM” and “TO” tag attributes. Edge **1490** is associated with a TO-type attribute of “region=us-east” while edge **1492** is associated with a TO-type attribute of “region=us-west.” Because the two sets of “FROM” and “TO” attributes are not identical, a separate count is maintained for each. The edge **1490** has an associated count of 2, while the edge **1492** has an associated count of 1. To determine the total number of requests or total count associated with the cross-service call from Service A to Service B, the number of counts for each set of “FROM” and “TO” tag attributes for an associated cross-service call may be summed up. In the example of FIG. 14 then, a total of 3 requests is computed to occur between Service A and Service B.

In one implementation, data sets for the metric events modality are stored as row of metrics extracted from the indexed tags in the service tiers, where each row is associated with either an edge or a node in the service graph. In an

implementation, the edges on the service graph (e.g., the edges **922** and **926** of FIG. 9) are rendered using both the “FROM” and “TO” tag attribute sets because rendering the edges requires information regarding directionality. The counts for the “FROM” and “TO” tag attribute sets for a given cross-service span pair are summed up to yield the number of requests made between the two services associated with the span pair. In other words, edges are rendered in the service graph by grouping “FROM” and “TO” tag attribute sets associated with a cross-service call and summing up the request counts associated with the cross-service call. In an implementation, this grouping may be performed using “group by” statements in a query language, e.g., SQL. In one implementation, the value of the number of requests between two services may be used to determine the thickness of the edge between the two services in the service graph.

In one implementation, the nodes (e.g., nodes associated with services **902**, **904**, **906**) on the service graph are also rendered using the aggregated cross-service “FROM” and “TO” tag attribute sets. However, rendering the nodes does not require directionality information and, therefore, the nodes may be rendered by collecting and extracting information from the “TO” type tag attributes. Stated differently, the nodes are rendered by grouping the “TO” tag attributes associated with a given service and summing up the request counts associated with the service. In an implementation, this grouping may be performed using “group by” statements in a query language, e.g., SQL. The “TO” tag attributes represent new services being called within the micro-services architecture. Accordingly, the counts associated with “TO” tag attributes for a given service may be summed up to determine the total number of requests made to the service. In one implementation, the value of the number of requests may also be used to determine the size of the node when rendering the service graph.

In an implementation, the “TO” type tag attributes for rendering the nodes may be aggregated separately from the “FROM” and “TO” tag attribute sets aggregated for rendering the edges. In the exemplary table of FIG. 14, information for Service B may be determined, for example, by analyzing the “TO” type tag attributes in the table.

FIG. 15A illustrates an exemplary on-screen GUI showing a visual representation of a portion of an exemplary trace illustrating a cross-service call, in implementations according to the present disclosure. As shown in FIG. 15A, front-end service **1539** makes a call to product catalog service **1538**. Accordingly, the front-end service **1539** and the product catalog service **1538** comprise a cross-service span pair. Note that spans **1540**, **1546** and **1547** may be part of the service tier for front-end service **1539**. Accordingly, even though the call is made by the span **1547** (“frontend:request/GetProduct”) to span **1545** (“productcatalogservice:/GetProducts”), indexed tags associated with the front-end service **1539** may also be extracted from the spans that are part of the service tier for the front-end service **1539**. In one implementation, the first matching tag within a service tier is extracted. For example, indexed tag “environment=prod” **1550** may be extracted from the span **1540**, even though it is repeated in the spans **1546** and **1547** because the span **1540** comprises the first matching instance of the tag **1550**. Assuming tags “environment” (referred to herein as “env”), “http.status_code” (referred to herein as “code”) and “k8s.io/pod/name” (referred to herein as “pod”) are indexed, then tags **1550**, **1551** and **1552** are extracted from the front-end service **1539** while tags **1560**, **1561** and **1562** are extracted from the product catalog service **1538**.

In an implementation, the extracted indexed tags are mapped to tag attributes. The extracted tags **1550**, **1551** and **1552** in the parent span (associated with the front-end service **1539**) may be mapped to a “FROM” tag attribute while the extracted tags **1560**, **1561** and **1562** in the child span may be mapped to a “TO” tag attribute. In one implementation, the mapped tags may be used to create node and edge data objects that are used to persist data for the metric events modality as shown in FIG. 15B.

FIG. 15B illustrates the manner in which data in the metric events modality is persisted using an edge data object comprising a memory-resident table of tag attributes with associated counts and using a node data object comprising a memory-resident table of tags with associated counts, in implementations according to the present disclosure. In one implementation of the monitoring service disclosed herein, a memory-resident table **1501** titled “Edge Health” may be maintained to keep track of the various dependencies in the application. The table **1501** may be stored in, for example, in the storage module **1366** (in FIG. 13). A memory-resident table **1500** titled “Node Health” may be maintained to keep track of the various service nodes in the application. Both tables comprise aggregated rows comprising metrics values. In one implementation, these rows are stored efficiently for fast aggregation.

For example, the table **1501** may comprise one or more exemplary rows related to the cross-service span pair discussed in connection with FIG. 15A. Row **1506** is one exemplary row that may be generated for the cross-service span pair of FIG. 15A. Note that for simplicity, only tag attributes “from pod” and “to pod” are illustrated in row **1506** in FIG. 15B, but row **1506** would typically also comprise information for tag attributes associated with indexed tags “code” and “env” that are also indexed. As discussed above, each row for the cross-service span pair of FIG. 15A will comprise a unique set of “FROM” and “TO” tag attributes. For example, if the front-end service (e.g., front-end service **1539** in FIG. 15A) makes multiple calls to the product catalog service (e.g., product catalog service **1538** of FIG. 15A), but any of the calls are associated with different values for the “pod” tag from the values shown in row **1506**, the information would be recorded in a new row. In other words, each row records a single unique combination of tag attributes and service names. If the value of either the “from pod” or “to pod” tag attribute changes, a new row is created to record the information. Accordingly, there may be multiple rows in the table **1501** for the cross-service call discussed in connection with FIG. 15A, where each row would comprise a unique combination of “FROM” and “TO” tag attributes for a given cross-service span pair.

Each row in the table **1501** comprises a count value for number of requests **1504**, errors **1505** and latency **1511**. The requests **1504** are incremented each time the same cross-service call with the same unique set of attributes for a respective row is observed on a trace. The errors **1505** are incremented each time a request associated with a respective row is observed on a trace that has an error. The latency **1511** metric relates to a histogram of the duration that a respective request took. Further, each row comprises a timestamp **1503** to record the time of the cross-service call.

Using the metrics associated with the requests **1504**, errors **1505** and latency **1511** and the timestamp **1503**, aggregations on the rows may be performed quickly and efficiently to determine SLIs for varying ranges of time. In response to a client query then, the numeric rows in the tables **1500** and **1501** may be summed into either timeseries buckets or into a single number depending on the query.

In one implementation, the metric events modality may maintain a separate memory-resident table **1500** titled “Node Health” in system memory associated with the service nodes in the application. Each row in the memory-resident table **1501** comprises a unique combination of service names and associated tags. For example, row **1508** is associated with the front-end service (e.g., service **1539** in FIG. 15A) and comprises corresponding tag values for “env,” “pod” and “code.” Similarly, row **1507** is associated with the product catalog service (e.g., product catalog service **1538** of FIG. 15A) and comprises corresponding tag values for “env,” “pod” and “code.”

Each unique combination of service name and corresponding tag values is associated with metrics that are maintained in the memory-resident table **1500**, e.g., request, error and latency (as discussed in connection with table **1501**). These metrics may be used to perform fast and efficient aggregations. For example, if the client queried the number of times “env=prod” in the application, assuming the two exemplary services illustrated in table **1500** are the only ones where “env=prod,” the request counts in each row would be aggregated to provide a result of 2.

Note that the memory-resident table **1500** may also comprise a “root cause” metric **1509** which tracks the number of times the corresponding service was the root cause of an error. For example, the “root cause” metric may be aggregated using the memory-resident table **1500** across multiple rows to determine the number of times each given service in an application was the root cause for an error.

In one implementation, a software tool may be employed to perform faster aggregations across the rows of tables **1500** and **1501**. For example, Apache Druid, which is an open-source data store designed for sub-second queries on real-time and historical data, may be used to perform the aggregations rapidly and efficiently. In different implementations, other tools may also be used to perform aggregations. In one implementation, the information in the memory-resident tables **1500** and **1501** may be used in the metric events modality to perform the metrics aggregations for rendering the service graph (e.g., graph **900** of FIG. 9) and computing the associated SLIs.

In one implementation, the metrics event modality may also store Trace IDs associated for each unique combination of cross-service span pairs and corresponding indexed tags. In one implementation, the aggregation module **1322** (of FIG. 13) of the monitoring platform aggregates across the cross-service span pairs by maintaining one or more exemplary Trace IDs for each unique set of “FROM” tag attributes (and their corresponding values) to “TO” tag attributes (and their corresponding values) for a cross-service pair. Accordingly, exemplary Trace IDs may be maintained for each unique cross-service call.

The exemplary Trace IDs stored with each unique set of “FROM” and “TO” tag attributes for a cross-service span pair may be used by the querying module **1382** to respond to queries requesting more particularized information pertaining to non-indexed tags associated with the spans. For example, if a client needs particularized information regarding span performance or span duration, the querying module **1382** may be able to use the aggregated rows of information stored in a database associated with the storage module **1366** to access one or more exemplary Trace IDs associated with the call. Using the Trace IDs then, the querying module may be able to access the sessionized traces **1308** and perform analytics on the retrieved exemplary traces to deliver the requisite span performance and span duration information. In one implementation, the full trace information may be

accessed from a storage set associated the full-fidelity modality, which stores the entire traces as ingested following sessionization. In a different implementation, however, the metric events modality may save full trace information for traces associated with the exemplary Trace IDs in a separate storage from the data set associated with the full-fidelity modality. In one implementation, because the metric events modality allows clients to retrieve raw trace data, it also allows clients to run an analysis on the retrieved data for an arbitrary set of tags (instead of being limited to the tags pre-indexed by indexing module 1307).
10

The metric events modality is particularly advantageous in circumstances where the client has identified a problem from the information provided by the metric time series. Having identified a problem either by manual monitoring of RED metrics or through an automatically generated alert, the client may be able to traverse deeper using the metric events data set and access relevant traces to receive more specific information regarding the problem. Also, the metric events modality allows the client to run an arbitrary analysis 15 on the traces, e.g., on a set of tags that has not previously been indexed, which provides the client with specific information that may be used to diagnose and resolve the problem.

FIG. 15C illustrates the manner in which data in the metric events modality is persisted using an edge data object 25 comprising a memory-resident table of extracted indexed tag attributes with associated Trace IDs and using a node data object comprising a memory-resident table of extracted tags with associated Trace IDs, in implementations according to the present disclosure. In one implementation of the monitoring service disclosed herein, a memory-resident table 1531 created to persist data is associated with the various dependencies in the application. Also, a memory-resident table 1530 created to persist data for the metric events 30 modality is associated with the various service nodes in the application. Note that table 1531 is created in a similar way to table 1501 in FIG. 15B and that table 1530 is created in a similar way to table 1500 of FIG. 15B. Instead of tracking RED metrics, however, the tables in FIG. 15C comprise a column for Trace IDs 1590 and Exemplar Type 1591. It should be noted that, in one implementation, memory- 35 resident table 1531 may be maintained in combination with memory-resident table 1501 and that memory-resident table 1530 may be maintained in combination with memory- 40 resident table 1500.

Row 1597 in table 1531 is one exemplary row that may be generated for the cross-service span pair of FIG. 15C. Note that for simplicity, only tag attributes “from pod” and “to pod” are illustrated in row 1597 in FIG. 15C, but row 45

1597 would typically also comprise information for tag attributes associated with indexed tags “code” and “env.” As discussed previously, each row for the cross-service span pair of FIG. 15A will comprise a unique set of “FROM” and “TO” tag attributes. Accordingly, there may be multiple rows in table 1531 for the cross-service call discussed in connection with FIG. 15A, where each row would comprise a unique combination of “FROM” and “TO” tag attributes for a given cross-service span pair. Each row in table 1531 comprises a Trace ID 1590, which keeps track of one or more Trace IDs associated with the unique combination of service names (and operation names) and tag attributes for the given row. In other words, the combination of service names (and operation names) and tag attributes in each row may comprise an index to access the associated Trace IDs.
55

In one implementation, the Exemplar Type column 1591 tracks the type of exemplary trace associated with the Trace

10 ID. Types of exemplars may be request, error, root cause errors or some latency bucket identifier. The Trace IDs in each row may be accessed to identify and retrieve the full trace associated with the ID for further analysis, e.g., an analysis on an arbitrary set of tags associated with the trace.

In one implementation, the monitoring system may maintain a separate table 1530 associated with the service nodes in the application. Rows 1595 and 1596 in table 1530 are two exemplary rows that may be generated for the cross-service span pair of FIG. 15A. Each row in table 1530 comprises a unique combination of service and associated tags. For example, row 1595 is associated with the front-end service (e.g., service 1539 in FIG. 15A) and comprises corresponding tag values for “env,” “pod” and “code.” Similarly, row 1596 is associated with the product catalog service (e.g., product catalog service 1538 of FIG. 15A) and comprises corresponding tag values for “env,” “pod” and “code.”
15

20 Each unique combination of service name and corresponding tag values is associated with a Trace ID and Exemplar type that is maintained in table 1530.

As noted above, in one implementation, metrics event data may be persisted in tables that consolidate the data shown in FIG. 15B and FIG. 15C. For example, table 1501 may comprise an additional column to track Trace IDs and similarly table 1500 may comprise an additional column to track Trace IDs.
25

30 The Trace IDs may be used in metrics events modality to retrieve full traces for more detailed analysis. In one implementation, full traces associated with the exemplary Trace IDs may be maintained in a dedicated storage associated with the metric events. In a different implementation, the full traces may be accessed from a data set associated with the 35 full-fidelity modality.

It should be noted that the metric events modality can 40 comprise higher-cardinality metrics information because a higher number of tags may be indexed for the metric events data set as compared to the dimensions associated with the metric time series. However, the metric time series modality may provide higher-fidelity information because it retains metadata associated with incoming spans (e.g., service name, operation name, count values, etc.) that are not collected in the metric events modality. Further, the metric time series modality also allows clients to configure alerts against one or more time series to monitor incoming data in real-time. Because metric events are generated from post-sessionized traces, the metrics data associated with metric events may not be computed as rapidly as compared with the metric time series modality.
45

4.3 Full-Fidelity Modality

In one implementation, the full-fidelity module 524 of FIG. 5 stores all the incoming trace data from the sessionization process in real time. Unlike the prior two modalities, the full-fidelity modality stores the trace data in its raw form. In one implementation, the data is stored in parquet-formatted batches of full traces in an unstructured format (e.g., blob storage) along with some metadata. The metadata may comprise the tags associated with the trace (both indexed and unindexed) and other properties such as service name and operation for more efficient querying. In one implementation, the format of the metadata may comprise a map of a service name to a map of tag names, wherein each tag name may be mapped to a list of tag values. The batches of full traces in unstructured format and the metadata are queried in the full-fidelity modality using a robust data engine to search for any tag across the traces. For example, PRESTO is an
50
55
60
65

open source distributed SQL query engine that may execute queries against data sources of varying sizes.

FIG. 16 is an exemplary on-screen GUI showing the manner in which a client may submit a query to be executed against the full set of traces stored in connection with the full-fidelity modality, in accordance with implementations of the monitoring service disclosed herein. The full-fidelity modality, in one implementation, allows a client to execute a query against arbitrary tags to receive a set of traces that matches the query. For example, in the GUI of FIG. 16, the client enters a query 1604 for traces where “Request Type=user,” “Service=adservice” and the tag “ad_size_bytes=4092.” In response, the platform returns a list 1602 of the traces matching the client-entered filters and, further, provides information about the traces, e.g., the Trace ID, duration, start time, root operation, root cause error status code and associated spans. As mentioned previously, the traces retrieved in response to a query may be analyzed to determine performance summaries for the spans comprised therein. Alternatively, the span performance analysis can be computed on all the traces stored as part of the full-fidelity data set.

In one implementation, the monitoring platform has the ability to run a full trace search (as shown in FIG. 16), and feed the traces collected into other modalities of analysis to get more detailed information about an arbitrary set of traces and an arbitrary set of attributes associated with the set of traces.

5.0 Multiple Modalities for Performing Real User Monitoring (RUM)

RUM is the practice of using data from an application or website’s real-life users to monitor and understand application performance. RUM tracks metrics such as DNS timing, time-to-first-byte, full page load time, JavaScript errors and the time it takes to load specific elements. These metrics are collected by monitoring actual user sessions. By monitoring real-user data across a variety of end-user configurations, browser versions, operating systems, feature flags, user status, locations, etc., software delivery teams can identify problems that undercut the user’s digital experience and user satisfaction. RUM is a specific type of application monitoring that relies on the passive collection of data produced by real users to identify application availability or performance issues. RUM provides insights that are difficult to achieve through other performance monitoring techniques because it synthesizes and reports on data collected from actual human users. While APM is used to monitor backend services and the interaction between them, RUM may be used to monitor activity and provide visibility all the way from the browser through the network down to the backend services.

There are several challenges associated with implementing an observability platform (e.g., monitoring service 306 of FIG. 3) that can perform both APM and RUM-related computations within a heterogeneous distributed system. One of the challenges associated with computing metrics for both RUM and APM, for example, is efficiently ingesting and aggregating significant amounts of span and trace data generated by a website or application. Conventional tracing and monitoring systems are simply unable to ingest vast amounts of span and tracing data and, therefore, have to resort to sampling the data intelligently to reduce the volume of stored trace data. Using sampling exclusively, however, results in data loss and, as a result, conventional monitoring tools do not allow clients access to all the spans and traces generated by real user interactions with a website or appli-

cation. Furthermore, conventional monitoring tools may calculate real-user metrics based on the sampled set of data and, accordingly, the calculations may be approximate at best and inaccurate at worst.

Implementations of the monitoring service (e.g. monitoring service 306) disclosed herein advantageously allow clients of the monitoring service the ability to ingest up to 100% of both RUM and APM-related spans and to generate metric data using the ingested spans. For RUM-related spans, for example, streams of metric time series data may provide clients with valuable real-time information pertaining to webpages (e.g. metrics related to accessing a particular endpoint provider) and also allow alerts to be configured to manage anomalous behavior associated with the webpages. Ingesting up to 100% of all span data also allows clients of the monitoring platform to retrieve responses to queries requiring a high degree of resolution, e.g., specific queries pertaining to specific user interactions with an application or browser over specific periods of time. Conventional monitoring tools that sample data simply are not able to provide accurate responses to queries requiring a high degree of resolution because they do not save all the generated span or trace data.

Note that as used herein, “users” refers to real-life users of an application or website, whereas “client” refers to a frontend developer of the application or website or site reliability engineer (SRE) (associated with the application or website) using a monitoring platform (e.g. monitoring service 306) to monitor the interactions of the real-life users with the application or website.

In addition to ingesting and storing up to 100% of the APM-related spans, implementations of the monitoring service disclosed herein also sessionize and store up to 100% of the RUM-related spans (e.g., spans generated as a result of a real user interacting with a website or application) received from the client in real time. Implementations of the monitoring service comprise an ingestion streaming pipeline that can ingest and consolidate the incoming spans into traces, and is further able to use advanced compression methods to store the traces. Within the RUM instrumentation, implementations of the monitoring service may use traces to organize units of activity and may also extract all necessary metrics based on the trace data. In one or more implementations, the monitoring service may use spans to organize units of activity and may also extract all necessary metrics based on the span data. In one implementation, the monitoring platform may also be able to compute certain metrics associated with the entirety of a user session. Because incoming trace and span information may be efficiently ingested and aggregated in real time, the monitoring platform is able to advantageously convey meaningful and accurate information (without the need for sampling) regarding the frontend interactions of a user with a website or an application, e.g., page load times, HTTP requests, time-to-first-byte, etc. High-cardinality metrics may be calculated with a high degree of accuracy because all incoming data is accounted for and there is no data loss as a result of sampling.

Implementations of the monitoring service disclosed herein further allow a client to store and analyze the RUM data using multiple modalities of analysis (similar to the modalities for APM data discussed in connection with FIG. 5). In one implementation, a first modality comprises converting incoming RUM-related spans (or traces) from one or more clients into a plurality of metric data streams (also referred to as metric time series) prior to sessionizing the spans. The plurality of metric data streams computed for

RUM data are similar to the metric data streams created for APM data discussed above. Each metric time series is associated with a single span identity, where a base span identity comprises a tuple of information corresponding to an associated type of span. Each metric time series in this modality (referred to herein as “metric time series modality”) represents a plurality of tuples, with each tuple representing a data point. Key performance metrics (KPIs) can be extracted directly from the metric time series in real-time and reported. Because the metric time series are created without paying a time penalty associated with sessionization, they can be used to perform real-time monitoring with sub-second resolution and to generate alerts within two to three seconds if a condition is violated.

In one or more implementations, a second modality of analysis sessionizes the incoming RUM-related spans and supports deriving higher-cardinality metrics (as compared with metric time series data) for a selected set of indexed tags, e.g., client-selected tags, global span tags, etc., over selected time durations (referred to herein as the “metric events modality”). The metric events computed from RUM data are similar to the metric events computed for APM data discussed above. In one implementation, however, the higher-cardinality metrics for the metric events modality for RUM data is generated directly from span data without needing to wait for the traces to be fully sessionized. In other words, generating metric events and extracting higher-cardinality metrics from spans for RUM may be more efficient than for APM because the monitoring service does not need to wait to store and consolidate the incoming spans into traces in order to generate the metric events or compute the higher-cardinality metrics. Because RUM data typically comprises spans generated by a user interacting with a browser or application, the metrics can be computed directly from the span data, as compared with APM data where the spans need to be consolidated into the traces to provide a complete picture of the manner in which the spans traverse through the various services in an application. Accordingly, generating metric events for RUM data is typically faster than generating metric events for APM data.

This modality is particularly useful for clients that need accurate SLI information for a larger set of high-value indexed tags. The metric events modality enables developers to aggregate metrics that have been pre-generated using the RUM span data to efficiently respond to queries submitted by a client. The aggregated real-user metrics help a client monitor end-user experience by providing visibility into the performance of a website or an application. Note that the metric events modality may not provide real-time streaming metrics and data retention that the metric time series modality provides.

In one or more implementations, the metric events modality track exemplary RUM-related spans associated with a pre-configured set of indexed tags (similar to the manner discussed in FIG. 15C for APM-related traces). The tags to be indexed may be pre-selected by the client or the monitoring platform. The Span IDs or the Trace IDs may be used to retrieve the associated spans or traces and analysis on the spans or traces may be performed to generate more particularized information regarding an end-user experience of a website or application. In one implementation, once the spans or traces are retrieved, an analysis may be run on an arbitrary set of tags (in addition to the pre-configured indexed tags).

Additionally, in one or more implementations, a third modality of analysis may comprise a “full-fidelity” modality where a full-fidelity analysis may be conducted on any

dimension or attribute of RUM data to gauge the performance of services in the microservices-based application (similar to the manner discussed in connection with FIG. 16). The full-fidelity modality allows clients to search most or all of the incoming span data (including all the tag data) that was ingested by the monitoring platform without relying on sampling. The full-fidelity modality may sacrifice speed for accuracy, and may be used by clients that need a more thorough analysis of the services across every dimension or attribute.

In an implementation, the three modalities associated with analyzing RUM-related data may be simultaneously supported by the monitoring platform by storing ingested trace data using three different formats, where each format corresponds to one of the three available modalities of analysis. Note that implementations of the monitoring service disclosed herein are not restricted to three discrete data sets. The data sets for the different modalities may overlap or may be saved as part of a single data set. When a client submits a query, the monitoring platform may determine which of the data sets is most suitable for addressing the query. Thereafter, the monitoring platform executes the query against the selected data set to deliver results to the client of the monitoring platform. By comparison, conventional monitoring systems typically focus on a single modality and do not provide clients the ability to seamlessly navigate between different modalities. Conventional monitoring systems also do not provide the ability to automatically select the most appropriate modality based on the content, structure, syntax or other specifics pertaining to an incoming query.

FIG. 17 is a flow diagram that illustrates an exemplary computer implemented method of ingesting and aggregating span information to support multiple modalities of analysis for RUM, in accordance with implementations of the monitoring service disclosed herein. FIG. 17 illustrates the manner in which RUM ingest module 588 (discussed in FIG. 5) ingests and aggregates spans associated with RUM data. The RUM ingestion engine is similar to the APM ingestion engine, but may use a separate deployment from the APM ingestion engine.

In one implementation, RUM-related spans are received at the monitoring service 306 of FIG. 3 from the beacon 1767 (which performs substantially the same functions as beacon 567 of FIG. 5). The ingested spans may be sharded by session ID and organization ID. The spans received from the beacon 1767 are directed to an ingest routing module 1729 which may comprise different components, e.g., gateway services, load balancer, etc. In an implementation, ingest routing module 1729 may comprise a queue (not shown in FIG. 17) in which spans are stored prior to being sessionized by the RUM sessionization module 1706.

In one implementation, the incoming spans are converted into a plurality of metric data streams prior to consolidating the spans into traces. The metric data streams are generated by module 1720 prior to the spans being sessionized. Because the metric time series are created without paying a time penalty associated with sessionization, they can be advantageously used to perform real-time monitoring and alerting.

The incoming spans for RUM may also be sessionized where the span information is combined into traces in a process called sessionization. The RUM sessionization module 1706 is responsible for stitching together or combining the traces 1708. The traces associated with RUM may be used to organize units of activity and the necessary metrics may be extracted based on the trace data. Note that, in one

implementation, the sessionized traces may also be input to the module **1720** to create metric time series to track traces (separately from the time series created to track spans).

In addition to a Trace ID, each trace also comprises a time-stamp; using the time-stamps and the Trace IDs, the RUM sessionization module **1706** creates traces **1708** from the incoming spans in real time and sessionizes them into discrete time windows. For example, the sessionization process may consolidate traces (from spans) within a first time window (associated with time window Y **1780**) before transmitting the traces to modules **1720**, **1722** or **1724**. Note that in one implementation, the ingested RUM-related spans may be transmitted to modules **1720**, **1722** and **1724** without consolidating them into traces, wherein metrics and other information may be extracted directly from the span data.

Subsequent to consolidating traces for the first time window, the sessionization process may consolidate traces within the subsequent time window (associated with time window “Y+M” **1785**) before transmitting those traces to the modules **1720**, **1722**, or **1724**. It should be noted that the time windows associated with each of the modules **1720**, **1722**, and **1724** may be different. In other words, the metric time series data may be collected over short time windows of 10 seconds each. By comparison, traces for the metric events modality (associated with the module **1722**) may be collected over longer time duration windows. Note that while in certain implementations, each modality can persist at different resolutions, in other implementations, the time window durations and resolutions for all modalities may be the same.

In some implementations of the monitoring service disclosed herein, the RUM sessionization module **1706** for RUM-related spans is able to ingest, process and store all or most of the spans received from the beacon **1767** in real time. By comparison, conventional monitoring systems do not accept all of the incoming spans or traces; instead, they sample incoming spans (or traces) to calculate SLIs at the root level of a trace before discarding the spans. Implementations of the monitoring service disclosed herein, by comparison, comprise an ingestion streaming pipeline that is able to ingest all the incoming spans into traces in real time. Further, in one implementation, the monitoring service is also able to consolidate the spans into traces in real time, and is further able to use advanced compression methods to store the data. Further, implementations of the monitoring service disclosed herein are able to generate metric time series from the span data (prior to sessionizing the spans) to provide real-time monitoring and alerting of certain KPIs.

As noted above, the RUM sessionization module **1706** can collect all the traces within a first time window Y **1780** using the time-stamps for the traces. Subsequently, the sessionized traces are fed to the modules **1722** and **1724**, for the respective modalities (metric events and full-fidelity) for extraction and persistence. Note that the tags analyzed for all three modalities in connection with RUM may be different than the tags analyzed for APM. In other words, each of the modules **1720**, **1722** and **1724** may be configured to perform a RUM-focused tag analysis.

In one implementation, the incoming span or trace data is indexed by an optional tag indexing module **1707**, which indexes one or more tags in the data. The tags may be client-selected tags or tags that the monitoring platform is configured to index by default. In a different implementation, tag indexing may be performed as part of data aggregation, e.g., by the modules **1720**, **1722** and **1724**. Note that in the implementation of FIG. 17, the tag indexing module

1707 will be configured to index tags that are needed to perform a RUM-focused tag analysis.

In an implementation, data sets associated with each of the modalities may be persisted in one or more databases **1717**. It should be noted that while modules **1720**, **1722** and **1724** perform substantially the same computations on RUM-related spans, as the corresponding modules **520**, **522** and **524** perform on APM-related spans, the schema for persisting the RUM data in database **1717** may be different from the schema for persisting APM data in one or more databases **555** of FIG. 5. For example, RUM data may be aggregated at the session level (rather than at the trace level). Accordingly, even though incoming RUM-related spans may be sessionized into traces and trace data may be ingested for the full-fidelity modality (e.g., into module **1724**), the data is consolidated and persisted based on session IDs.

As noted previously, the spans associated with RUM (ingested, for example, from the beacon **1767**) may need to be treated differently from APM-related spans. For example, the spans related to RUM may need to be ingested and sharded by a session identifier (session ID) (and, optionally, an organization identifier) instead of using the Trace ID. A session ID is an identifier that connects a series of spans or traces. RUM data is typically organized into page views (which show details of a page visit) and sessions (which group all the page views by a user in a single visit). A session ID is typically used to filter for all the views in a specific session. For RUM, a developer is typically more interested in the behavior of a user over the course of a session, e.g., a user session interacting with a particular website or application. Spans associated with RUM are usually sharded and tracked using a session identifier (or session ID). Accordingly, the way full-fidelity data for RUM data is persisted in database **1717** may be different from the way full-fidelity data is persisted for APM data.

For RUM, data is persisted in shards (or chunks) based on time and session ID rather than by traces where each shard or chunk is associated with a window of time in a user session. The full-fidelity modality for RUM data stores the chunks of data in raw form, e.g., in parquet formatted batches of spans in an unstructured format with some metadata. The metadata may comprise the tags associated with the trace (both indexed and unindexed) and other properties such as service name and operation for more efficient querying. In one implementation, the format of the metadata may comprise a map of a service name to a map of tag names, wherein each tag name may be mapped to a list of tag values. The batches of spans in unstructured format and the metadata are queried in the full-fidelity modality using a robust data engine to search for any tag across the traces. For example, PRESTO is an open source distributed SQL query engine that may execute queries against data sources of varying sizes.

At query time, data is retrieved in chunks from the memory to be analyzed to generate a response to the query. In one implementation, indexes are created within the full-fidelity data set in database **1717** to be able to efficiently find relevant data within the data set.

Referring back to FIG. 17, the data sets for the respective modalities may be separate data sets, overlapping data sets or a single data set that supports all the modalities. Note that the databases **1717** may be a single database that stores data sets corresponding to all three modalities. Alternatively, the databases **1717** may represent different respective databases for each of the three modalities. Furthermore, the databases

1717 may also represent distributed databases across which relevant information for each of the three modalities is stored.

In one implementation, a RUM analysis engine **1795** retrieves information regarding backend traces from APM module **1796**. APM module **1796** may extract APM trace information received from, for example, traces **508** in FIG. 5. The RUM analysis engine **1795** receives APM trace information and forms connections between the fronted RUM traces **1708** and the backend APM traces **508**. This allows a client to monitor the manner in which errors or problems arising at the backend propagate to the frontend and vice versa. By connecting the frontend and backend traces, the monitoring platform is able to provide complete visibility into any transaction all the way from a user browser, through the network, and to any backend service.

In one implementation, data associated with each of the three modalities is generated at the time of ingestion and stored separately from each other. The structure, content, type or syntax of query submitted by a client will typically dictate which of the three modalities and corresponding data set will be selected. In one implementation, an interface through which the query is submitted may also determine which of the three modalities and corresponding data set is selected. In an implementation, there may be some commonality in the data for the three modalities in which case the storage for the data may overlap. An alternative implementation may also comprise one or two of the three modalities (instead of all three) described above.

A client may send in a request to retrieve information pertaining to a website or application through query interface **1782**. Note that query interface **1782** may, in one implementation, be a common interface for querying both APM and RUM data. The underlying querying engine (e.g., the query engine and reporting system **324** from FIG. 3) will analyze the structure, content, type and/or syntax of the query, and also the interface through which the query is submitted, to determine to which of the modalities and respective data set to route the query for servicing.

In one implementation, the monitoring service may be able to consolidate both RUM and APM data at query time in order to respond to a user query. The query interface **1782** may be able to access modalities for both the APM and RUM-related data and provide a client with an appropriate response based on the query. In other words, in one implementation, the query can apply constraints to both the frontend RUM metadata and backend APM metadata, which allows a client to target both RUM and APM data using a single unified query. As noted above, the RUM analysis engine **1795** may form connections between the frontend and backend traces (as will be further explained below), which allow the query interface **1782** to target both the RUM and APM metadata with a single unified query. For example, at query time when a user submits a query through the query interface **1782**, the monitoring service may retrieve RUM related data from the database **1717** and further retrieve APM data from the APM module **1796** (via the RUM analysis engine **1795**). This way the frontend RUM span data may be aggregated with the corresponding APM backend traces in order to provide a response to the user query. Implementations of the monitoring service disclosed herein are, therefore, able to create a query time join of the RUM and APM data in response to a user submitting a query through the query interface **1782**.

In an implementation, the data sets corresponding to the modalities are structured in a way that allows the querying engine to navigate between them fluidly. For example, a

client may submit a query through the query interface **1782**, which may potentially result in the query engine accessing and returning data associated with the metric events modality for RUM data. Thereafter, if the client requires more in-depth information, the querying engine may seamlessly navigate to data associated with a different modality (e.g., full-fidelity) to provide the client with further details.

By way of further example, a client may submit a query through the query interface **1782**, which may potentially result in the query engine accessing and returning metric events associated with both RUM and APM data (using the linkage information from the RUM analysis engine **1795**). Conventional monitoring systems, by comparison, do not provide more than a single modality or the ability to navigate between multiple modalities of data analysis. Further, conventional monitoring systems do not provide developers the ability to query both APM and RUM data using the same interface, or provide an aggregate analysis of the manner in which the APM and RUM data are connected, for example, by rendering a service graph that shown both frontend browser data and backend microservice data.

5.1 End-to-End Visibility of a Real User Session

In certain instances, a frontend developer or site reliability engineer may need an overview of an entire session (e.g., the entire time duration that a user is session-interacting with a particular website or application) of a real user interacting with a website, application or web interface to gain insights into the end-user experience. A session typically groups all successive page views by a single user in a single visit and can be upwards of an hour. The developer may also need end-to-end visibility of a user session for troubleshooting purposes, e.g., to determine where the user experienced lags or had problems with navigation. As noted previously, conventional monitoring tools are unable to provide developers with end-to-end visibility into a user's session nor do they offer the ability to intelligently and more thoroughly explore areas of interest within the session.

As noted previously, implementations of the monitoring platform disclosed herein are able to ingest, store and analyze significant amounts of data (e.g., span data, trace data, etc.) generated by both the frontend (e.g., a website) or the backend (e.g., a service on the backend) of an application owner's architecture. Further, implementations of the monitoring platform disclosed herein use the data efficiently to provide the developer with insights into the performance of a website or application or to detect problematic conditions associated with, for example, browser performance, network performance, erroneous processes, failing services, etc. In particular, implementations of the present monitoring platform are able to construct an end-to-end representation of an entire user session by ingesting up to 100% of the incoming spans from the client (e.g., clients' browser, etc.) into the SaaS backend of the monitoring platform and analyzing them.

FIG. 18 illustrates an exemplary on-screen GUI comprising a service graph illustrating an access of one or more endpoints by a page with which a user is interacting, in accordance with implementations of the monitoring service disclosed herein. The service graph **1800** may comprise one or more nodes (e.g. node **1810**) that correspond to a page or a view (e.g., associated with a website URL or application) with which a user is interacting.

As shown in service graph **1800**, the page represented by node **1810** makes calls to several endpoints (e.g., endpoints associated with nodes **1820**, **1830**, **1840**, etc.). The endpoints correspond to resources that the page (represented by node **1810**) is attempting to access. Several different types of

endpoints may be displayed in service graph **1800**, e.g., endpoints associated with static resources, endpoints associated with third party providers, etc. This allows a client to gain insight into the manner in which different endpoints and endpoint providers (e.g., third party providers) are influencing the end user experience. In one implementation, the size of a node associated with either a page or an endpoint provider conveys the amount of traffic related to the node (as was discussed in connection with FIG. 9).

In one implementation, an application name and/or activity name for a node may be displayed alongside the node. For example, the URL for the page associated with node **1810** (<http://robotshop.k8s.int.lab0.signalfx.com>) may be displayed alongside the node **1810** in the GUI. As will be explained further below, in one implementation, the URL displayed alongside the node may be normalized. The application name or activity name may be extracted from span tags of spans associated with the node.

Note that in one implementation, the connections (e.g., edge **1808**) shown in service graph **1800** may comprise metric information regarding the respective access as was discussed in detail in connection with FIG. 9. For example, the access from node **1810** to the endpoint provider associated with node **1820** takes 223 ms as shown on the edge **1808**. In one implementation, the connection may also comprise information pertaining to an error rate.

The resources or endpoints may be either be internal or external with respect to a client of the monitoring platform. In one implementation, the endpoints may relate to external resources, e.g., an external service such as a payment processor, a content delivery network (CDN), etc. Alternatively, in one implementation of the monitoring platform, the resources may be part of a backend owned by the client. More specifically, the client may own existing backend infrastructure that supports one or more of the endpoints and can, therefore, exercise control over those endpoints. For example, nodes **1820** and **1840** may correspond to endpoints that a client's backend infrastructure supports. Because nodes **1820** and **1840** correspond to endpoints that a client controls, the client may be able to glean additional information regarding the behavior of those endpoints from its own backend, where the additional information may provide a client further insight into the performance of the endpoints.

Note that in one implementation, the service graph **1800** may be scoped by several constraints as discussed in connection with FIG. 10. To scope a service graph or similar visualization entails filtering and displaying the service graph or visualization across one or more constraints. In one implementation the GUI of FIG. 18 may include drop-down menus (not shown) that allow service graph **1800** to be filtered in accordance with different constraints, e.g., tags associated with environment, incident, etc. Further, the service graph **1800** may be filtered to display a particular type of endpoint or a specific view (or page).

In one implementation, the monitoring platform also provides information regarding spans that may be of particular interest to a client. For example, the GUI of FIG. 18 provides a list of exemplar spans (and the associated sessions in which they originate) associated with service graph **1800** in an adjacent panel **1801**. The list of spans, in one implementation, may provide information regarding a session ID **1850**, the span ID **1855** of the exemplary span, a timestamp **1860** for the respective span and any errors **1870** associated with the edge, page or endpoint which with the respective span is associated (e.g., HTTP errors, Javascript errors, etc.). As explained previously, the session ID **1850** is

associated with a specific session that a user is actively participating on a platform provided by a client.

Note that a span ID **1855** is displayed in the panel **1801** as opposed to a Trace ID because for RUM (as compared with APM), the spans provide a higher level of resolution and convey more meaningful information to a client. One of the differences between browser-emitted spans (associated with RUM) and backend spans (associated with APM) is that each browser span has all the metadata needed to analyze it, e.g., sessionId, location.href, activity, all tags, etc. As explained above, in order to analyze events for RUM, a client need not wait for a sessionized browser trace to be able to analyze individual spans or any propagation of metadata. The client can extract necessary information directly from the spans. Accordingly, for example, metric events may be generated for RUM faster than they are generated for APM data because the monitoring platform does not need to wait for the spans to be sessionized into traces.

Note that the spans displayed in the panel **1801** may be exemplars that capture representative spans associated with the service graph **1800**. Alternatively, the panel **1801** may list all the spans associated with the service graph **1800**. Because the monitoring platform ingests all the spans associated with user activity, the monitoring platform has the option of either displaying up to 100% of the spans associated with the service graph **1800** or exemplar representative spans that provide an overview of the activity in the service graph **1800** (as shown in FIG. 18).

In one implementation, the panel **1801** may list one or more exemplar spans (for each session ID **1850**) that capture activity in which a client may be interested, e.g., spans associated with a certain threshold number of errors, spans associated with different types of errors, spans associated with specific client selected endpoints or pages, spans associated with a key metric, etc. The monitoring platform may, therefore, be configured to either display one or more representative spans for a session that would be of particular interest to a client.

The list in the panel **1801** may comprise spans that are generated as a result of the calls made by one or more pages (or views) during a user session to various endpoints or resources. Each of the spans in panel **1801** may be associated with an access to a particular resource or endpoint as depicted in service graph **1800**. Each span may be the result of an activity that is either user-generated or was generated for the user by a browser (or application) with which the user is interacting. Note that not all the spans shown in the list are the result of direct user-interactions with a browser. Some of the calls depicted in the service graph **1800** or spans shown in the panel **1801** may, for example, be generated by background threads or by other processes that support the browsing session.

In one implementation, the spans in the panel **1801** may be categorized to provide a user better insight into the types of resources with which the spans are associated. For example, spans may be categorized, for example, based on whether the resources being accessed are internal or external to a client of the monitoring platform. There may be other criteria for organizing the spans in the panel **1801** as well, e.g., spans related to CDN traffic, spans categorized by type of endpoint or resource accessed, spans relevant to a key metric, etc. Categorizing the spans advantageously allows a client of the monitoring platform to better understand traffic patterns and user behavior. In one implementation, the

service graph **1800** may also be categorized so a client can visually distinguish between the various types of pages, endpoints and connections.

FIG. 19 illustrates an exemplary on-screen GUI comprising a service graph illustrating an access of an endpoint by multiple pages, in accordance with implementations of the monitoring service disclosed herein. The service graph **1900** depicts certain pages (e.g., pages associated with nodes **1910**, **1920**, etc.) in an application or website accessing a particular endpoint provider (e.g., the/cart endpoint **1940**). As shown in FIG. 19, the service graph **1900** can be filtered by a particular endpoint, e.g., /cart endpoint **1940** using drop-down menu **1922**. While FIG. 18 illustrates a page accessing several different endpoint providers, FIG. 19 illustrates that the service graph **1900** may also show several pages accessing a single endpoint provider.

In one implementation, similar to the service graph discussed in connection with FIG. 9, the service map may indicate the health of the various nodes. For example, certain endpoint nodes such as the/cart endpoint **1940** may comprise a solid-filled circular region **1930** indicating errors associated with the particular endpoint, where the size of the region **1930** indicates an approximate percentage of errors returned when the node is accessed. For example, the endpoint **1940** returns errors approximately 50% of the time when called. A page node, e.g., node **1920** with a solid-filled circular region, e.g., region **1923** indicates that the associated view had an error, e.g., a console error.

In one implementation selecting a particular node (e.g., endpoint **1940**) may highlight corresponding spans associated with the node in an adjacent panel **1901**. The client may then be able to further examine errors (e.g., HTTP errors, JavaScript errors, etc.) associated with various calls made from or to the node. In a different implementation, the panel **1901** may list exemplary spans of interest associated with the service graph **1900** (similar to the panel **1801** of FIG. 18) regardless of node or edge selection by a client.

In one implementation, aggregated metrics associated with a selected node (e.g. the /cart endpoint **1940**) may be displayed in a side-panel **1902**. For example, requests and errors **1931** and latency **1933** associated with the/cart endpoint **1940** may be displayed in the side-panel **1902**. It should be noted that metrics shown in the side-panel **1902** may be aggregated and computed using the metric events modality (e.g., using aggregated metrics from the metric event aggregation module **1722**). In a different implementation, real-time metrics may also be computed using the metric time series modality (e.g., using metrics aggregated from the metric time series module **1720**).

In one implementation, the side-panel **1902** may also provide a tag breakdown **1932** with values of all the indexed tags. The tag breakdown and analysis surfaces problems that a client may be interested in and prevents the client from going through all the data manually. For example, the tag value **1957** associated with the browser tag informs a client that most of the tags associated with errors for the/cart endpoint **1940** correspond to the Chrome browser. In other words, most of the errors for the/cart endpoint **1940** resulted from an access through a Chrome browser. In one implementation, clicking the full analysis option **1989** provides the client a full tag analysis. The full tag analysis allows a client to access a breakdown of the errors by the various tags.

FIG. 20A illustrates an exemplary on-screen GUI comprising an end-to-end view of a user session, in accordance with implementations of the monitoring service disclosed herein. As noted earlier, implementations of the present

monitoring platform construct an end-to-end representation of an entire user session by ingesting up to 100% of the incoming spans from the client (e.g., clients' browser, etc.) into the SaaS backend of the monitoring platform and analyzing them. Further, implementations of the monitoring platform disclosed herein also provide end-to-end visibility of a single user session (e.g. a user session interacting with a particular website or application) with the ability to perform more in-depth investigation of specific chunks of time associated with a user session. The chunks of time may either be client-selected or automatically selected by the monitoring platform based on a determination of the types of behavior a client may be interested in examining. In one implementation, the service graphs illustrated in FIGS. 18 and 19 are constructed using spans collected and analyzed during a single user session.

The GUI of FIG. 20A, in one implementation, conveys high-level metrics and information about the session including the start time **2010**, the session duration **2020**, the session ID **2031**, the agent **2032** and the location **2035**. As noted above, the session duration may be upwards of an hour depending on how long the user is interacting with an application or browser. The session illustrated in FIG. 20A, for example, is 6.7 hours. In one implementation, the field for the agent **2032** may comprise information about the browser and operating system (OS) used by the user. Agent information may enable a client to identify browsers, devices and platforms most used by users which can be beneficial in making informed optimizations to an application or website. In one implementation, the location **2035** may comprise information about the location of the browser, user, etc. As noted above, a session ID is an identifier that connects a series of spans or traces.

The graphic visualization **2000** is meant to capture the entire user session at an aggregate level along a time axis. In particular, the visualization graphically displays aggregated events and metrics computed for the user session. In one implementation, for example, the visualization graphically displays events **2041** (e.g., page load events), errors **2042**, JavaScript errors **2043** and requests **2044** associated with the session. Note that the errors **2042** are aggregated separately from the JavaScript errors **2043** to provide the client some insight into where the errors occurred, e.g., to distinguish between a frontend JavaScript error and an error that may have surfaced from the backend.

The aggregated metrics may comprise metrics aggregated for the duration of the entire session or a selected portion thereof. The graphic visualization **2000** provides a client with an efficient overview regarding the most active segments of time within the session. In one implementation, the graphic visualization provides visual or other types of indicators to direct the client's attention to the portions of the visualization in which the client would be most interested, e.g., page transitions, errors, etc., that a user experienced during the session.

In one implementation, a client of the monitoring platform may zoom into select parts of the user session. For example, a client may select region **2005** based on some activity of interest during the selected region. The client may be interested in the region **2005** based on a visual indication of a high number of events, errors, JavaScript errors or requests. For example, a client may select the region **2005** based on the several spikes in the number of events occurring during that region.

In one implementation, the graphic visualization **2000** may also display an aggregate number of events, errors, or requests associated with a selected region. For example, a

pop-up box **2045** indicates the aggregate number of events associated with the region **2005**. In one implementation the GUI may provide additional event activity metrics **2046** pertaining to the transactions in the selected region **2005**, e.g., the number of document loads, route changes, console errors, document resource loads, clicks, etc.

Note that in one implementation, the aggregated metrics related to events, errors and requests may be computed using the full-fidelity module **1724** from FIG. 17. Because visualization **2000** represents a single user session, the events **2041**, errors **2042**, JavaScript errors **2043** and requests **2044** may, in one implementation, all be computed using the set of traces associated with the single user session. These set of traces may be available in the full-fidelity data set corresponding to module **1724**. The set of traces associated with a single user session may also be persisted in module **1724** for efficient queries thereafter. Note, however, that as discussed above, the full-fidelity modality persists RUM data in shards based on a session ID (rather than a Trace ID).

In one implementation, a client may select a region of interest based on time values. For example, a client may be interested in user activity between the second and third hour of the user session. A client can then select the desired time bounds to define the region of interest and collect aggregated metrics associated with events, requests and errors for that region.

In one implementation, the GUI may provide a waterfall view of spans associated with all the events in the user session. In this implementation, all the events in the user session would be viewable in the panel **2050**. In a different implementation, the waterfall view displayed in the panel **2050** is scoped to the client selected region **2005**. In other words, the spans shown in the panel **2050** are associated with all the events encapsulated within client selected region **2005**. In one implementation, only exemplar spans associated with the client selection region **2005** are displayed. In a different implementation, all the spans associated with the client selected region **2005** are listed in the panel **2050**.

The spans displayed may be root-level or parent spans (e.g., span **2052**) that can be expanded out to reveal other child spans (e.g., spans **2053**, **2054**, etc.). For example, the document load event associated with parent span **2052** may be a combination of other sub-events, e.g., a document fetch associated with the child span **2053** and other different resource accesses such as the one associated with child span **2054**. Although not shown in FIG. 20A, in one implementation, each of the spans shown in the waterfall view of the panel **2050** is displayed adjacent to an icon indicating whether one or more errors are associated with the span. Displaying an error status of a span enables a client to visually identify whether a particular span needs to be explored further. Note that each of the spans (including both the parent and children spans) may be expanded out to get further information regarding the various attributes associated with the spans, as shown in connection with FIG. 20B.

FIG. 20B illustrates an exemplary on-screen GUI illustrating the manner in which a span may be expanded within the waterfall view to get further information regarding the attributes of the span, in accordance with implementations of the monitoring service disclosed herein. As mentioned above, any of the spans shown in the waterfall view of the panel **2050** may be further expanded to view the span attributes. For example, span **2064** shown in FIG. 20B may be expanded to display its various associated tags. The expanded view for the span **2064** shows the Span ID **2063** and the Parent ID **2062** for the span **2064**. Also, the expanded view shows all the related tags **2061**. For example,

the related tags may include the HTTP URL **2088** associated with the page that generated the span **2064**, the session ID **2089**, and links to a backend span and trace associated with the span **2064** (e.g., links **2067** and **2068**). In one implementation, additional information extracted from the attributes for one or more spans may be aggregated and conveyed to a client in the GUI of FIG. 20A.

As mentioned previously, some of the resources accessed by a page may be part of a backend owned by the client. In such cases, the frontend spans, e.g., spans **2053**, **2054**, etc. displayed in the GUI of FIG. 20A may be linked to backend APM traces to allow the client further insight into the performance of related endpoint providers. Where a frontend span associated with the user session can be linked to a backend APM trace, an APM icon **2051** is provided next to the frontend span in the waterfall view to indicate that linkage information exists for the respective span. The APM icon **2051** may also comprise a hyperlink to the linked backend trace. The linked trace is a backend representation of a user request for the resource associated with the respective span. A client may then be able to expand the span in the GUI to retrieve the link to the backend trace, where the link is provided as one of the attributes of the respective span. As shown in FIG. 20B, the attributes of the span **2064** comprise a backend trace link **2068** and a backend span link **2067**. In one implementation, when the links are accessed, the monitoring platform directs the client to a GUI that provides further information regarding the backend spans or traces. In one implementation where the span **2064** is associated with an error, the backend trace link **2068** and the backend span link **2067** may link directly to the trace or span respectively in the backend where the error originated.

In one implementation, the RUM analysis engine **1795** of FIG. 17 performs the analysis necessary to link the frontend span (e.g., the span **2064**) with the backend trace. For example, for a given span comprising a link to a backend trace, the RUM analysis engine **1795** may first check the corresponding APM deployment to determine if a backend trace exists that has the corresponding Trace ID included in the span attributes. Once the RUM analysis engine **1795** determines that the Trace ID exists in the backend deployment and that it is accessible, the monitoring platform can add a hyperlink to the backend trace next to the span in the panel **2050**.

FIG. 20C illustrates an exemplary on-screen GUI illustrating the manner in which hovering a cursor over a backend trace link provided for a span in the waterfall view conveys further information regarding a backend trace in a pop-up window, in accordance with implementations of the monitoring service disclosed herein. In one implementation, a client may hover a cursor over the APM icon **2051** of FIG. 20A (or even over one of the links **2067** and **2068** in FIG. 20B) to receive further information regarding a backend trace (or span). For example, hovering over an APM icon **2072** shown in FIG. 20C may result in a pop-up window **2071** on-screen and displaying summary information regarding the backend trace that connects to the associated frontend span **2073**. The pop-up window may further convey meaningful information regarding the linked backend trace, e.g., a performance summary **2078**, associated services **2079**, a Trace ID **2077** and a service/operation name **2075** related to the trace.

The pop-up window **2071** provides a client a preview of the trace so the client knows what to expect if the backend trace is accessed and loaded. The pop-up window **2071** also provides a hyperlinked Trace ID **2077** to be able to fast-track a client to the actual APM trace on the backend. Because the

monitoring platform has access to full-fidelity data, the client can click the hyperlink to directly access the backend trace, which may, for example, be stored in a data set associated with the full-fidelity modality discussed in connection with FIG. 17. In one implementation, the client may also access the Workflow link 2099 to access the aggregate behavior in the system for the transaction related to the span 2073. Accessing the Workflow link 2099 may direct the client to a service graph associated with the backend services monitored by APM, which allows the client to perform an in-depth investigation.

Implementations of the monitoring platform therefore, advantageously provide a client with an end-to-end view of an entire user session graphically while indicating regions of interest in which a user may be interested. Implementations of the monitoring platform, for example, may provide a graphical representation of the entire user session which tracks the number of events, errors, and requests in the session. Conventional monitoring methods, for example, did not provide the ability to provide clients with an overview of an entire session while also providing the ability to gain further resolution into specific client-selected portions of the session.

In one or more implementations, the monitoring platform allows a client to more intelligently navigate around the user session to discover portions of the session in which a client may be interested. In one implementation, for example, the visualization 2000 is automatically segmented into various chunks with certain vital statistics provided in the GUI for each chunk of time so a client can assess which regions are of particular interest. Thereafter, a client may be able to select a pre-segmented chunk or region of interest to receive information regarding the spans associated with the respective selected chunk in the panel 2050. For example, the region 2005 may be associated with a chunk of time that the monitoring platform automatically flags as being of interest to a client based on the spikes in event activity during that time period. A client is then able to review the metrics (e.g., event activity metrics 2046) associated with the pre-segmented chunk of time to determine if further exploration is warranted and also review the spans associated with the chunk in the panel 2050.

In one implementation, the monitoring platform is able to intelligently pre-select a region of interest for a client based on spans in which the client is interested. For example, a client may want to investigate the edge 1808 in FIG. 18 associated with a call to the endpoint or node 1820. Alternatively, a client may want to investigate a node (e.g., a page node or an endpoint node) and explore exemplar spans associated with calls from or to the respective node. In either case, a client may start by selecting an exemplar span associated with either the relevant edge or node from the panel 1801.

Upon selecting the relevant span, the client, in one implementation, may be directed to the graphic visualization 2000 of FIG. 20A with a region of interest (e.g., region 2005) pre-selected for the client, where the region of interest would include information regarding the selected endpoint or edge. In this way, a client needing to explore a particular endpoint, page or edge in the GUIs of FIGS. 18 and 19 may be directed to the GUI of FIG. 20A with the relevant portion of the graphic visualization 2000 scoped to the portion in which the client is most interested. This allows the client to not only investigate the segment of interest but also advantageously provides the client with an overview of the relative location of the region of the interest within the overall session. Thereafter, the client can inspect the seg-

ment of interest more closely while having an overall idea of other proximate events that were taking place during the user session.

By way of further example, a client may need to explore errors generated by a page associated with node 1920. The client may identify exemplar spans within the panel 1901 associated with node 1920 and select one of those spans. Upon selection of the relevant span, the monitoring platform would automatically bring up the GUI of FIG. 20A with a portion of the graph pre-selected where the selected portion comprises information regarding the span selected through the GUI of FIG. 19. The monitoring platform therefore automatically takes a client to a specific segment on graphical visualization 2000 associated with the user selection in FIG. 19. The client can then visually analyze where the error occurs within the session in relation to other events in the session. For example, the client is able to determine whether the user continued to navigate around the page after experiencing the error or if the user chose to leave the page and transition to a different page.

In one implementation, upon selection of an exemplar span in the panel 1801 of FIG. 18 or the panel 1901 of FIG. 19, in addition to pre-selecting a region of interest within the graphic visualization 2000, the monitoring platform may also present the specific span of interest to the client within the waterfall view of the panel 2050 of the GUI. In other words, the monitoring platform can direct the client to the actual instance of the span that exhibited the behavior in which the client expressed an interest. The client then has the option of expanding the span and analyzing the relevant attributes of interest associated with the span. In one implementation, the span within the waterfall view of the panel 2050 may either be highlighted or already expanded to allow a client to easily review the span attributes. For example, the span of interest may be expanded similar to the manner shown in FIG. 20B so the client can visually inspect the various attributes associated with the span. In one implementation, the GUI of FIG. 20A may be configured to visually indicate to the client that a span displayed in the panel 2050 is related to the specific edge, endpoint or span selected in the interfaces of FIGS. 18 and 19.

Note that a client may also be able to select (or interact in some way with) a node or an edge in the service graph 1800 of FIG. 18 (or service graph 1900 of FIG. 19) and be directed to the graphical visualization 2000 of FIG. 20A with a region of interest pre-selected and a specific instance of a span of interest highlighted (or expanded) in the waterfall view of the panel 2050. In this way, a client is provided not only specific information regarding an error span of interest in the waterfall view but also an overview of where in relation to other events the error occurs on the session timeline.

An exemplary manner in which a client may approach troubleshooting latency problems experienced by one or more users may start with the client investigating a user's interactions with a page or application at the service graph level. Viewing the service graph of FIG. 19, for example, the client may observe that the/cart endpoint 1940 is returning errors when invoked in response to a call. The client may select the endpoint 1940 and view information regarding an associated error span in the panel 1901. Subsequently, the client may select or double-click on the entry corresponding to the error span in the panel 1901 and be directed to the session view of FIG. 20A, where a region (e.g., region 2005) associated with the selected error span may be pre-selected for a client. Event activity metrics 2046 for the pre-selected segment may also be displayed to provide a client with some aggregated metrics associated within the segment. The cli-

ent's attention may also be directed to the actual error span of interest within the panel **2050**, for example, by having the error span expanded to display its attributes. If the client controls the backend services to which the/cart endpoint **1940** maps, the error span will typically contain a link to the backend trace associated with the error span. The client can then access the link to be directed to the backend trace. As noted earlier, implementations of the monitoring platform are, therefore, able to map endpoint provider (or page provider) nodes on the frontend to the backend microservices comprised within a software architecture.

In one implementation of the monitoring platform, the spans within the waterfall view of the panel **2050** may be organized by page access. Instead of laying out the spans in a chronological order, in this implementation, the spans may be organized and grouped by page view. This allows clients to have clearer insight into the manner in which a user navigated around a particular website or application. Any time a user transitions to a new page, either through a document load or a route change, the waterfall is reset and any activity subsequent to the transition is grouped separately from the prior activity.

By way of example, selecting the region **2005** of FIG. **20A** allows a user to view a list of all the spans associated with the region **2005** in the panel **2050** grouped by page view. This allows the client to easily determine if the event spikes within the region **2005** are all associated with a single page view or multiple different page views. The client is also able to easily receive insight into a user's journey, in particular, the manner in which the user navigated from one page to the next.

Grouping the spans by page view also allows the client to conveniently determine which of the pages had the most associated errors. For example, if the number of errors within the region **2005** were high, grouping the activity by page view allows a client to easily determine which of the page views is the most problematic. As noted previously, each of the spans are displayed adjacent to an icon indicating whether the span comprises an error. Accordingly, grouping the spans by page view may also allow a client to visually determine which of the pages were associated with the most error spans.

FIG. **21** illustrates an exemplary on-screen GUI displaying aggregate metrics for a specific page, in accordance with implementations of the monitoring service disclosed herein. In one implementation, the monitoring platform may be able to provide the client with aggregated metrics for a specific page or even an endpoint selected, for example, from the service graphs displayed in FIGS. **18** and **19**. FIG. **21** illustrates aggregated metrics associated with a particular page or endpoint for a select time period. The aggregated metrics can, in one implementation, be calculated using the metric events modality for RUM-based spans discussed in connection with FIG. **17**. Alternatively, in a different implementation, more real-time metrics may be calculated using the metric time series modality also discussed in FIG. **17**. In another implementation, the full-fidelity modality may also be used to compute the metrics.

The aggregated metrics, among other things, include end user response time **2152** and a graph **2153** conveying page load times over a selected time period. Other metrics **2154**, e.g., network time, response time, server time, requests, load time, etc. may also be displayed for the convenience of a client in the GUI of FIG. **21**. In one implementation, the client may be provided additional details regarding how much of the latency associated with a page load was due to network latency versus latency attributable to a client's

backend. In one implementation, the metrics (e.g., network time, server time, etc.) may be used to provide a client a way to visualize network time and/or server time across the service graphs shown in the GUIs of FIGS. **18** and **19**. Providing metrics for specific pages and/or endpoints through an interface such as the one shown in FIG. **21** allows the monitoring platform to advantageously provide a client with more targeted information, especially, in instances where the service graph view (in FIGS. **18** and **19**) or the session view (in FIG. **20A**) may be particularly crowded with information.

Session exemplars associated with the page or endpoint associated with the GUI may be displayed in a panel **2155**. The session exemplars may comprise details regarding the session ID, a timestamp associated with the session, the duration of the session, the agent (e.g., a browser, platform, OS) used during the session, a location of the user and a number of errors encountered during the session. In one implementation, where aggregated metrics for a specific page are displayed (as shown in FIG. **21**), information extracted from the location field **2167** associated with each session may be used to construct a geo-map allowing a client to visualize page views broken down by location.

Location information, among other uses, helps a client understand regional performance of a website. In one implementation, the location information may also be used to visualize page load times broken down by location (e.g., by city, by country, by zip code, etc.). For example, an exemplary GUI may provide a client a map of the world allowing a client to hover over any country (or region) to see average load times or average page views for that country.

In one implementation, page views and other metrics may also be filtered based on URL, browser type and version, operating system, user id, etc. The GUI of FIG. **21** may provide visualizations that enable a client to see a breakdown of the page views based on URL, browser type and version, user id or operating system.

In one implementation, for each page, a geo-map may be constructed to analyze the endpoint traffic for the respective page, e.g., the map may be able to visually indicate to a client the physical location of the various endpoints. This allows a client to obtain a better understanding of the proximity of the various resources being accessed by any particular page.

In one implementation, where aggregated metrics for a particular endpoint are displayed, information from the location field **2167** may be used to construct a geo-map of pages accessing the particular endpoint. This allows a client to obtain a better understanding of where the pages or related users accessing a particular endpoint or resource are located.

FIG. **22** illustrates an exemplary on-screen GUI displaying a geo-map associated with a particular website or application, in accordance with implementations of the monitoring service disclosed herein. The geo-map **2205** provides, among other things, a bird's eye view of where the traffic for a particular site or application is coming from. In addition to understanding the regional performance of a site, the geo-map can be used to visualize page load times broken down by location, to analyze resource access for a respective page by location, etc. The GUI of FIG. **22** may also provide other relevant information to a client, e.g., high-level metrics **2210** for an application that are easily accessible, errors by view **2215**, top endpoints with errors **2220**, or visit counts by browser **2225**.

5.1.1 Aggregating Metrics for Workflows Associated with a Real User Session

There are several challenges associated with implementing an observability platform (e.g., monitoring service 306 of FIG. 3) that can perform RUM-related computations within a heterogeneous distributed system. One of the challenges is associated with providing clients meaningful information regarding where work related to a particular user interaction with a website (or a client process) is occurring and, also, regarding application-specific use cases that cannot be easily captured with basic instrumentation techniques.

In addition to enabling workflows for APM (as discussed in connection with FIGS. 12A and 12B previously), implementations of the monitoring platform disclosed herein (e.g., monitoring service 306) allow workflows (also known as “custom events”) to be used in connection with RUM to provide clients with meaningful information regarding work performed in relation to a frontend user interaction with a website (or application) or a client process. An example of a user interaction may be a user selecting an “Add to Cart” option on a website of an online retailer. Conventional tracing and monitoring tools do not have the ability to track and consolidate related frontend spans (associated with a workflow) generated in response to a user interaction or a client process during a real user session. Further, conventional tracing tools do not provide clients the ability to extract and aggregate metrics from the related frontend spans associated with the workflow.

Implementations of the monitoring platform disclosed herein advantageously allow users to identify and track one or more related RUM spans associated with a workflow that are generated during a real user session in response to a particular user-interaction or a client process. Further, implementations of the monitoring platform are able to aggregate metrics from the one or more related RUM spans associated with the workflow. For example, a user may need to monitor a cumulative duration or error count associated with a particular user interaction, e.g., selecting an “Add to Cart” option on a website of an online retailer. Using metadata associated with a workflow attribute (and corresponding value) extracted from the tags of one or more related spans generated in response to the selection of the “Add to Cart” option, implementations of the monitoring service can identify other spans related to the user action. Subsequently, implementations of the monitoring service are able to compute the error or duration metrics associated with the selection of the “Add to Cart” option using the entire set of related spans.

Extracting a workflow attribute from one or more spans associated with a particular user-interaction or a client process on the frontend advantageously allows implementations of the monitoring service to efficiently identify other spans associated with the respective user-interaction or client process. As a result, implementations of the monitoring service are able to identify and track spans ingested during a RUM session that are not explicitly tagged with a workflow attribute but may be related to a particular workflow (and to the spans which are explicitly tagged with a workflow attribute associated with the particular workflow). By associating other spans that may not be explicitly tagged with a workflow attribute to a particular workflow, implementations of the monitoring service disclosed herein are able to more accurately compute metrics for the particular workflow and provide clients with meaningful information about the entire set of spans that may be generated from a given user interaction or a client process.

As noted earlier, software developed by a client (e.g., clients A and B in FIG. 3) may be instrumented in order to monitor different aspects of the software. Instrumentation results in each span being annotated with one or more tags that provide context about the execution, e.g., a client process related to the span. The workflow dimension is one of the attributes (or tags) with which a span may be annotated through instrumentation to provide contextual information regarding the user interaction or client process to which the span relates. In an implementation, instrumentation for RUM may include the workflow dimension within a set of tags (e.g., global tags) attributed to a span known to relate to a particular user interaction or client process. Referring back to FIG. 20B, as shown in the waterfall view therein, a span may be tagged with a workflow attribute 2021 (“Workflow Name”) and a corresponding workflow value 2022 (“Chartload”). The workflow attribute and its respective value indicates that the span shown in FIG. 20B may be associated with loading a chart on a particular website.

In RUM, the user may be viewed as a single service, one which interacts with other external services and performs some internal-only work. At this level, clients are typically interested in aggregate activities that might include many different external resource calls working in collaboration. This can make it difficult to narrow down at an aggregate level what the state of the system is when an issue or problem occurs and specifically where in the overall flow an error or performance degradation occurs. By using workflows in connection with RUM, implementations of the monitoring service use workflows to provide clients meaningful information regarding where in an overall flow a particular error or performance degradation occurs.

One of the differences between browser-emitted spans (associated with RUM) and backend spans (associated with APM) is that in order to analyze events for RUM, a client need not wait for a sessionized browser trace to be able to analyze individual spans or any propagation of metadata. Further, instrumentation for RUM makes tracking and associating spans to a single causal source challenging. Because of this challenge, the RUM spans are not easily organized into causality-related traces. Implementations of the monitoring service disclosed herein address this challenge by identifying and organizing associated spans for secondary analysis into workflows. This aggregation may be performed at ingest time or at query time. In order to extract meaningful information for a client using workflows, however, implementations of the monitoring service first need to identify the entire set of spans related to a particular workflow.

In contrast to APM, where an application can be instrumented so that each span in a trace associated with a particular workflow is tagged with the corresponding workflow attribute, not all frontend (or browser-related) spans generated for RUM that are associated with a particular user interaction or client process may be tagged with the appropriate workflow attribute. For example, while instrumentation for RUM may allow one or more spans generated in response to a user attempting to load a chart on a website to be tagged with the workflow attribute 2021 (and the associated workflow value 2022 as shown in the example of FIG. 20B), it may not be possible or convenient for the instrumenter to identify all the spans generated in response to the user-initiated activity based on instrumentation alone.

FIG. 23 presents a flowchart 2300 illustrating a computer-implemented process for identifying the entire set of RUM-

related spans that are associated with a particular workflow, in accordance with implementations of the present monitoring service.

At block 2302, a tagged span is identified, wherein the tagged span is tagged with an attribute associated with the particular workflow.

At block 2304, any unassociated span is filtered out. Unassociated spans may be filtered out, for example, by filtering out any span that does not have the same session ID as the tagged span. In other implementations, any spans that do not have the same instance identifier instance ID) may also be filtered out. Filtering out spans that do not have the same instance ID filters out span that may not be part of the same window instance (within a given session) as the tagged span.

At block 2306, the remaining spans are ordered by their start time.

At block 2308, any spans that have a start time occurring between the start time and the end time of the tagged span are designated as associated with the particular workflow.

At block 2310, from the remaining spans, any span that has a start time occurring after the end time of the tagged span is considered associated with the particular workflow provided that the start time of the respective span occurs prior to a predetermined duration of span inactivity. In other words, a span would not be associated with the particular workflow if a predetermined duration of time with no span activity has elapsed prior to the start time of the span. In one implementation, the predetermined duration of time is configurable by a client of the monitoring platform. In a different implementation, the predetermined duration may be a static constant programmed into the monitoring platform.

At block 2312, a determination is made if the predetermined duration has been exceeded. If the predetermined duration has been exceeded, then the process ends at block 2314. If not, the process repeats for each of the remaining spans that were ordered at block 2306.

FIG. 24 illustrates an exemplary manner in which spans related to a particular workflow may be identified and grouped, in accordance with implementations of the monitoring service disclosed herein. As noted above, to start identifying spans that may be related to the particular workflow, implementations of the monitoring service will first filter out any unassociated or unrelated spans. Accordingly, spans that are unrelated to the tagged span (Span A 2402 tagged with a workflow attribute value of “Chartload” in FIG. 24) will first be filtered out, e.g., either because they do not share the same session ID or instance ID as the tagged span or both.

Thereafter, any span in the example of FIG. 24 with a start time between the start time of the tagged span (e.g., Span A 2402) and the end time of the tagged span is designated as related to the “Chartload” workflow. Accordingly, Span B 2404, Span C 2406 and Span D 2408 are designated as related to Span A 2402 and to the “Chartload” workflow. Subsequently, any spans that have a start time occurring after the end time of Span A 2402 are considered associated with the “Chartload” workflow provided that the start time of the respective span occurs prior to a predetermined duration of span inactivity. In the example of FIG. 23, the predetermined duration of span inactivity is 100 ms. Accordingly, Span E 2410 is designated as part of the “Chartload” workflow because the start time of Span E occurs prior to a 100 ms period of inactivity. Because Span E 2410 starts after a 25 ms period of inactivity following Span D 2408, Span E is still considered to be associated with the “Chartload”

workflow. Span F 2412, however, starts more than 100 ms of inactivity following Span 2410. As a result, Span F 2412 is not included in the group of spans associated with the “Chartload” workflow.

In one implementation, other heuristics besides the time-based approach illustrated in FIG. 24 may also be used to associate spans related to a particular workflow. For example, other attributes or tags of a span (besides the workflow attribute) may be compared to determine whether a particular span is related to or associated with a particular workflow.

Implementations of the monitoring service disclosed herein allow a client to define and instrument a “workflow” dimension and to tag one or more spans of interest with the defined workflow attribute. In an implementation, the monitoring service 306 (referenced in FIG. 3) will extract the tags related to the workflow dimension and their respective values from the ingested tagged spans and will analyze the tagged spans (and other associated spans) to determine the paths associated with each respective workflow. The path for a workflow in RUM comprises the various endpoint dependencies that are invoked by a particular user-interaction or client process during a real user session. By allowing a client to define the workflow dimension and track particular spans of interest, implementations of the monitoring service provide a client with more control over conducting dimensional analysis. In other words, clients can slice (or scope) the workflow dimension across other attributes (as discussed in connection with FIGS. 12A and 12B). Further, implementations of the monitoring service allow clients to identify spans related to the tagged spans and to compute metrics for all related spans.

In one implementation, metrics, e.g., a count value, error count, duration, etc. can be aggregated for a particular workflow. The metrics may be computed based on the entire set of spans associated with a particular workflow. Alternatively, the metrics may be computed based on the spans that are explicitly tagged with a particular workflow attribute. Further, any metrics extracted for a real user session can then be filtered to compute metrics associated specifically with a particular workflow (as also discussed in connection with FIGS. 10, 12A and 12B).

In one or more implementations, the monitoring service 306 may be programmed to execute a complex set of rules to extract further information pertaining to the workflows and their associated values from the tags of ingested spans (and their associated spans). Further, pattern matching may also be employed to extract workflow information from ingested spans.

FIG. 25 illustrates an exemplary on-screen RUM monitoring GUI showing an aggregate view of instrumented workflows, in accordance with implementations of the monitoring service disclosed herein. As shown in FIG. 25, the GUI may present an aggregate view over many different signals (associated with user interactions). For example, the GUI may provide information about document load latencies 2525 and endpoint latencies 2526.

The GUI may also include information about instrumented workflows. For example, the GUI may display the top workflows 2510 (sorted by duration or error) indicating which workflows need the most attention. Another section in the GUI may display the various workflow categories 2529. Further, other sections of the GUI may also illustrate workflow errors 2512 and workflow performance 2515. In one implementation, each of the workflows shown in the GUI would be displayed next to a count, error and latency (or duration) metric aggregated over a selected time period (e.g.,

over the course of a single session). Further, the GUI may allow users the ability to break down the workflows across multiple different attributes using a drop down menu (similar to the manner discussed in FIGS. 10, 12A and 12B). The computations for each of the break-downs may be efficiently determined using the metrics data aggregated for the metric events mode.

In one implementation, each of the workflow names in the various sections shown in FIG. 25 may be linked to another GUI illustrating further details about a selected workflow. FIG. 26 illustrates an exemplary on-screen GUI showing a detailed view of a selected workflow, in accordance with implementations of the monitoring service disclosed herein. Because “workflow” is a tag that can be scoped or filtered on (similar to a page or resource/endpoint), the monitoring platform can support the selecting of a workflow and displaying the associated error and latency metrics for the given workflow. As shown in FIG. 26, panel 2610 illustrates a latency graph for a selected workflow.

Further, as mentioned previously, the GUI may allow users the ability to break down the workflows across multiple different attributes (or dimensions). In other words, the aggregated workflow metrics may be used by the monitoring platform for dimensional analysis. For example, panel 2620 allows clients to break down the behavior of the workflow across several different dimensions, e.g., application, component, operation, name, type, browser, etc. The workflow metrics may also be scoped across other dimensions, e.g., current page, geographical location, operating system, etc.

FIG. 27 illustrates an exemplary on-screen GUI showing an aggregated waterfall view of an instrumented workflow, in accordance with implementations of the monitoring service disclosed herein. While panel 2710 displays the latency graph for the selected workflow (similar to FIG. 26), selecting the waterfall view in panel 2720 displays exemplary spans associated with the workflow. Note that the metric events modality discussed in connection with FIG. 17 allows for exemplary RUM spans and traces to be stored. Displaying exemplary spans in the panel 2720 allows clients to better understand, at an aggregate level, the general type of work performed within the workflow.

In one implementation, the workflow analysis is performed on RUM spans, but where specific connections can be established to backend traces, the panel 2720 may contain links to backend APM traces so that a client can perform deeper end-to-end analysis. As mentioned in connection with FIG. 20A, where a frontend span associated with the user session can be linked to a backend APM trace, an APM icon (e.g., similar to the APM icon 2051) can be provided next to the frontend span in the waterfall view to indicate that linkage information exists for the respective span. The linked trace is a backend representation of a user request for the resource associated with the respective span. A client may then be able to expand the span in the GUI to retrieve the link to the backend trace, where the link is provided as one of the attributes of the respective span. In one implementation, the panel 2720 would allow the client to view the exemplary APM trace for any backend calls performed during the selected workflow (metrics for which are shown in the panel 2710). In one implementation, the linkage may also be workflow-specific. In other words, RUM workflows on the frontend may be linked to APM workflows on the backend.

FIG. 28 illustrates an exemplary on-screen GUI of the manner in which associated spans from an individual session may be grouped into workflows, in accordance with implementations of the monitoring service disclosed herein. FIG.

27 illustrates events from a user session that may be displayed in a panel adjacent to a graphical view of the user session, e.g., the panel 2050 shown in FIG. 20A. In one implementation, when viewing spans associated with a particular user session, the monitoring service groups associated spans by workflow so that it is easier to identify associations and patterns in the span data. As shown in FIG. 27, spans 2810 associated with a given workflow may be highlighted for better visual analysis.

FIG. 29 presents a flowchart 2900 illustrating a computer-implemented process for aggregating metrics associated with a user interaction during a real user session, in accordance with implementations of the present monitoring service.

At block 2902, a tagged span associated with a given workflow is identified from a plurality of spans ingested during the real user session. The tagged span, in one implementation, may be generated as a result of selective instrumentation by a client to include a workflow attribute associated with the given workflow in the span.

At block 2904, the monitoring platform identifies one or more other spans generated during the real user session that are associated with the tagged span. One exemplary method of identifying associated spans was discussed in connection with FIG. 23.

At block 2906, the monitoring platform groups together the tagged span and the one or more other spans as relating to the given workflow.

At block 2908, metrics are aggregated for the given workflow. In one implementation, the metrics may be aggregated using metadata extracted from the tagged span over a given duration of time. In another implementation, metrics can be aggregated from both the tagged spans and the one or more spans identified as associated spans.

The disclosed system advantageously addresses a problem in traditional data analysis of instrumented software tied to computer technology, namely, the technical problem of identifying spans and aggregating metrics that are associated with a given user interaction or client process generated during a real user session. The disclosed system advantageously solves this technical problem by providing a solution also rooted in computer technology, namely, by associating a span tagged with the workflow attribute with other spans related to the user interaction (or client process) that may not be tagged with the particular workflow attribute. Further, the disclosed system advantageously allows metrics to be computed from the entire set of spans associated with a particular workflow. The disclosed subject technology further provides improvements to the functioning aspects of the computer itself by identifying and organizing associated RUM-related spans for secondary analysis into workflows at query time.

While the principles of the invention have been described above in connection with specific apparatus, it is to be clearly understood that this description is made only by way of example and not as a limitation on the scope of the invention. Further, the foregoing description, for purpose of explanation, has been described with reference to specific implementations. However, the illustrative discussions above are not intended to be exhaustive or to limit the invention to the precise forms disclosed. Many modifications and variations are possible in view of the above teachings. The implementations were chosen and described in order to best explain the principles of the invention and its practical applications, to thereby enable others skilled in the

65

art to best utilize the invention and various implementations with various modifications as may be suited to the particular use contemplated.

What is claimed is:

1. A method of aggregating metrics associated with a user interaction during a real user session, the method comprising:

identifying a tagged span from a plurality of ingested spans associated with the real user session, wherein the tagged span is associated with a workflow, wherein further the workflow comprises spans generated in response to the user interaction, and wherein the tagged span comprises a tag associated with the workflow;

identifying at least one associated span from the plurality of ingested spans, wherein the at least one associated span is associated with the workflow;

consolidating the at least one associated span with the tagged span into a group of spans associated with the workflow; and

computing metrics for the workflow over a duration of time using the group of spans.

2. The method of claim 1, wherein the computing comprises aggregating metrics extracted from the tagged span in the group of spans over the duration of time.

3. The method of claim 1, wherein the computing comprises aggregating metrics extracted from the tagged span and the at least one associated span in the group of spans over the duration of time.

4. The method of claim 1, further comprising:

displaying the at least one associated span and the tagged span as grouped together in a graphical user interface.

5. The method of claim 1, wherein the tagged span is selectively instrumented responsive to client input to include the tag associated with the workflow.

6. The method of claim 1, wherein the identifying the at least one associated span from the plurality of ingested spans comprises:

filtering out spans from the plurality of ingested spans that are unassociated with the workflow;

ordering remaining spans from the plurality of ingested spans in accordance with respective start times of the remaining spans;

from the remaining spans: a) designating a span as an associated span if it has a start time occurring between a start time and an end time of the tagged span; and b)

designating a span as an associated span if it has a start time occurring after the end time of the tagged span responsive to a determination that a start time of the span occurs following the end time of the tagged span but prior to a predetermined duration of span inactivity.

7. The method of claim 1, wherein the identifying the at least one associated span from the plurality of ingested spans comprises:

filtering out spans from the plurality of ingested spans that are unassociated with the workflow by filtering out spans comprising a different session identifier than the tagged span;

ordering remaining spans from the plurality of ingested spans in accordance with respective start times of the remaining spans;

from the remaining spans: a) designating a span as an associated span if it has a start time occurring between a start time and an end time of the tagged span; and b) designating a span as an associated span if it has a start time occurring after the end time of the tagged span responsive to a determination that a start time of the

66

span occurs following the end time of the tagged span but prior to a predetermined duration of span inactivity.

8. The method of claim 1, wherein the identifying the at least one associated span from the plurality of ingested spans comprises:

filtering out spans from the plurality of ingested spans that are unassociated with the workflow by filtering out spans comprising a different session identifier and a different instance identifier than the tagged span; ordering remaining spans from the plurality of ingested spans in accordance with respective start times of the remaining spans;

from the remaining spans: a) designating a span as an associated span if it has a start time occurring between a start time and an end time of the tagged span; and b) designating a span as an associated span if it has a start time occurring after the end time of the tagged span responsive to a determination that a start time of the span occurs following the end time of the tagged span but prior to a predetermined duration of span inactivity.

9. The method of claim 1, wherein the identifying the at least one associated span from the plurality of ingested spans comprises:

filtering out spans from the plurality of ingested spans that are unassociated with the workflow; ordering remaining spans from the plurality of ingested spans in accordance with respective start times of the remaining spans;

from the remaining spans: a) designating a span as an associated span if it has a start time occurring between a start time and an end time of the tagged span; and b) designating a span as an associated span if it has a start time occurring after the end time of the tagged span responsive to a determination that a start time of the span occurs following the end time of the tagged span but prior to a predetermined duration of span inactivity, wherein the predetermined duration is configurable by a client.

10. The method of claim 1, further comprising: scoping the metrics across other attributes associated with the plurality of ingested spans.

11. The method of claim 1, wherein the metrics comprise a workflow count value, an error count and latency values associated with the workflow.

12. The method of claim 1, further comprising: displaying the metrics in a graphical user interface.

13. The method of claim 1, further comprising: displaying the at least one associated span and the tagged span in a graphical user interface; and displaying links to a backend trace associated with the at least one associated span or the tagged span in the graphical user interface.

14. The method of claim 1, further comprising: displaying the at least one associated span and the tagged span in a graphical user interface; and displaying a backend trace associated with the at least one associated span or the tagged span in the graphical user interface.

15. The method of claim 1, wherein the identifying the at least one associated span from the plurality of ingested spans comprises:

filtering out spans from the plurality of ingested spans that are unassociated with the workflow; from remaining spans of the plurality of spans, identifying the at least one associated span by comparing tags of the remaining spans with the tagged span.

67

16. A non-transitory computer-readable medium having computer-readable program code embodied therein for causing a computer system to perform a method of aggregating metrics associated with a user interaction during a real user session, the method comprising:

identifying a tagged span from a plurality of ingested spans associated with the real user session, wherein the tagged span is associated with a workflow, wherein further the workflow comprises spans generated in response to the user interaction, and wherein the tagged span comprises a tag associated with the workflow;

identifying at least one associated span from the plurality of ingested spans, wherein the at least one associated span is associated with the workflow;

consolidating the at least one associated span with the 15 tagged span into a group of spans associated with the workflow; and

computing metrics for the workflow over a duration of time using the group of spans.

17. The non-transitory computer-readable medium of claim **16**, wherein the computing comprises aggregating metrics extracted from the tagged span in the group of spans over the duration of time.

18. The non-transitory computer-readable medium of claim **16**, wherein the computing comprises aggregating 25 metrics extracted from the tagged span and the at least one associated span in the group of spans over the duration of time.

19. The non-transitory computer-readable medium of claim **16**, wherein the identifying the at least one associated 30 span from the plurality of ingested spans comprises:

filtering out spans from the plurality of ingested spans that are unassociated with the workflow;

68

ordering remaining spans from the plurality of ingested spans in accordance with respective start times of the remaining spans;

from the remaining spans: a) designating a span as an associated span if it has a start time occurring between a start time and an end time of the tagged span; and b) designating a span as an associated span if it has a start time occurring after the end time of the tagged span responsive to a determination that a start time of the span occurs following the end time of the tagged span but prior to a predetermined duration of span inactivity.

20. A system for performing a method of aggregating metrics associated with a user interaction during a real user session, the system comprising:

a processing device communicatively coupled with a memory and configured to:

identify a tagged span from a plurality of ingested spans associated with the real user session, wherein the tagged span is associated with a workflow, wherein further the workflow comprises spans generated in response to the user interaction, and wherein the tagged span comprises a tag associated with the workflow;

identify at least one associated span from the plurality of ingested spans, wherein the at least one associated span is associated with the workflow;

consolidate the at least one associated span with the tagged span into a group of spans associated with the workflow; and

compute metrics for the workflow over a duration of time using the group of spans.

* * * * *